

# QuantumNAS: Noise-Adaptive Search for Robust Quantum Circuits

Hanrui Wang<sup>1</sup>, Yongshan Ding<sup>2</sup>, Jiaqi Gu<sup>3</sup>, Yujun Lin<sup>1</sup>, David Pan<sup>3</sup>, Fred Chong<sup>2</sup>, Song Han<sup>1</sup>

<sup>1</sup>MIT, <sup>2</sup>University of Chicago, <sup>3</sup>University of Texas at Austin

# Outline

- Introduction
- Background
- QuantumNAS
  - SuperCircuit Training
  - Noise-Adaptive Co-Search
  - Searched SubCircuit Training
  - Gate Pruning
- Evaluation
- Conclusion

# Outline

- Introduction
  - Background
  - QuantumNAS
    - SuperCircuit Training
    - Noise-Adaptive Co-Search
    - Searched SubCircuit Training
    - Gate Pruning
  - Evaluation
  - Conclusion

# Introduction

- Quantum computing has potentials to bring exponential advantages
- Currently: NISQ era
  - Insufficient qubits for error correction
  - Imperfect qubits: quantum noise forms the bottleneck

# Introduction

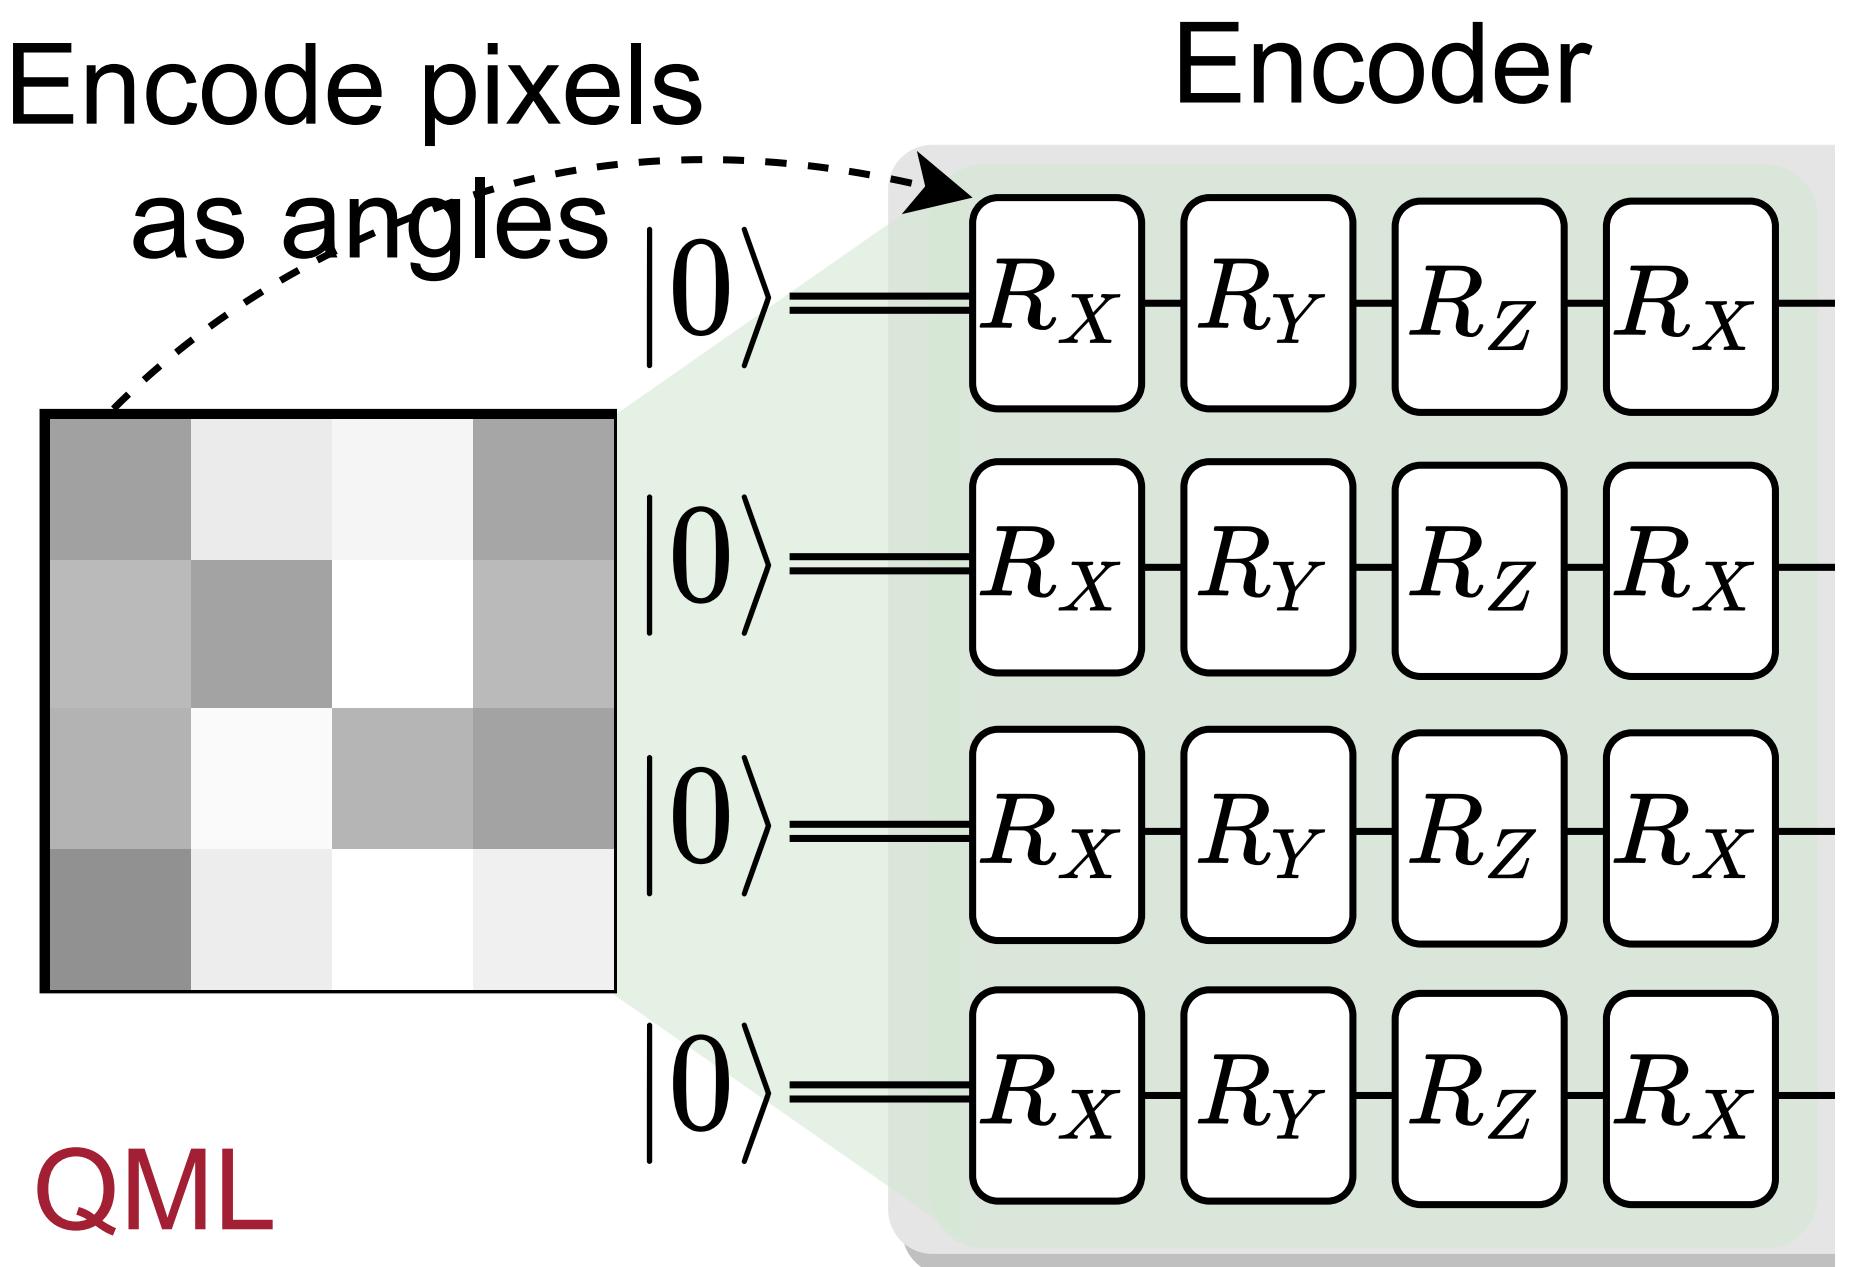
- Existing works focus on noise-adaptive quantum compilation:
  - Noise adaptive qubit mapping
  - Noise adaptive instruction scheduling
  - Noise adaptive crosstalk mitigation
  - ...
- They are optimizing the compilation process with a **fixed** quantum circuit
- QuantumNAS:
  - Target at **parameterized** quantum circuits
  - **Noise-adaptive co-optimization** of compilation (qubit mapping) and circuit (Ansatz)

# Outline

- Introduction
- Background
- QuantumNAS
  - SuperCircuit Training
  - Noise-Adaptive Co-Search
  - Searched SubCircuit Training
  - Gate Pruning
- Evaluation
- Conclusion

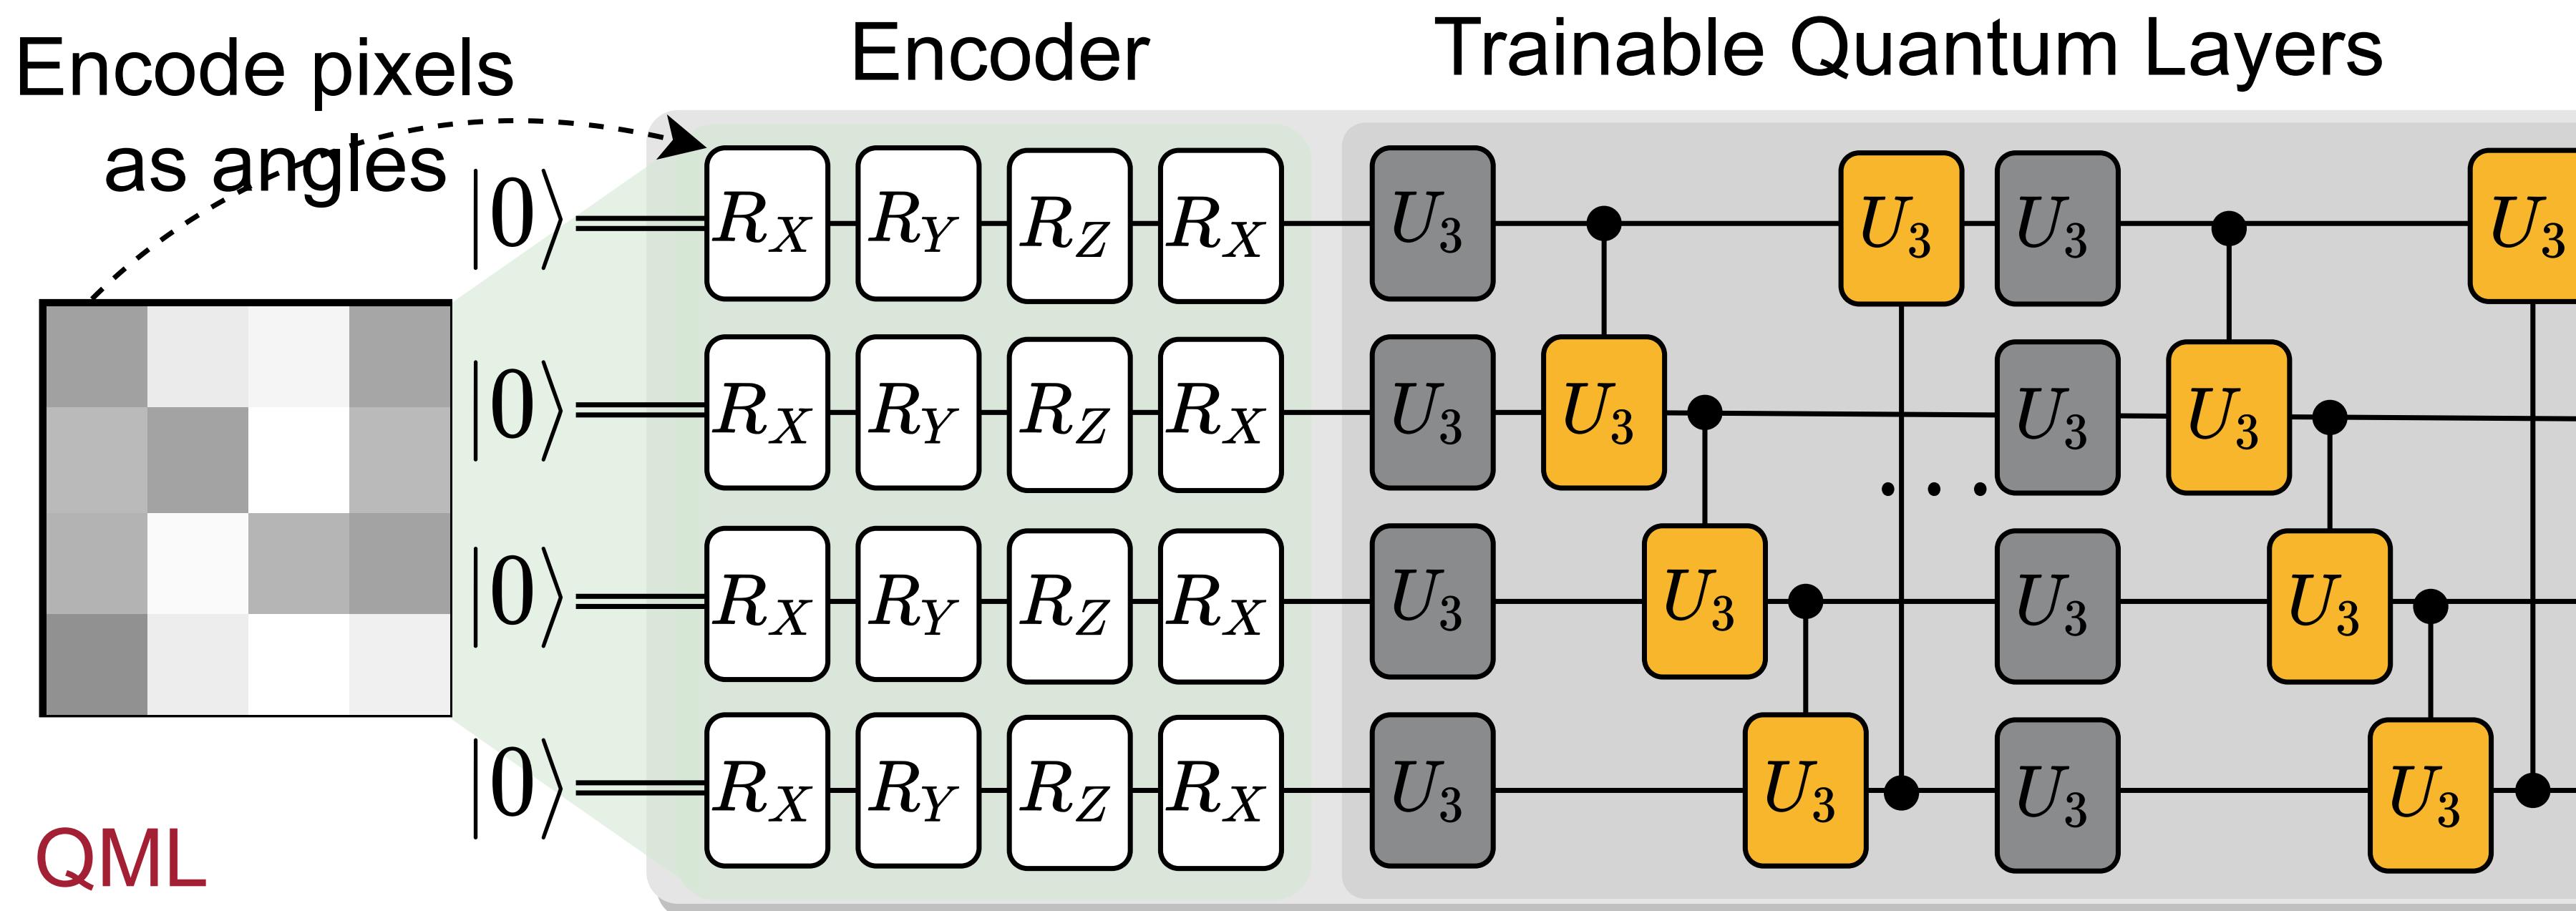
# Background

- Parameterized quantum circuits
  - Quantum Neural Networks



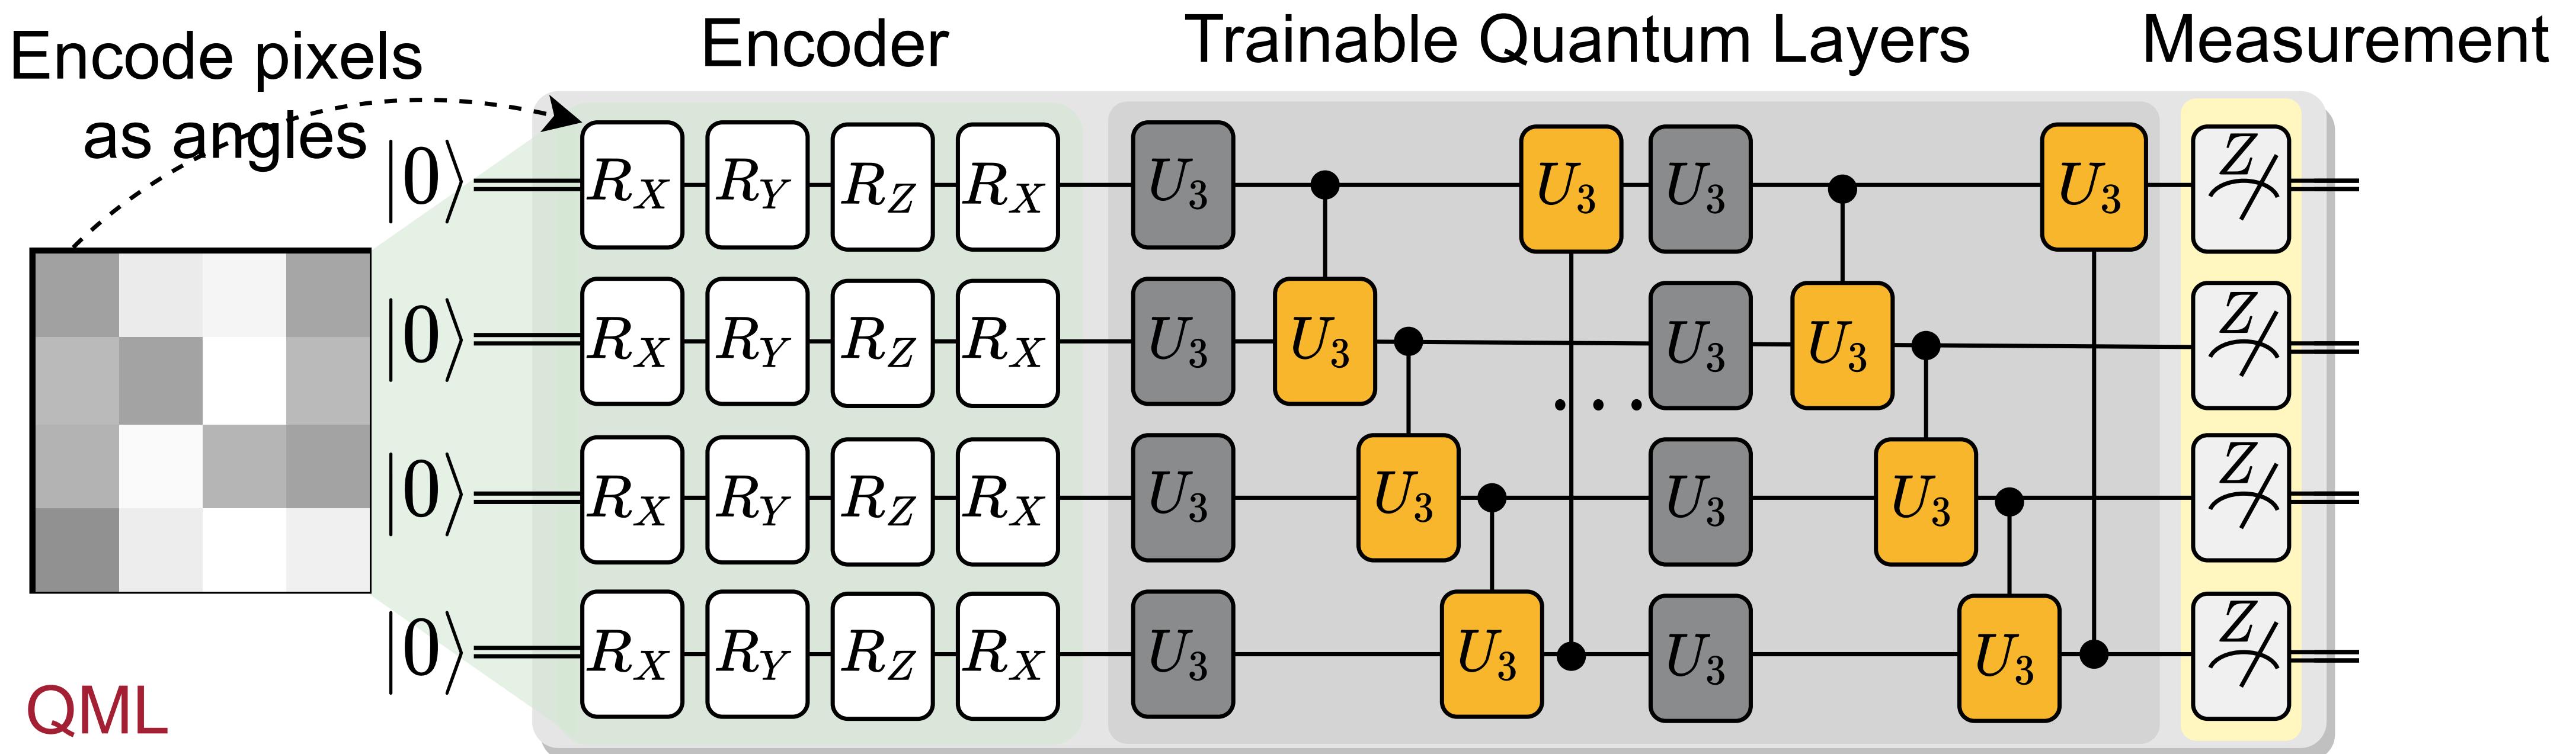
# Background

- Parameterized quantum circuits
  - Quantum Neural Networks



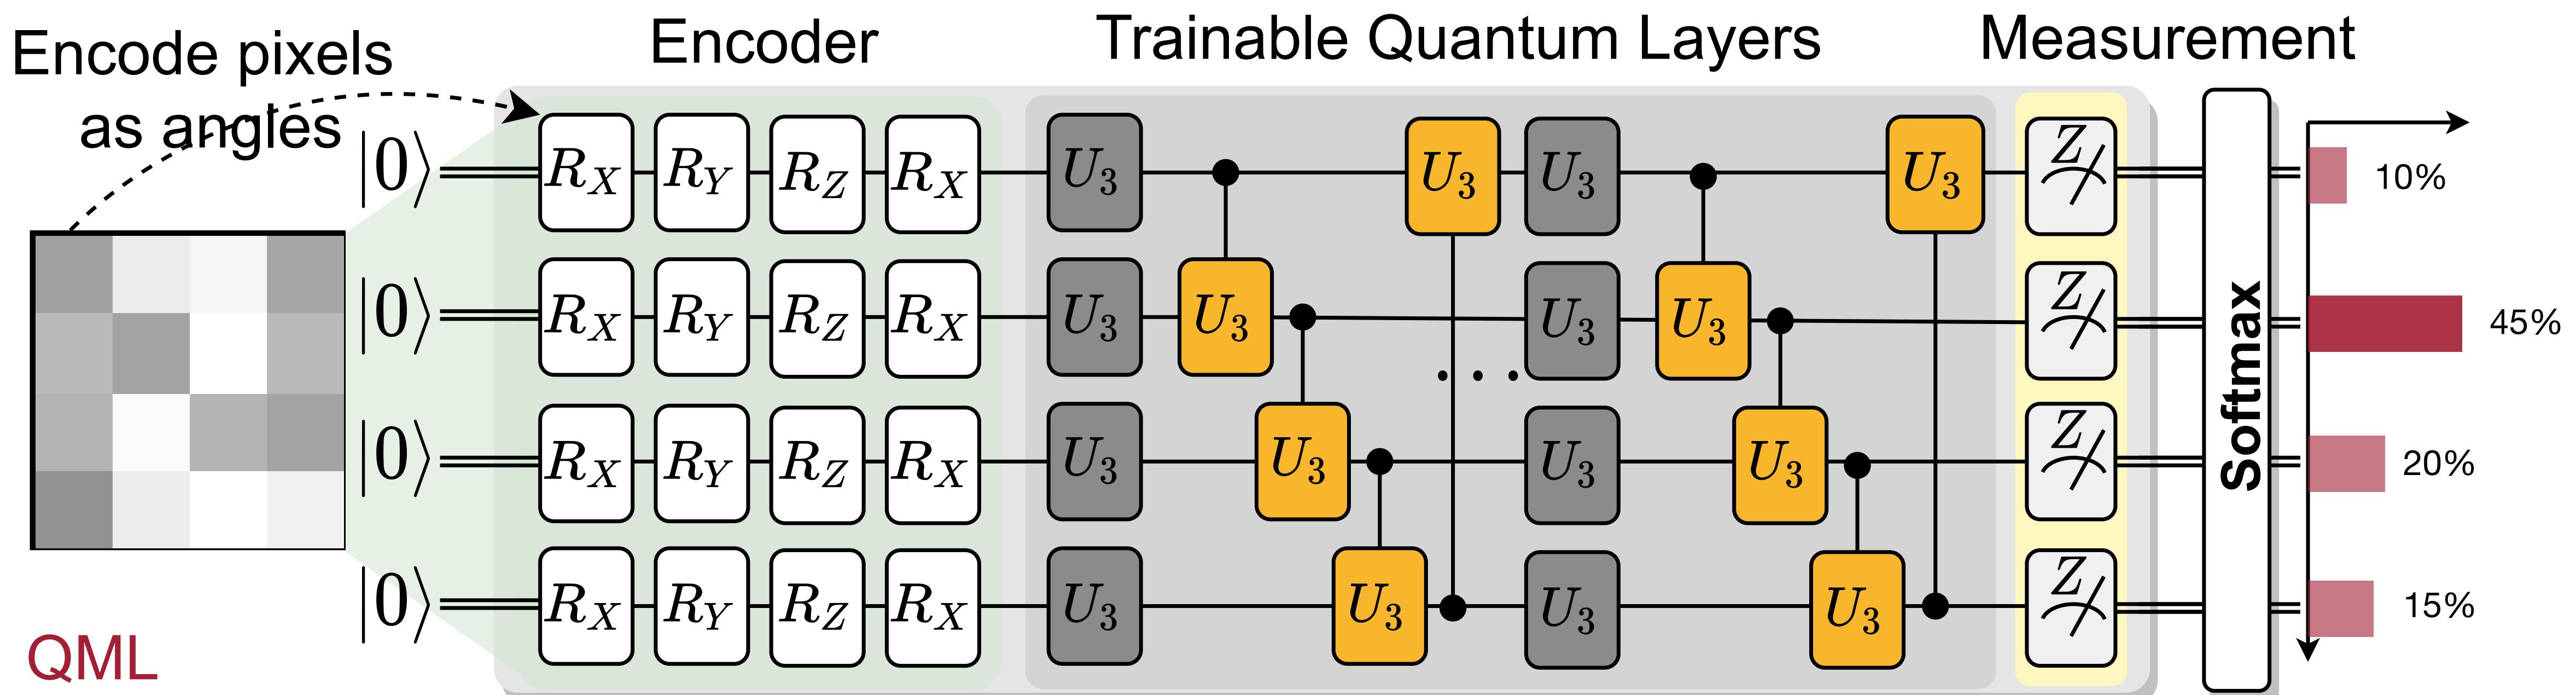
# Background

- Parameterized quantum circuits
  - Quantum Neural Networks



# Background

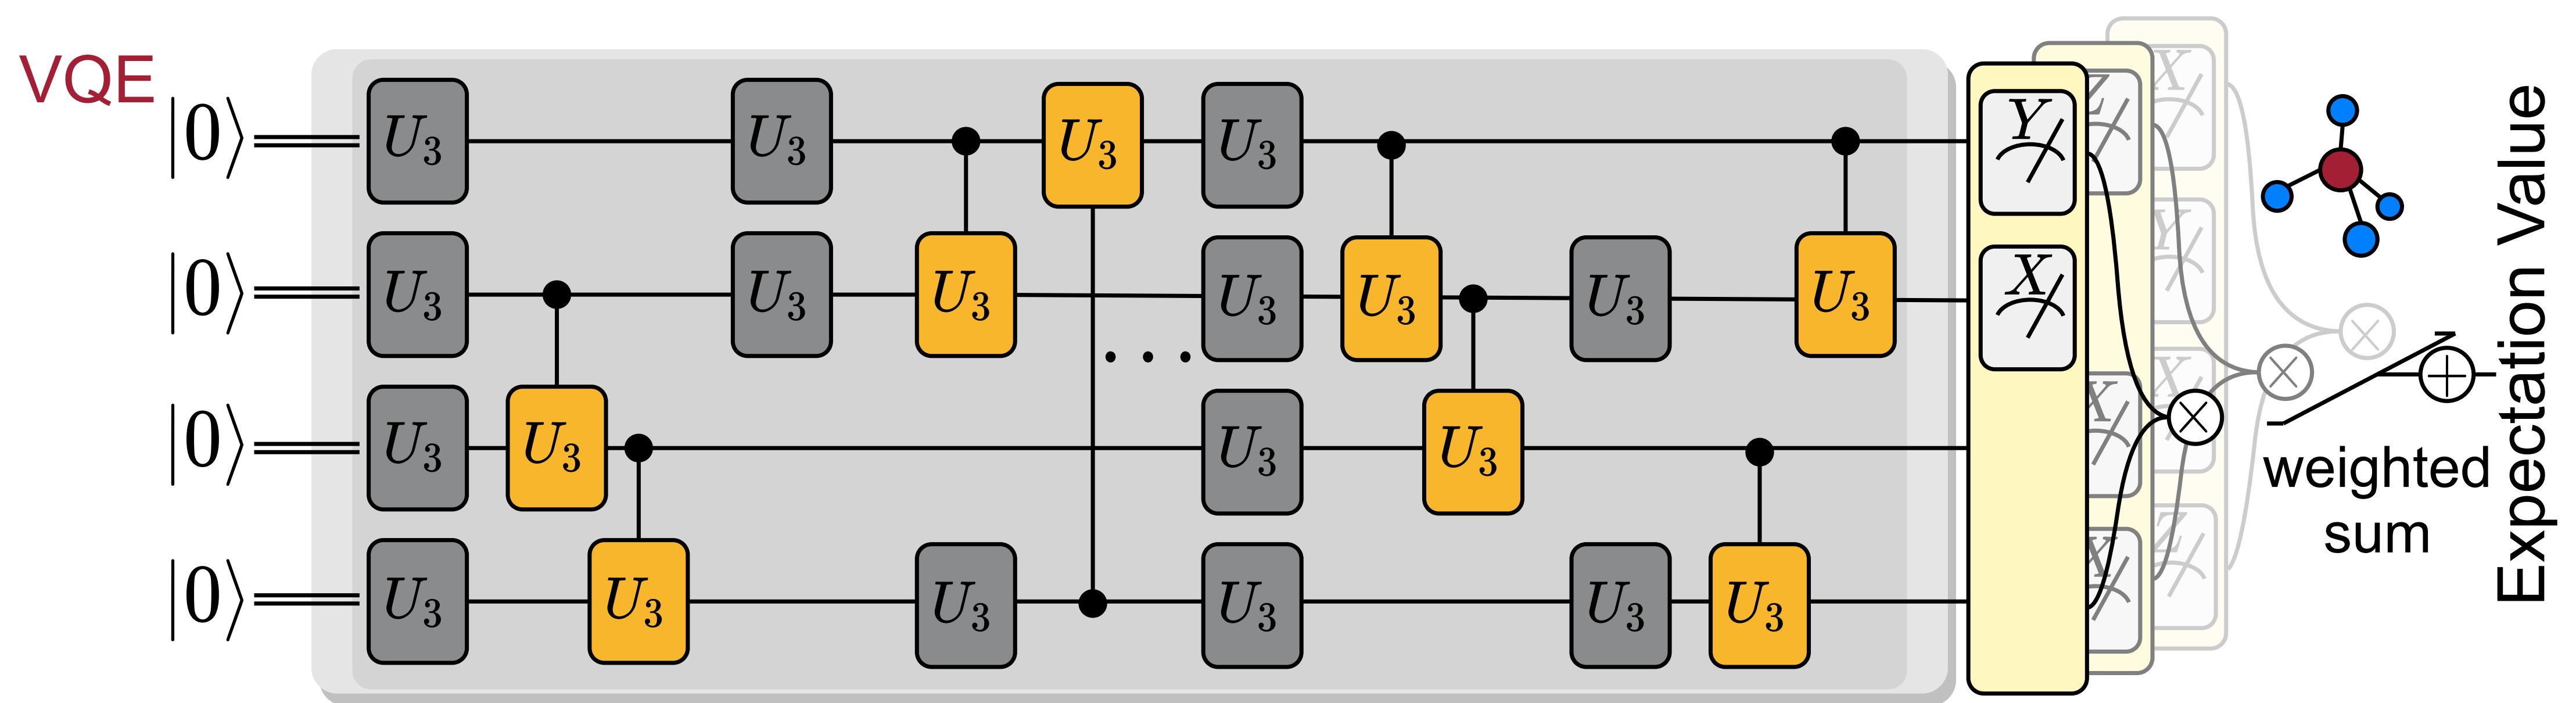
- Parameterized quantum circuits
    - Quantum Neural Networks



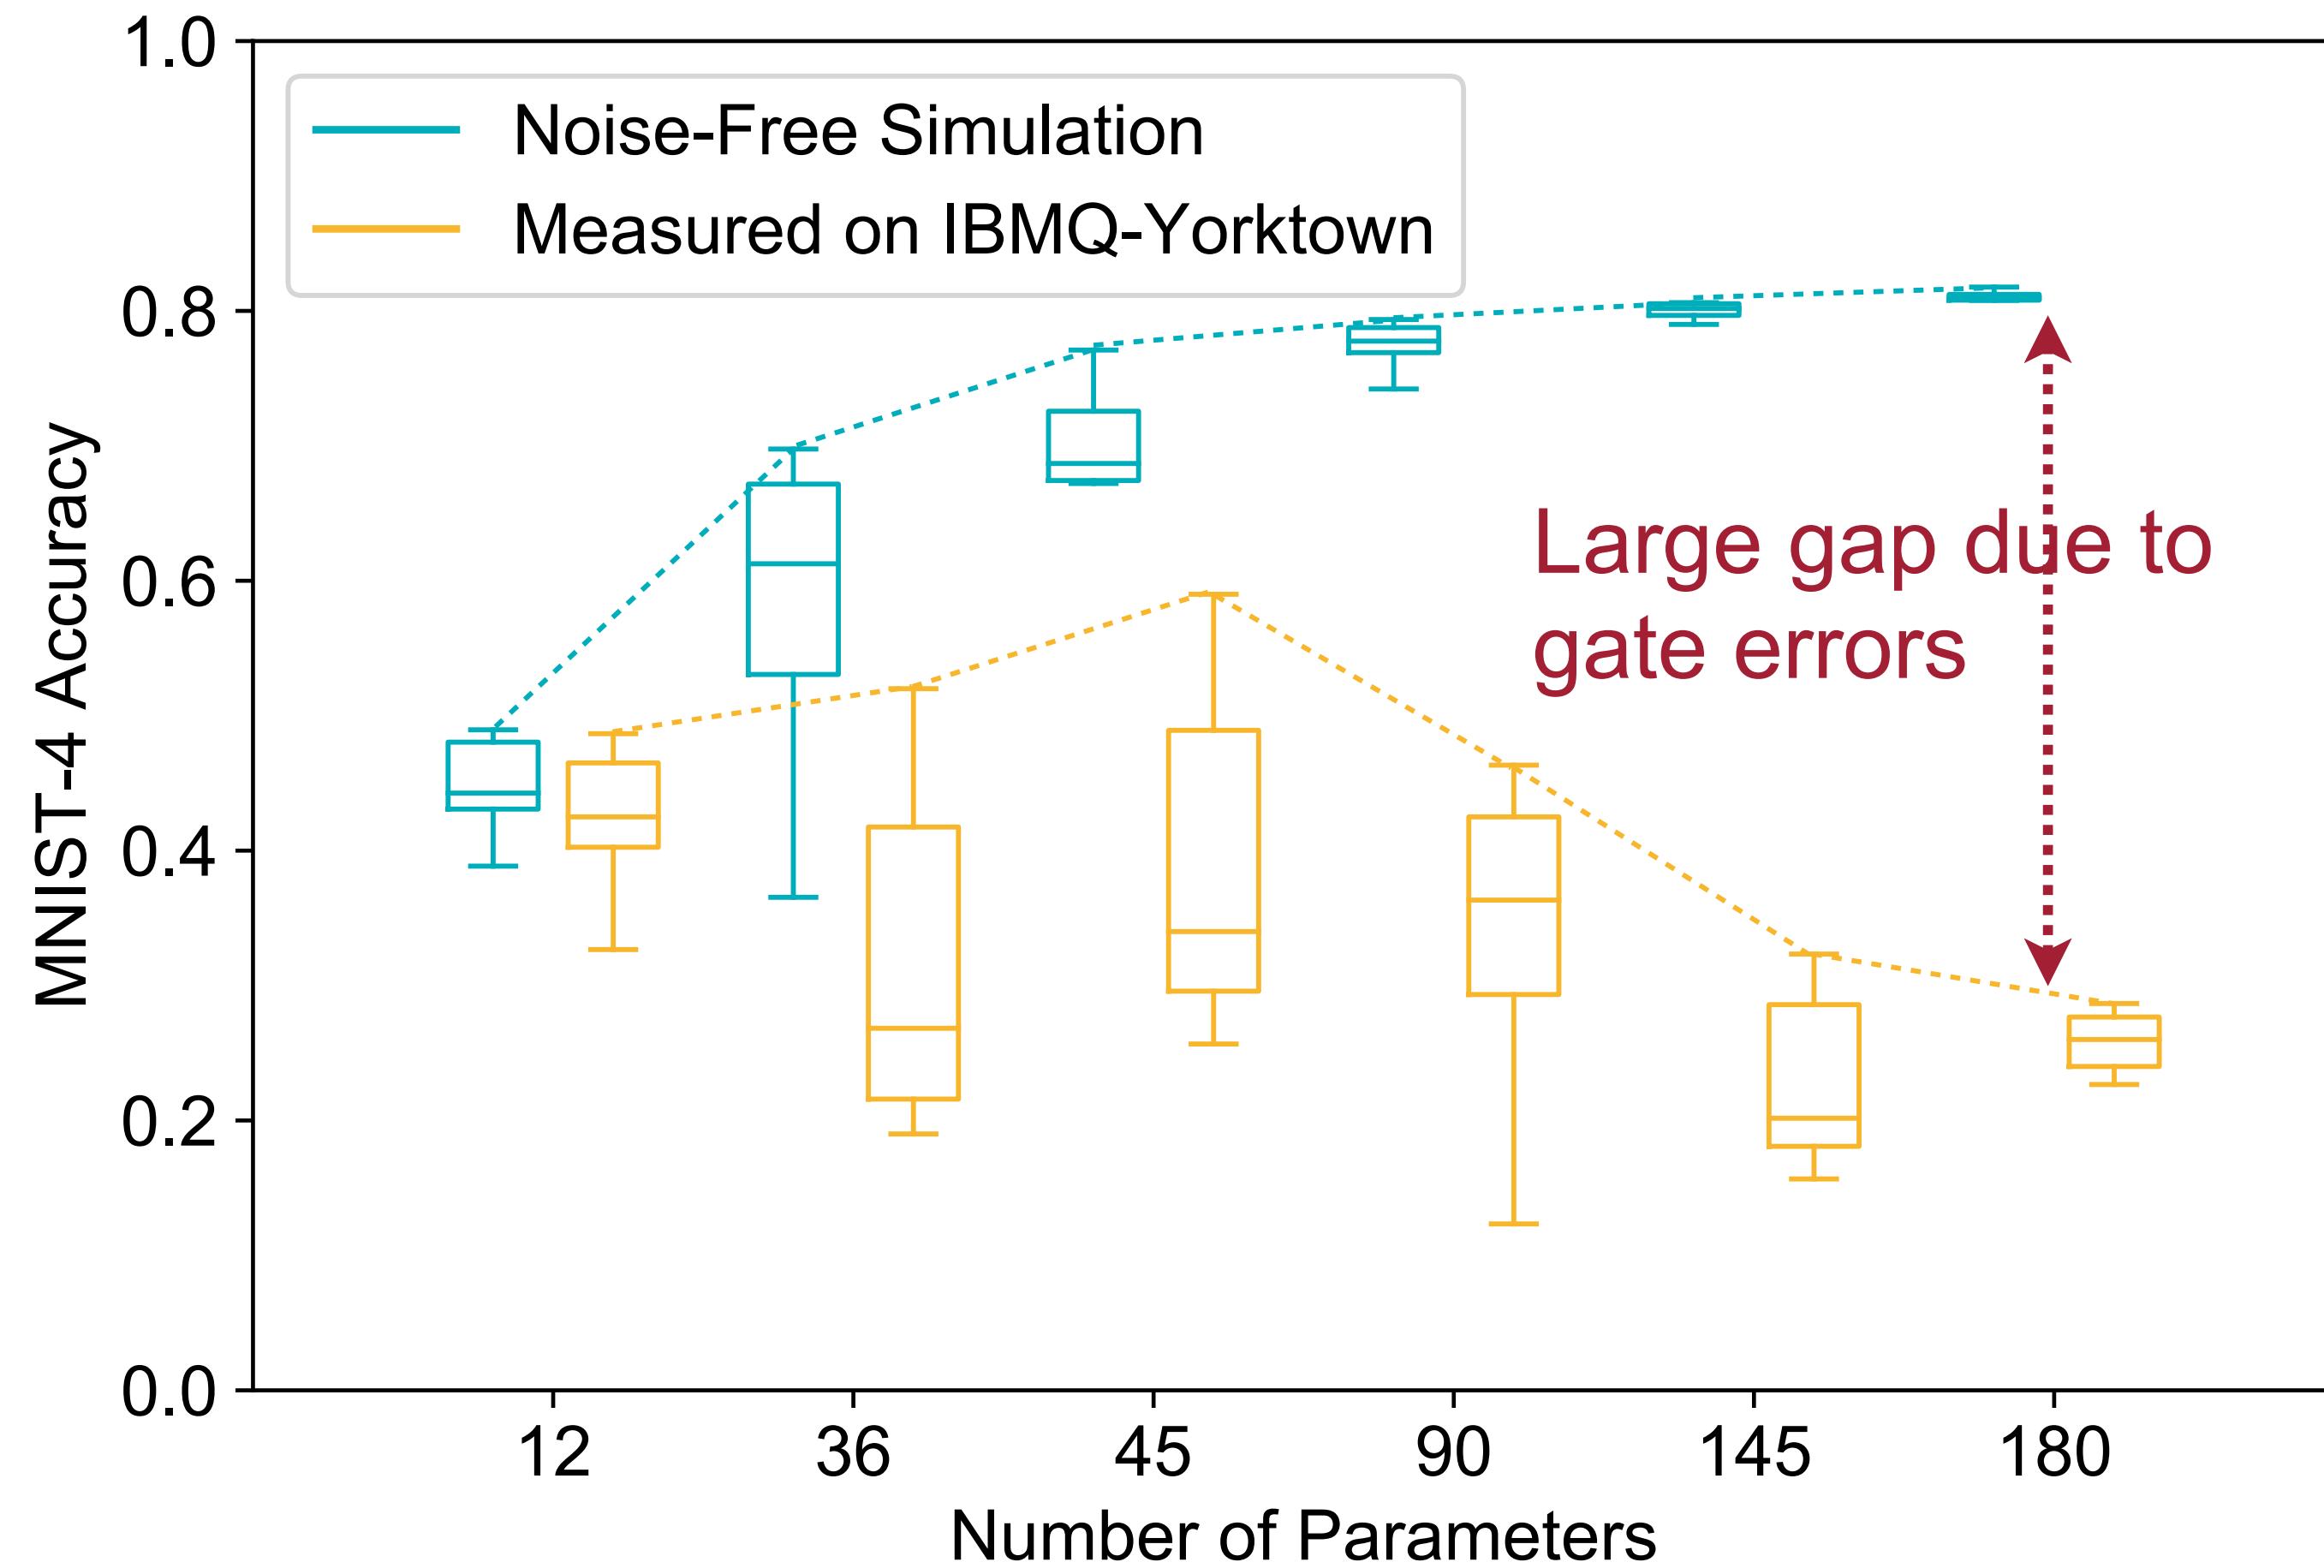
- Softmax 
$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

# Background

- Parameterized quantum circuits
  - Variational Quantum Eigensolver



# Background



- More parameters increase the noise-free accuracy but degrade measured accuracy
- Under same #parameters, measured accuracy of different circuit ansatzes varies a lot

# Background

- How to find the most robust circuit?
- Naive search for robust circuit
  - In one search step: sample many circuit ansatzes, train each to get **final** parameters, compare performance
  - Too costly
- QuantumNAS aims to find robust circuit in a more **scalable** way

# QuantumNAS: Decouple the Training and Search

## Naive Search

```
For q_devices:  
  For search episodes: // meta controller  
    For circuit training iterations:  
      update_parameters(); Expensive  
      If good_circuit: break;
```

# QuantumNAS: Decouple the Training and Search

Naive Search

```
For q_devices:  
  For search episodes: // meta controller  
    For circuit training iterations:  
      update_parameters(); Expensive  
    If good_circuit: break;
```

=>

QuantumNAS

```
For SuperCircuit training iterations: Expensive  
  update_parameters(); training  
  .....  
  decouple  
  For q_devices:  
    For search episodes:  
      sample from SuperCircuit; Light-Weight  
      If good_circuit: break;  
      //no training
```

# QuantumNAS

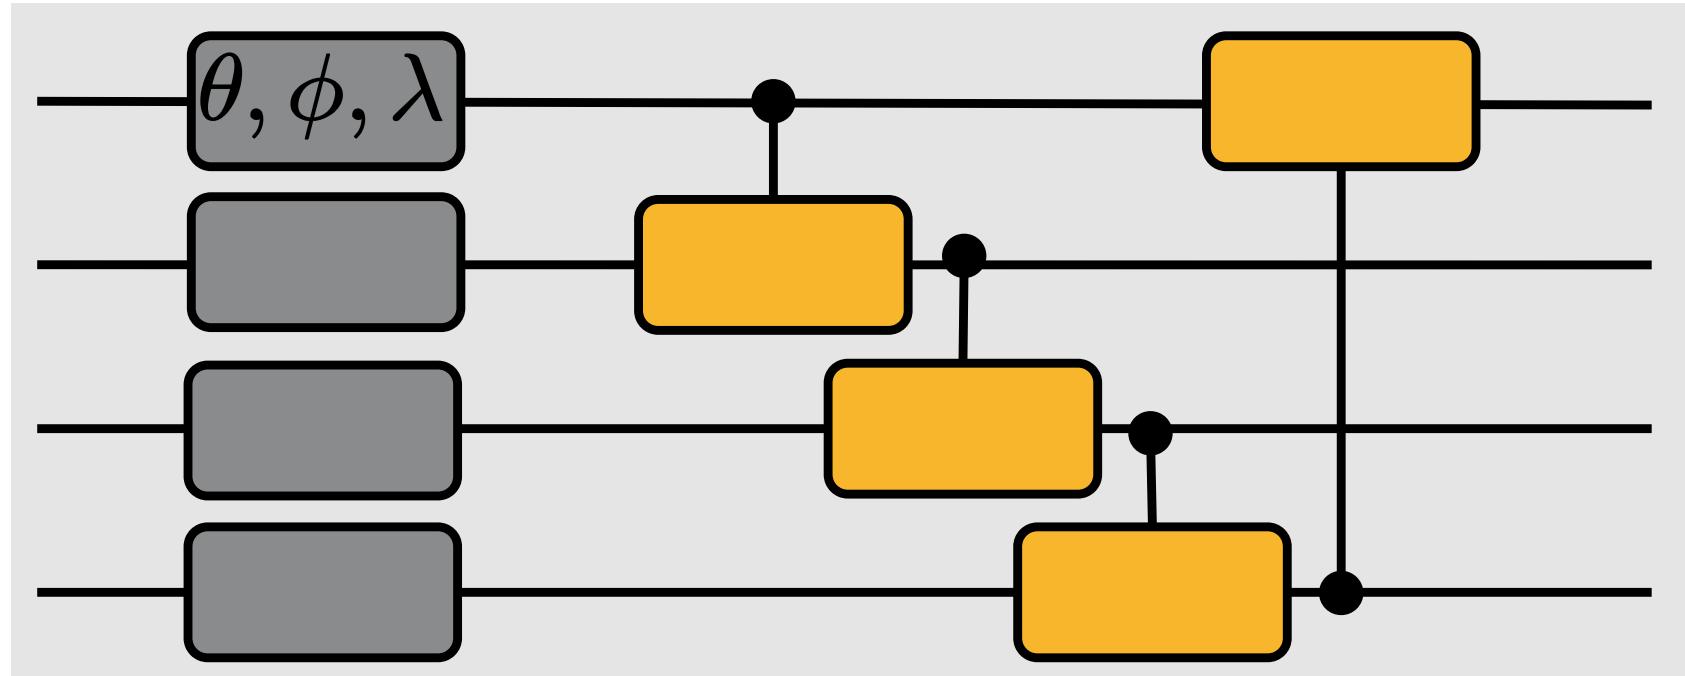
- SuperCircuit Construction and Training
- Noise-Adaptive Evolutionary Co-Search of SubCircuit and Qubit Mapping
- Searched SubCircuit Training
- Iterative Pruning

# QuantumNAS

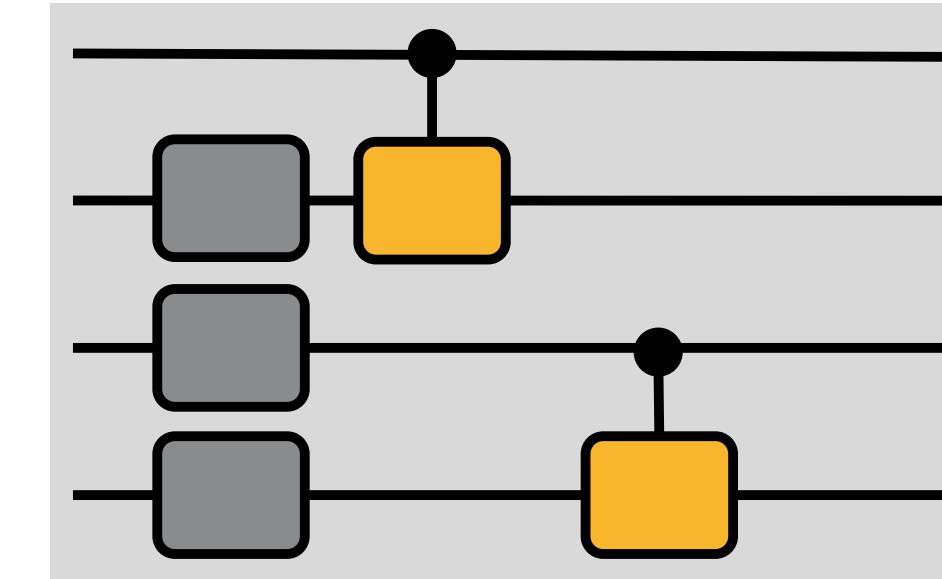
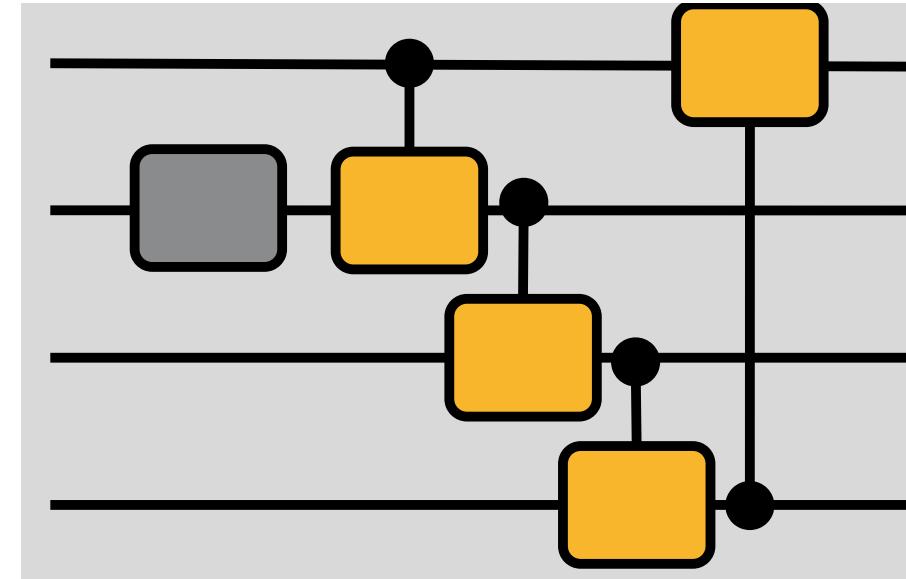
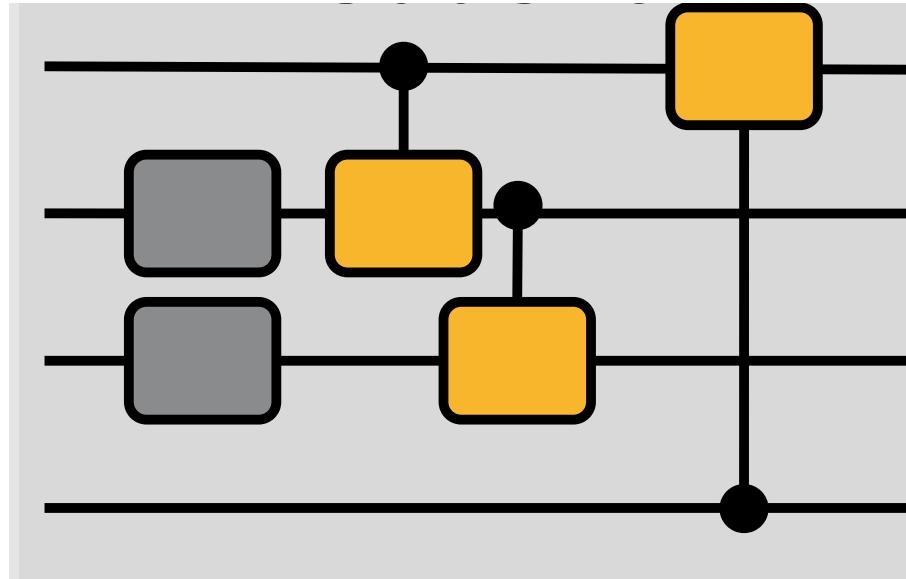
- SuperCircuit Construction and Training
  - Noise-Adaptive Evolutionary Co-Search of SubCircuit and Qubit Mapping
  - Searched SubCircuit Training
  - Iterative Pruning

# SuperCircuit Construction

- Why construct a SuperCircuit?
  - Use it for efficient search of circuit candidates with no need to train them to final
- SuperCircuit: the circuit with the largest number of gates in the design space
  - Example: SuperCircuit in U3+CU3 space

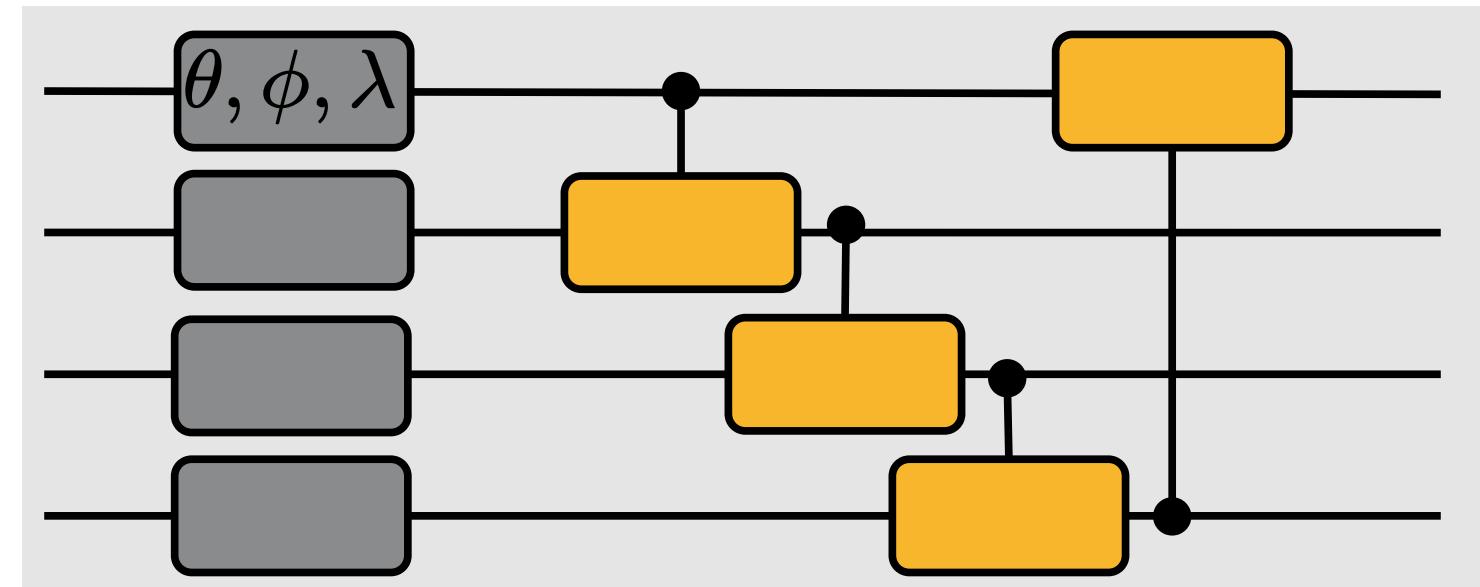


- Each candidate circuit in the design space (called SubCircuit) is a subset of the SuperCircuit

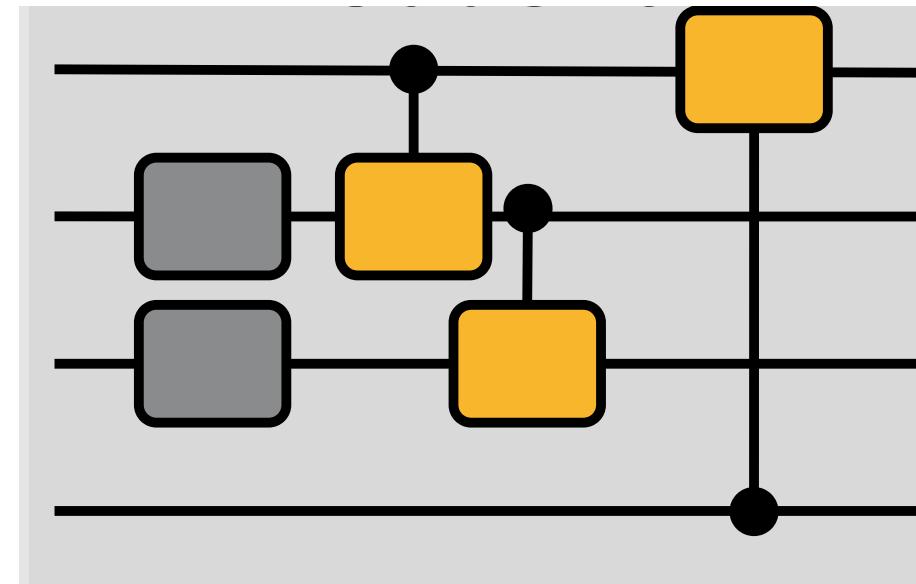


# SuperCircuit Training

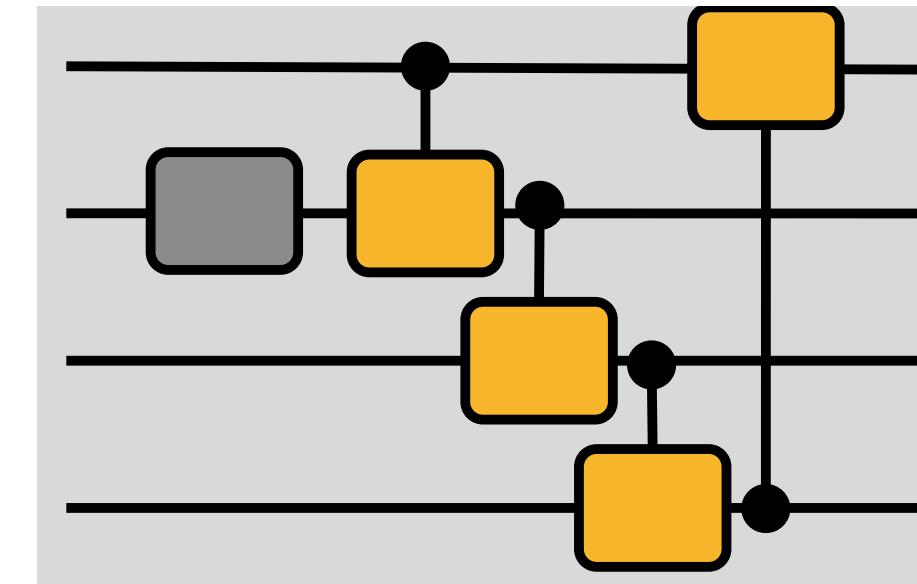
- In one SuperCircuit Training step:
  - Sample a gate subset of SuperCircuit (a SubCircuit)
    - Front Sampling, Progressive Shrinking and Restricted Sampling
  - Only use the subset to perform task and updates the parameters in the subset
  - Parameter updates are cumulative across steps



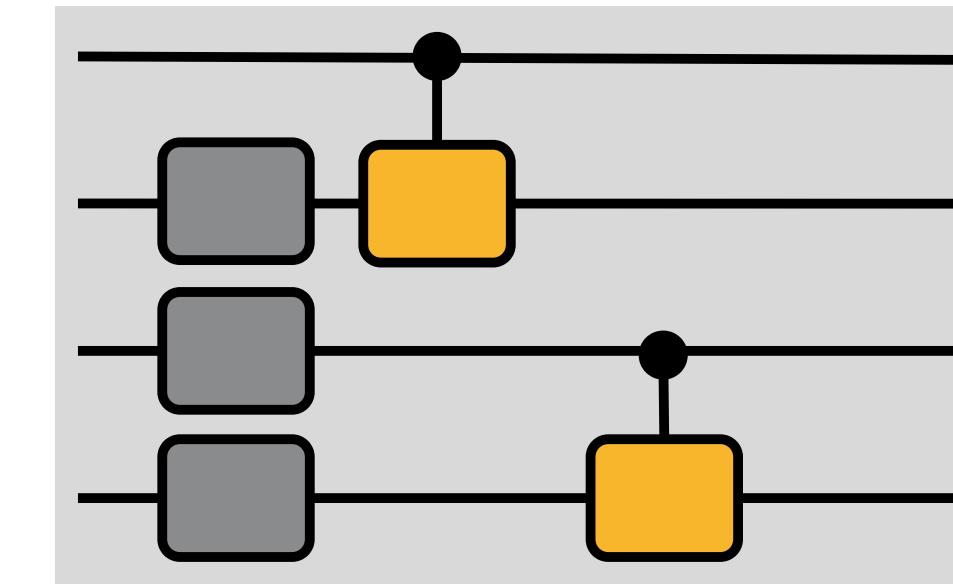
- Supercircuit:
- Sampled subsets (SubCircuits of different steps)



Step1



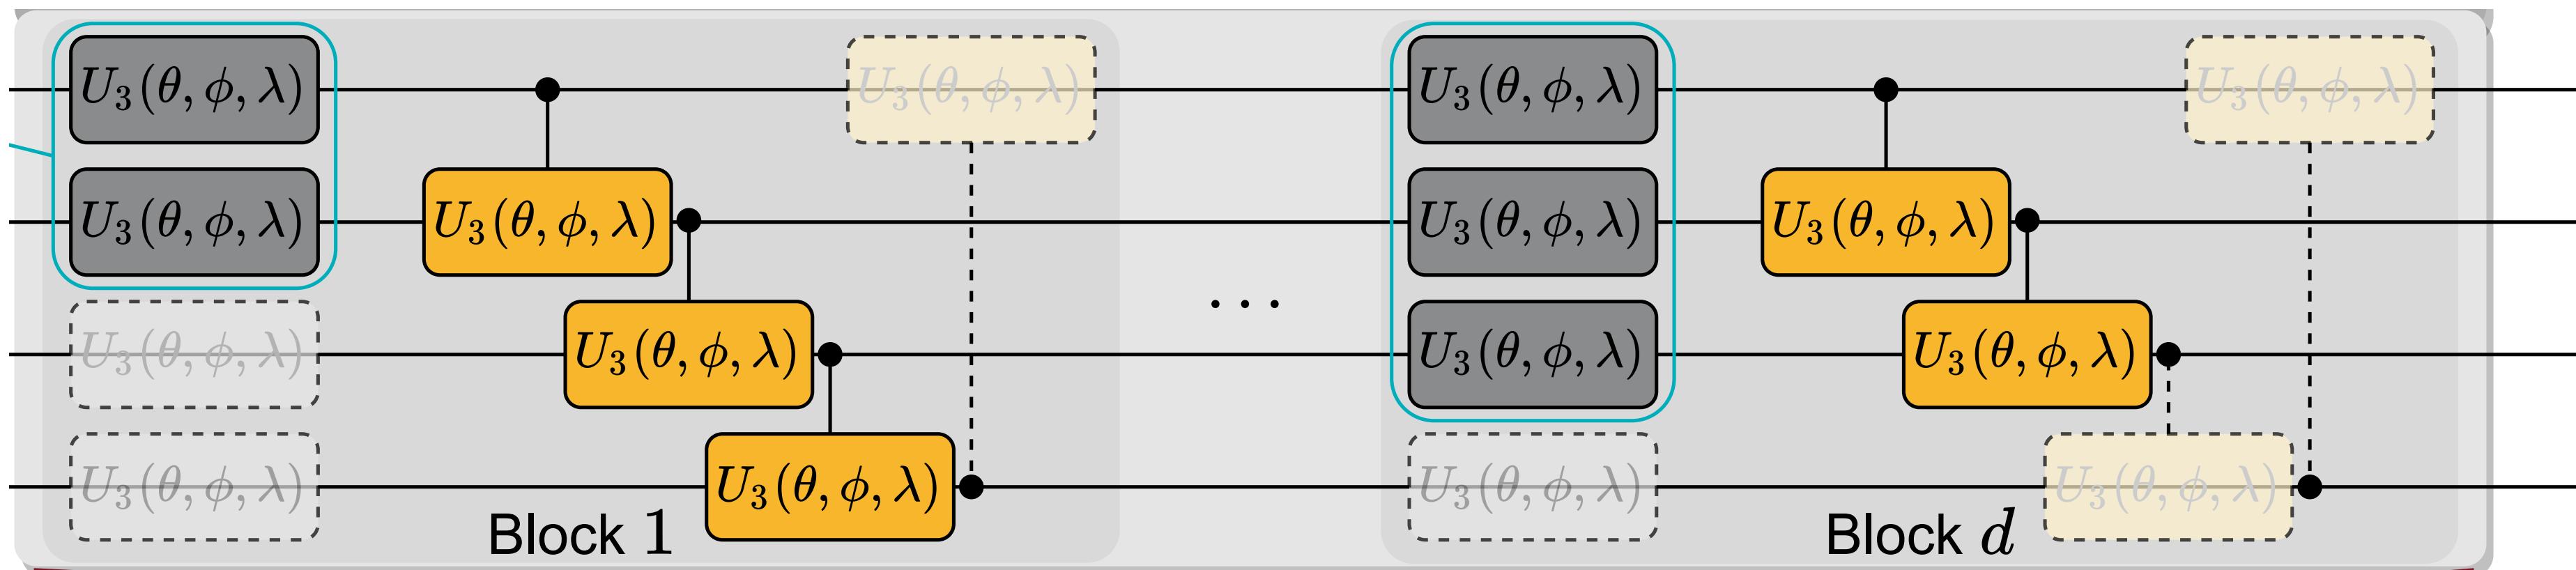
Step2



Step3

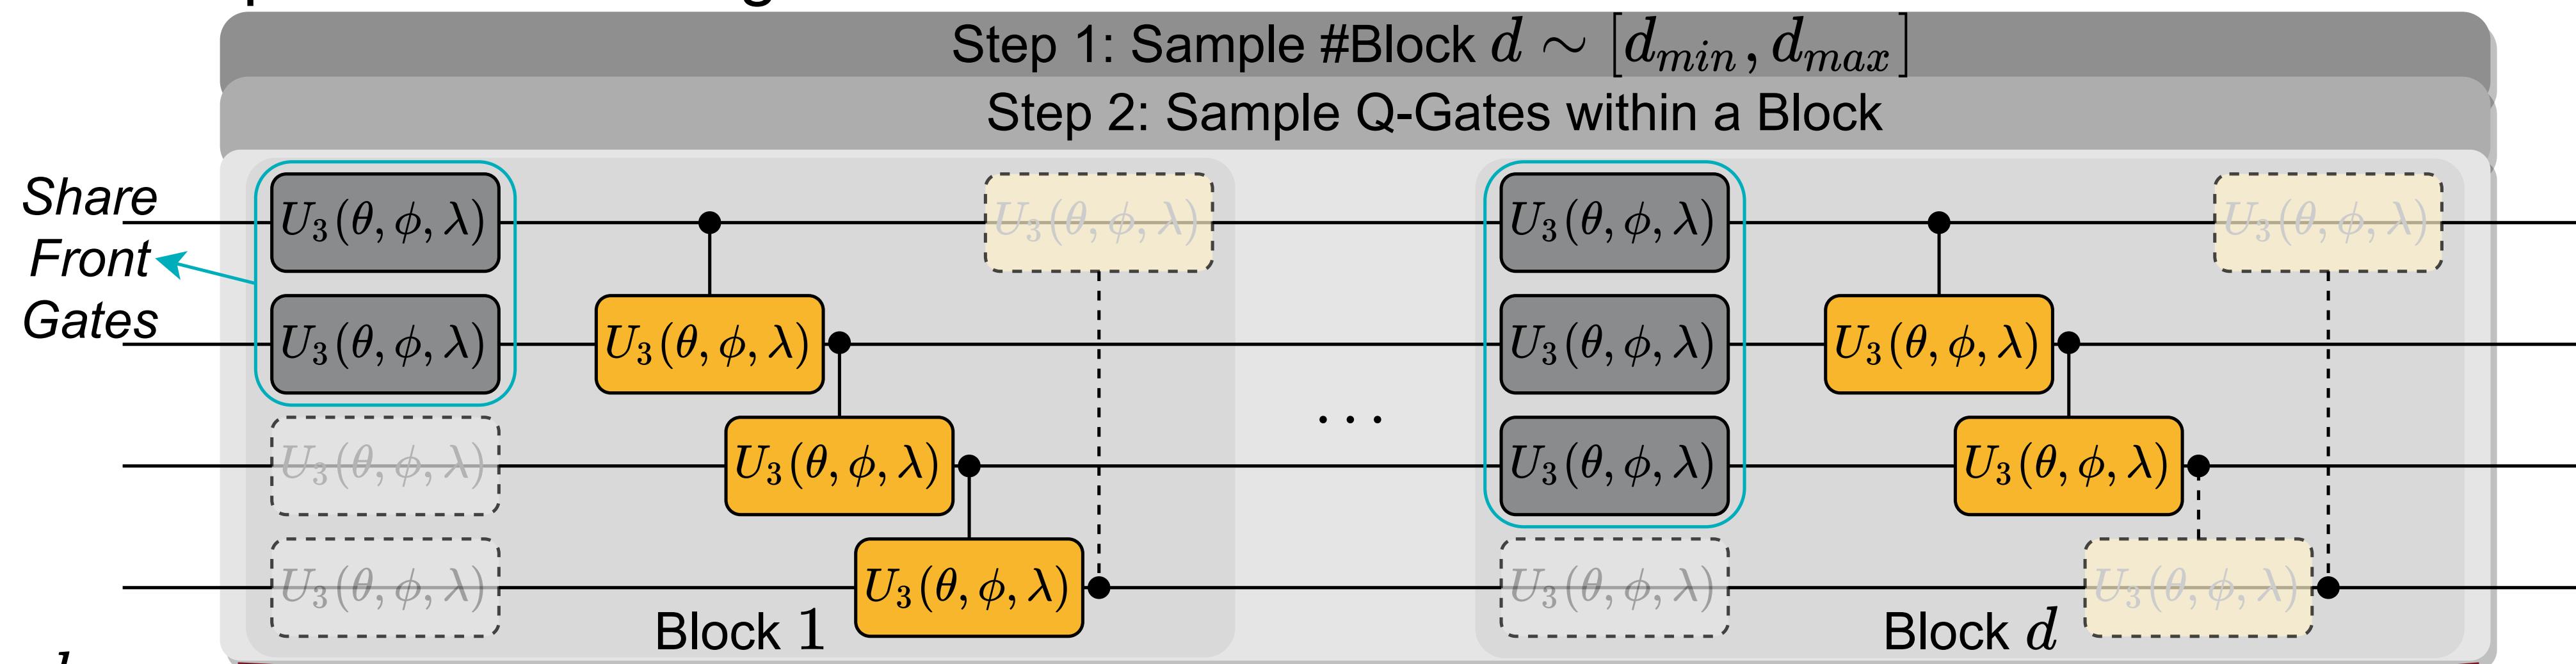
# More Details on Construction and Sampling

- In reality, we can stack multiple blocks of quantum layers to have enough parameters and a large design space



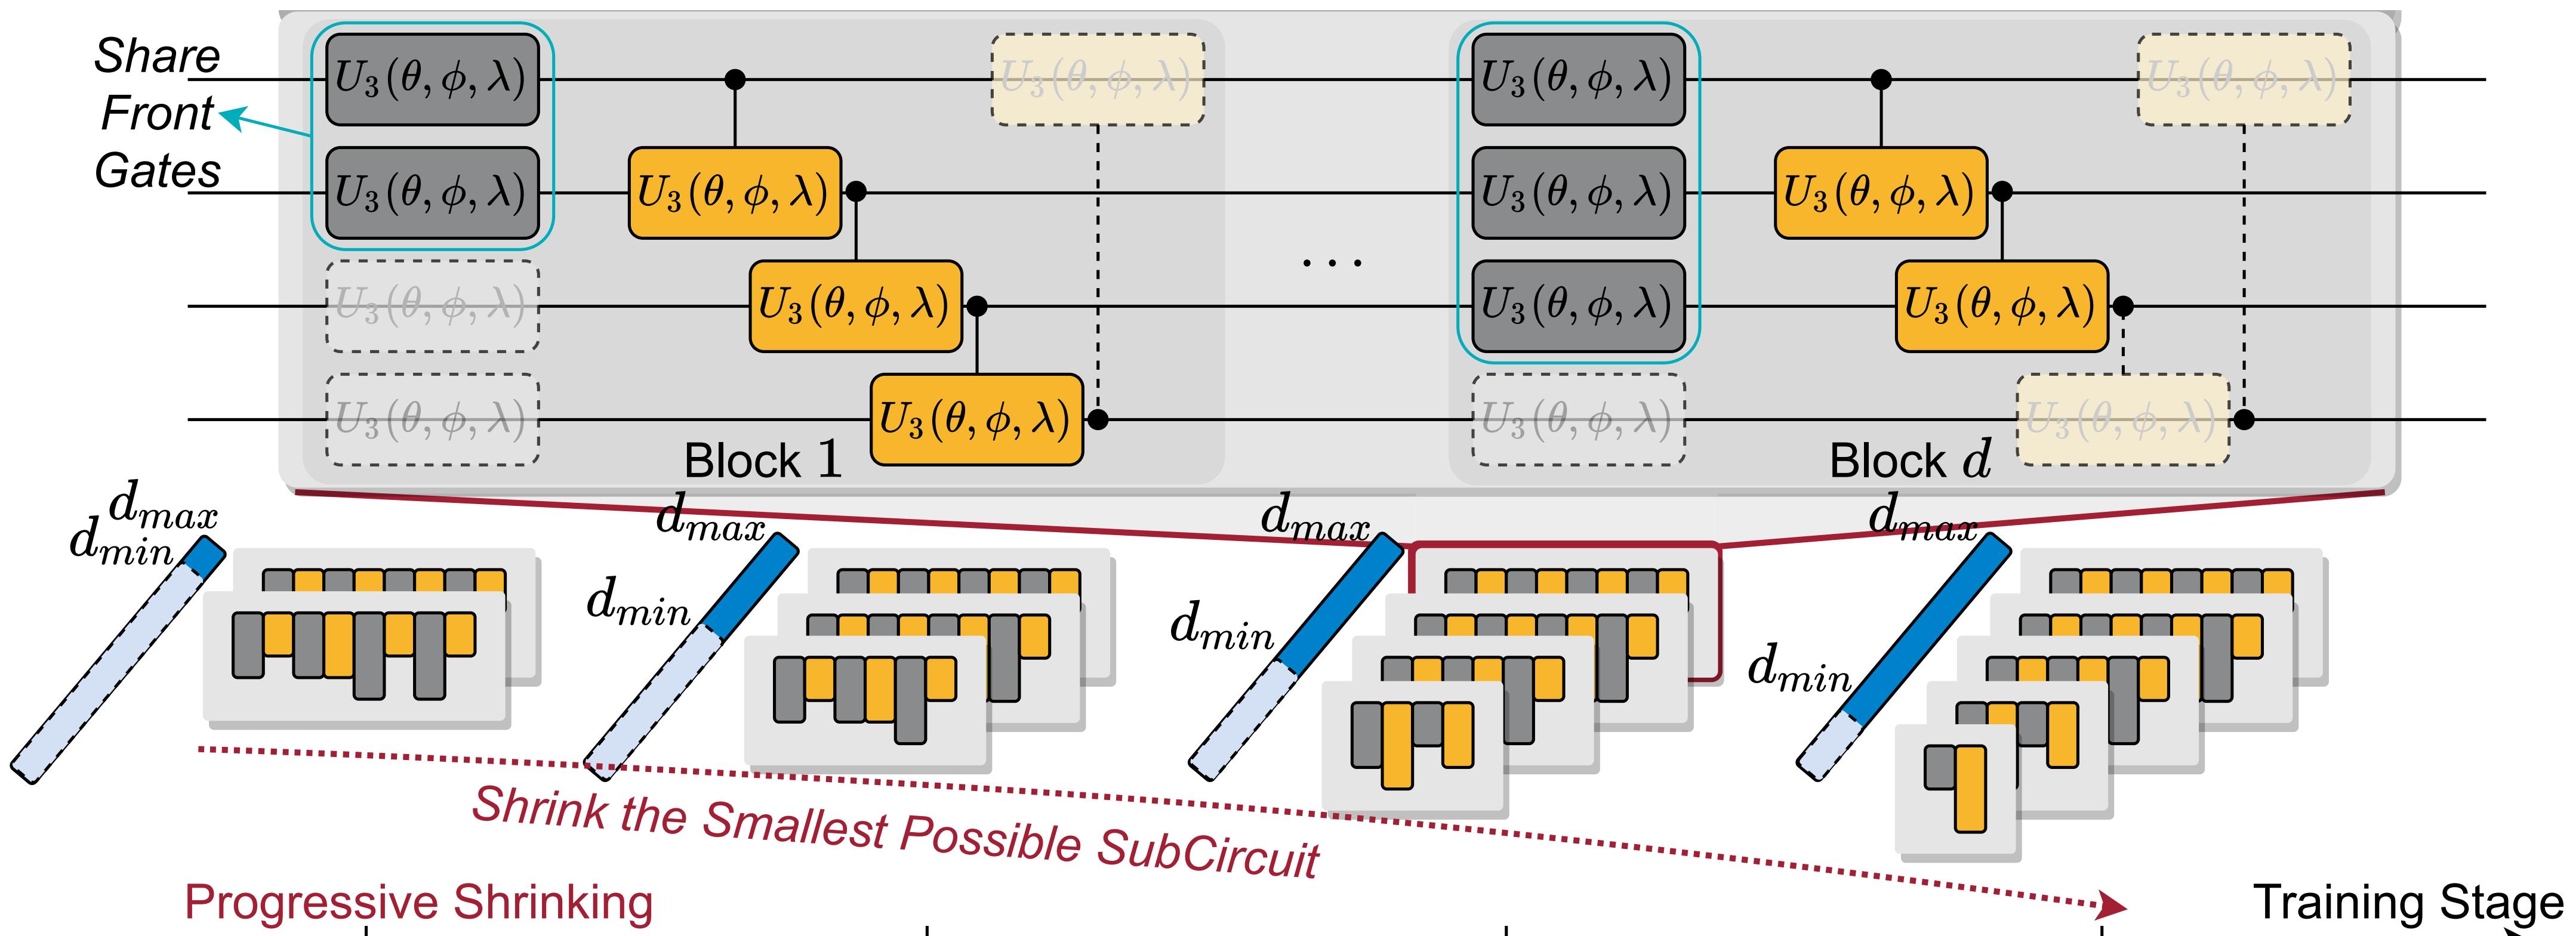
# Front Sampling

- In reality, we can stack multiple blocks of quantum layers to have enough parameters and a large design space
- During sampling, we first sample total number of blocks, then sample gates within each block
  - Front sampling: Only the **front several** blocks and front several gates can be sampled to make SuperCircuit training more stable



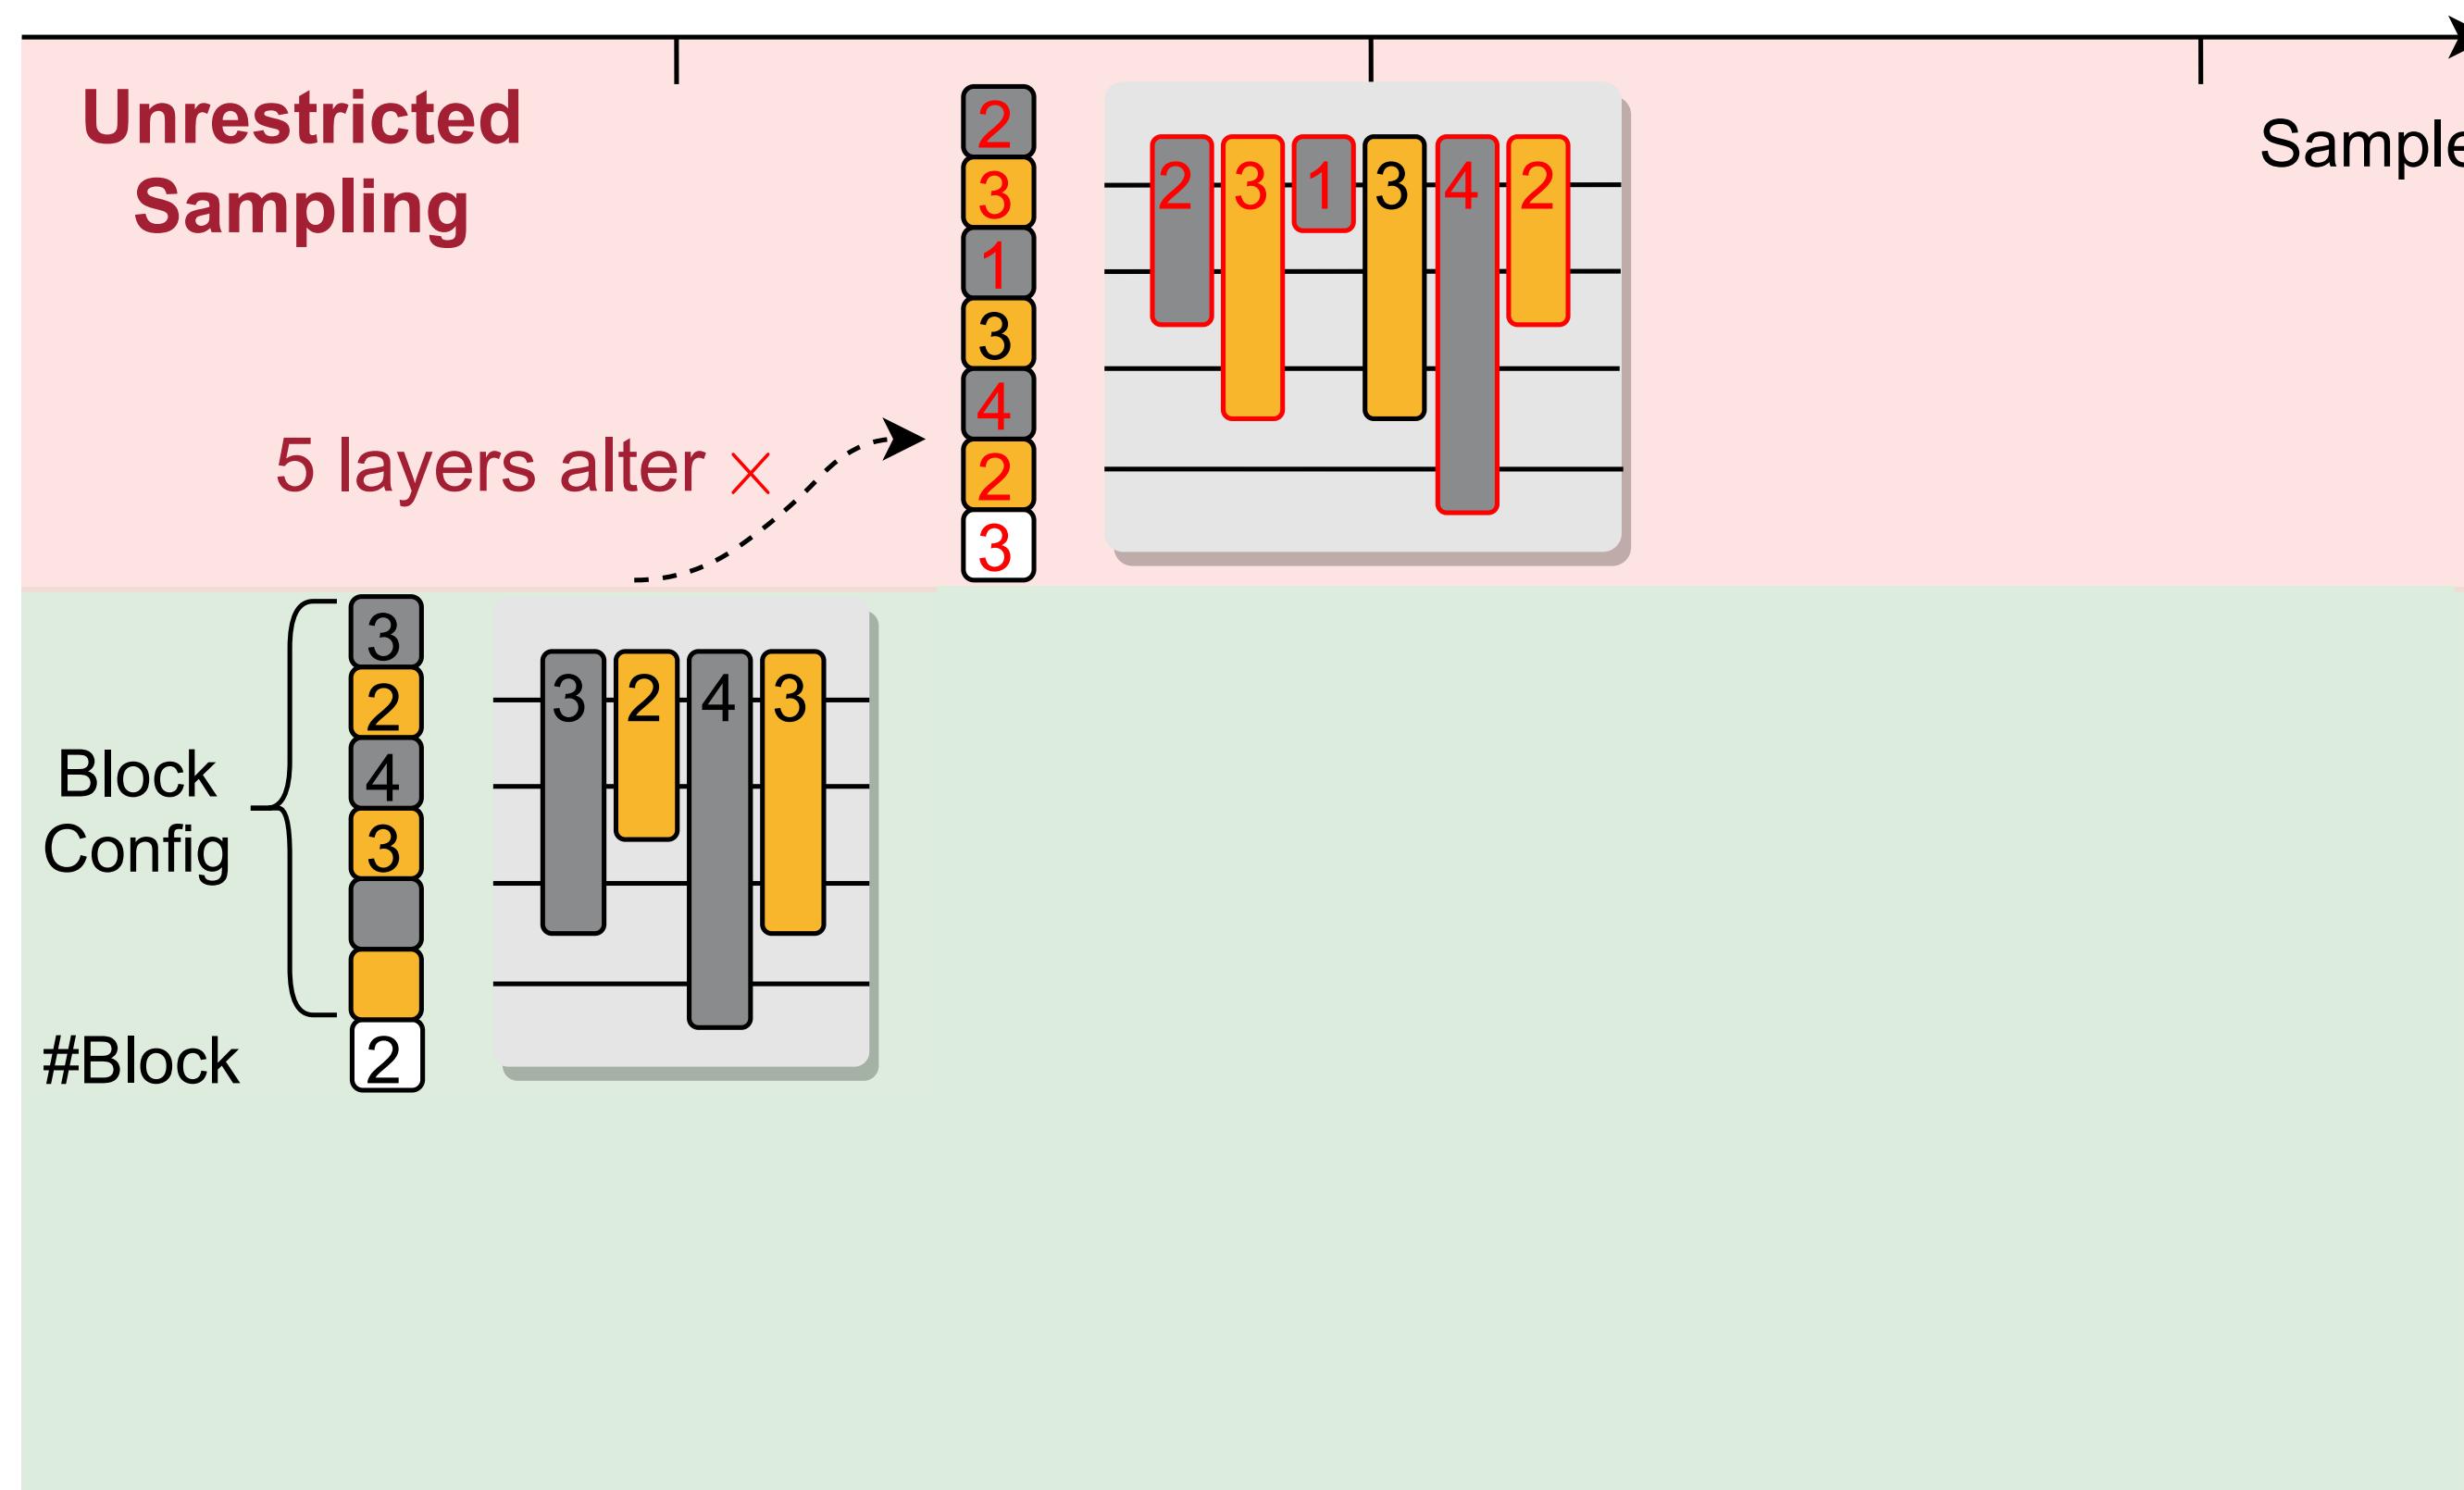
# Progressive Shrinking

- Progressive Shrinking:
  - Make parameters in large index block well-trained
  - Progressively reduce the lower bound of possible block number



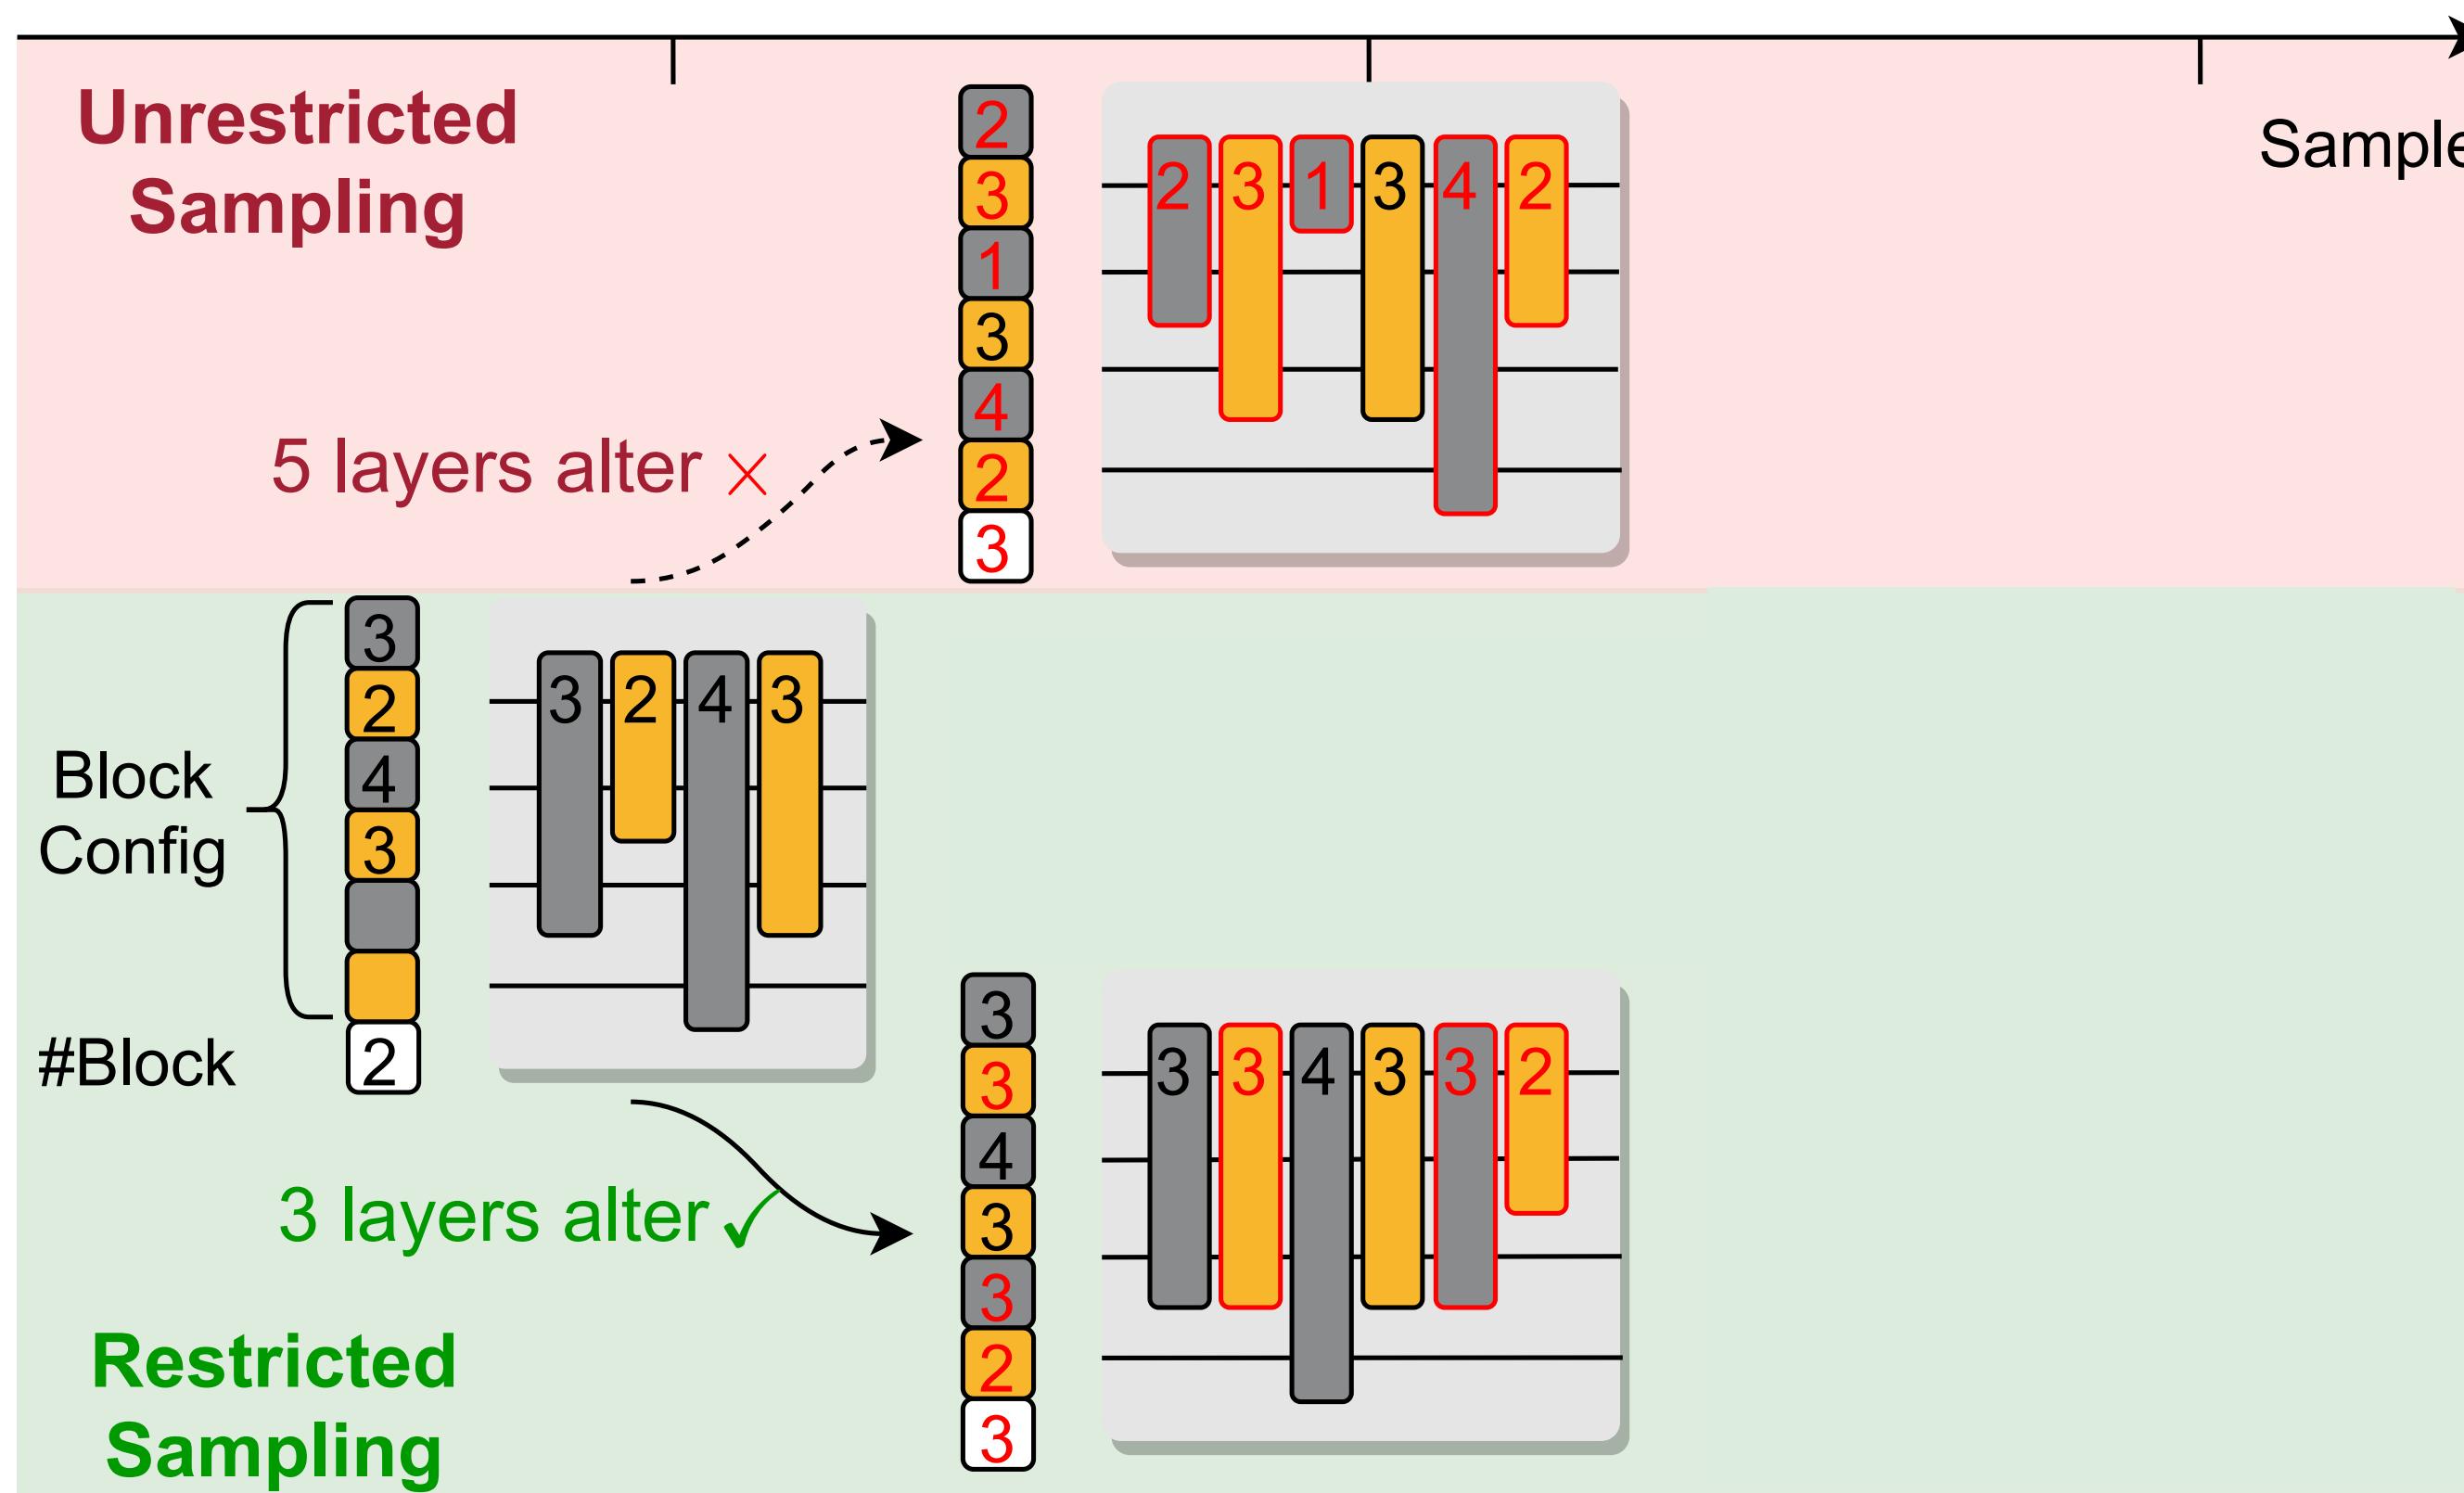
# Restricted Sampling

- Restricted Sampling:
  - Restrict the difference between two consecutively sampled SubCircuits



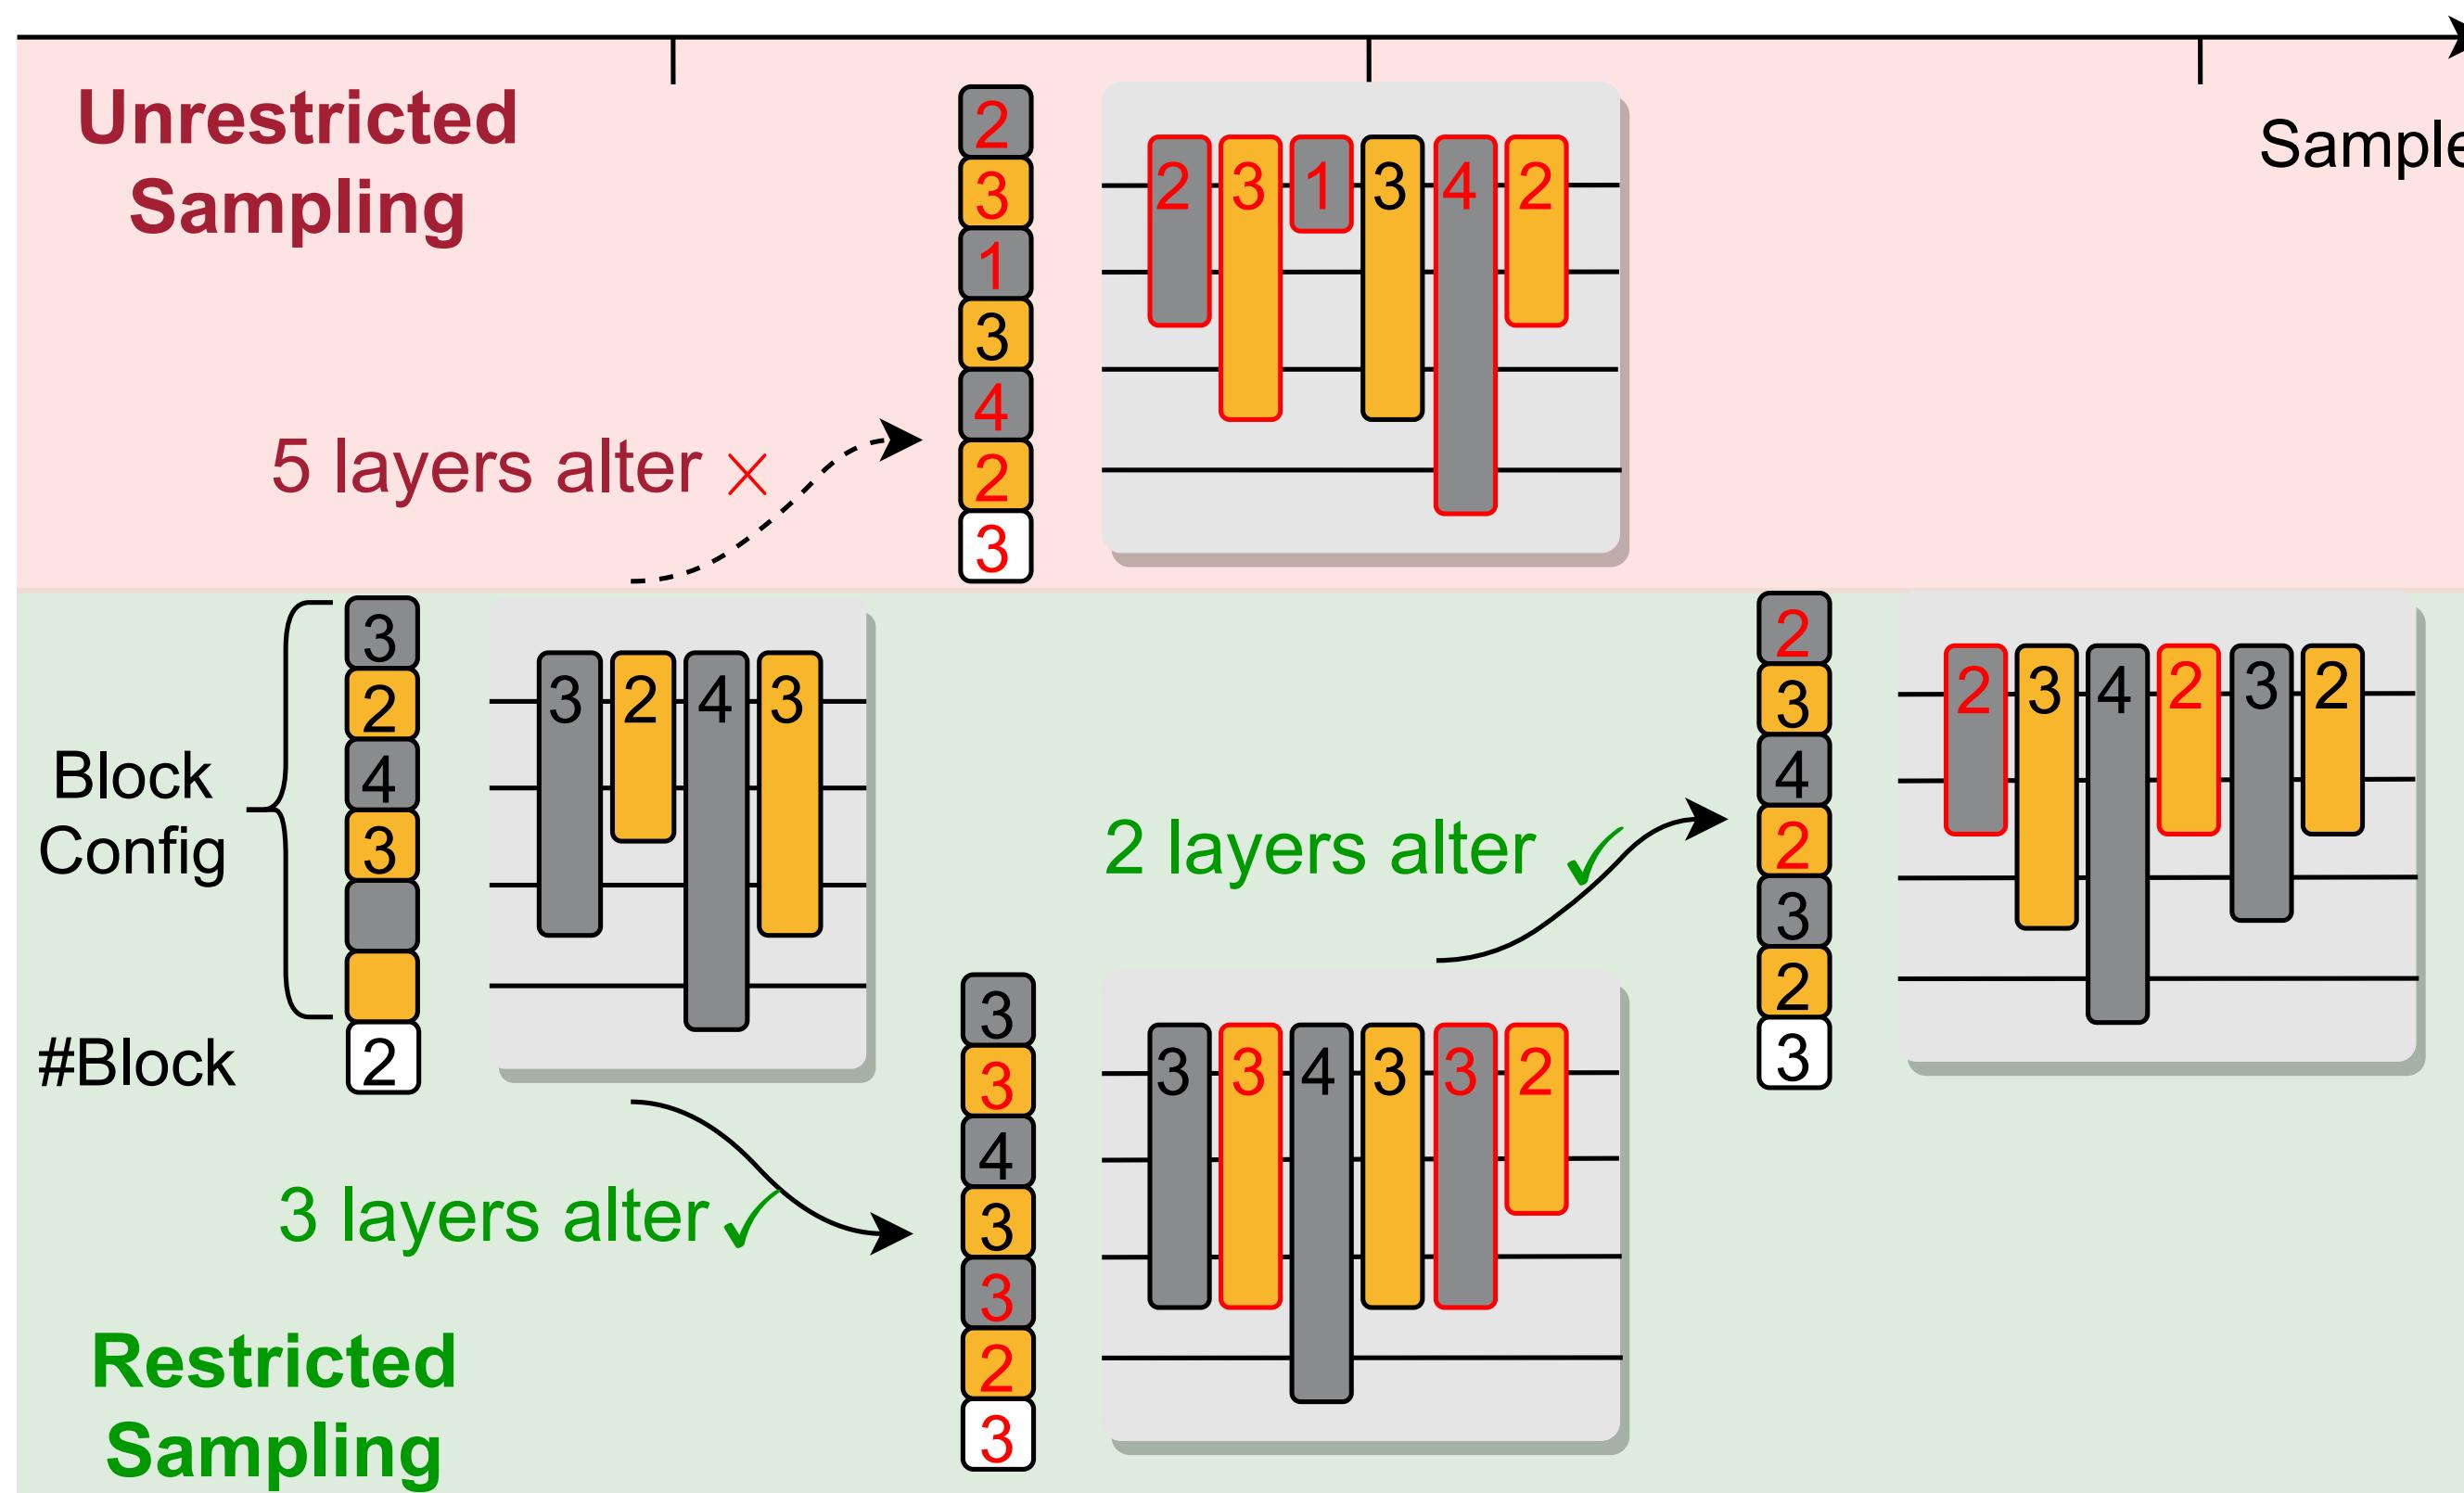
# Restricted Sampling

- Restricted Sampling:
  - Restrict the difference between two consecutively sampled SubCircuits



# Restricted Sampling

- Restricted Sampling:
  - Restrict the difference between two consecutively sampled SubCircuits

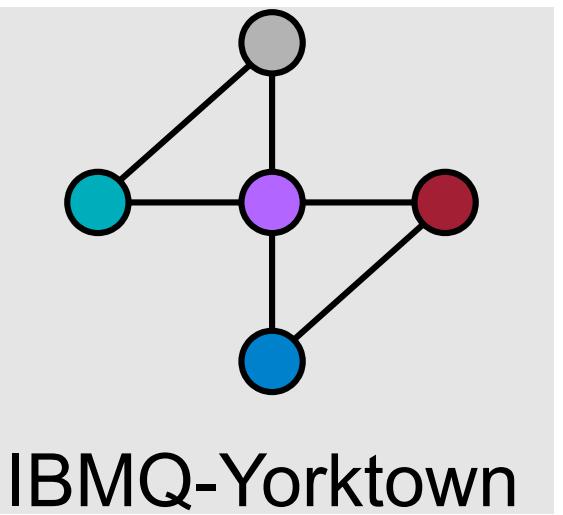


# QuantumNAS

- SuperCircuit Construction and Training
- Noise-Adaptive Evolutionary Co-Search of SubCircuit and Qubit Mapping
- Searched SubCircuit Training
- Iterative Pruning

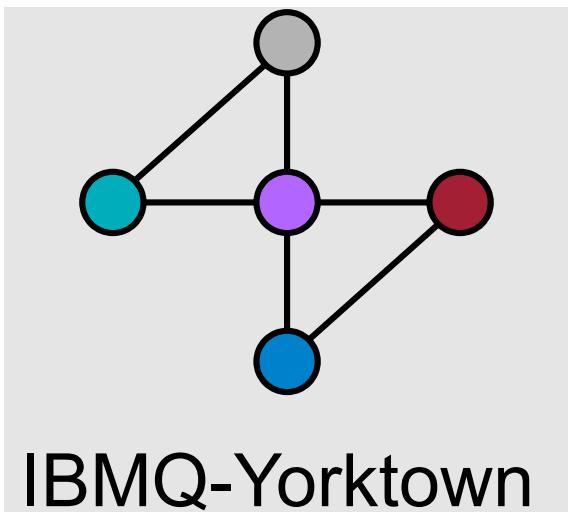
# Noise-Adaptive Evolutionary Co-Search

- Search the best SubCircuit and its qubit mapping on target device

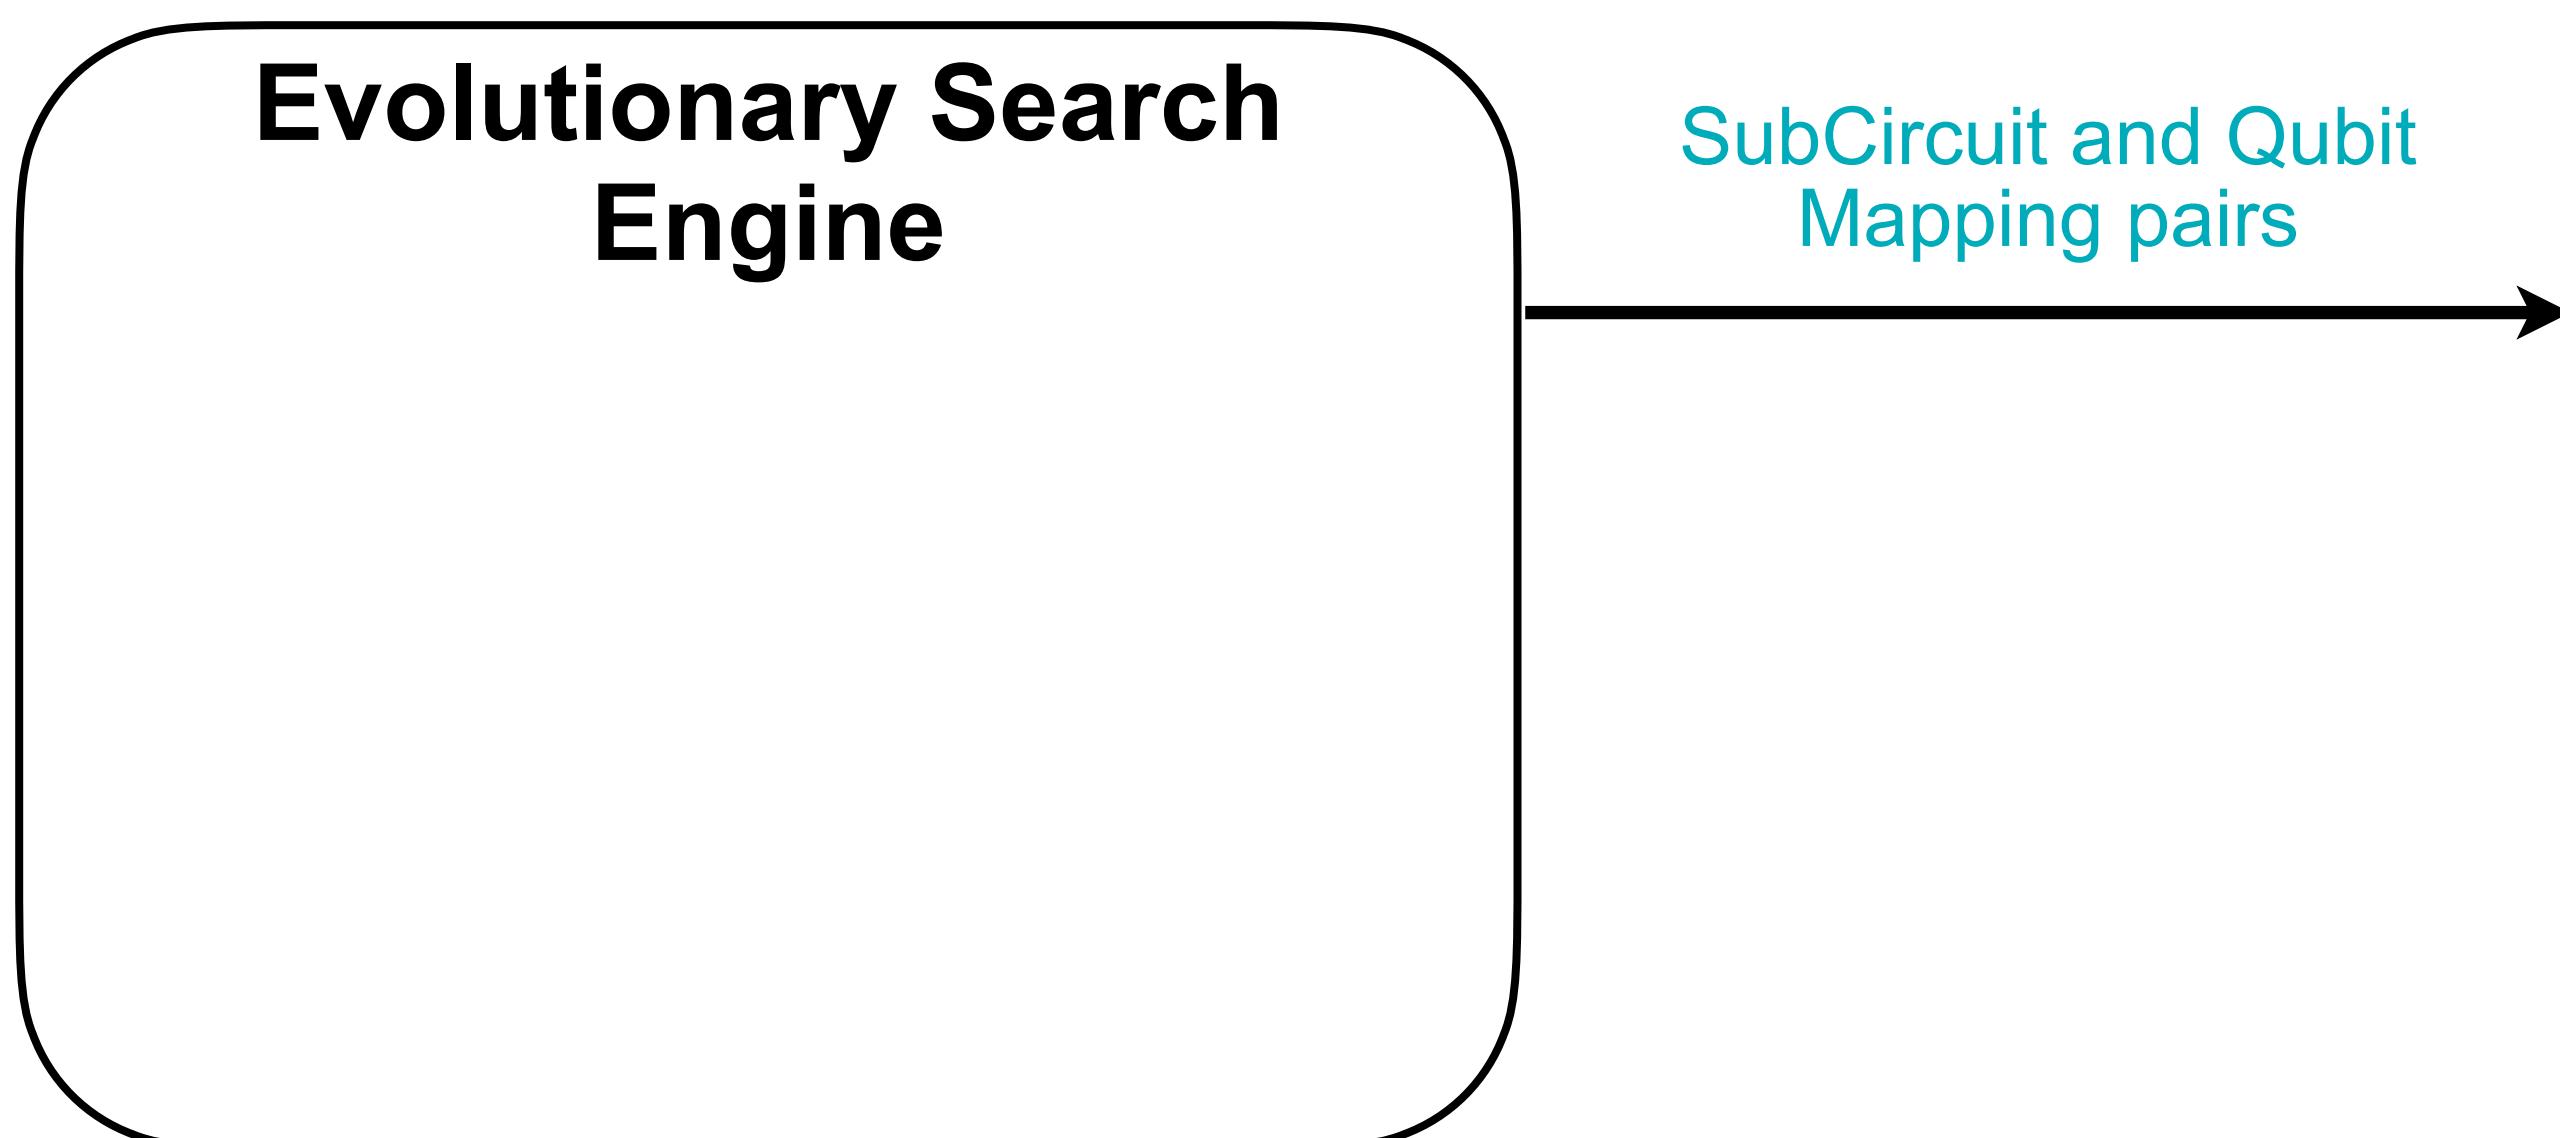
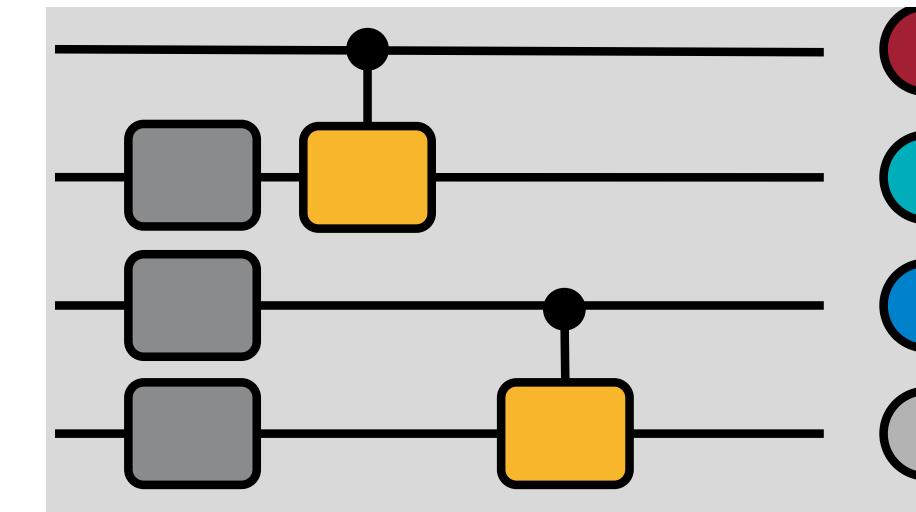
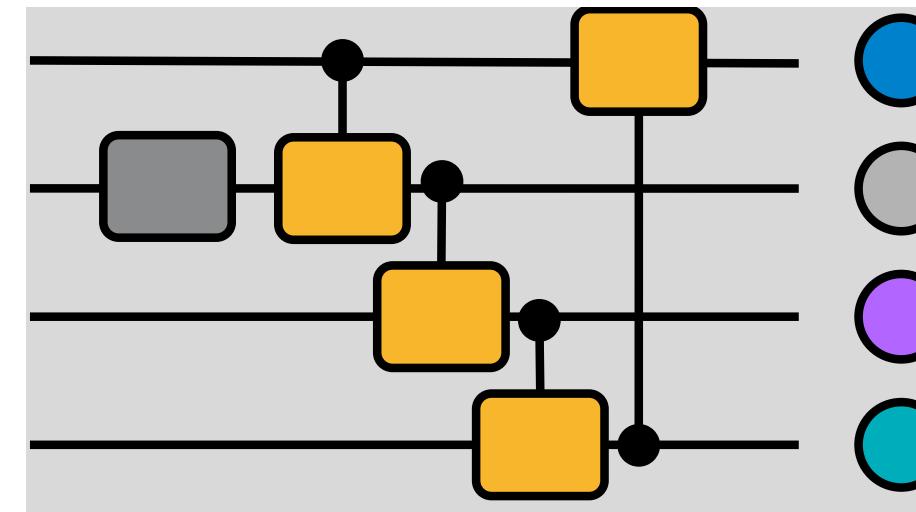
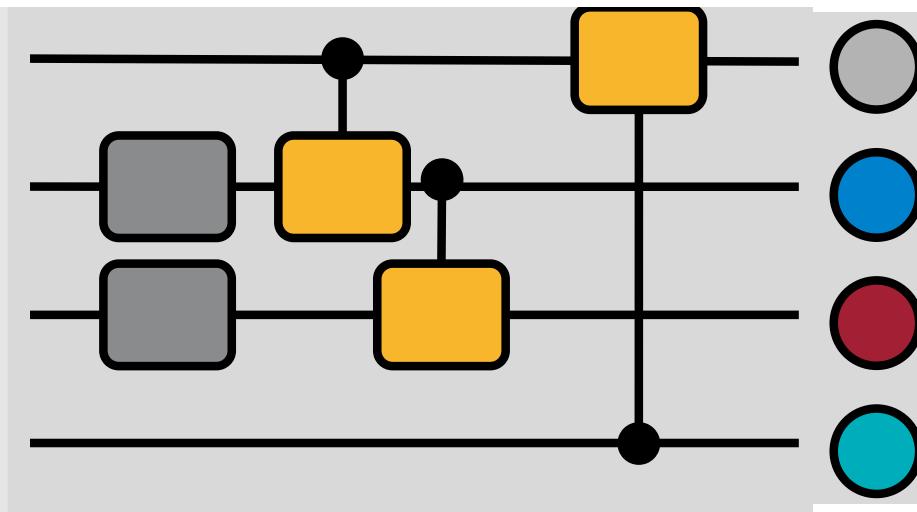


# Noise-Adaptive Evolutionary Co-Search

- Search the best SubCircuit and its qubit mapping

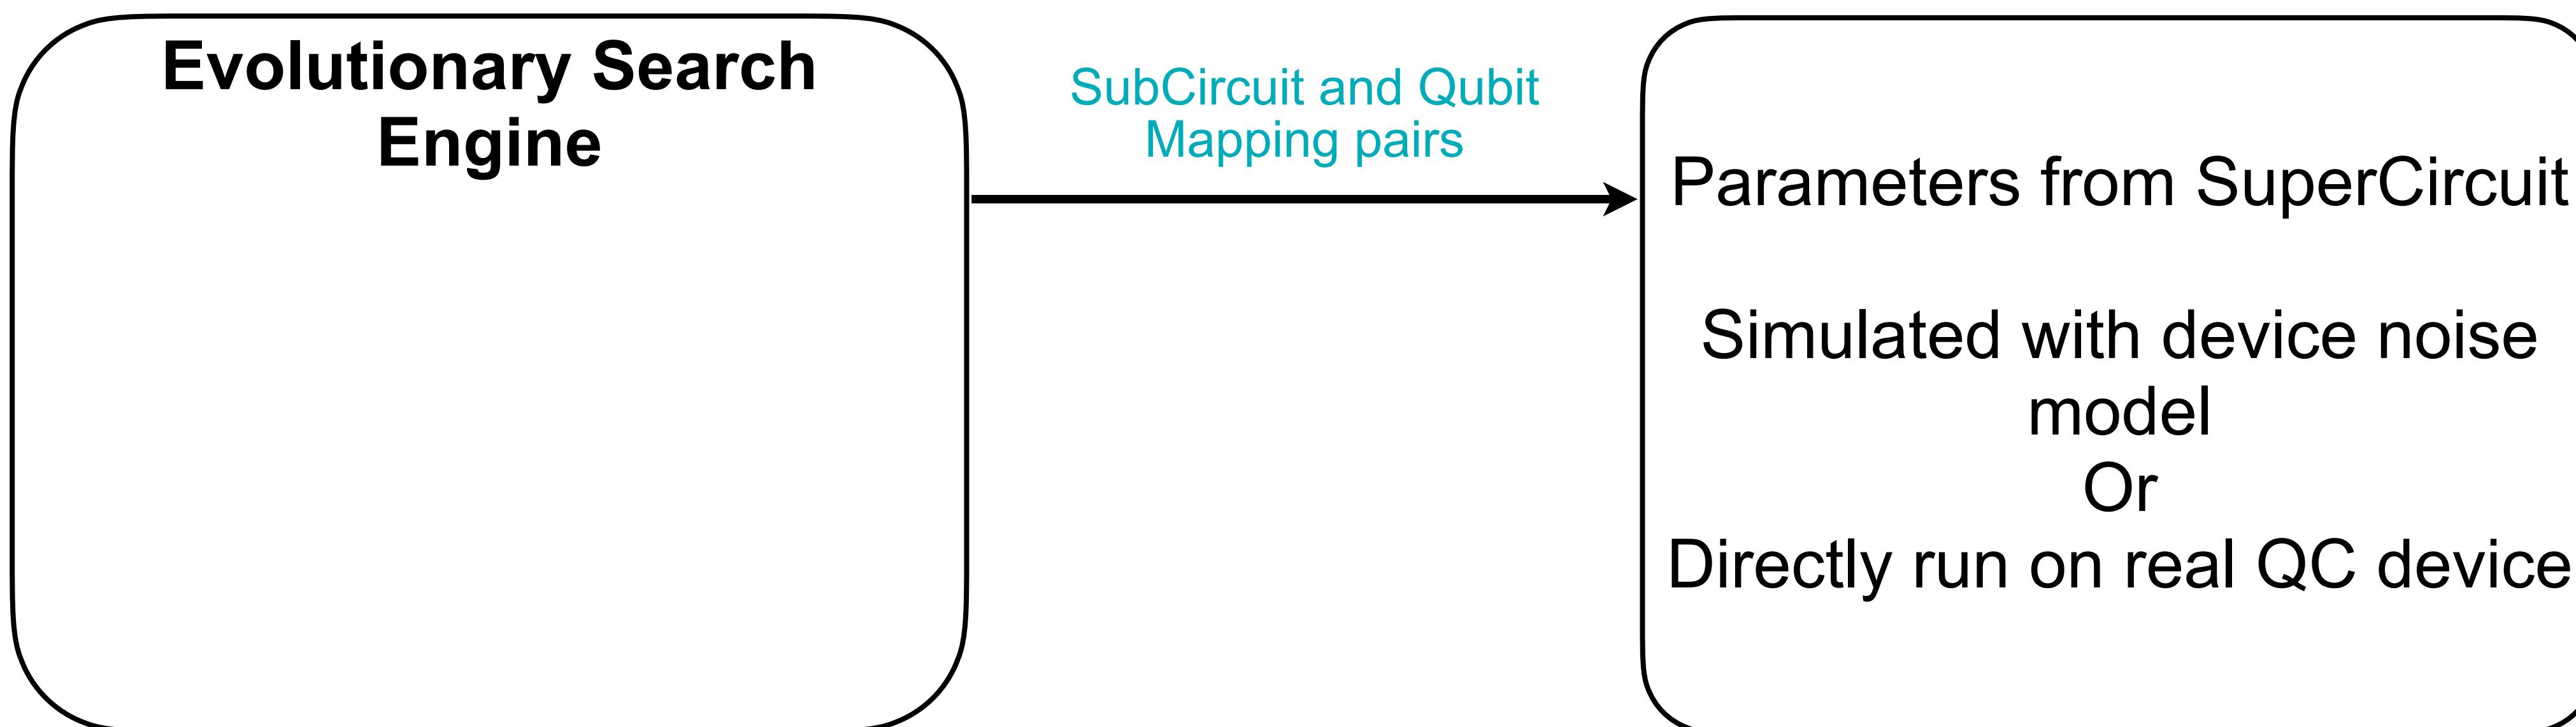
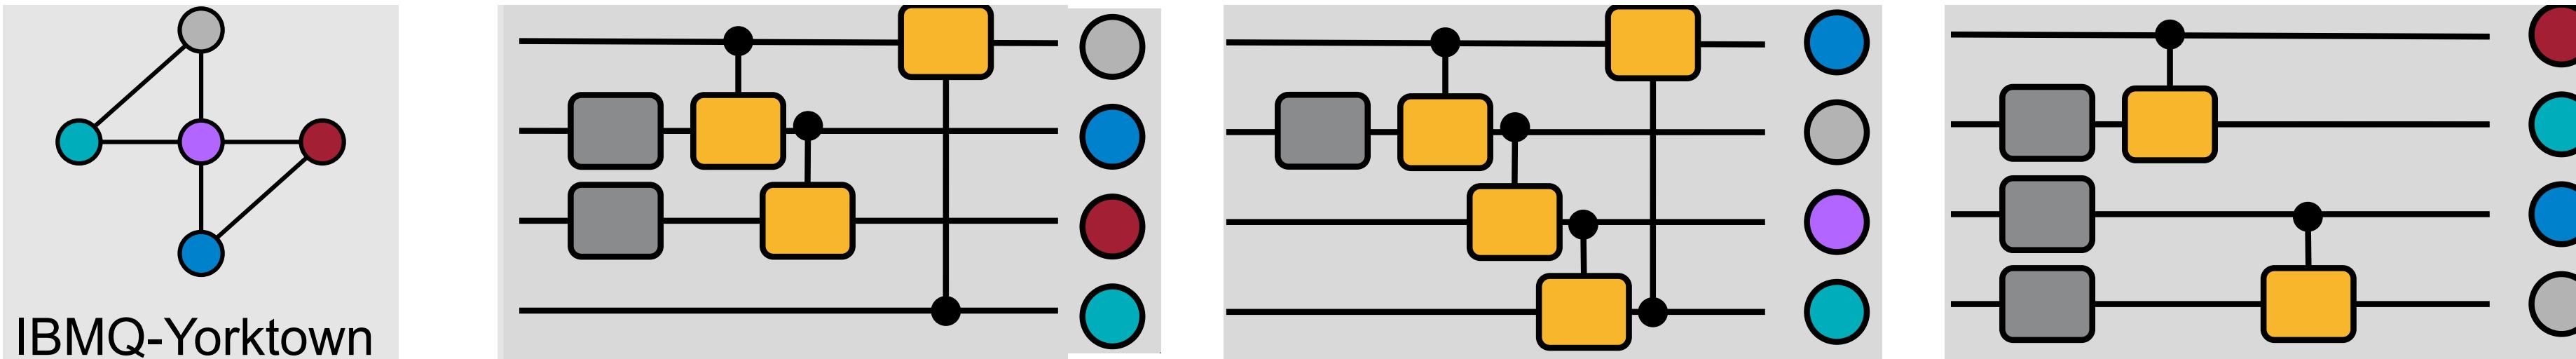


IBMQ-Yorktown



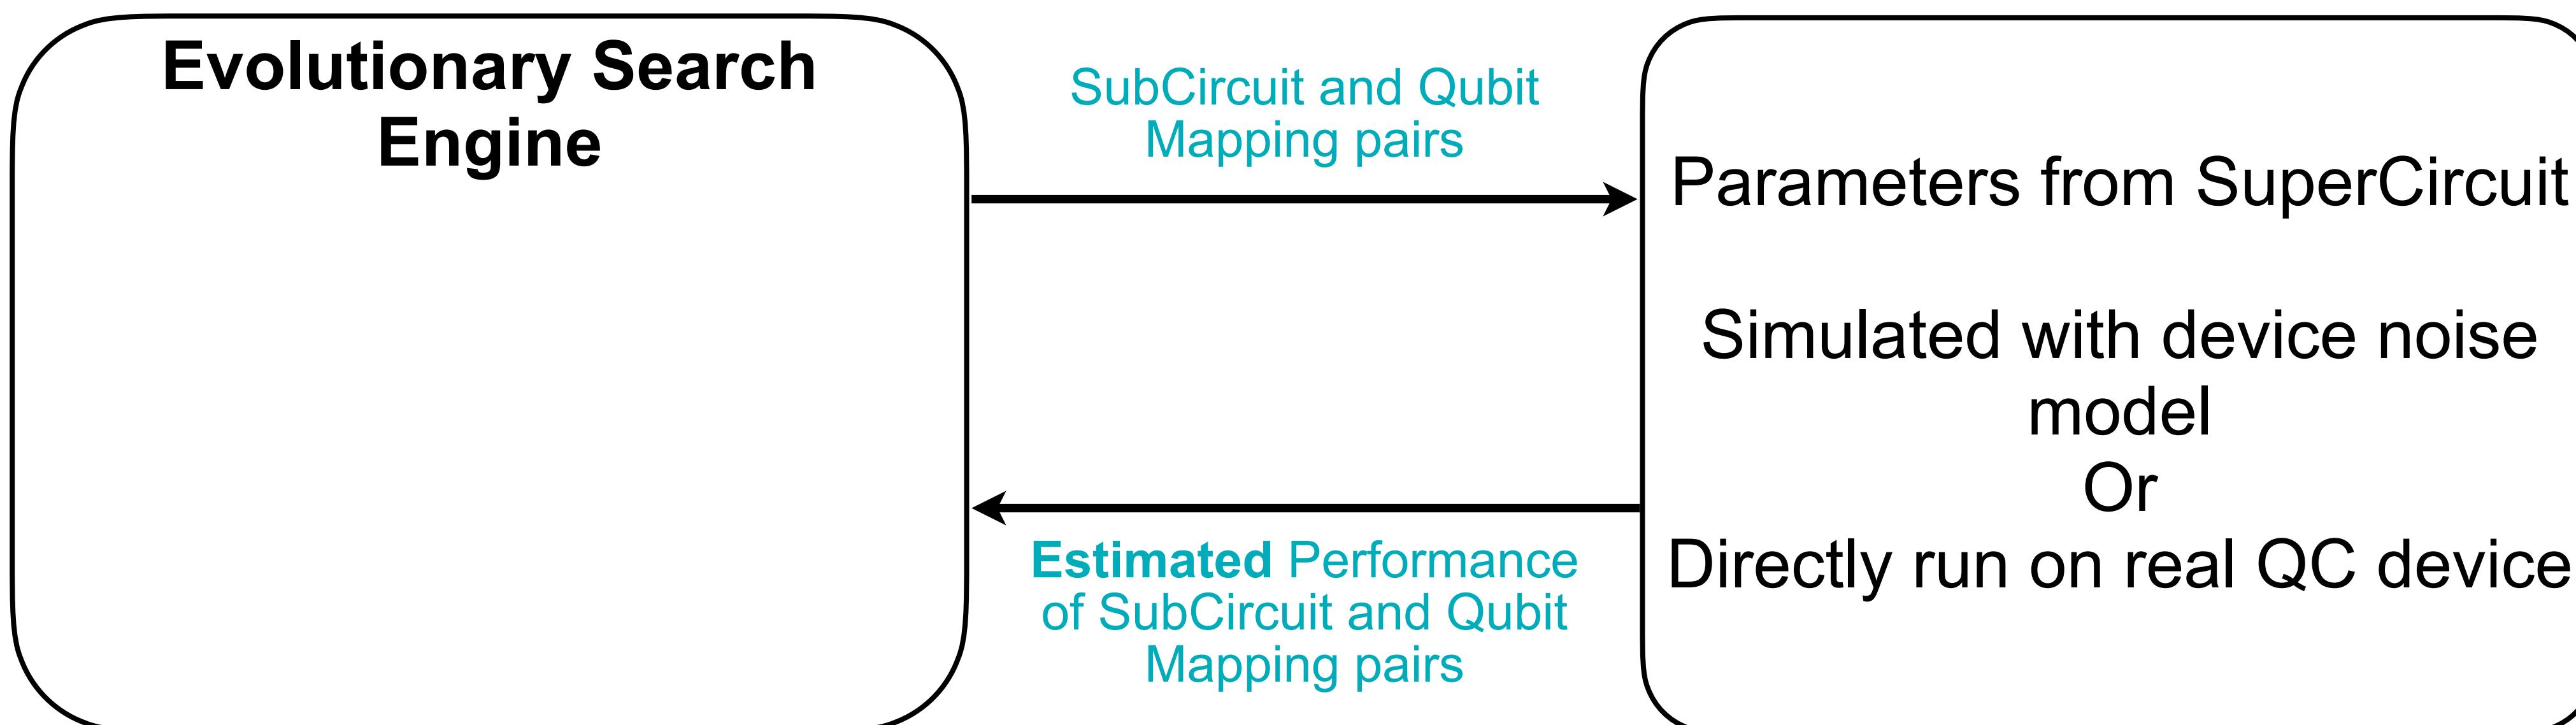
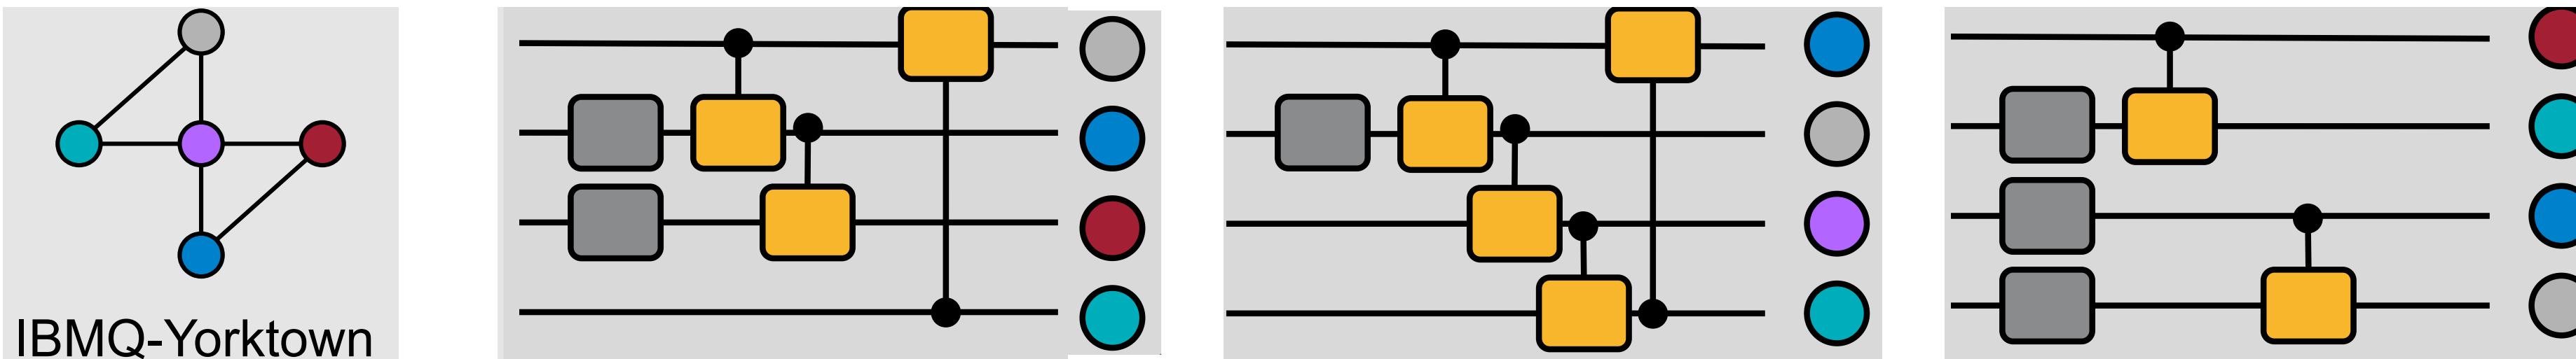
# Noise-Adaptive Evolutionary Co-Search

- Search the best SubCircuit and its qubit mapping



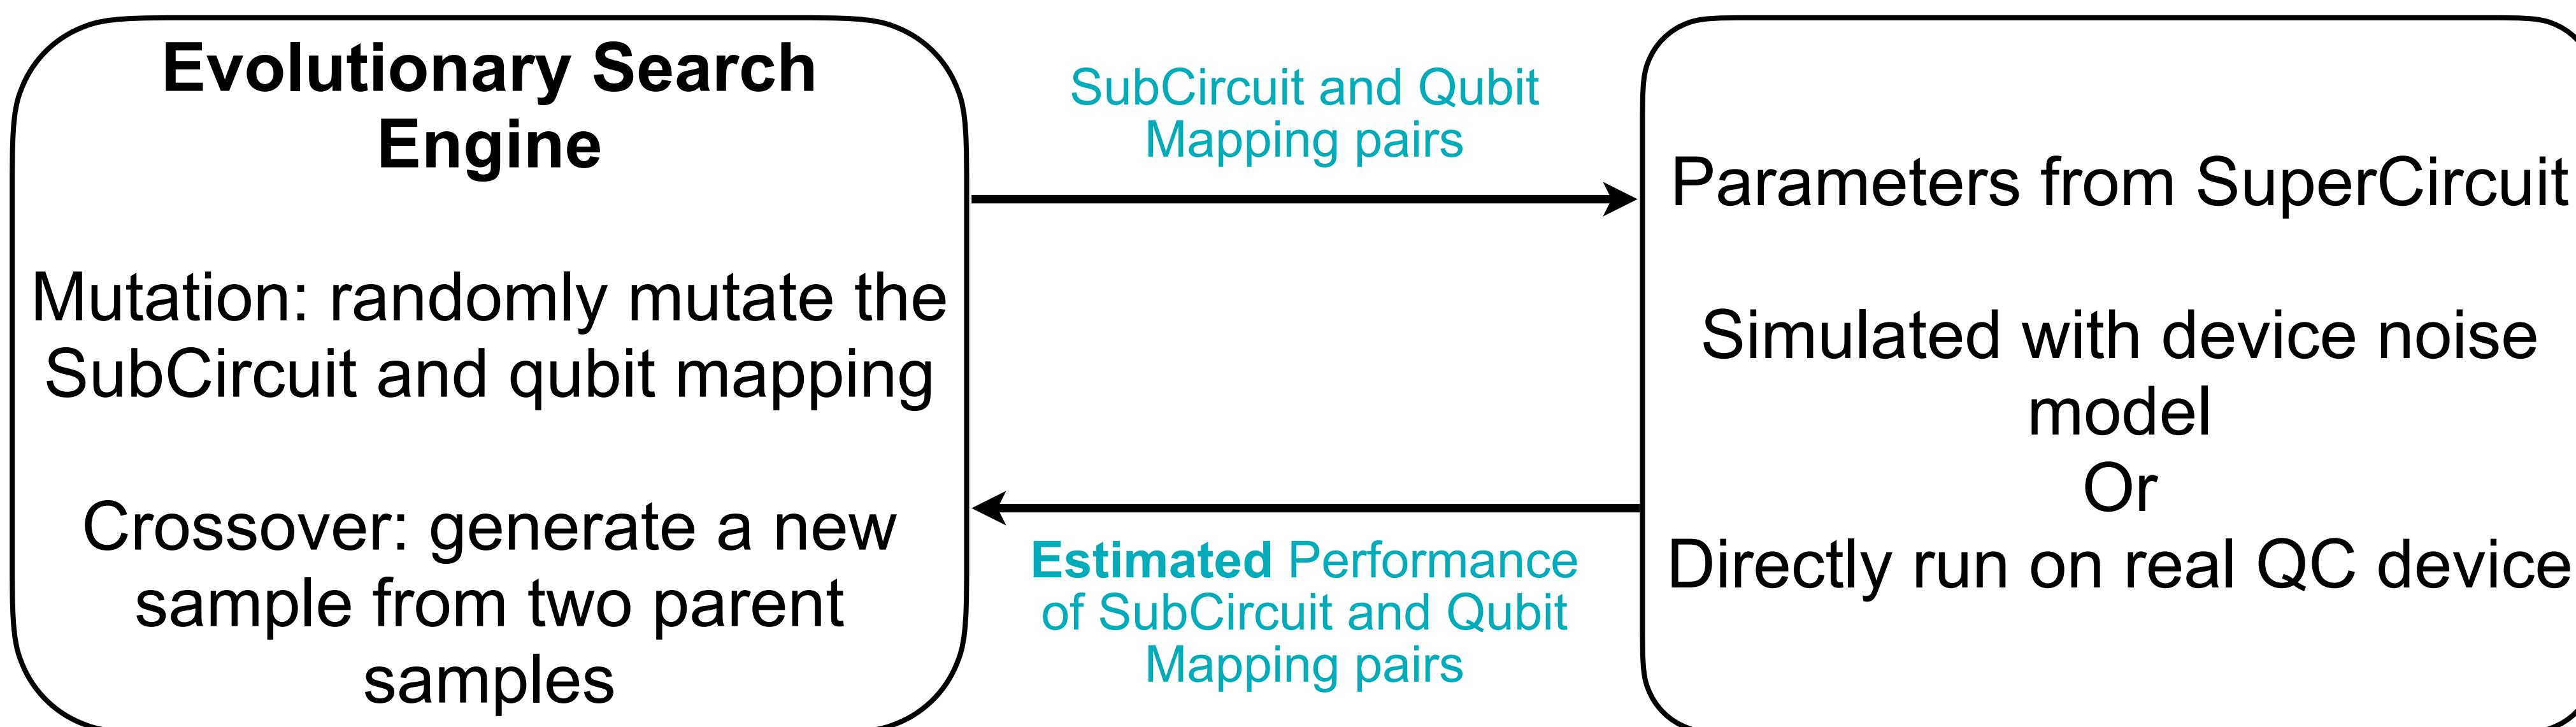
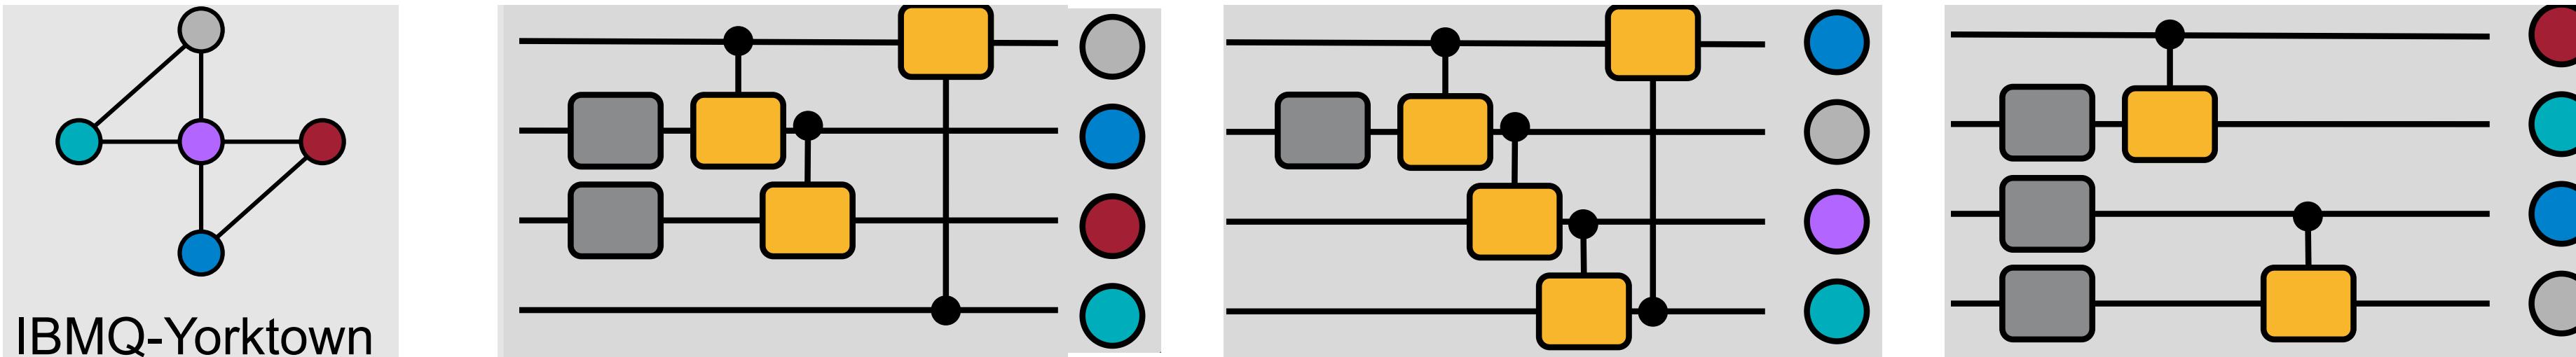
# Noise-Adaptive Evolutionary Co-Search

- Search the best SubCircuit and its qubit mapping



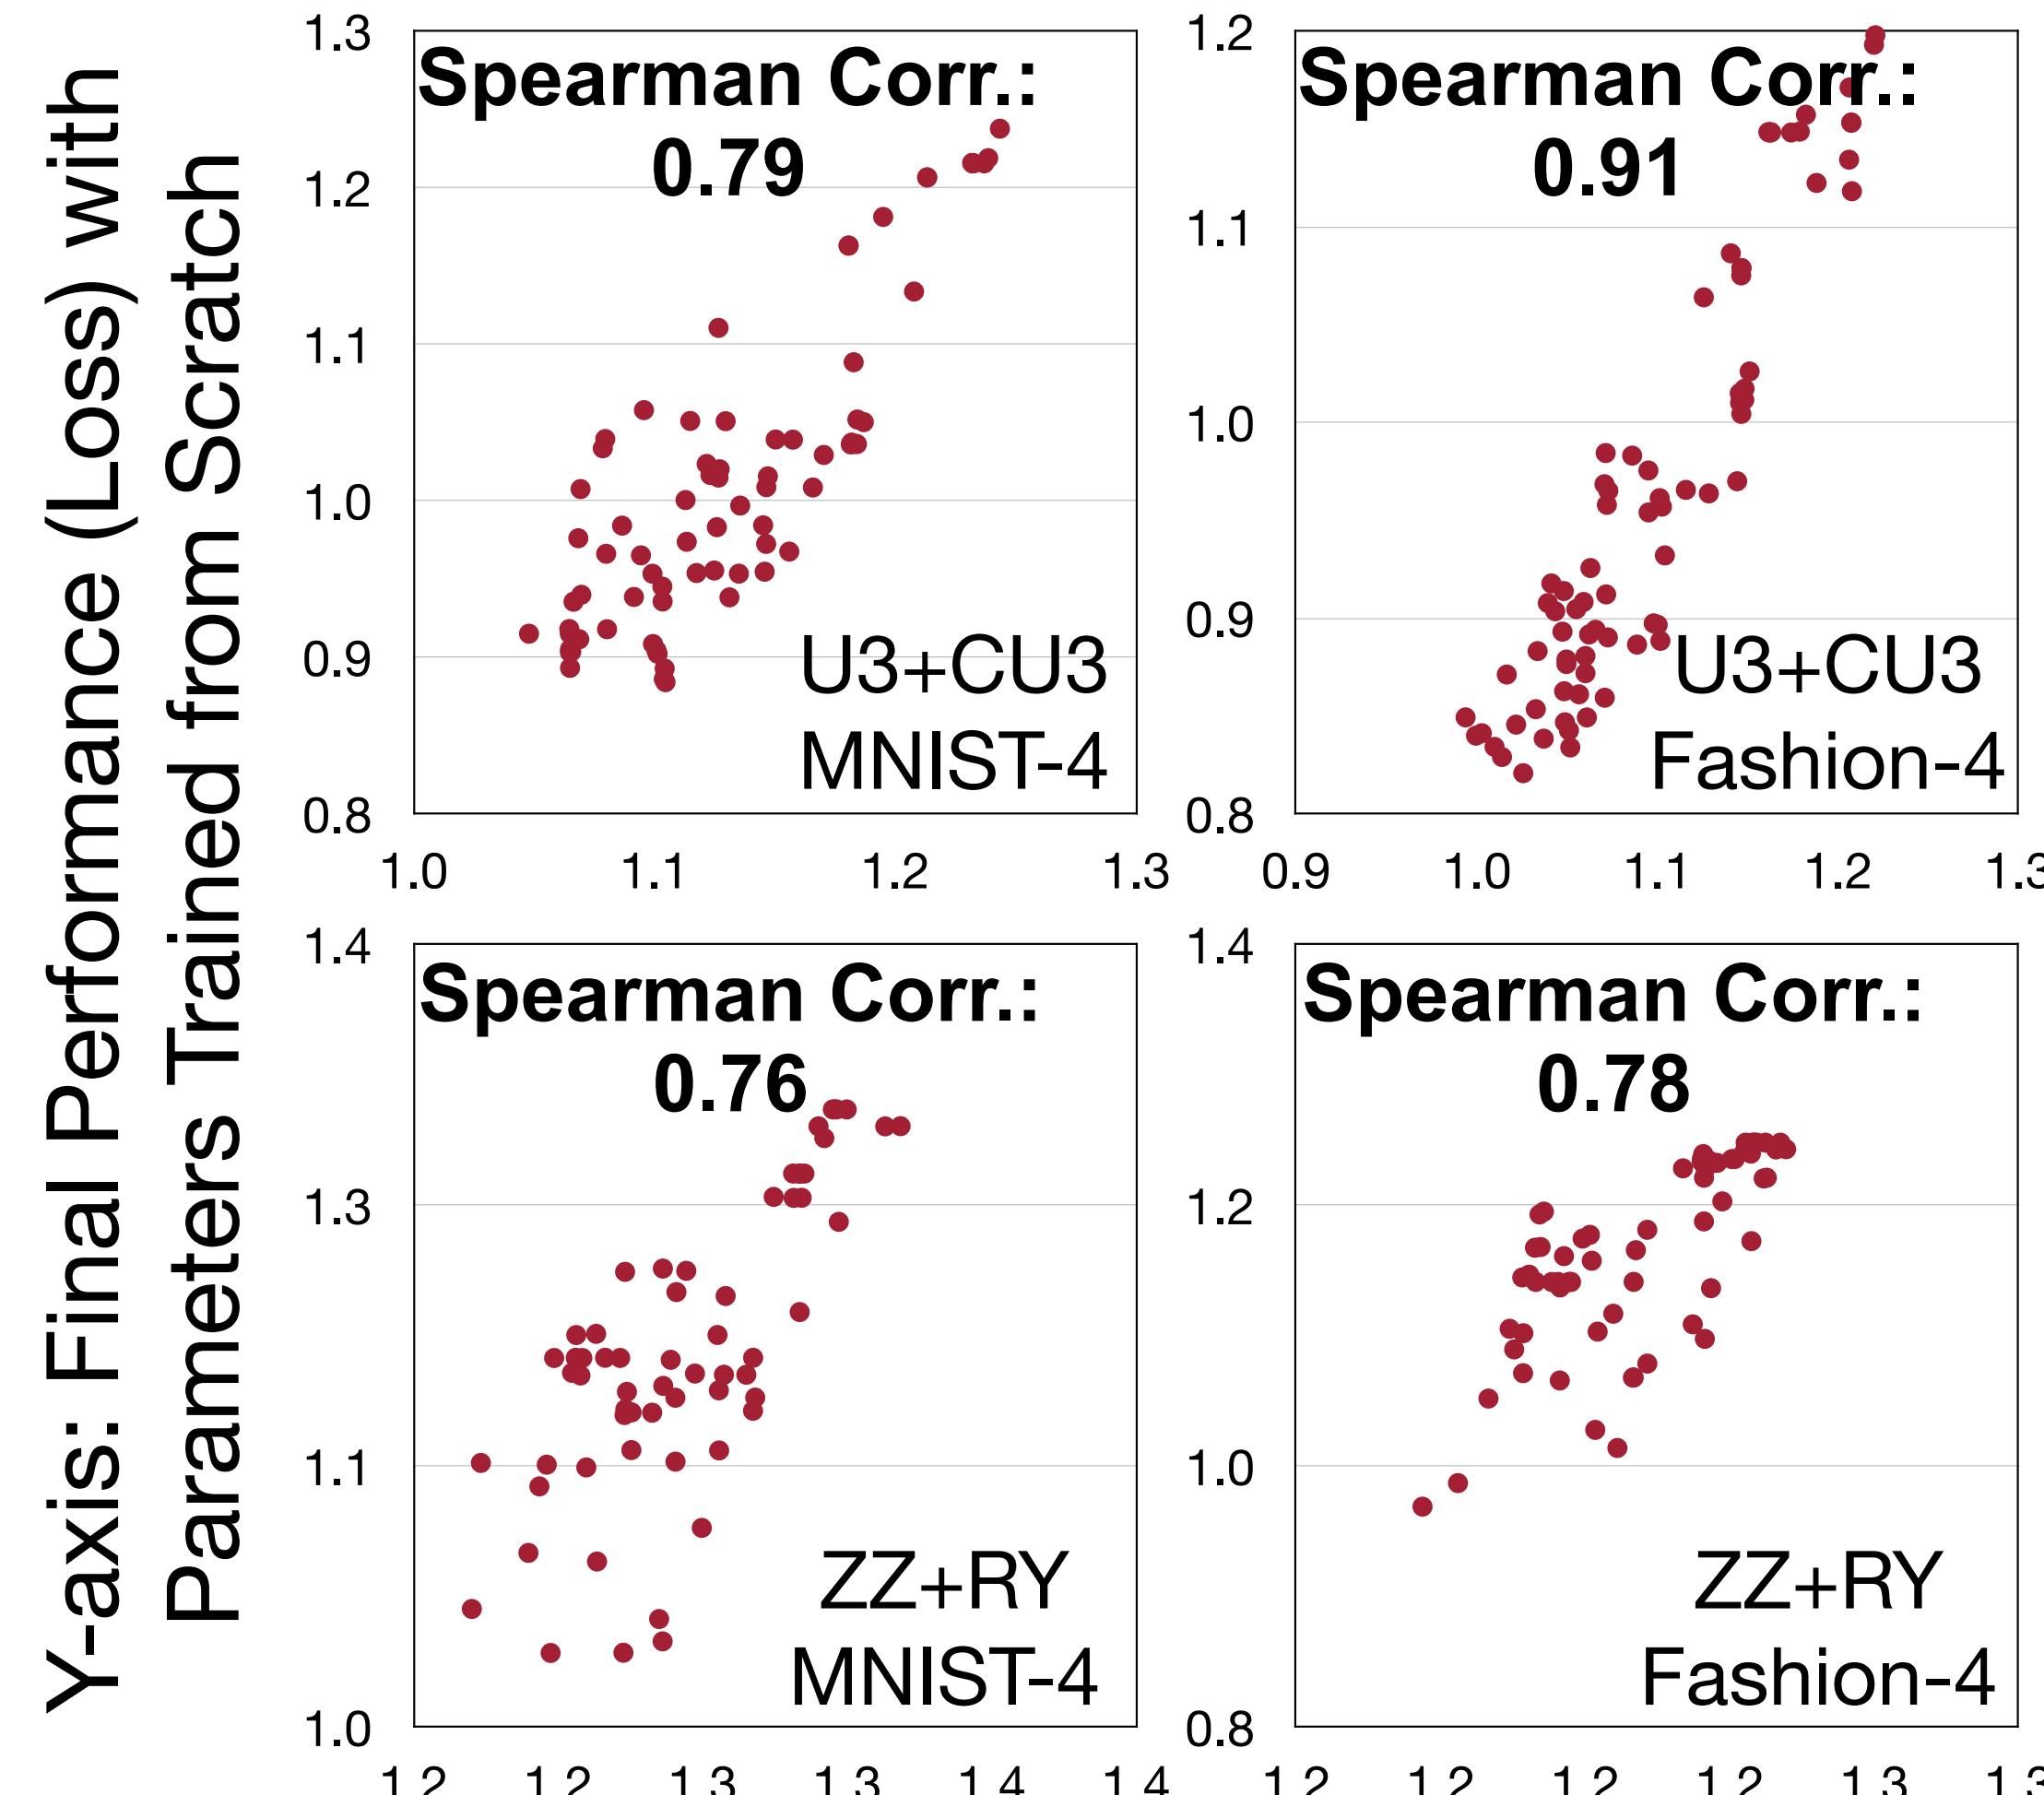
# Noise-Adaptive Evolutionary Co-Search

- Search the best SubCircuit and its qubit mapping



# How Reliable is the SuperCircuit?

- Inherited parameters from SuperCircuit can provide accurate relative performance

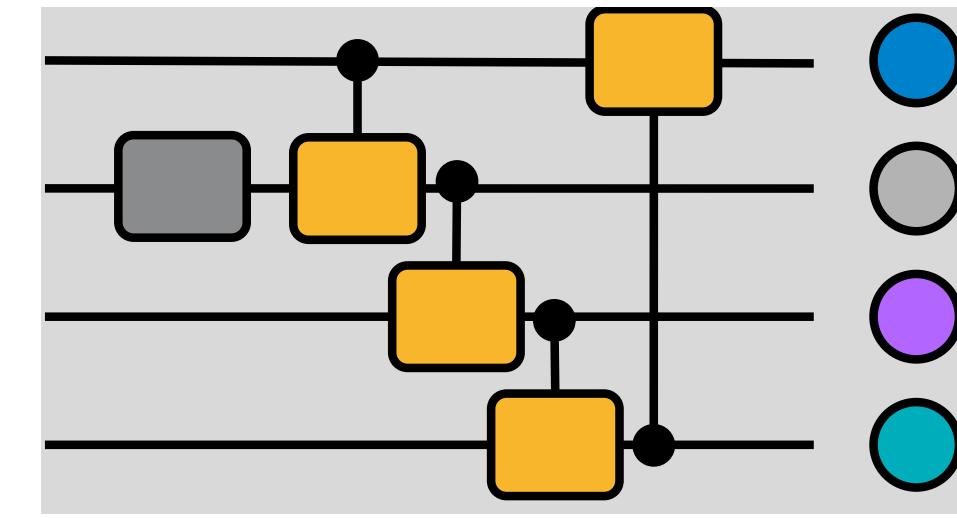
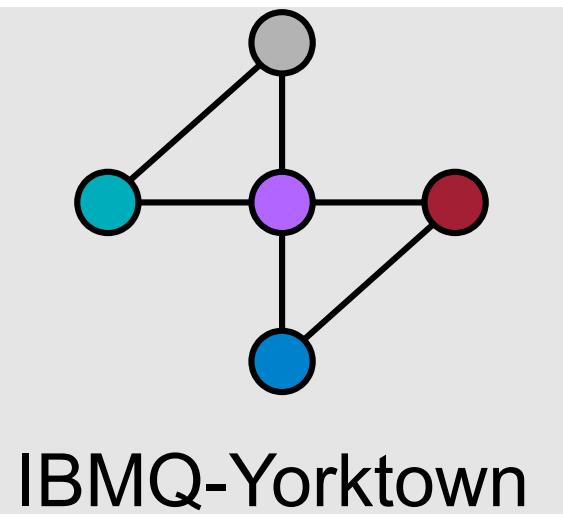


# QuantumNAS

- SuperCircuit Construction and Training
- Noise-Adaptive Evolutionary Co-Search of SubCircuit and Qubit Mapping
- **Searched SubCircuit Training**
- Iterative Pruning

# Searched SubCircuit Training

- Train the SubCircuit parameters from scratch and get final parameters

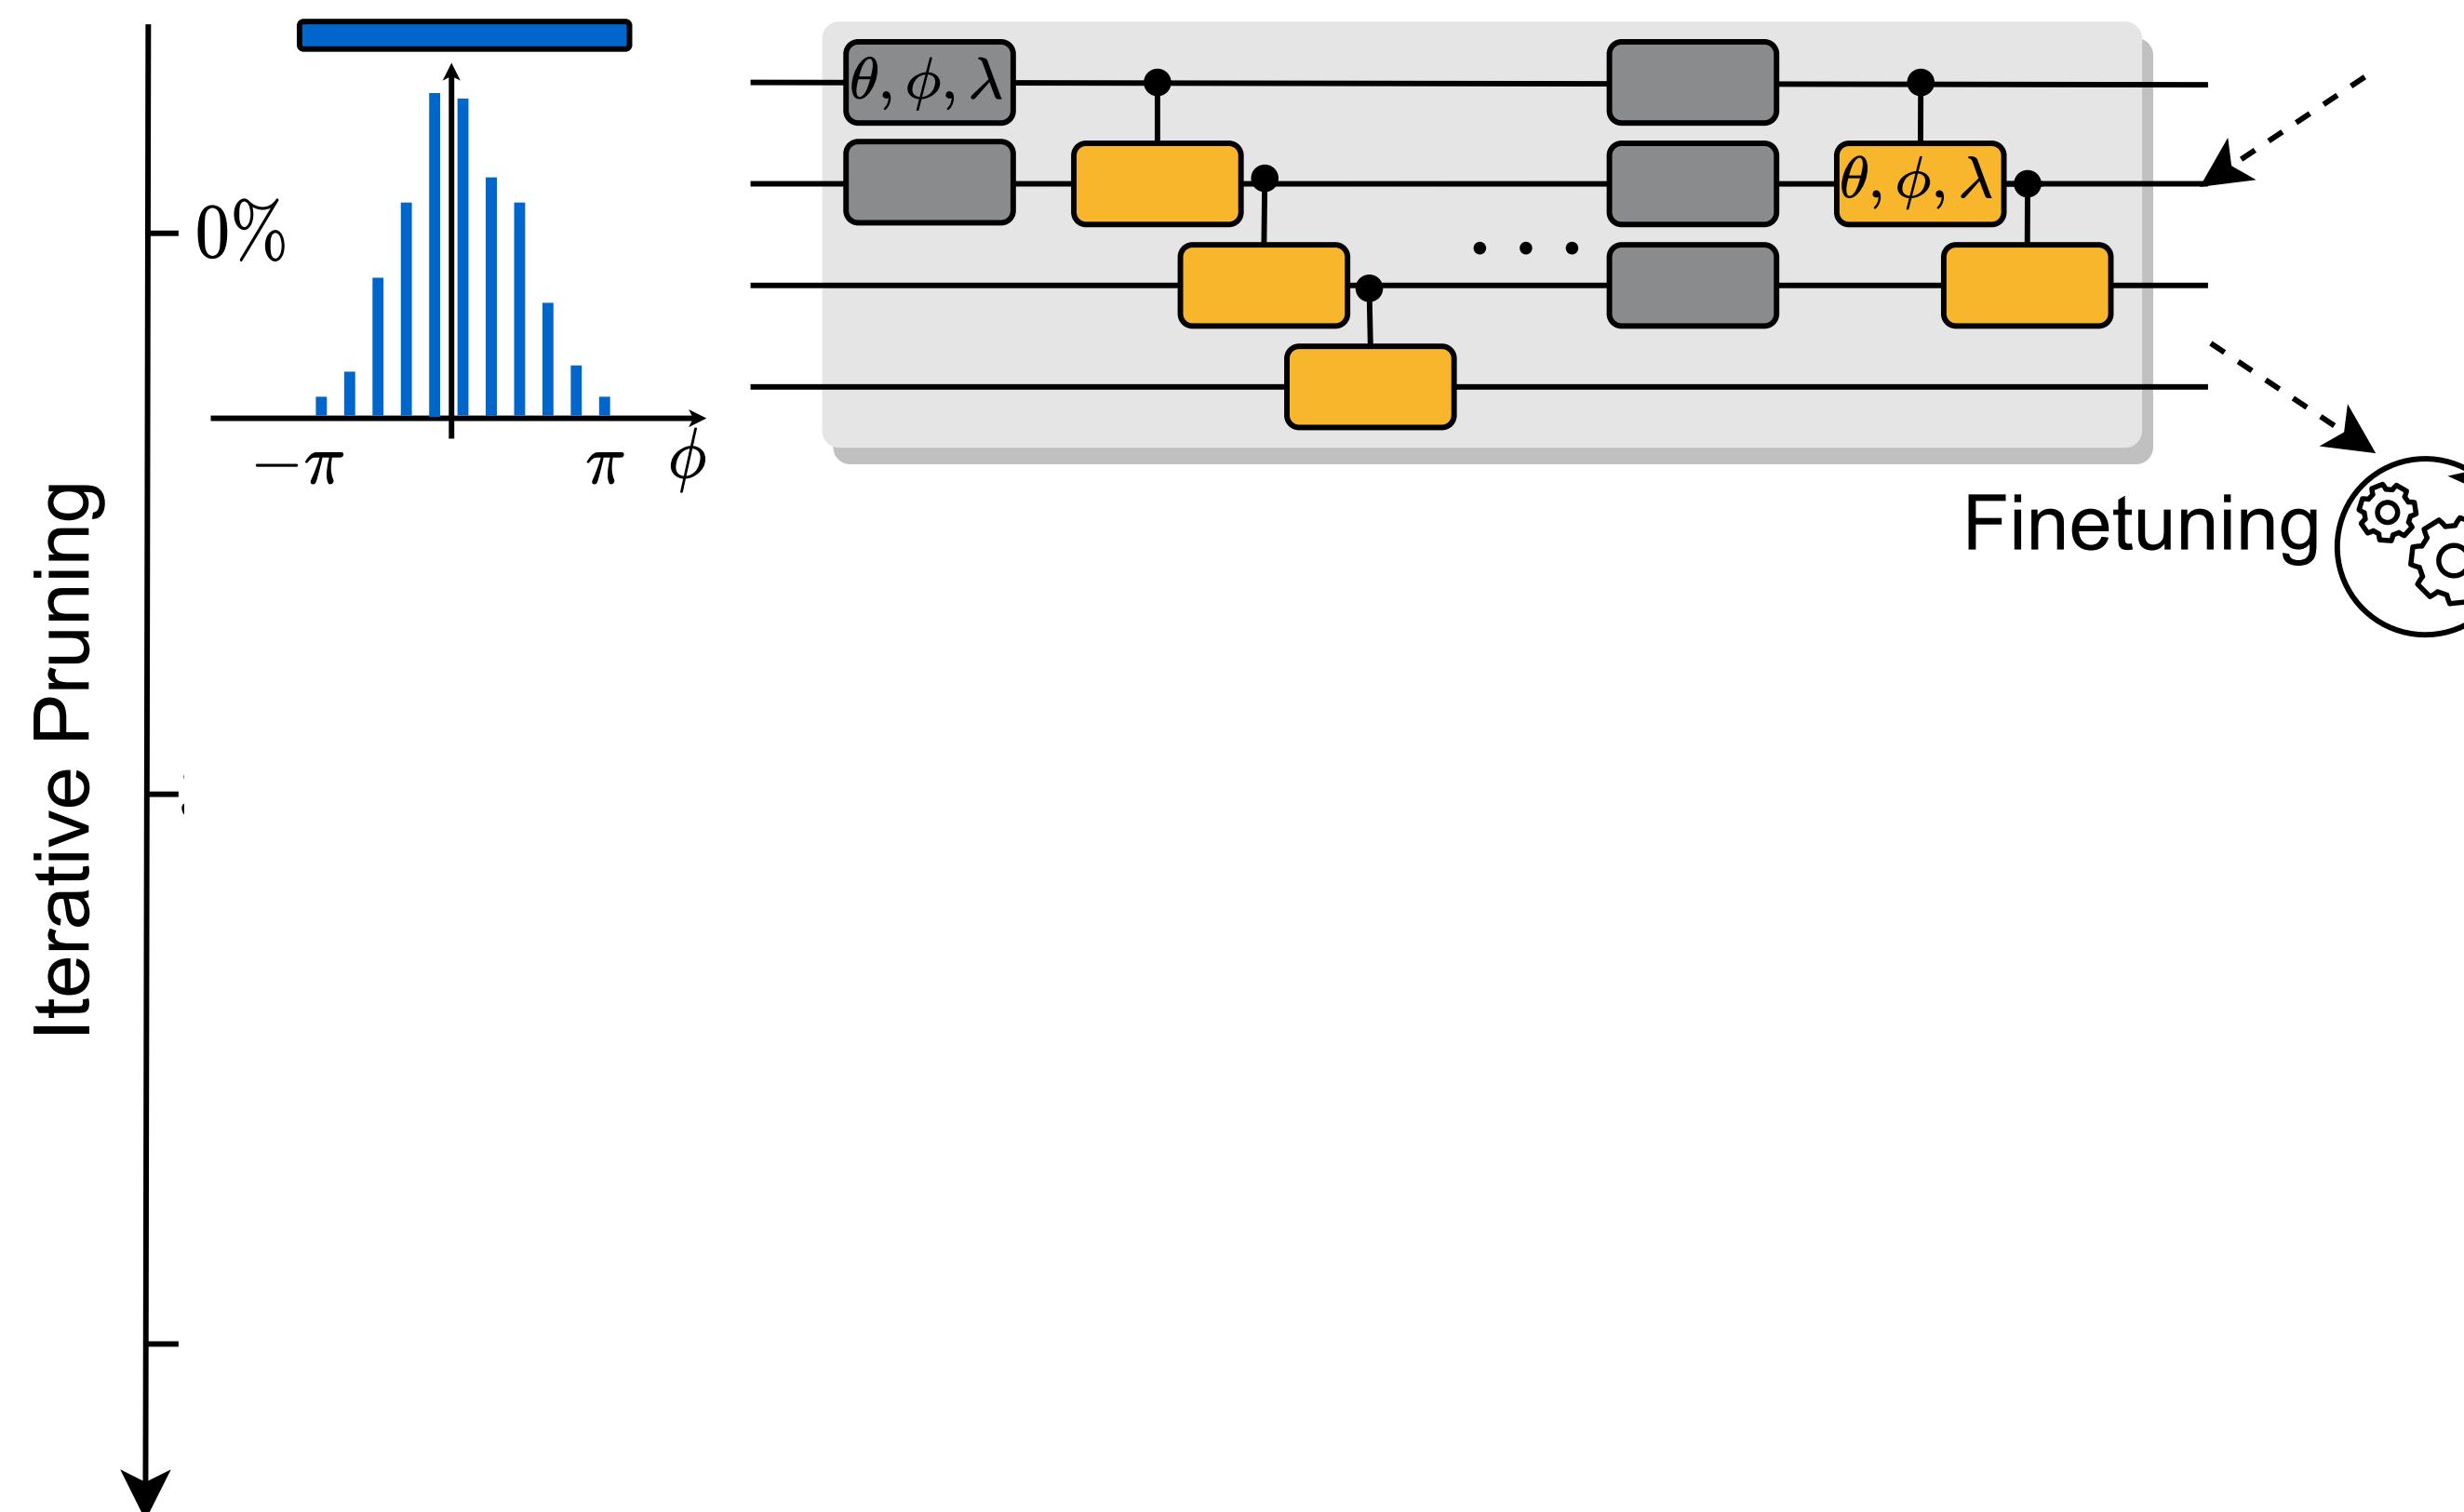


# QuantumNAS

- SuperCircuit Construction and Training
- Noise-Adaptive Evolutionary Co-Search of SubCircuit and Qubit Mapping
- Searched SubCircuit Training
- Iterative Pruning

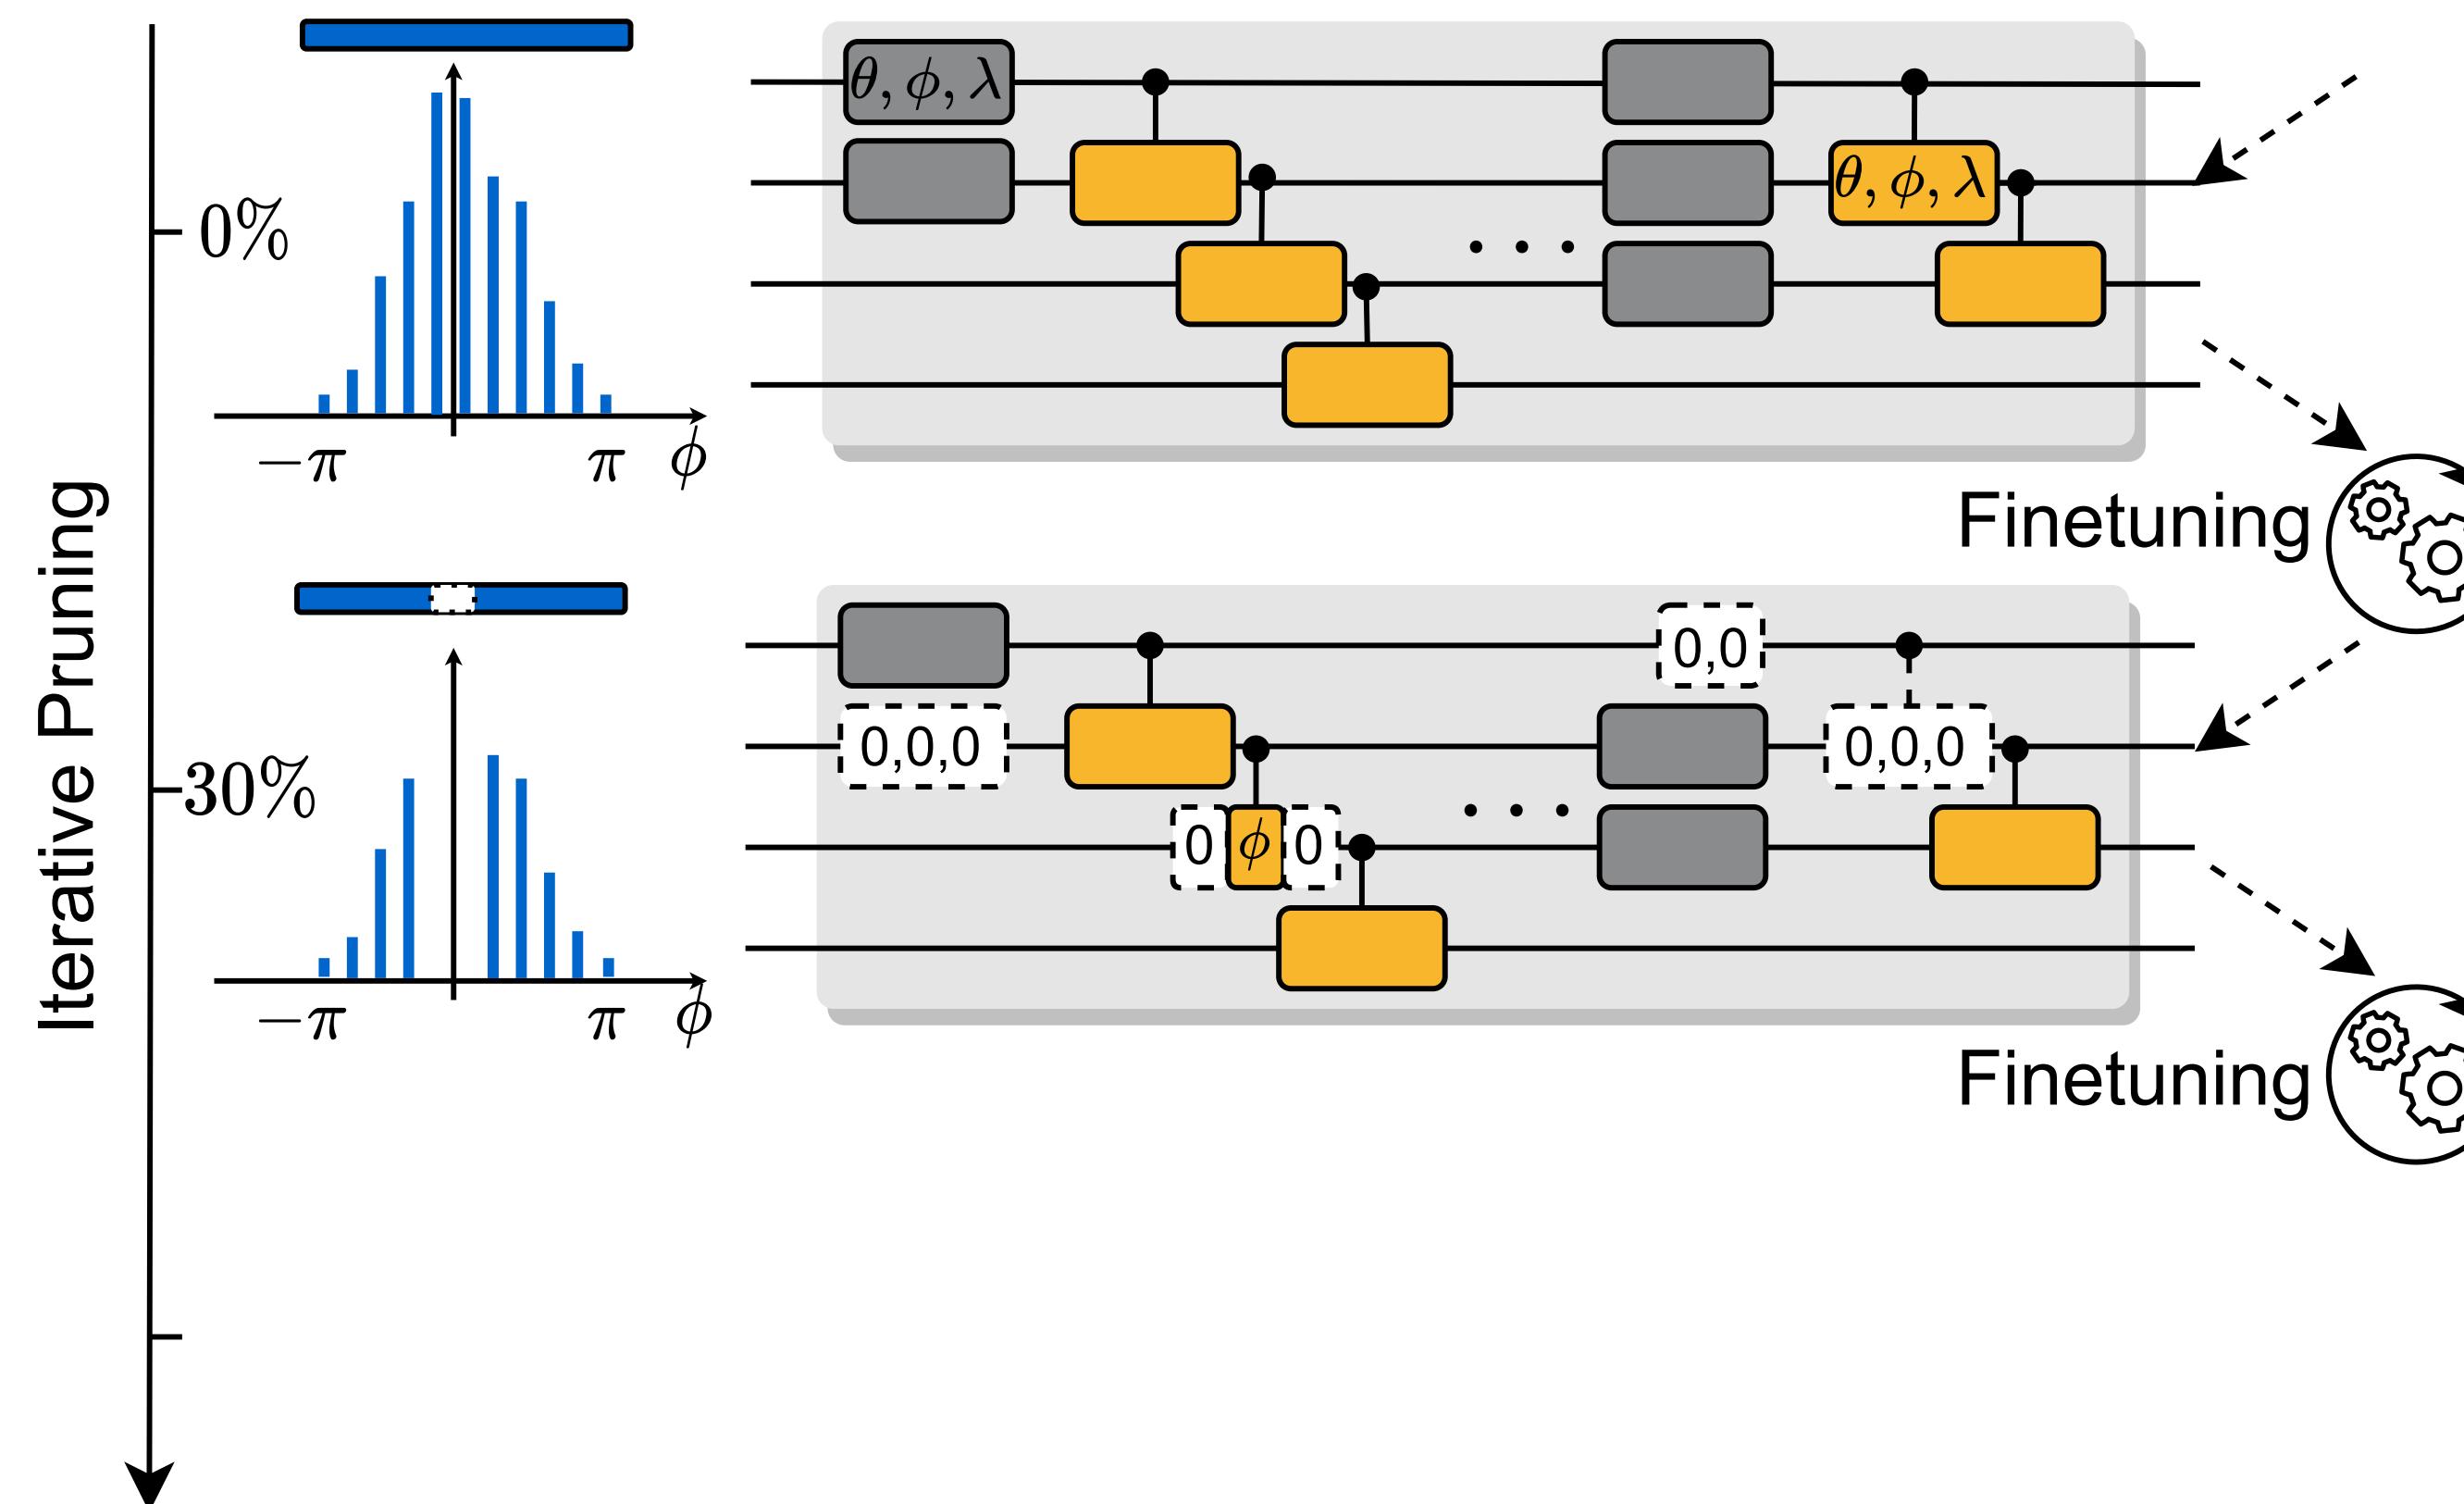
# Iterative Pruning

- Some gates has parameters close to 0
  - Rotation gate with angle close to 0 has small impact on the state
  - Iteratively remove small-magnitude gates and fine-tune the remaining parameters



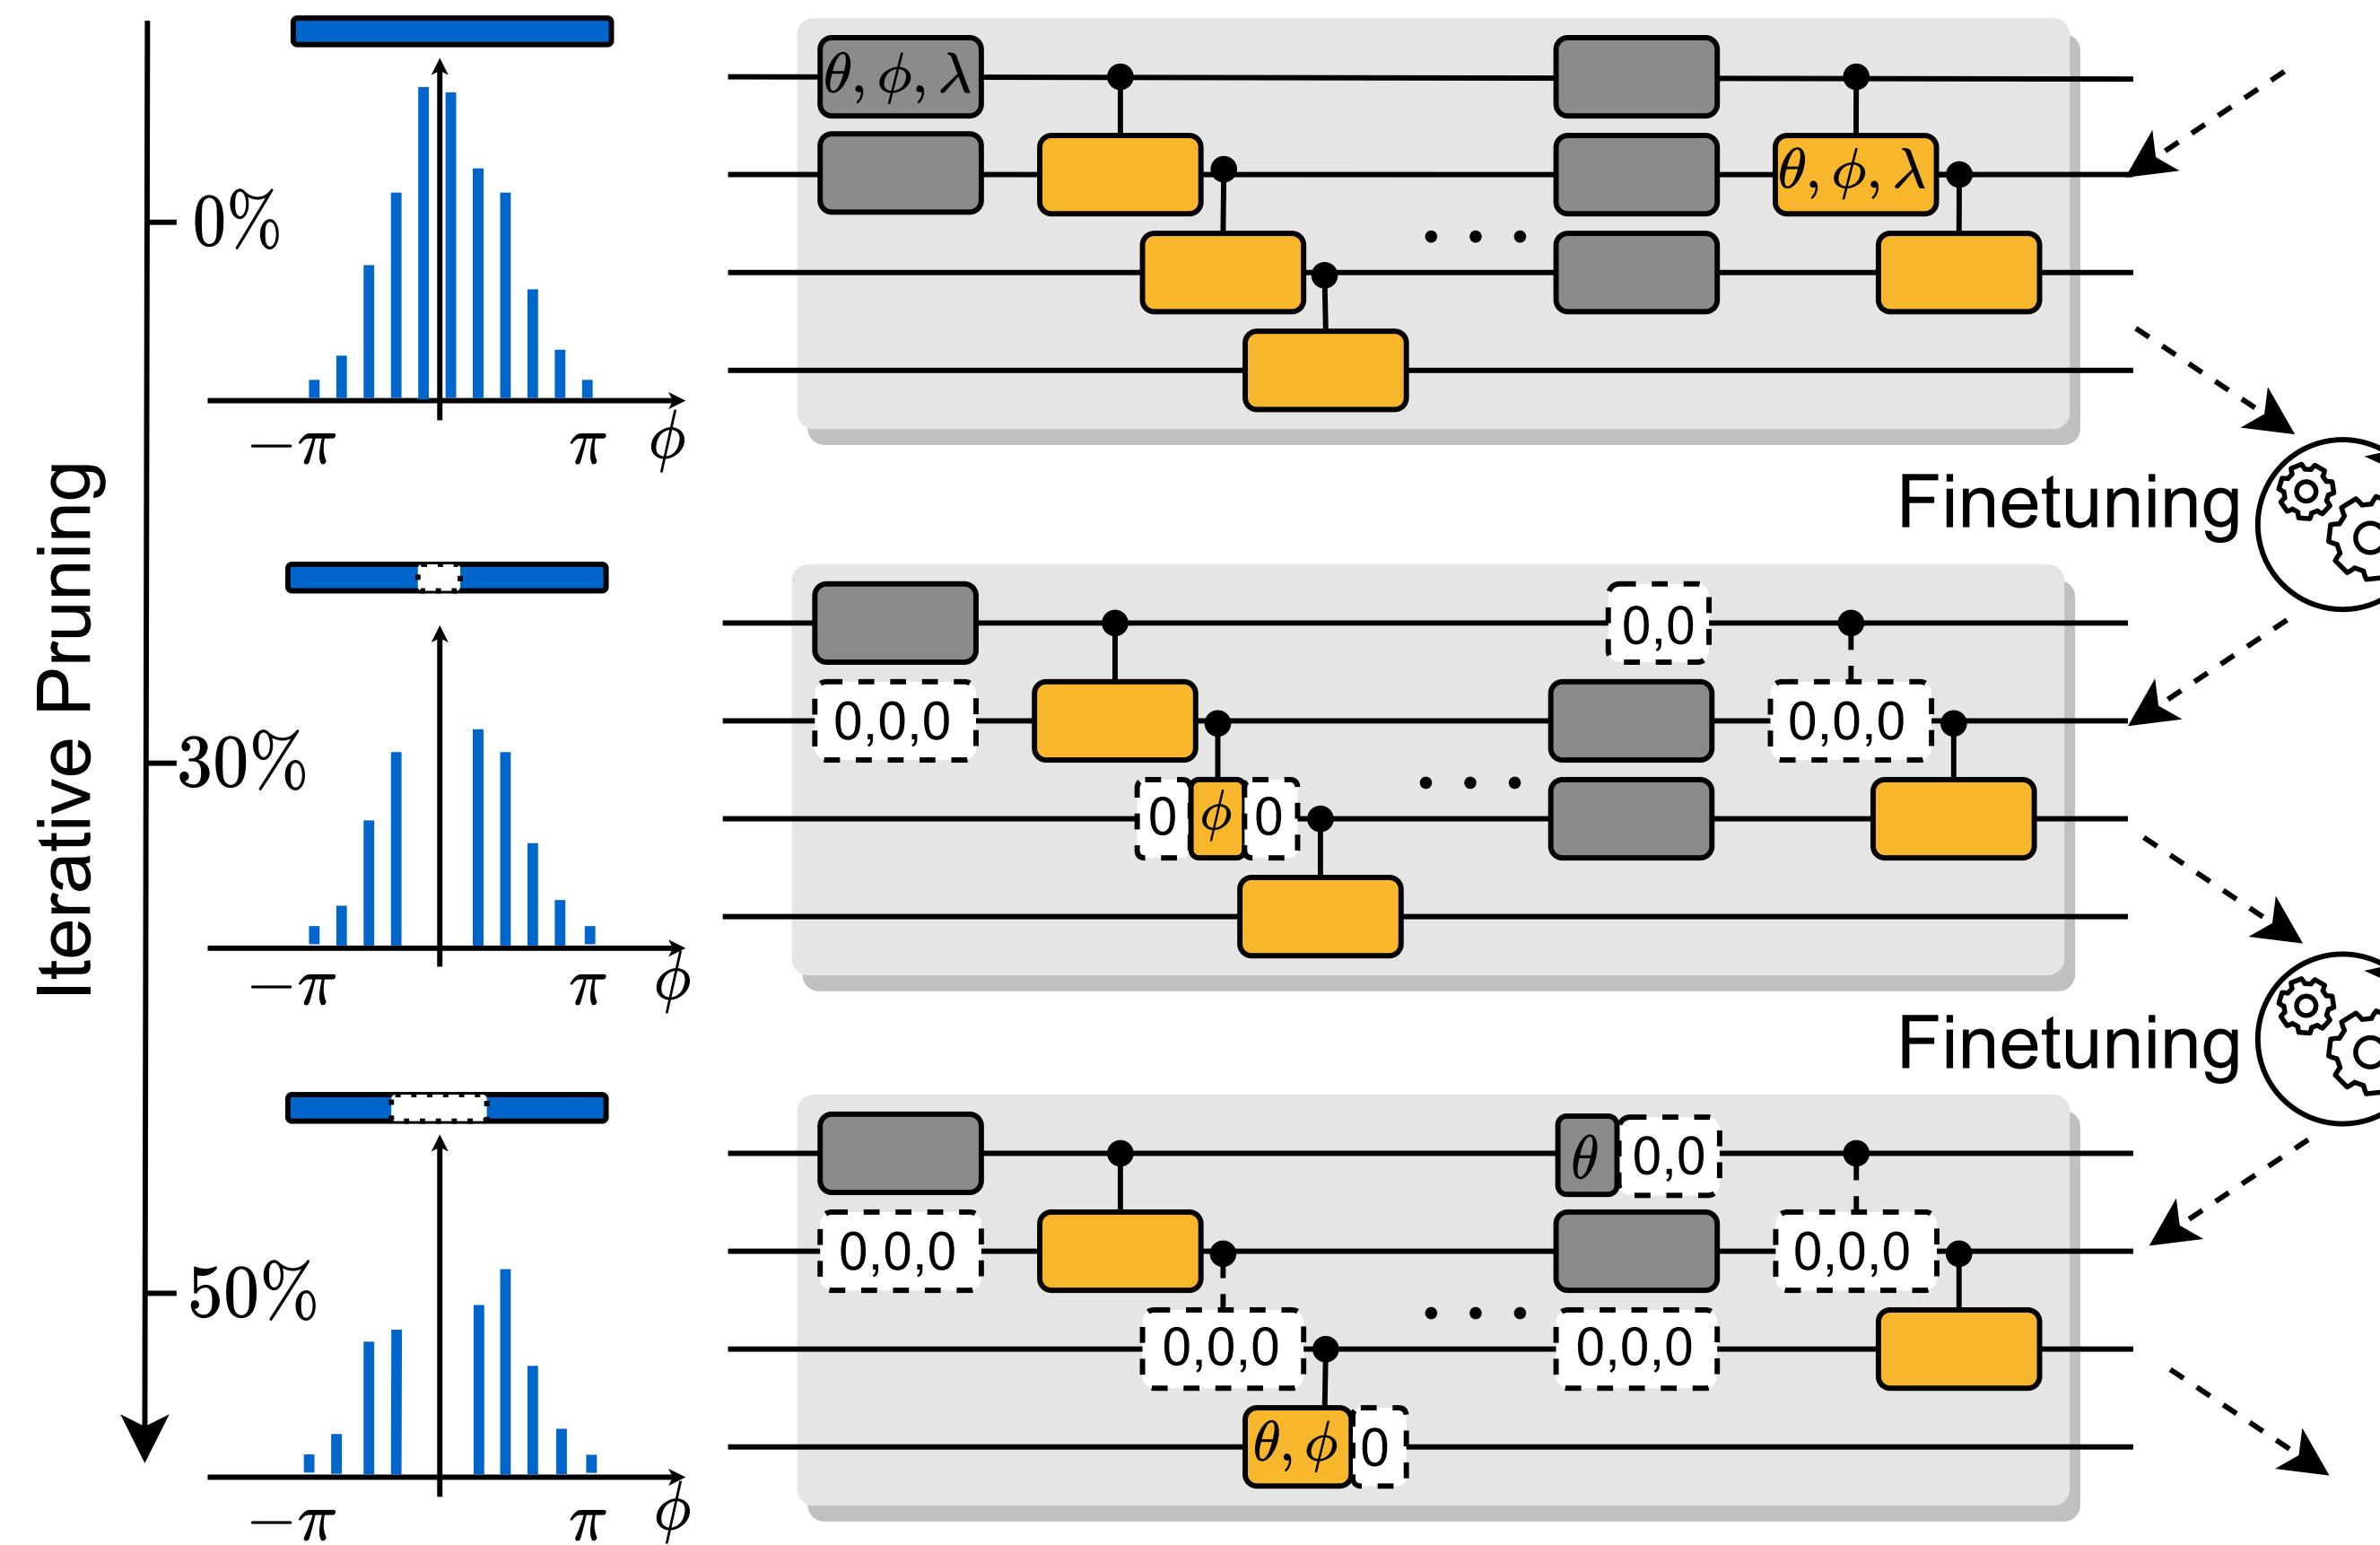
# Iterative Pruning

- Some gates has parameters close to 0
  - Rotation gate with angle close to 0 has small impact on the state
  - Iteratively remove small-magnitude gates and fine-tune the remaining parameters



# Iterative Pruning

- Some gates has parameters close to 0
  - Rotation gate with angle close to 0 has small impact on the state
  - Iteratively remove small-magnitude gates and fine-tune the remaining parameters



# Parameterized QC Library — QuantumEngine

- Easy construction of parameterized quantum circuits in PyTorch.
- Support **batch** mode inference and training on CPU/GPU.
- Support both **static and dynamic** computation graph for easy debugging.
- Support **easy deployment** on real quantum devices such as IBMQ.

# Parameterized QC Library — QuantumEngine

```
import torch.nn as nn
import torch.nn.functional as F
import torchquantum as tq
import torchquantum.functional as tqf

class QFCModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.n_wires = 4
        self.q_device = tq.QuantumDevice(n_wires=self.n_wires)
        self.measure = tq.MeasureAll(tq.PauliZ)

        self.encoder_gates = [tqf.rx] * 4 + [tqf.ry] * 4 + \
                            [tqf.rz] * 4 + [tqf.rx] * 4
        self.rx0 = tq.RX(has_params=True, trainable=True)
        self.ry0 = tq.RY(has_params=True, trainable=True)
        self.rz0 = tq.RZ(has_params=True, trainable=True)
        self.crx0 = tq.CRX(has_params=True, trainable=True)
```

```
def forward(self, x):
    bsz = x.shape[0]
    # down-sample the image
    x = F.avg_pool2d(x, 6).view(bsz, 16)

    # reset qubit states
    self.q_device.reset_states(bsz)

    # encode the classical image to quantum domain
    for k, gate in enumerate(self.encoder_gates):
        gate(self.q_device, wires=k % self.n_wires, params=x[:, k])

    # add some trainable gates (need to instantiate ahead of time)
    self.rx0(self.q_device, wires=0)
    self.ry0(self.q_device, wires=1)
    self.rz0(self.q_device, wires=3)
    self.crx0(self.q_device, wires=[0, 2])

    # add some more non-parameterized gates (add on-the-fly)
    tqf.hadamard(self.q_device, wires=3)
    tqf.sx(self.q_device, wires=2)
    tqf.cnot(self.q_device, wires=[3, 0])
    tqf.qubitunitary(self.q_device0, wires=[1, 2], params=[[1, 0, 0, 0],
                                                           [0, 1, 0, 0],
                                                           [0, 0, 0, 1j],
                                                           [0, 0, -1j, 0]])

    # perform measurement to get expectations (back to classical domain)
    x = self.measure(self.q_device).reshape(bsz, 2, 2)

    # classification
    x = x.sum(-1).squeeze()
    x = F.log_softmax(x, dim=1)

    return x
```

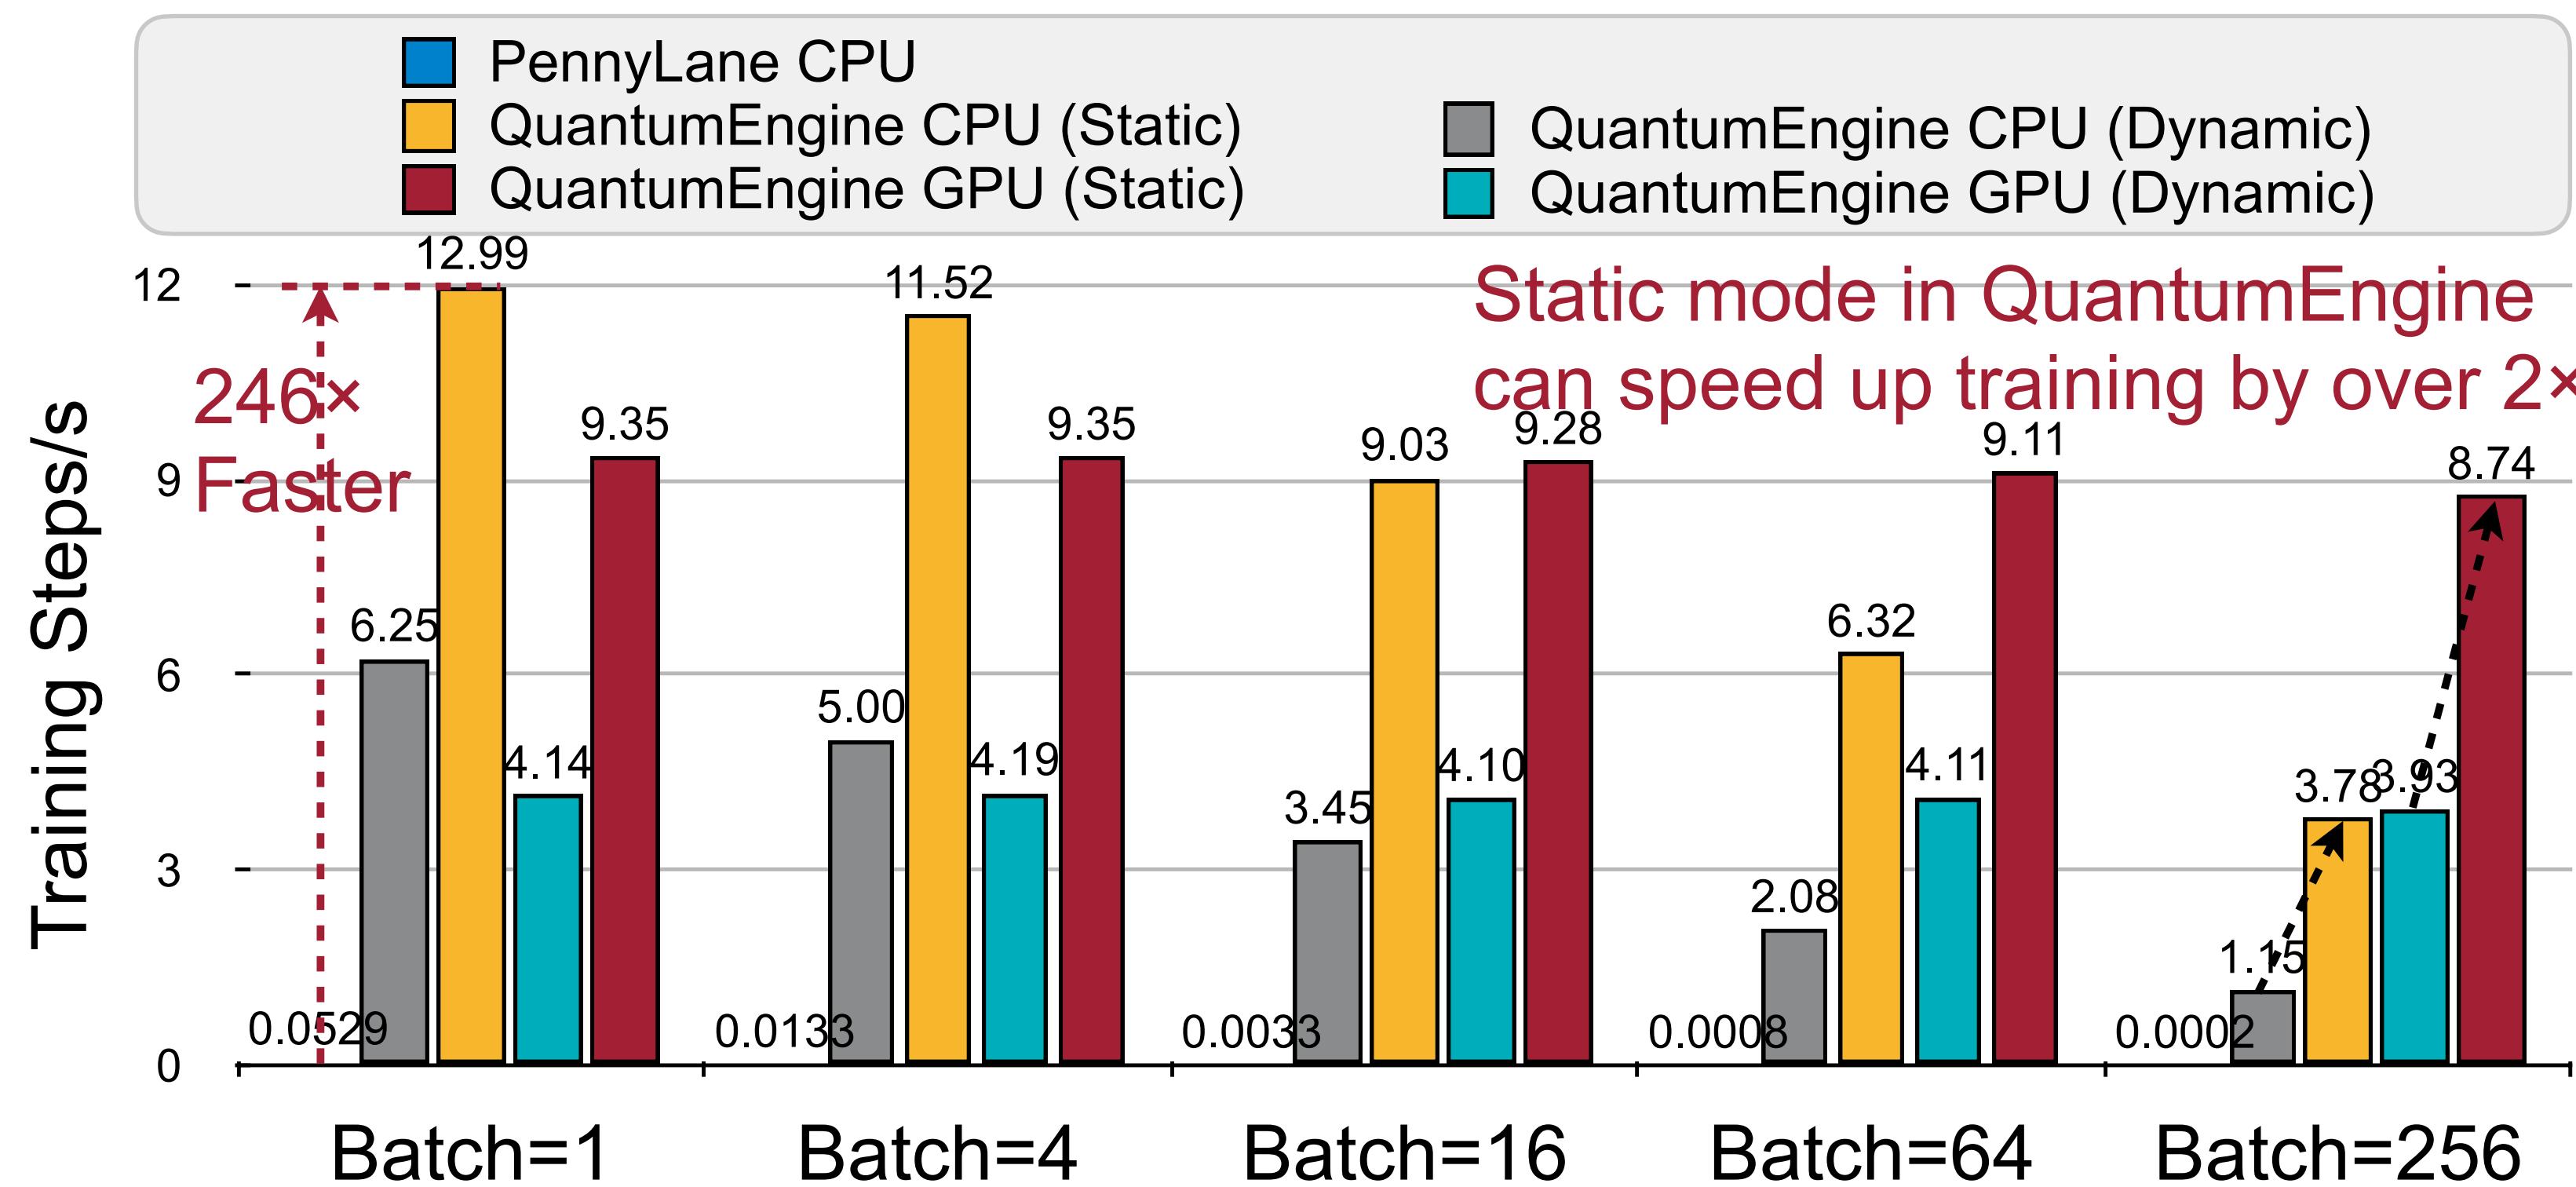
# Parameterized QC Library — QuantumEngine

- Easy deployment:

```
# then try to run on REAL QC
backend_name = 'ibmqx2'
print(f"\nTest on Real Quantum Computer {backend_name}")
processor_real_qc = QiskitProcessor(use_real_qc=True,
                                     backend_name=backend_name)
model.set_qiskit_processor(processor_real_qc)
valid_test(dataflow, 'test', model, device, qiskit=True)
```

# Parameterized QC Library — QuantumEngine

- Fast training speed:



# Outline

- Introduction
- Background
- QuantumNAS
  - SuperCircuit Training
  - Noise-Adaptive Co-Search
  - Searched SubCircuit Training
  - Gate Pruning
- Evaluation
- Conclusion

# Evaluation Setups: Benchmarks and Devices

- Benchmarks
  - QML classification tasks: MNIST 10-class, 4-class, 2-class, Fashion 4-class, 2-class, Vowel 4-class
  - VQE task molecules: H<sub>2</sub>, H<sub>2</sub>O, LiH, CH<sub>4</sub>, BeH<sub>2</sub>
- Quantum Devices
  - IBMQ
  - #Qubits: 5 to 65
  - Quantum Volume: 8 to 128

# Evaluation Setups: Design Spaces

- Circuit Design Spaces
  - ‘U3+CU3’: one block has one U3 layer with a U3 gate on each qubit, and one CU3 layer with ring connections, 8 blocks
  - ‘ZZ+RY’: one block has one ZZ layer, and one RY layer, 8 blocks
  - ‘RXYZ’: one block has one RX, RY, RZ and CZ layer, 8 blocks
  - ‘ZZ+XX’: one block has one ZZ layer, and one XX layer, 8 blocks
  - ‘RXYZ+U1+CU3’: one block has 11 layers: RX, S, CNOT, RY, T, SWAP, RZ, H,  $\sqrt{\text{SWAP}}$ , U1, and CU3, 4 blocks
  - ‘IBMQ Basis’: use IBMQ basis gate set: one block has 6 layers: RZ, X, RZ, SX, RZ, CNOT, 20 blocks

# Evaluation Setups: Baselines

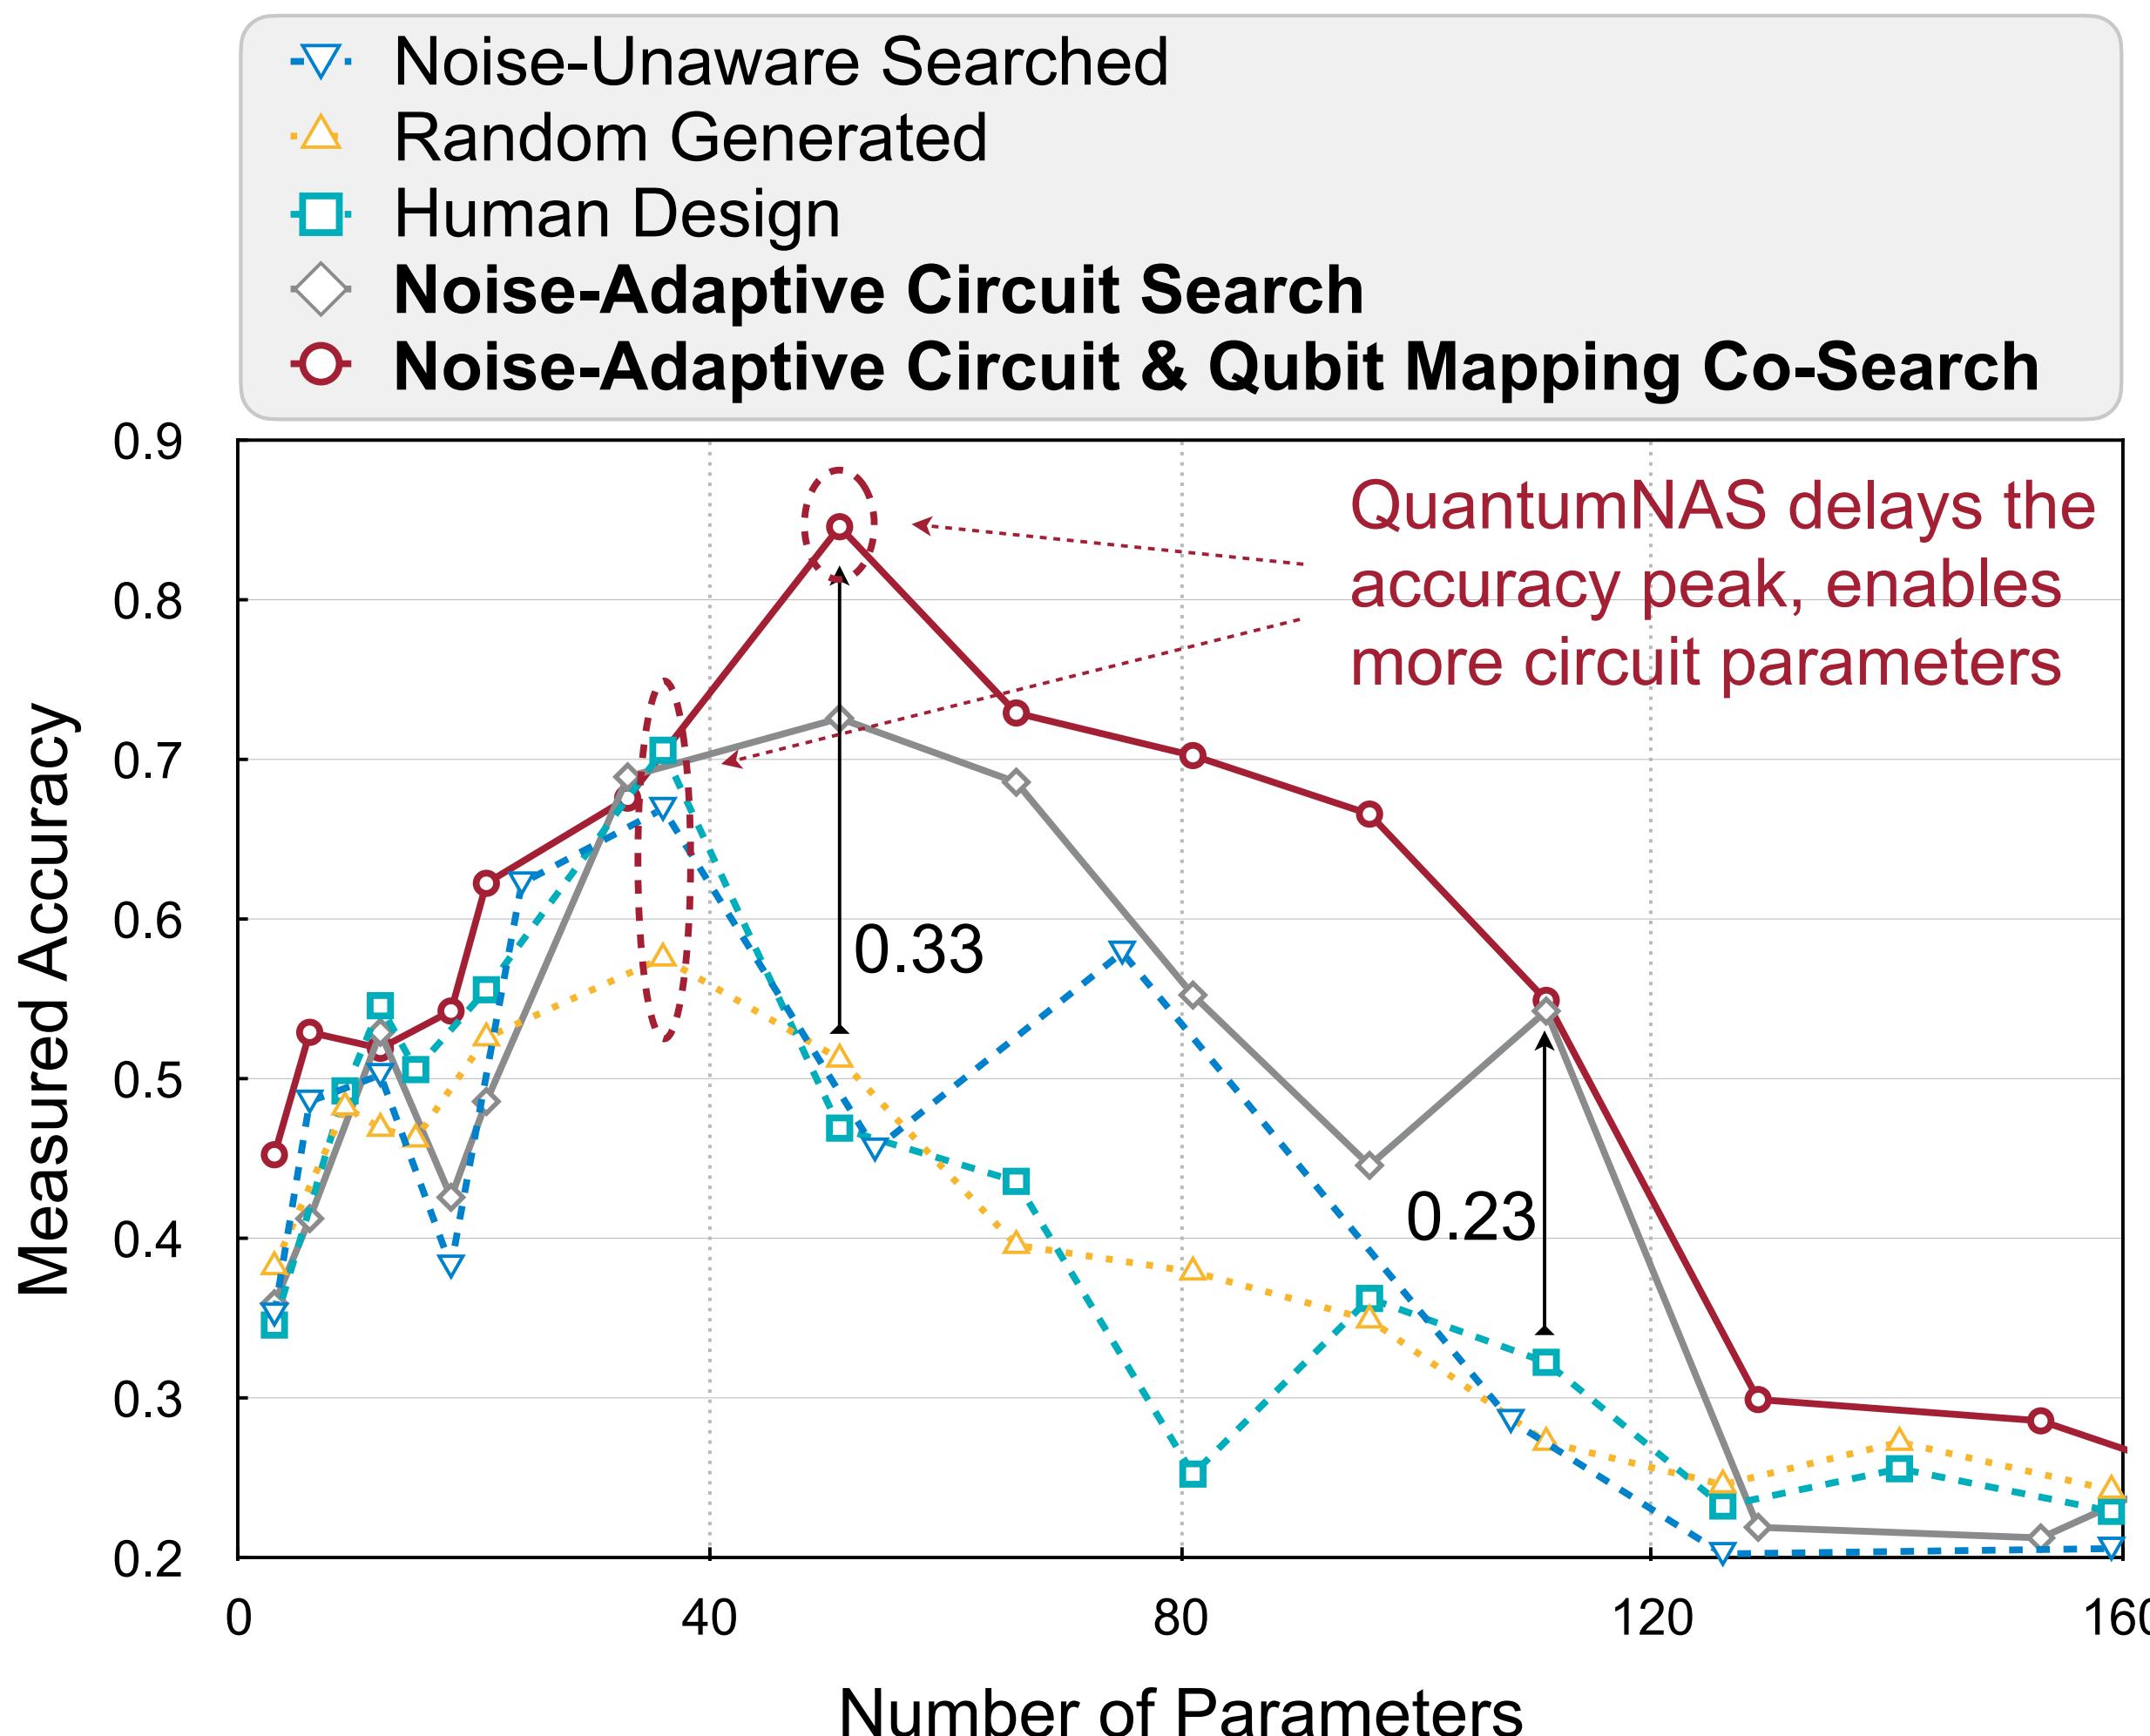
- 6 Baselines for QML/QVE
  - *Noise-Unaware Search*: no noise information is involved during search
  - *Random generation*: same gate set, same #parameters with QuantumNAS
  - *Human Design*: same #parameters with QuantumNAS, full width in the front blocks
  - *Human+noise-adaptive mapping*: human design with noise-adaptive qubit mapping<sup>1</sup>
  - *Human+Sabre mapping*: human design with Sabre qubit mapping<sup>2</sup>
  - *Human(1/2 #params)+Sabre mapping*: 1/2 #parameters with QuantumNAS with Sabre qubit mapping
- UCCSD for VQE

<sup>1</sup>Murali, Prakash, et al. "Noise-adaptive compiler mappings for noisy intermediate-scale quantum computers." ASPLOS. 2019.

<sup>2</sup>Li, Gushu, Yufei Ding, and Yuan Xie. "Tackling the qubit mapping problem for NISQ-era quantum devices." ASPLOS. 2019.

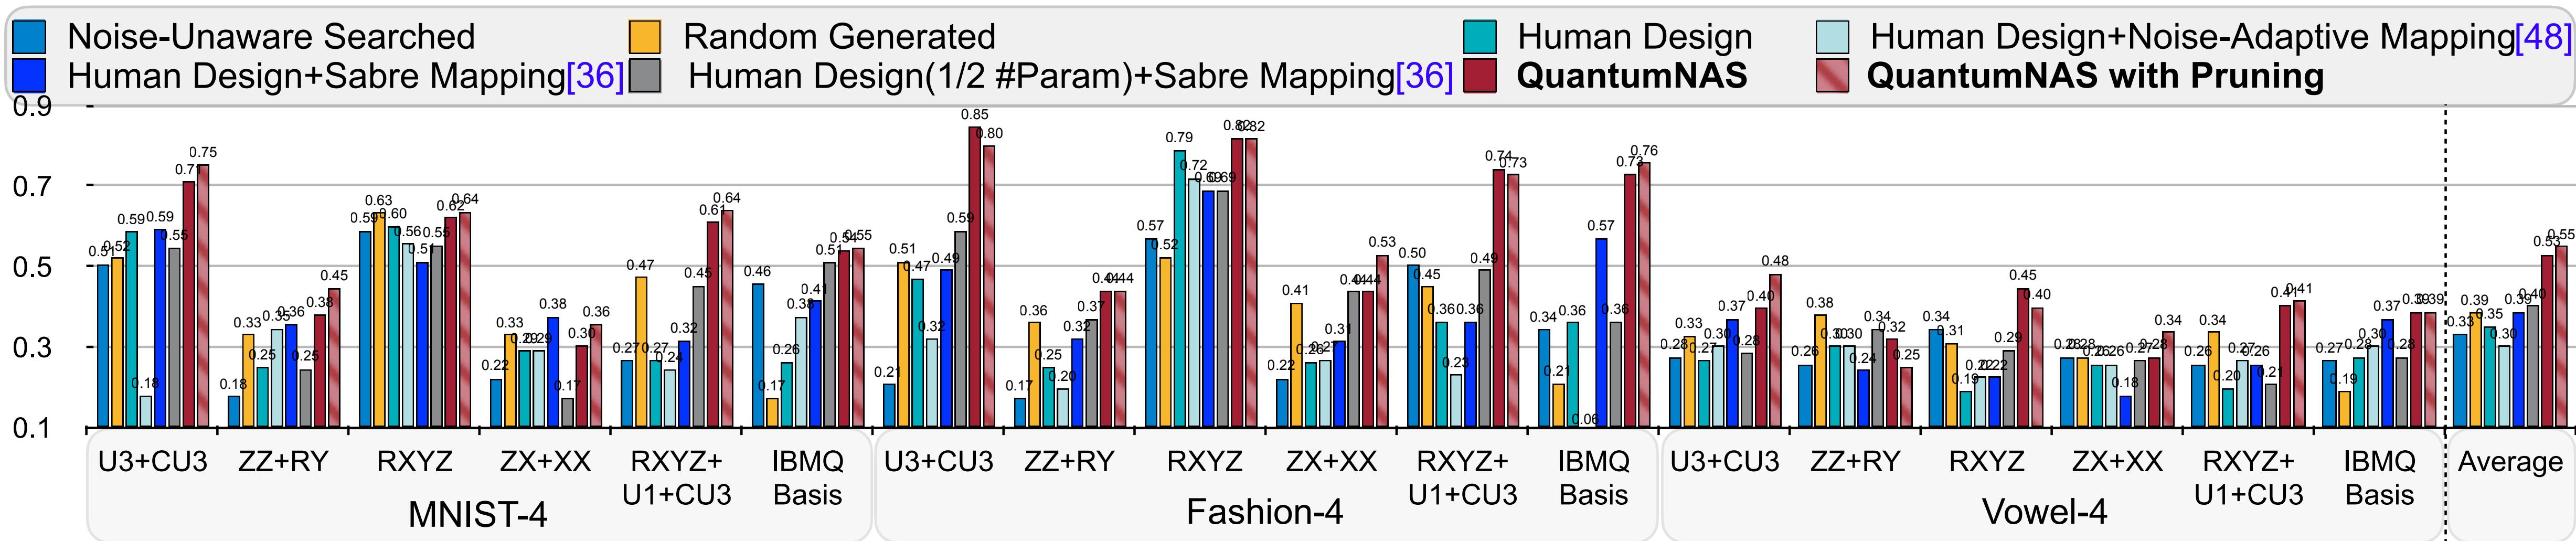
# QML Results

- 4-classification: MNIST-4 U3+CU3 on IBMQ-Yorktown



# QML Results

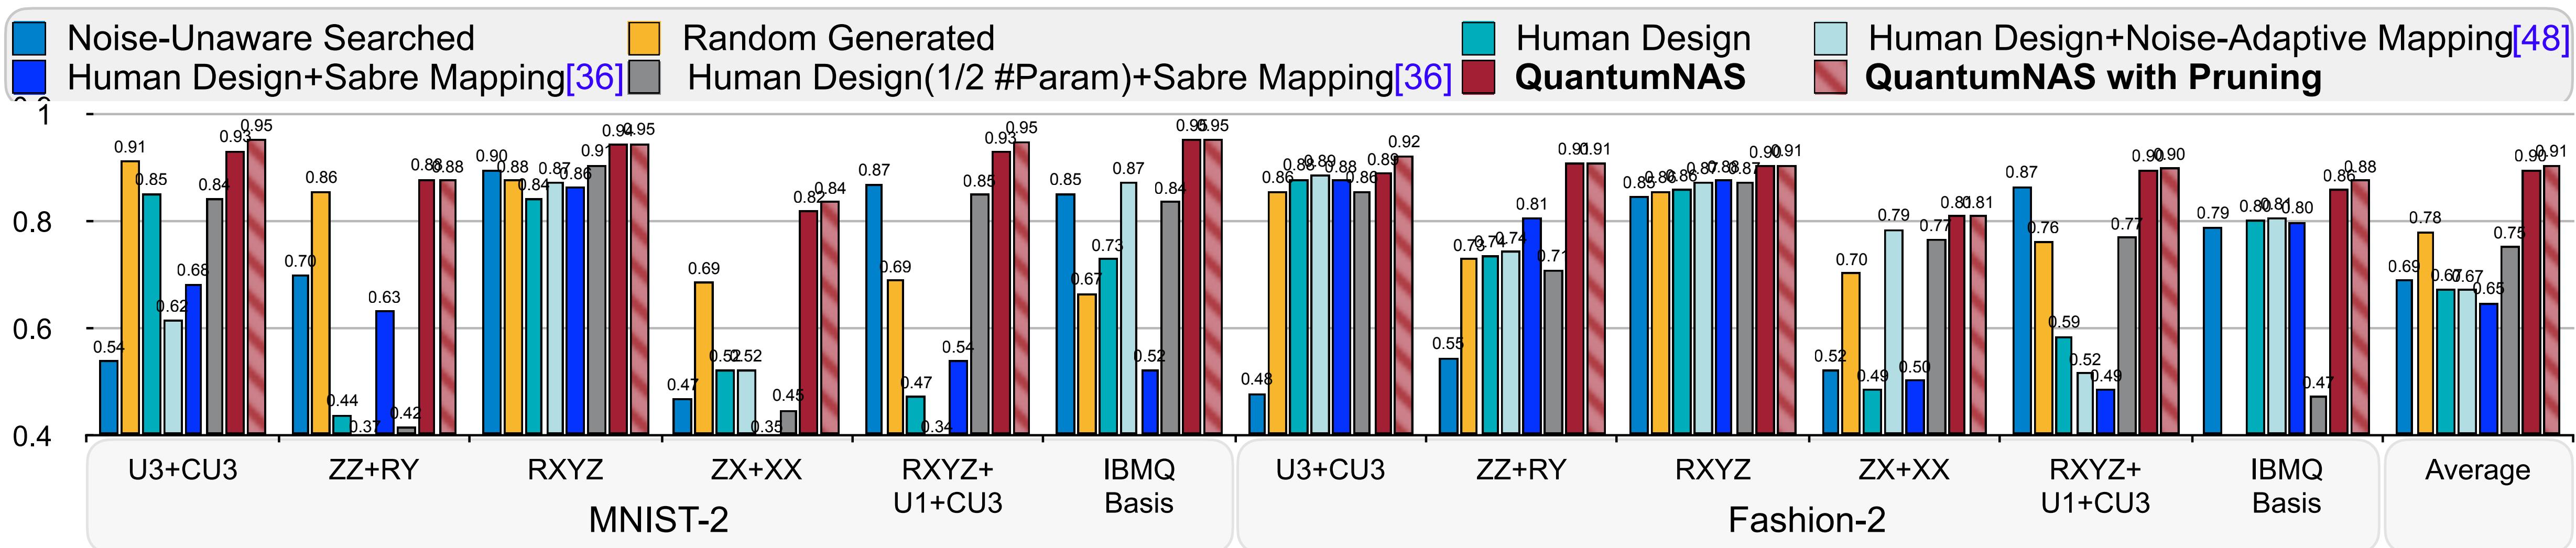
- 4-classifications: MNIST-4, Fashion-4, Vowel-4 in 6 design spaces on IBMQ-Yorktown



| Method           | Noise-Unaware Searched | Random | Human | Human+Noise-adaptive mapping | Human+Sabre mapping | Human(1/2params)+Sabre mapping | QuantumNAS | QuantumNAS w/ Pruning |
|------------------|------------------------|--------|-------|------------------------------|---------------------|--------------------------------|------------|-----------------------|
| Average Accuracy | 0.33                   | 0.39   | 0.35  | 0.30                         | 0.39                | 0.40                           | 0.53       | 0.55                  |

# QML Results

- 2-classifications: MNIST-2 and Fashion-2 in 6 design spaces on IBMQ-Yorktown



| Method           | Noise-Unaware Searched | Random | Human | Human+Noise-adaptive mapping | Human+Sabre mapping | Human(1/2params)+Sabre mapping | QuantumNAS | QuantumNAS w/ Pruning |
|------------------|------------------------|--------|-------|------------------------------|---------------------|--------------------------------|------------|-----------------------|
| Average Accuracy | 0.69                   | 0.78   | 0.67  | 0.67                         | 0.65                | 0.75                           | 0.90       | 0.91                  |

# Results on Different Devices

- On different 5-Qubit devices
- MNIST-4, Fashion-4, Vowel-4, MNIST-2, Fashion-2 averaged

| Method              | Noise-Unaware<br>Searched | Random | Human | QuantumNAS  |
|---------------------|---------------------------|--------|-------|-------------|
| Belem (5Q, 16QV)    | 0.47                      | 0.50   | 0.67  | <b>0.76</b> |
| Quito (5Q, 16QV)    | 0.73                      | 0.68   | 0.74  | <b>0.79</b> |
| Athens (5Q, 32QV)   | 0.50                      | 0.63   | 0.68  | <b>0.77</b> |
| Santiago (5Q, 32QV) | 0.74                      | 0.73   | 0.75  | <b>0.80</b> |

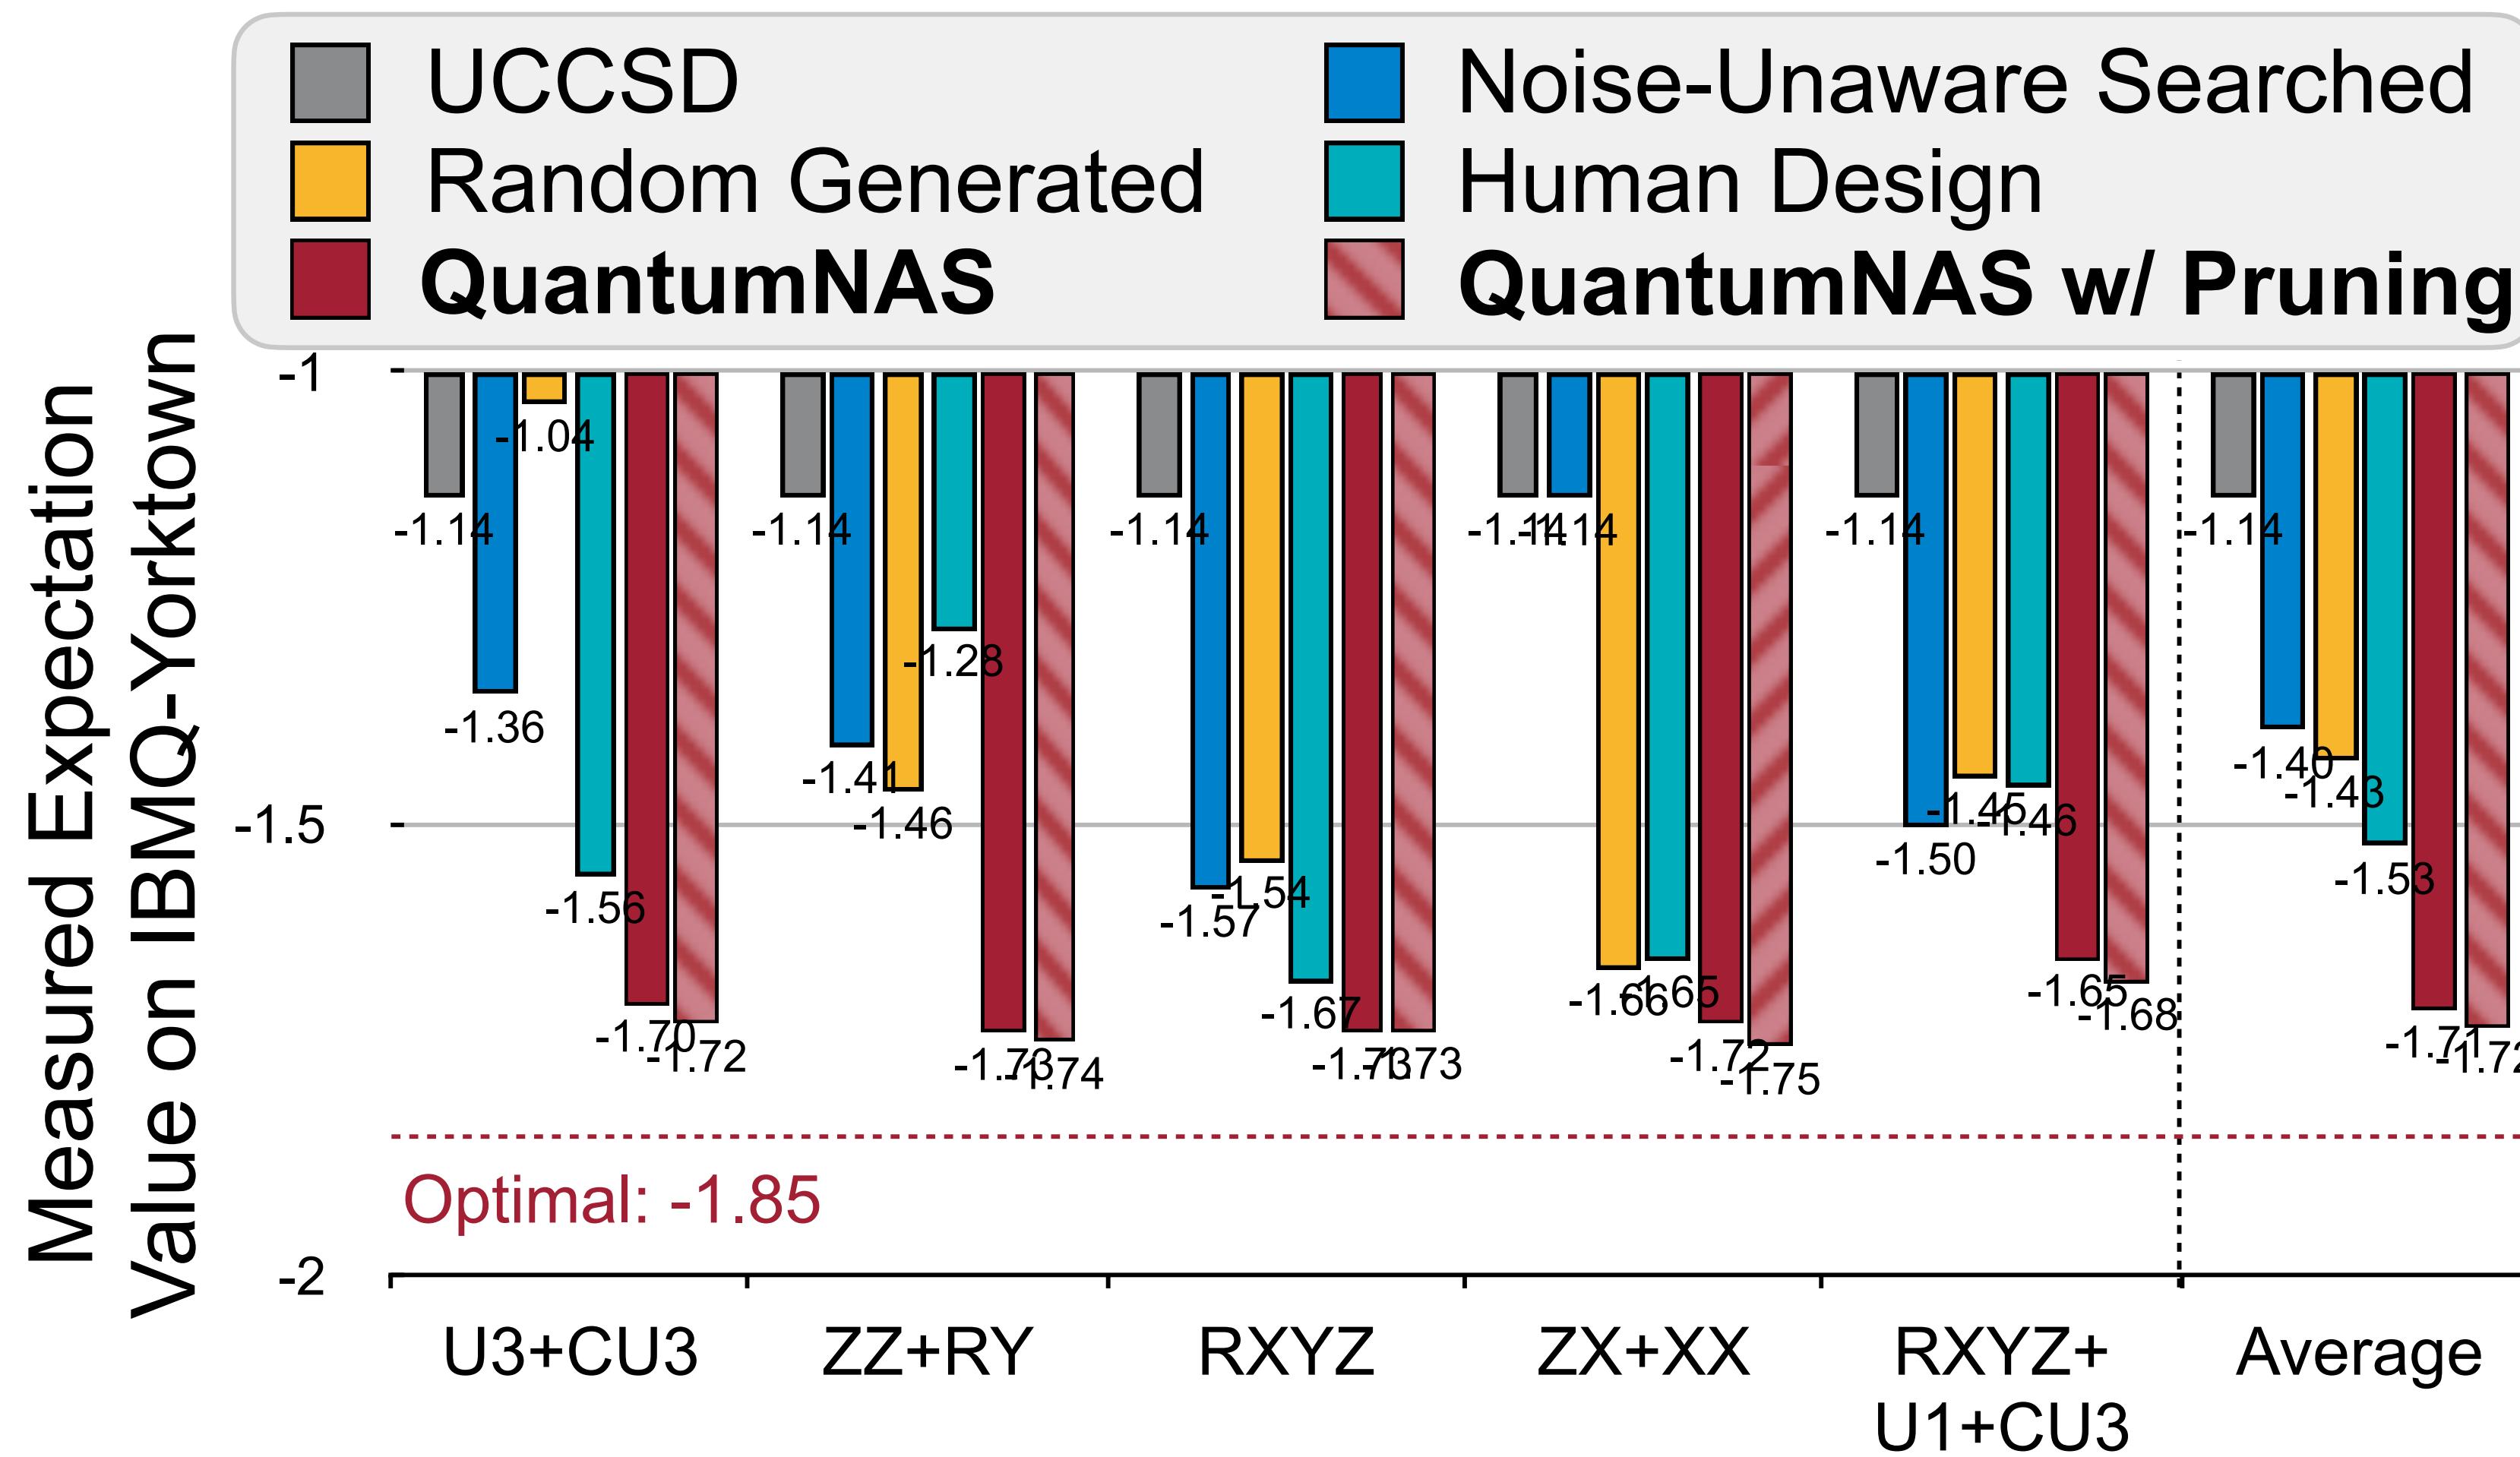
# On Large Devices

- On large devices
- MNIST-4, Fashion-4, Vowel-4, MNIST-2, Fashion-2, MNIST-10 averaged

| Method                         | Noise-Unaware<br>Searched | Random | Human | QuantumNAS  |
|--------------------------------|---------------------------|--------|-------|-------------|
| Melbourne (15Q, 8QV, use 15Q)  | 0.31                      | 0.42   | 0.53  | <b>0.60</b> |
| Guadalupe (16Q, 32QV, use 16Q) | 0.39                      | 0.42   | 0.58  | <b>0.62</b> |
| Montreal (27Q, 128QV, use 21Q) | 0.57                      | 0.48   | 0.61  | <b>0.65</b> |
| Manhattan (65Q, 32QV, use 21Q) | 0.52                      | 0.53   | 0.62  | <b>0.66</b> |

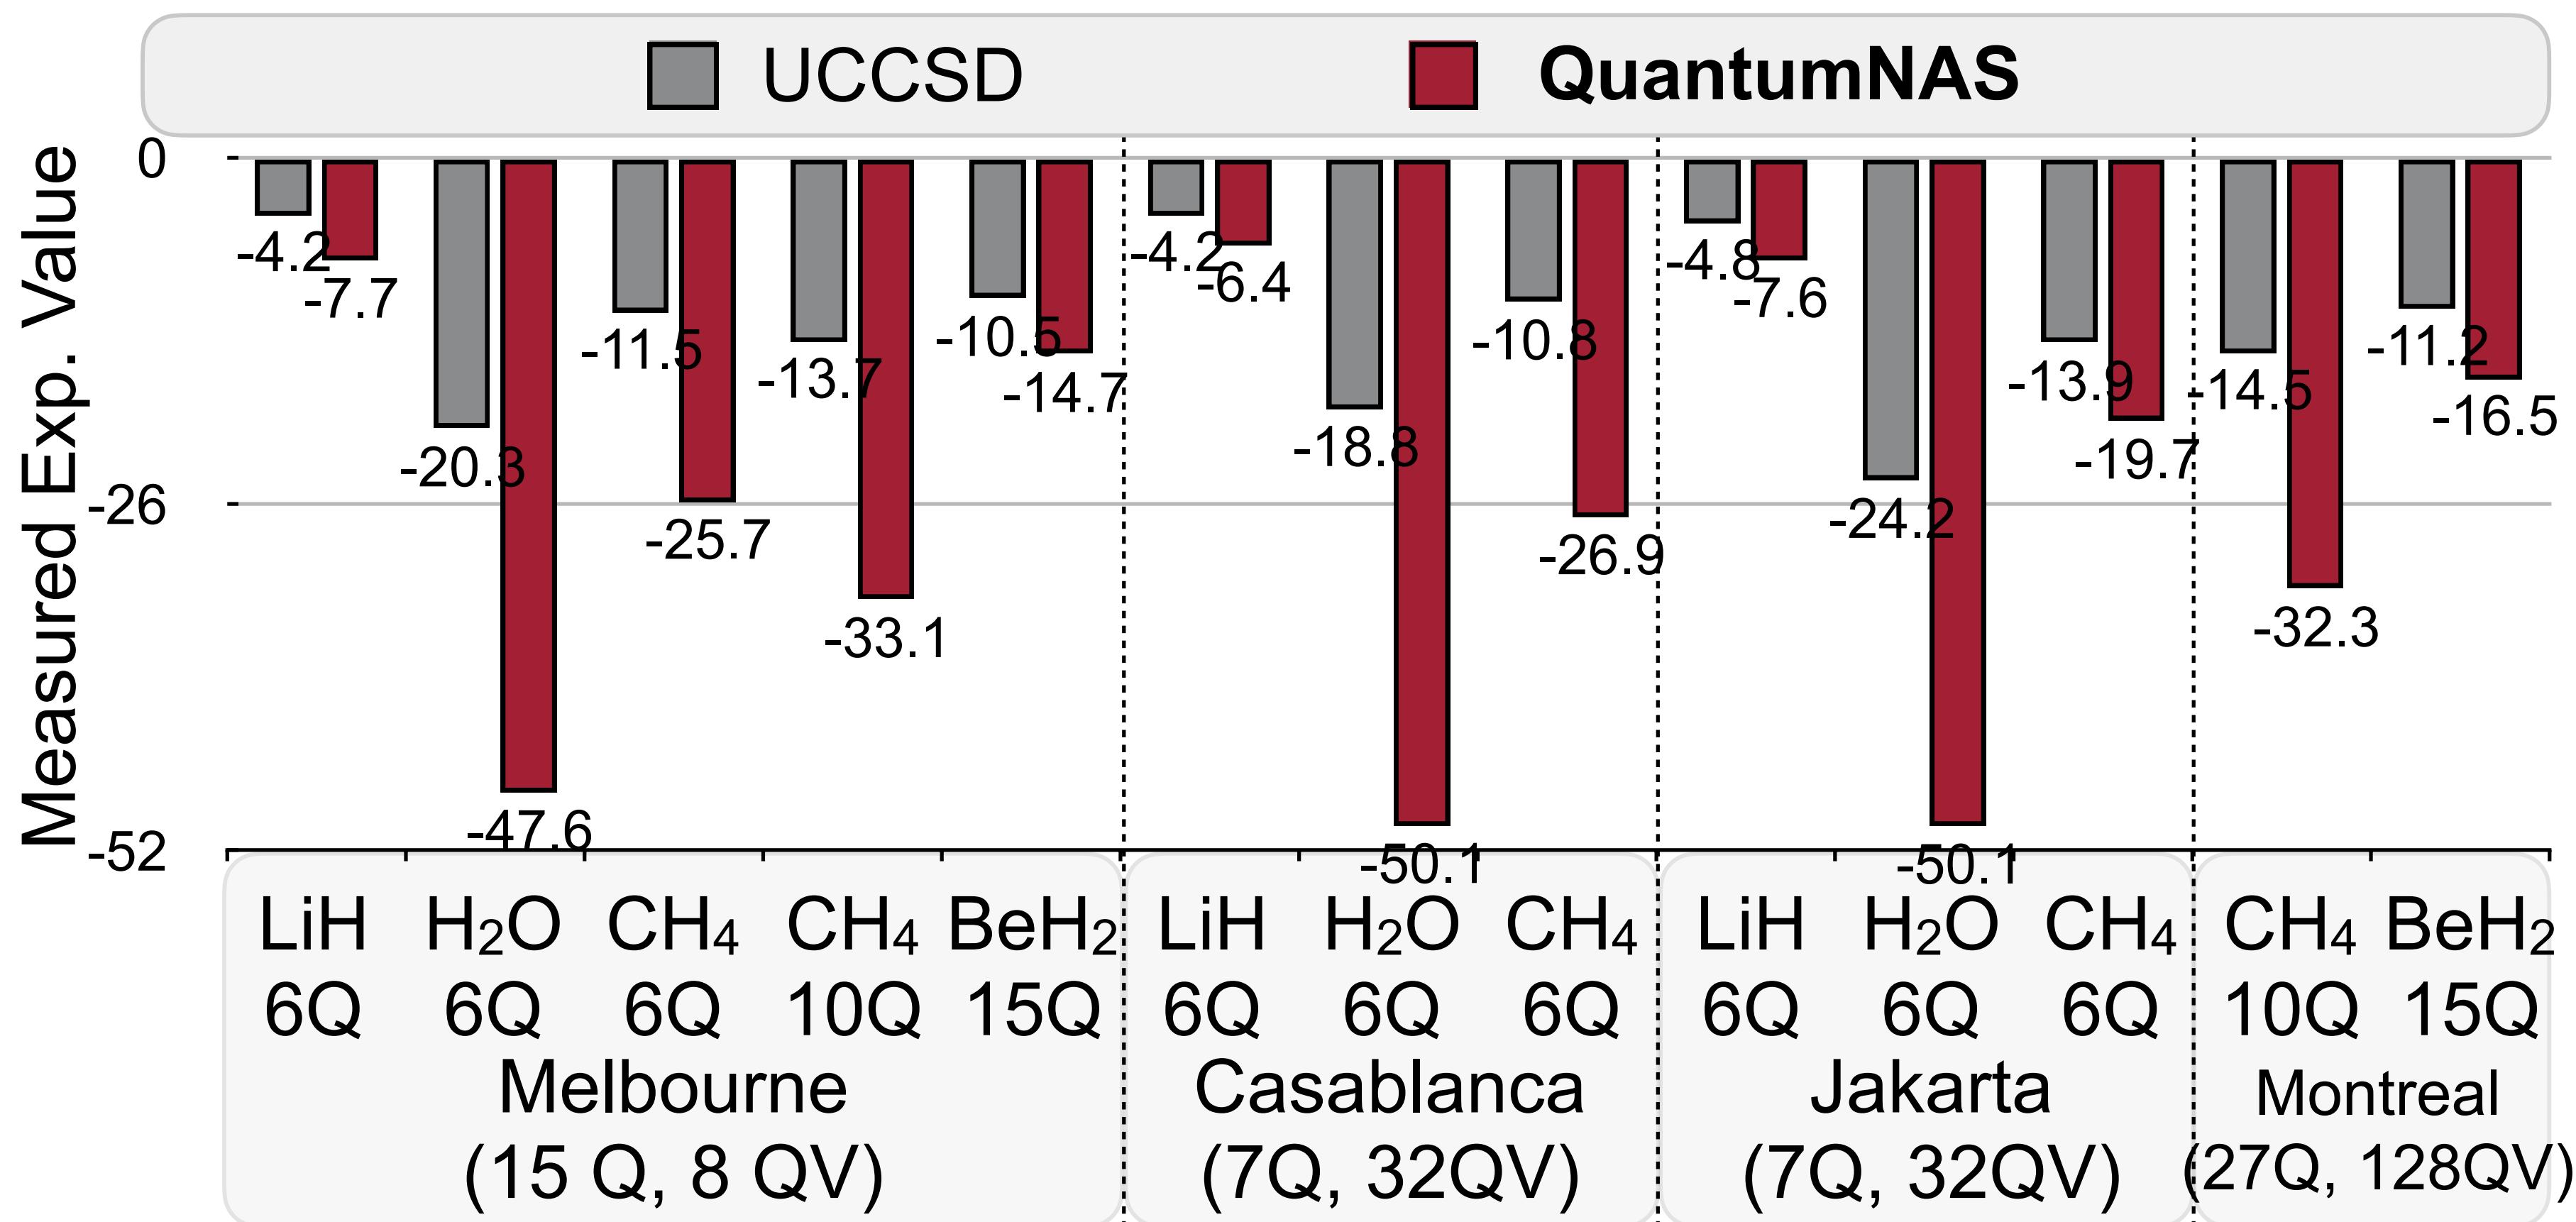
# H2 VQE Results

- H2 in different design spaces on IBMQ-Yorktown



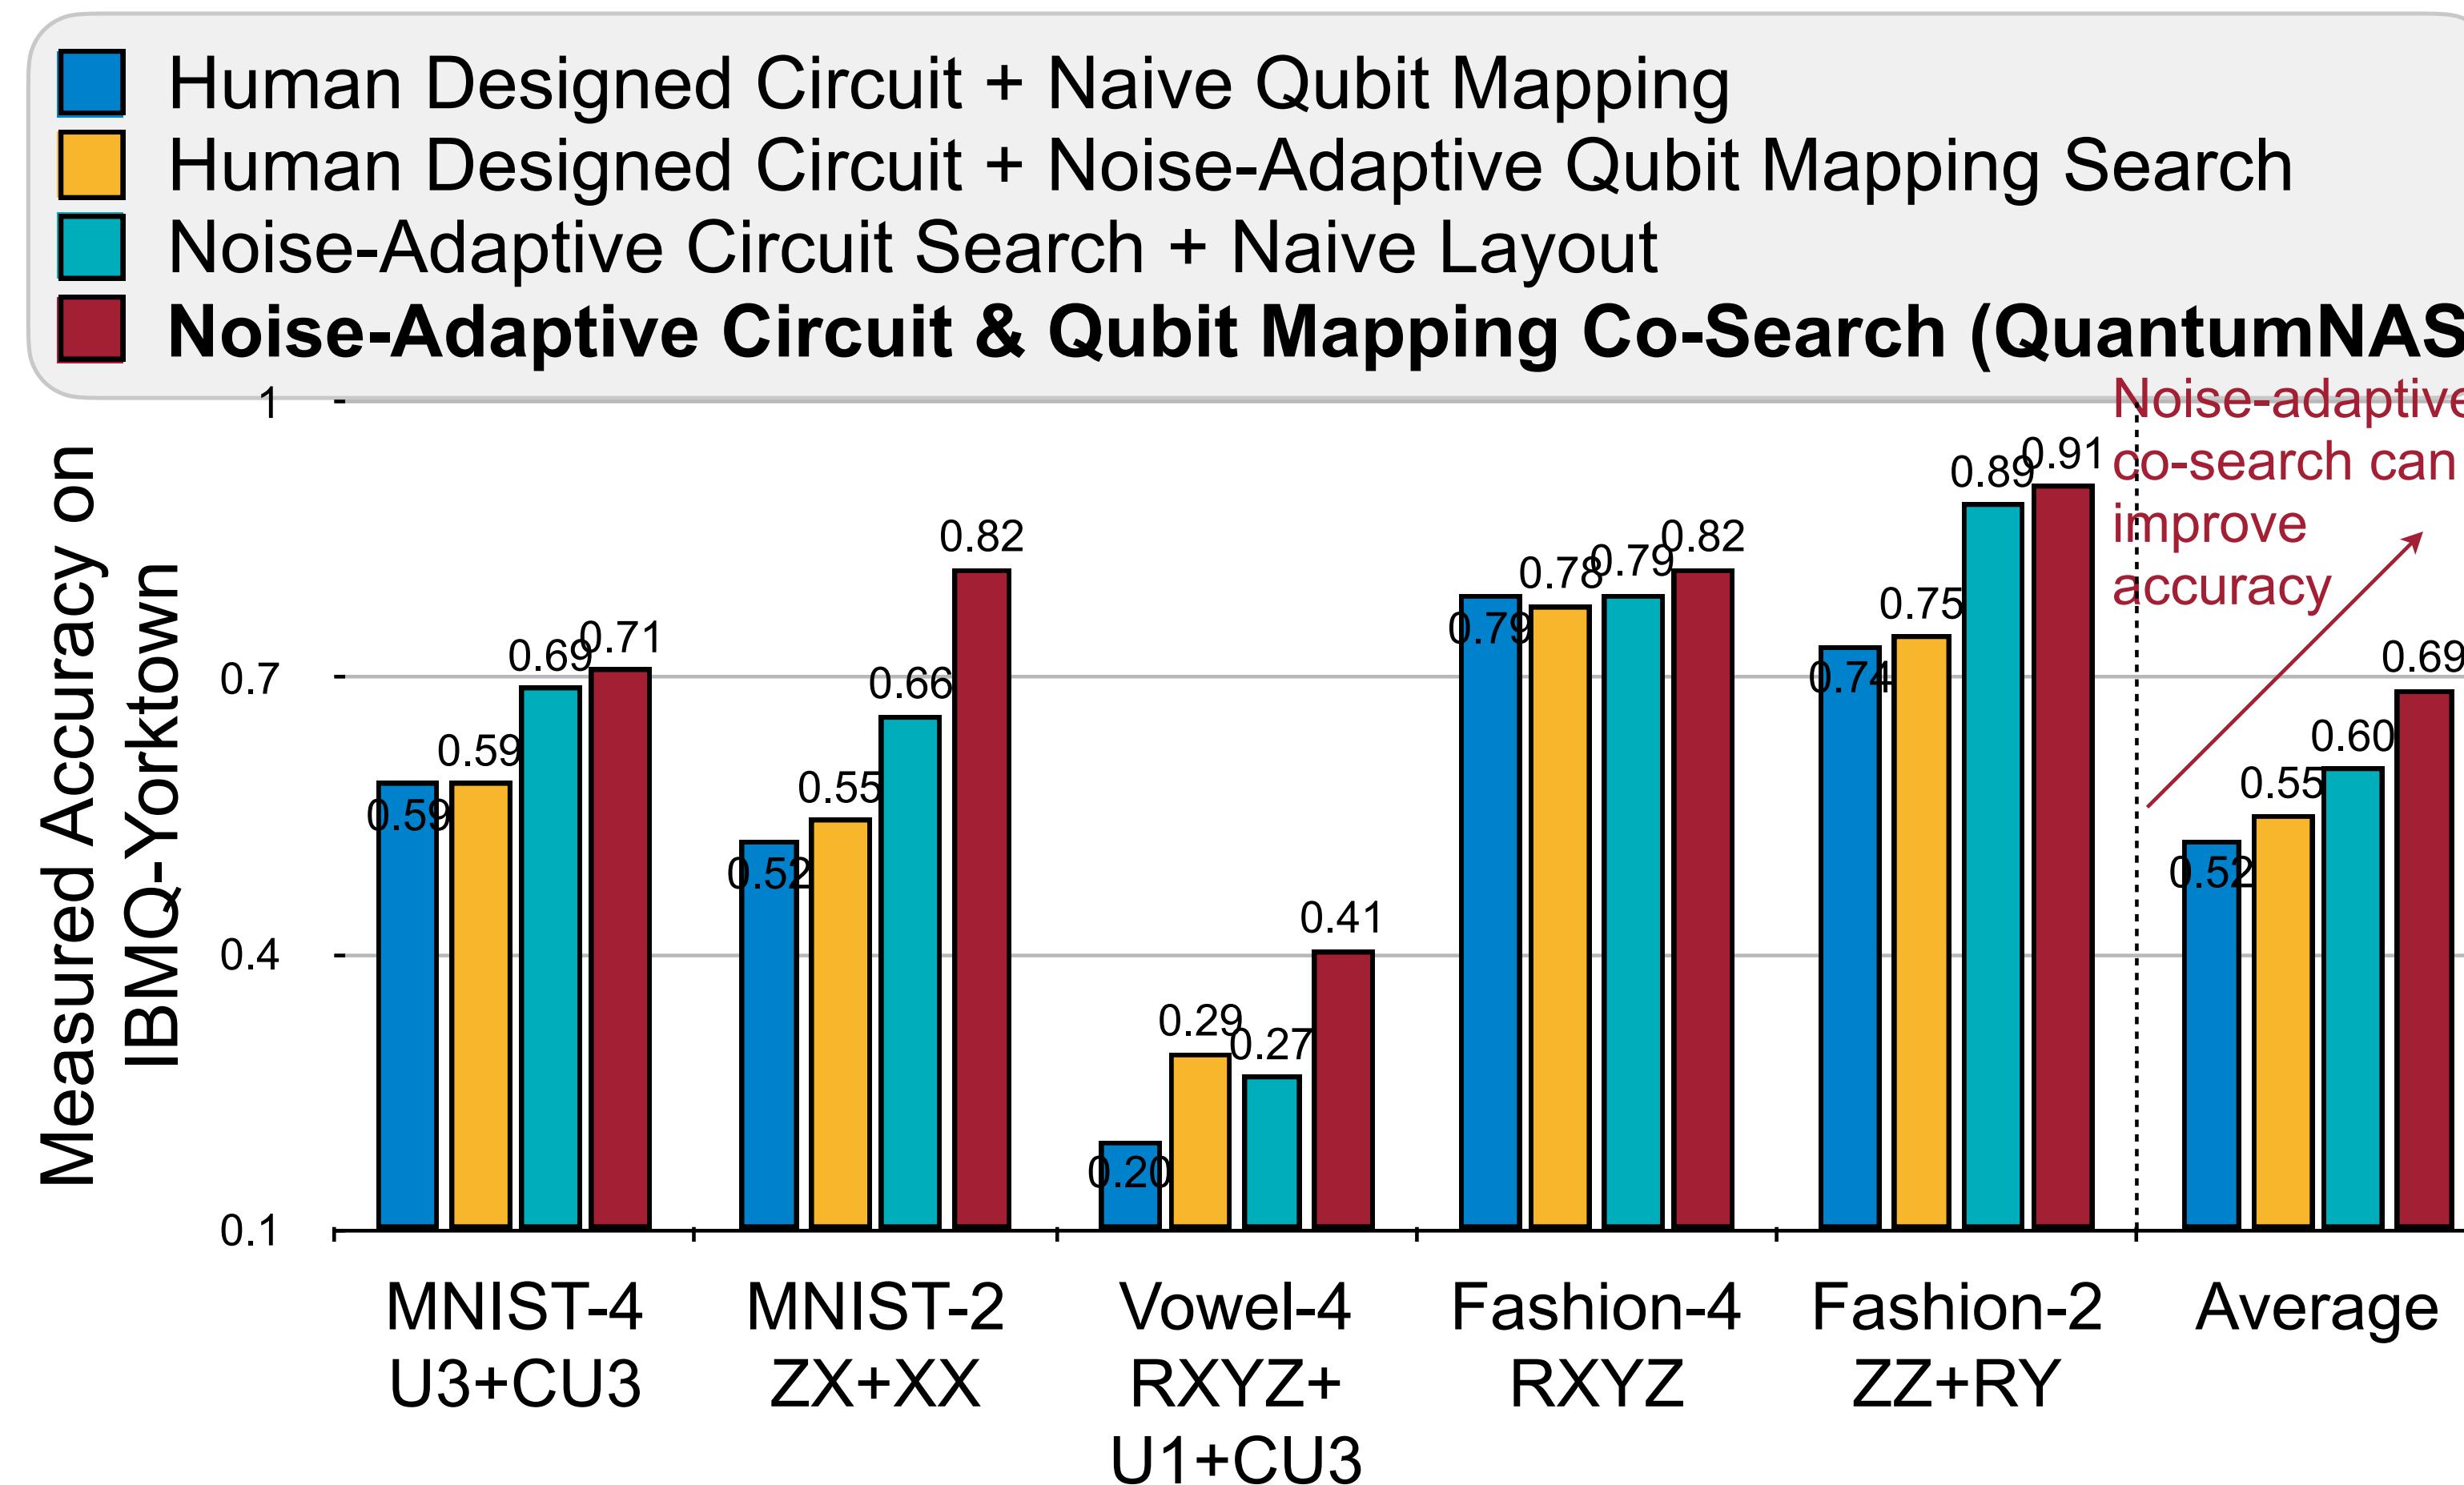
# VQE Results of Various Molecules

- VQE on different devices



# Effect of Circuit & Qubit Mapping Co-Search

- Co-search is better than only search ansatz/mapping



# Scalability

- The whole process can be done on quantum machines
  - SuperCircuit, SubCircuit training with parameter shift on QC device (no harder than training a normal variational quantum circuit)
  - Search with evaluation results on QC device
- Even with whole pipeline on Quantum, QuantumNAS still reduces search cost, because only one SuperCircuit is trained instead of many circuit candidates.
- Different devices can reuse the same SuperCircuit

# Time Cost

- On 1 Nvidia Titan RTX 2080 ti GPU

| Stage | SuperCircuit Training | Noise-Adaptive Co-search | SubCircuit Training | Deployment on Real QC |
|-------|-----------------------|--------------------------|---------------------|-----------------------|
| 4Q    | 0.5h                  | 3h                       | 0.5h                | 0.5h                  |
| 15Q   | 5h                    | 5h                       | 5h                  | 1h                    |
| 21Q   | 20h                   | 10h                      | 15h                 | 1h                    |

# Outline

- Introduction
- Background
- QuantumNAS
  - SuperCircuit Training
  - Noise-Adaptive Co-Search
  - Searched SubCircuit Training
  - Gate Pruning
- Evaluation
- Conclusion



# Conclusion

- SuperCircuit based co-search for most robust circuit Ansatz and qubit mapping
- Iterative pruning to further remove redundant gates
- Achieved over 95% 2-class, 85% 4-class and 32% 10-class classification accuracy
- QC library open-sourced: <https://github.com/Hanrui-Wang/pytorch-quantum>
- Outlook:
  - A general method applicable to all parameterized QC program
  - Search for Ansatz to alleviate Barren Plateau
  - Search for best embedding/readout

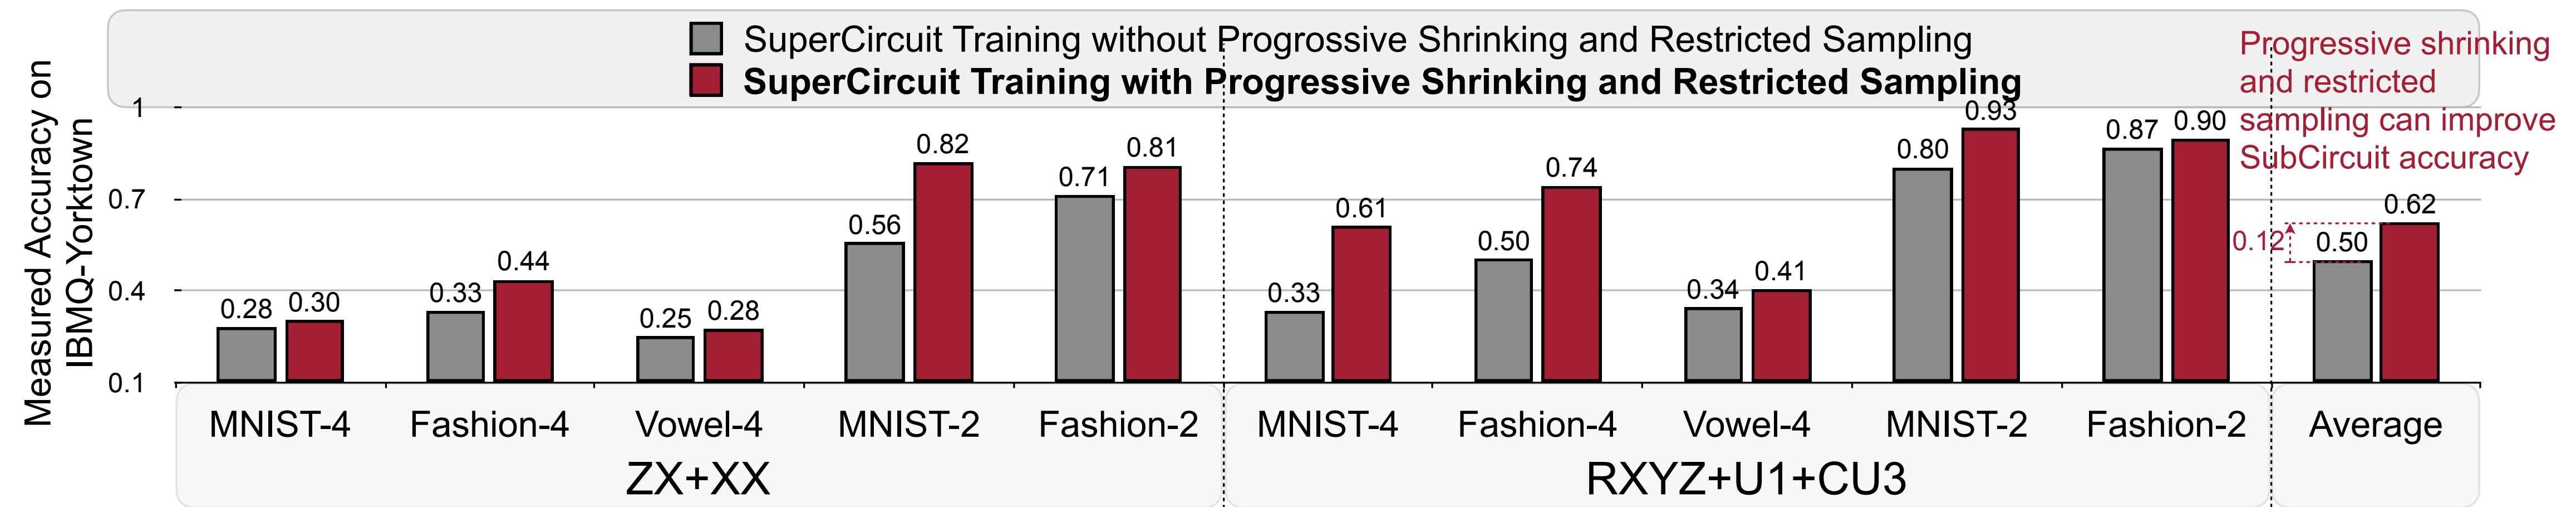
**Thank you!**

# Backup Slides

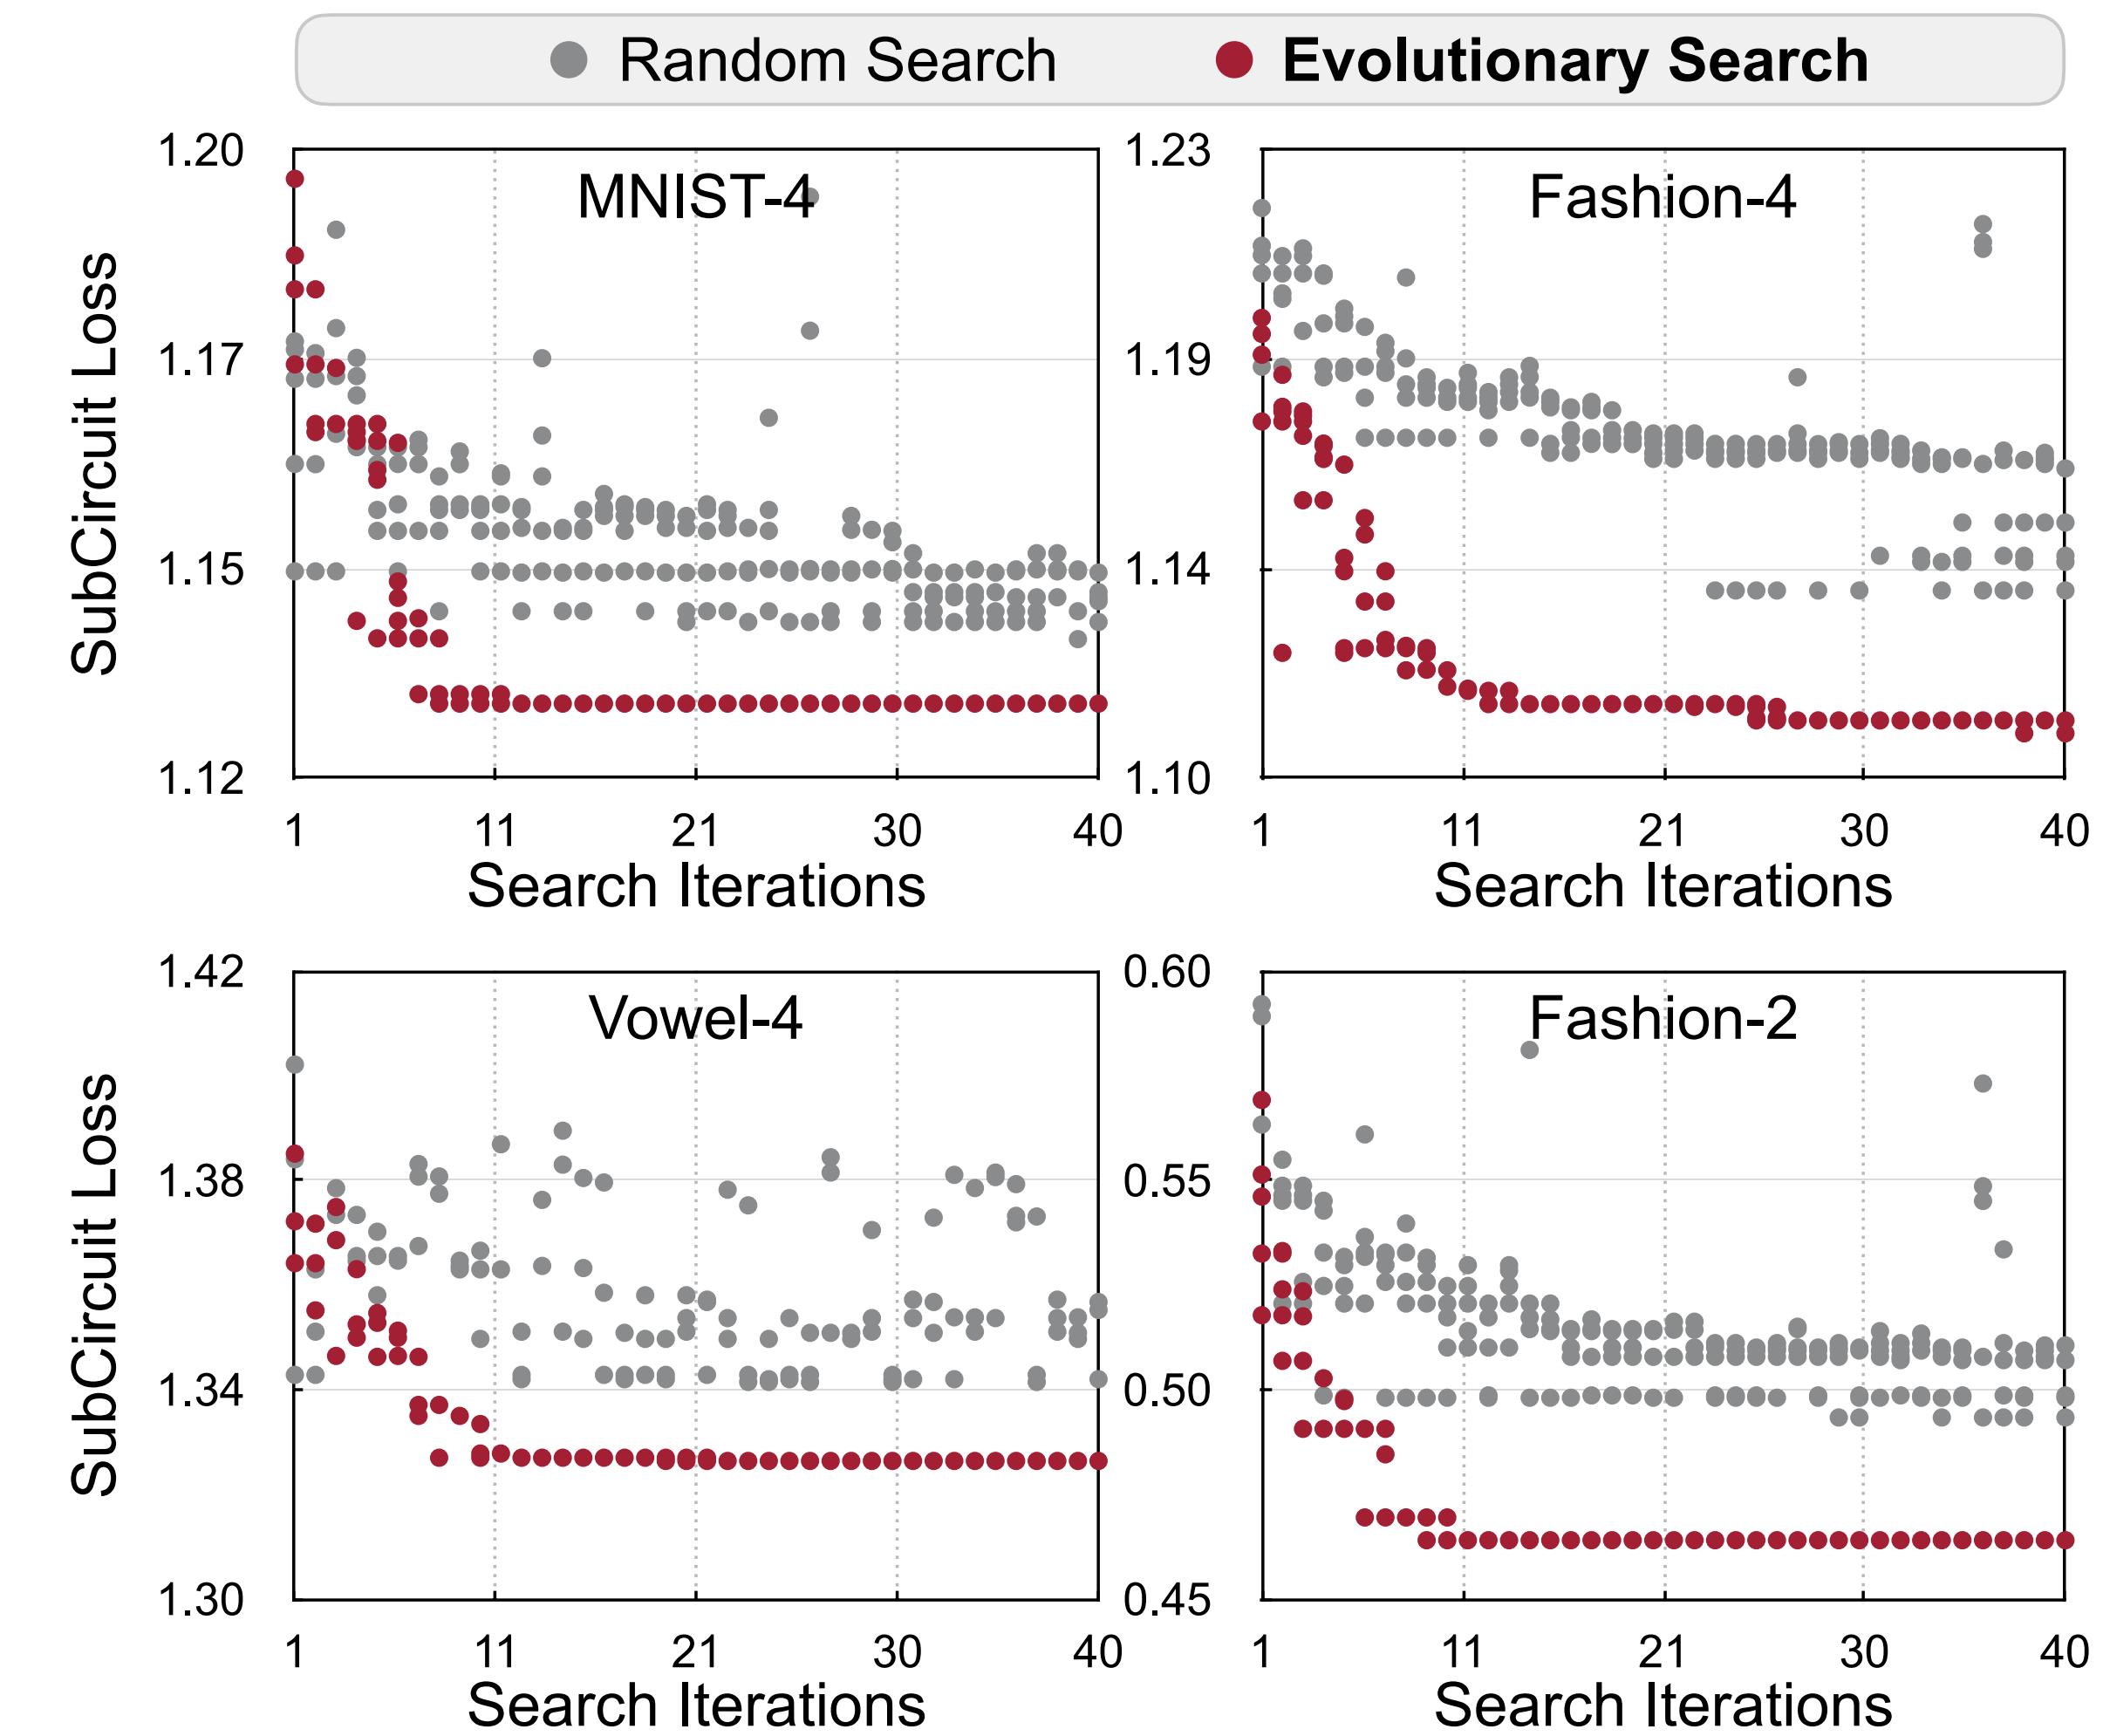
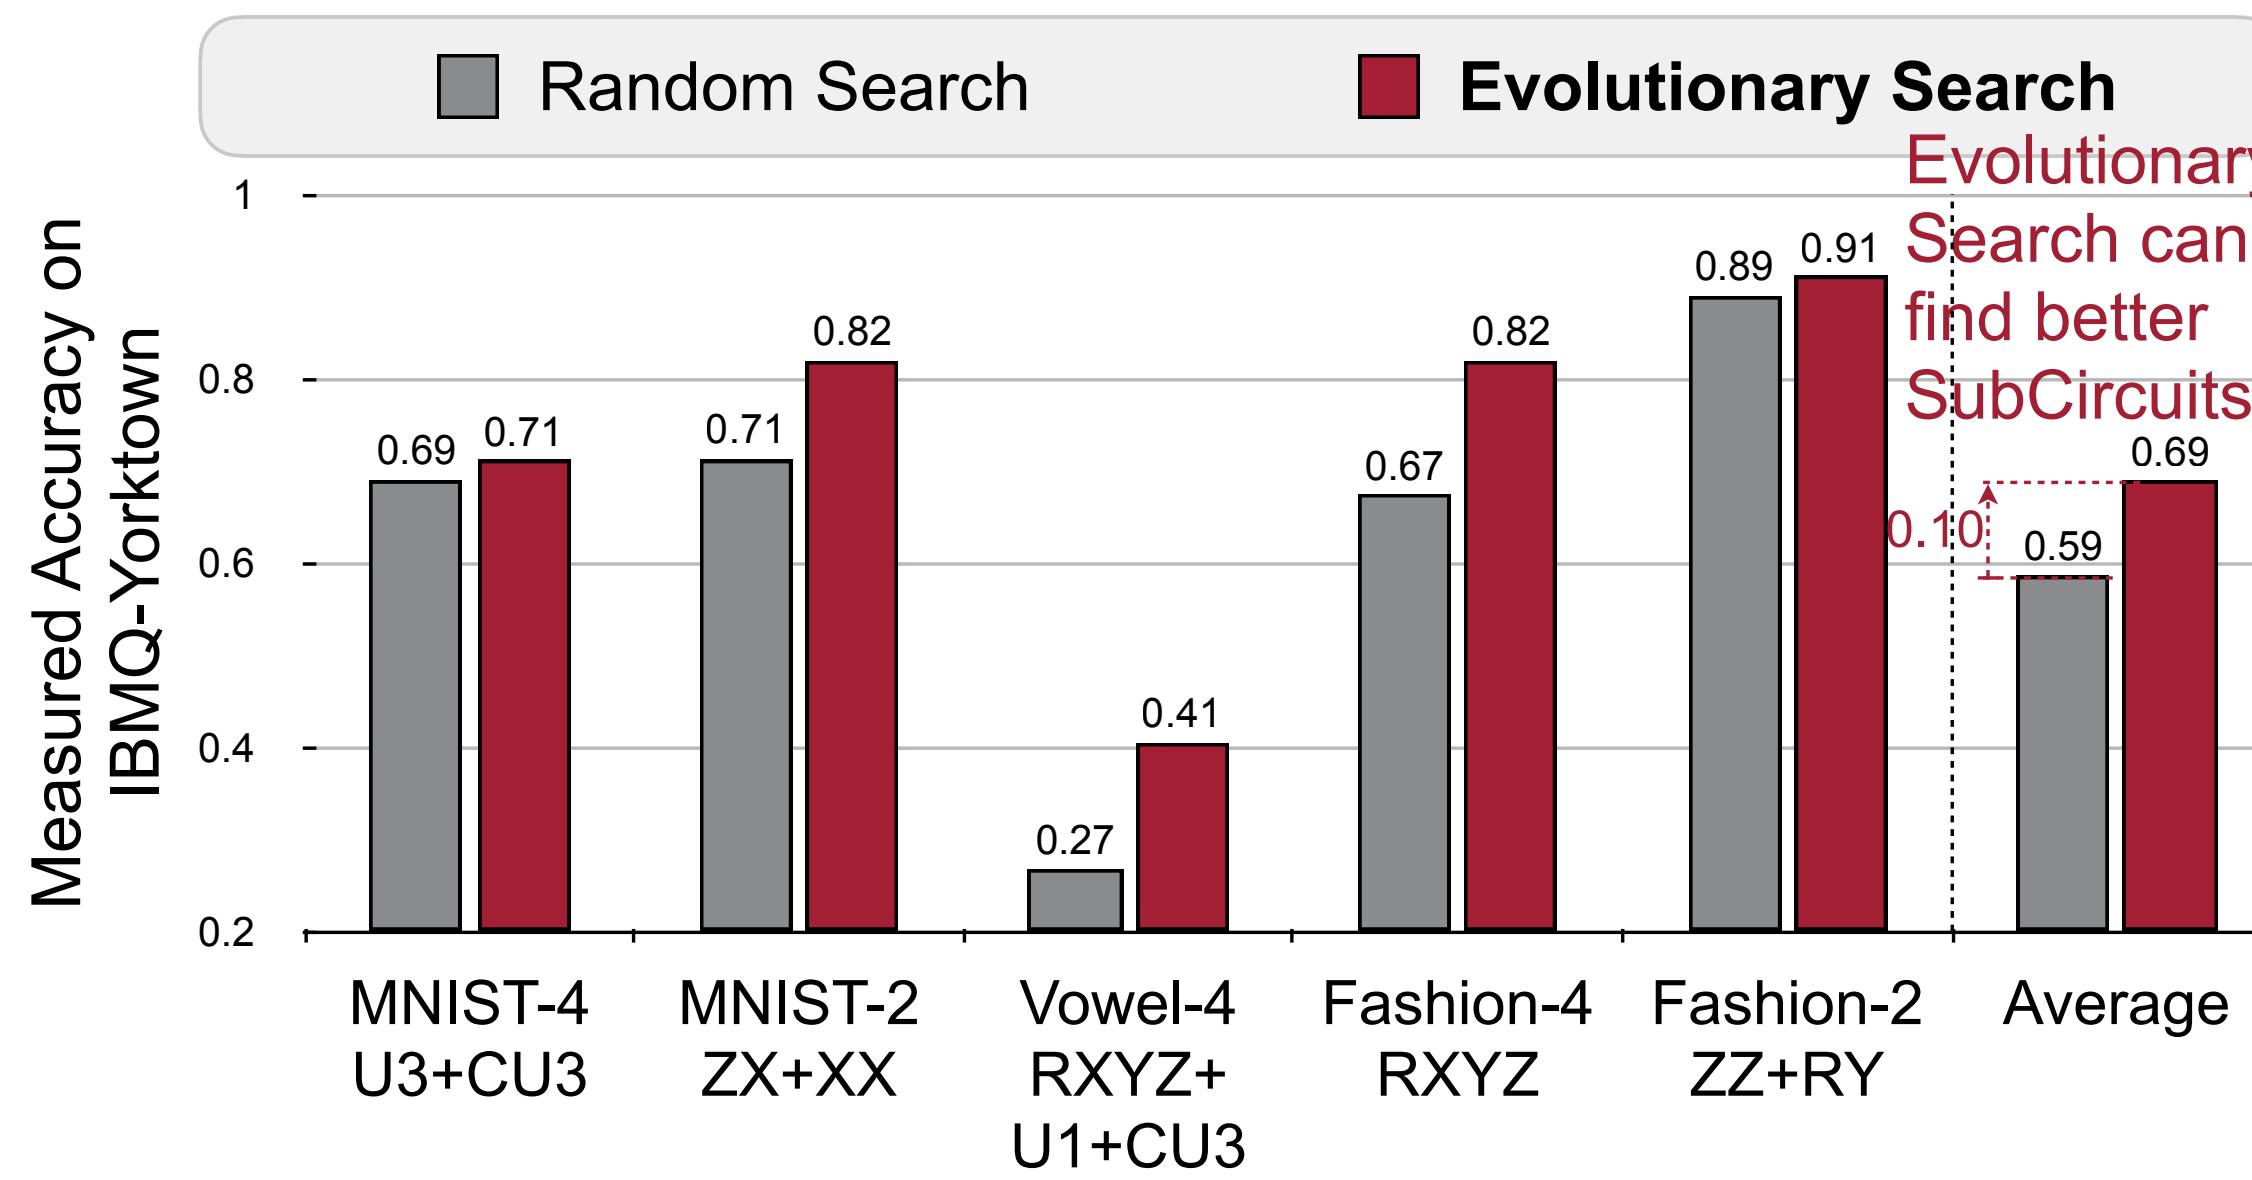
# Iterative pruning

| U3              | $ (\theta, \phi, \lambda) $ | $(\mathbf{0}, \phi, \lambda)$ | $ (\theta, \phi, \mathbf{0}) $ | $(\theta, \mathbf{0}, \mathbf{0})$ | $ (\mathbf{0}, \phi, \mathbf{0}) $ | $(\mathbf{0}, \mathbf{0}, \lambda)$ |
|-----------------|-----------------------------|-------------------------------|--------------------------------|------------------------------------|------------------------------------|-------------------------------------|
| #Compiled Gates | 5                           | 1                             | 4                              | 4                                  | 1                                  | 1                                   |

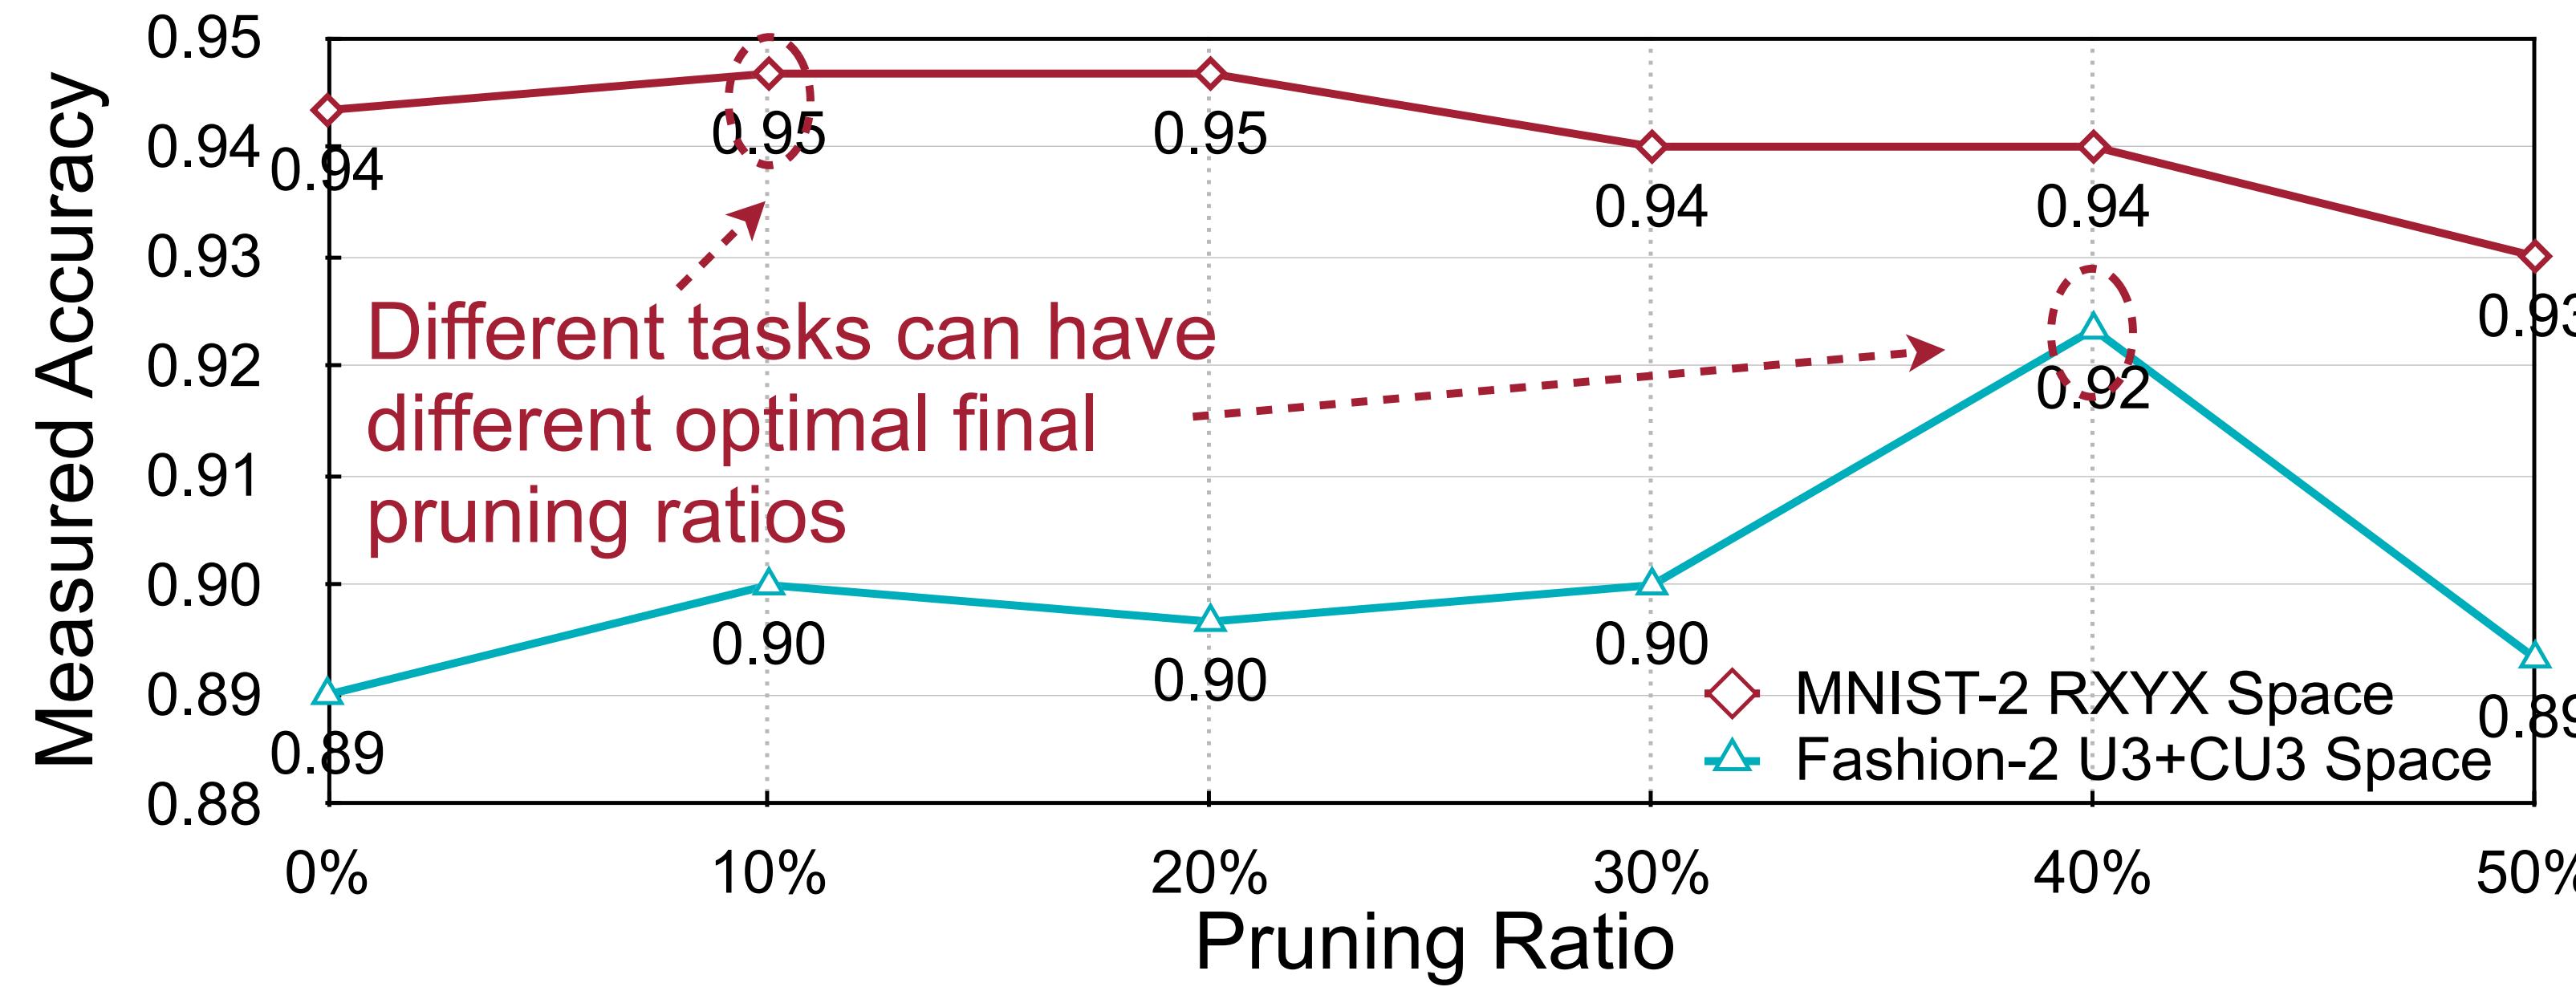
# Progressive shrinking



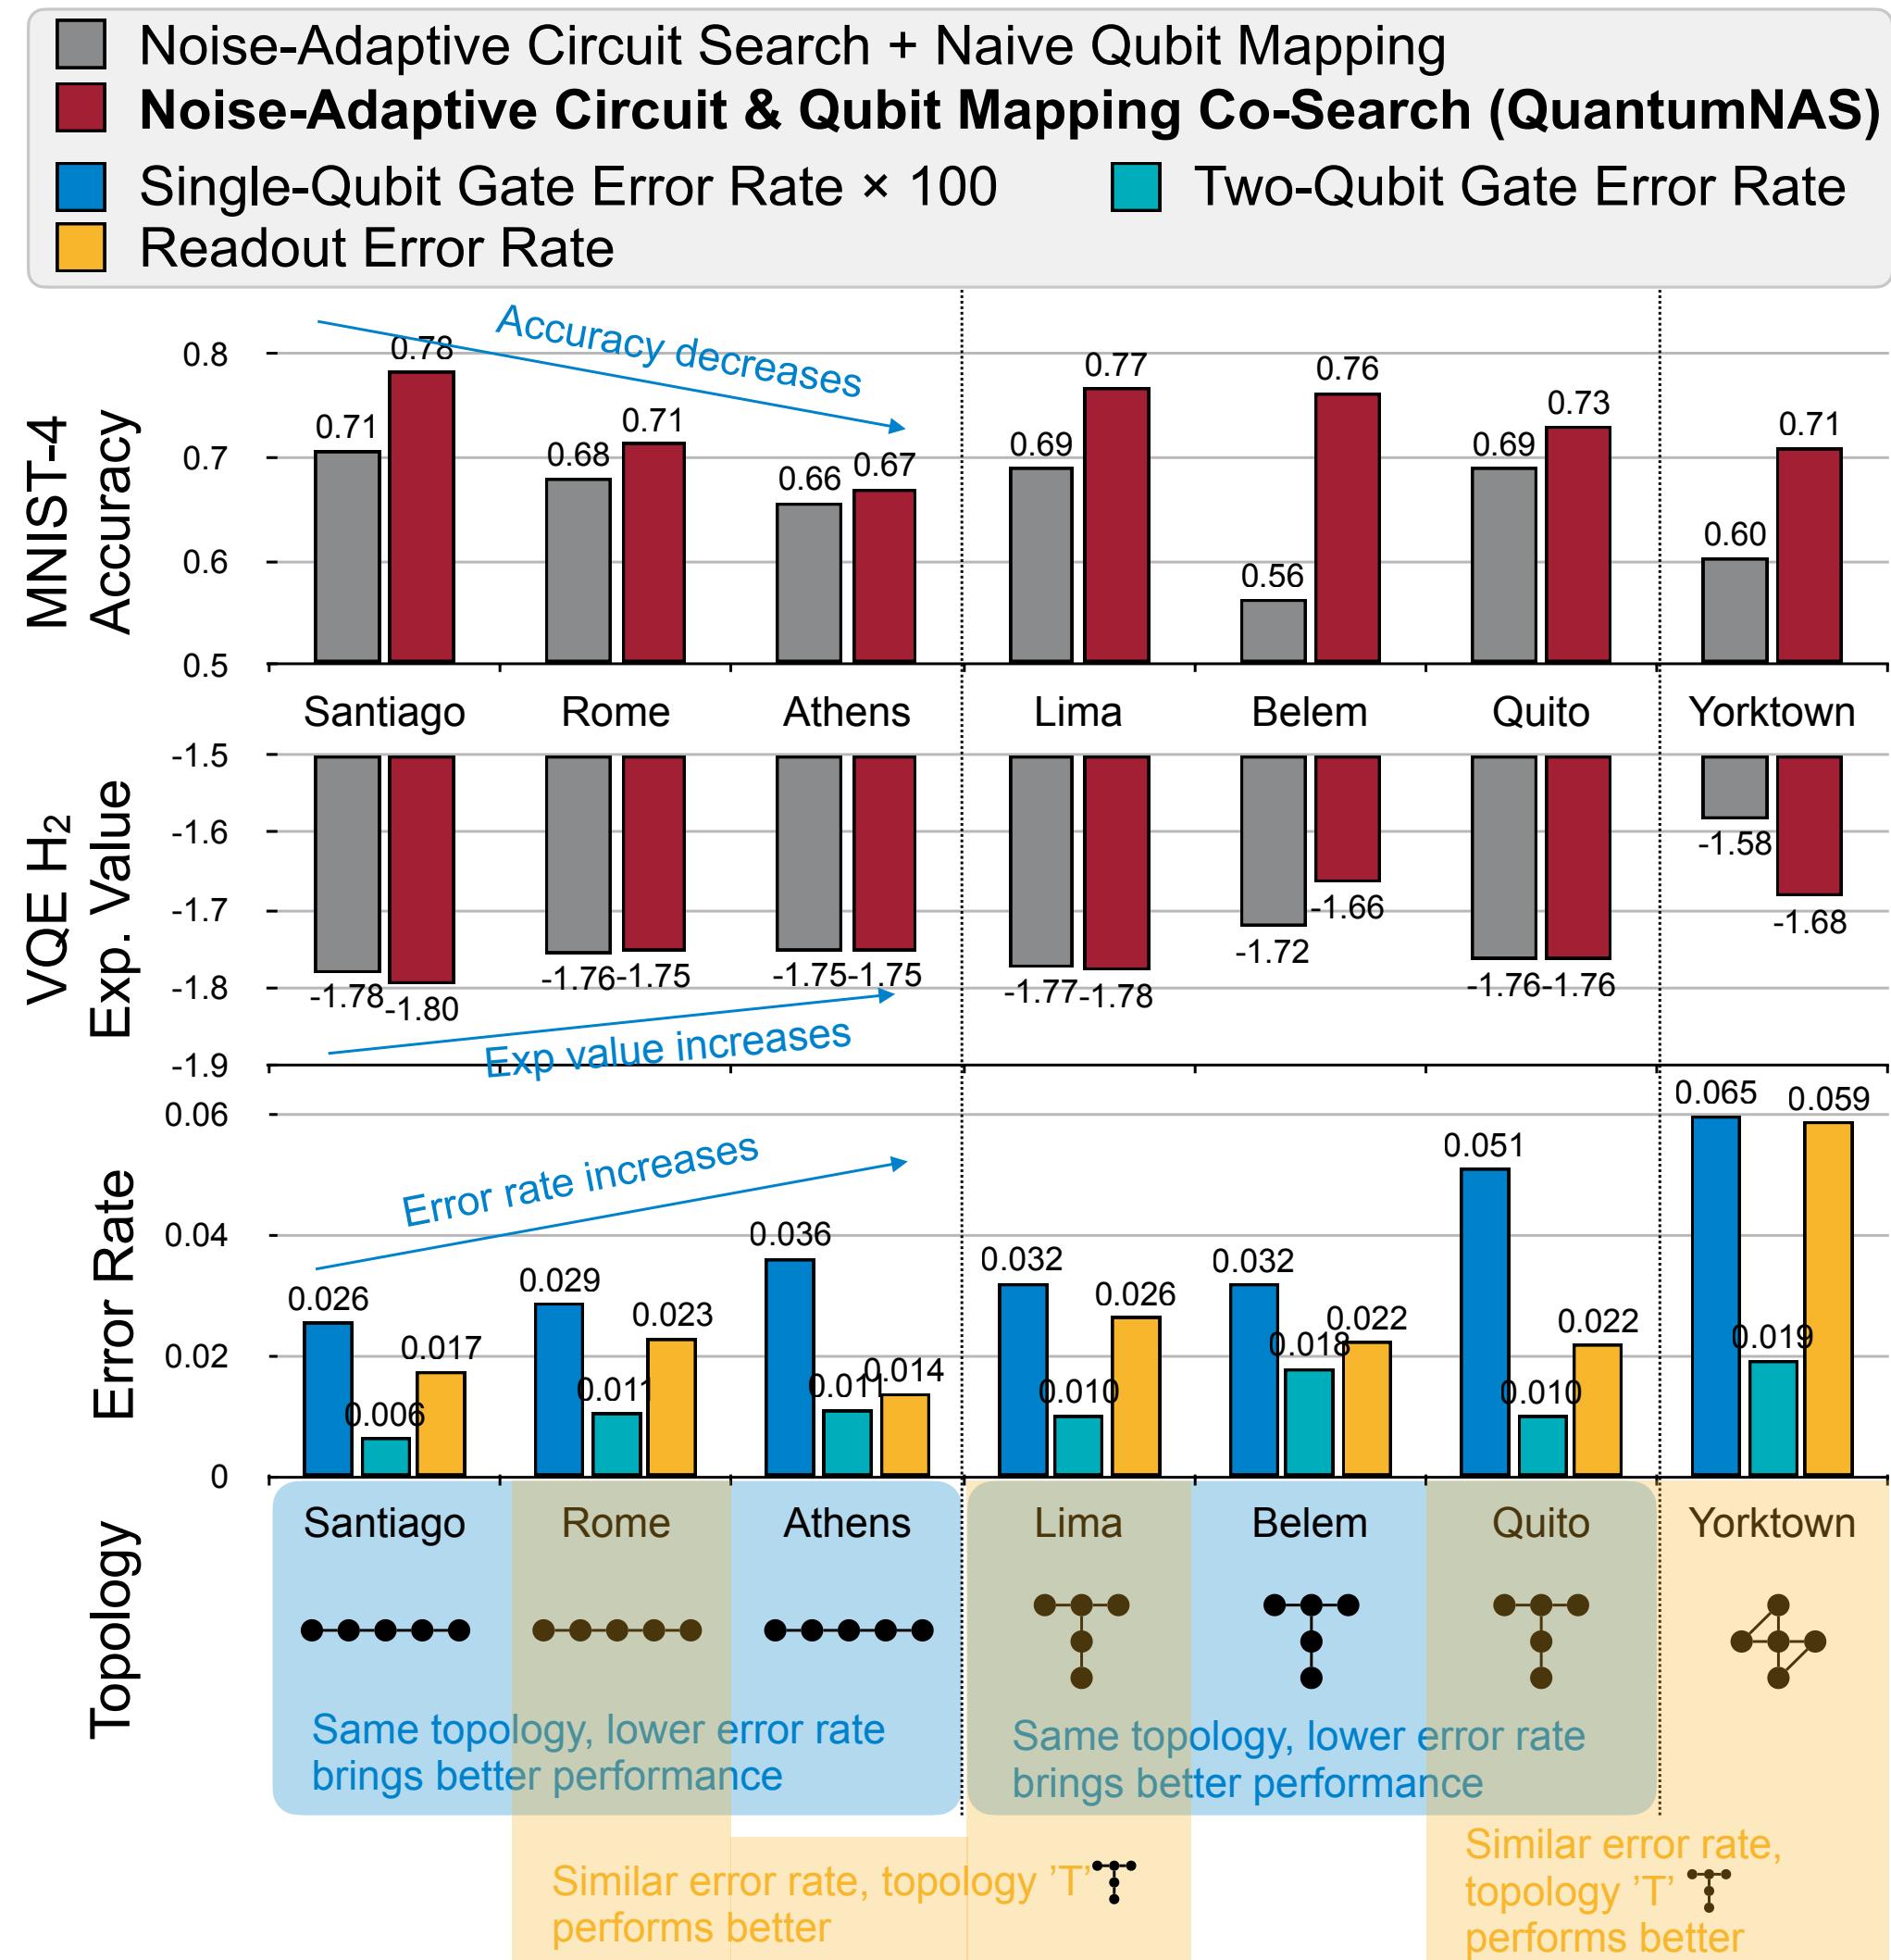
# Random search vs evolutionary search



# Pruning



# Topology/Error rate



# Device-specific

| Run on ↓ Searched for → | Yorktown    | Belem       | Santiago    |
|-------------------------|-------------|-------------|-------------|
| Yorktown                | <b>0.85</b> | 0.60        | 0.54        |
| Belem                   | 0.67        | <b>0.77</b> | 0.43        |
| Santiago                | 0.82        | 0.81        | <b>0.85</b> |