



qmlsys.mit.edu

QuantumNAS: Noise-Aware Search for Robust Quantum Circuits and the TorchQuantum Library

Hanrui Wang
MIT HAN Lab

[HPCA'22] QuantumNAS: Noise-adaptive search for robust quantum circuits

[DAC'22] QuantumNAT: Quantum Noise-Aware Training with Noise Injection, Quantization and Normalization

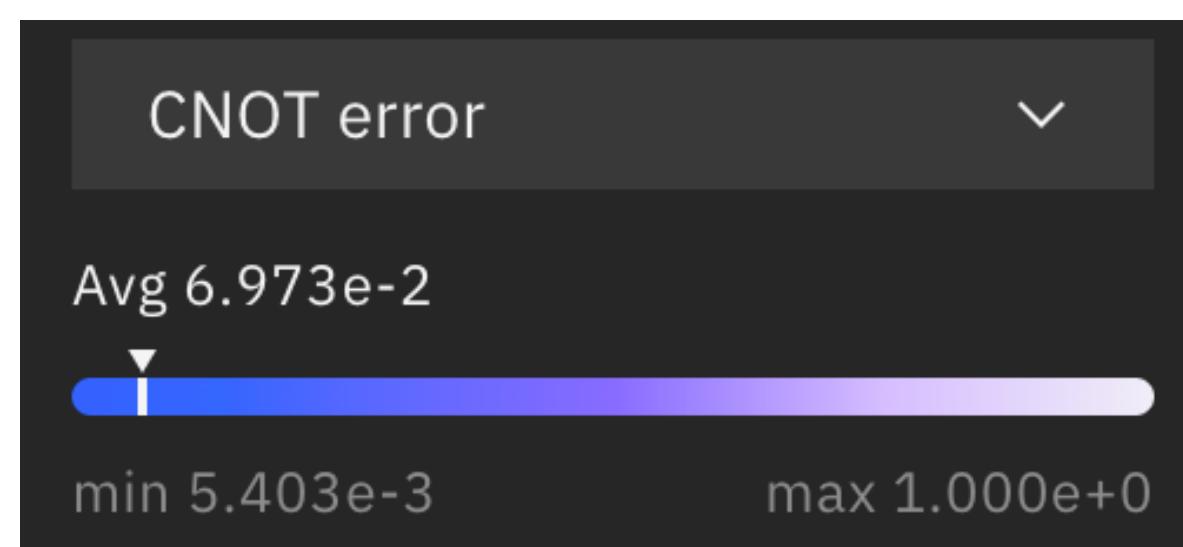
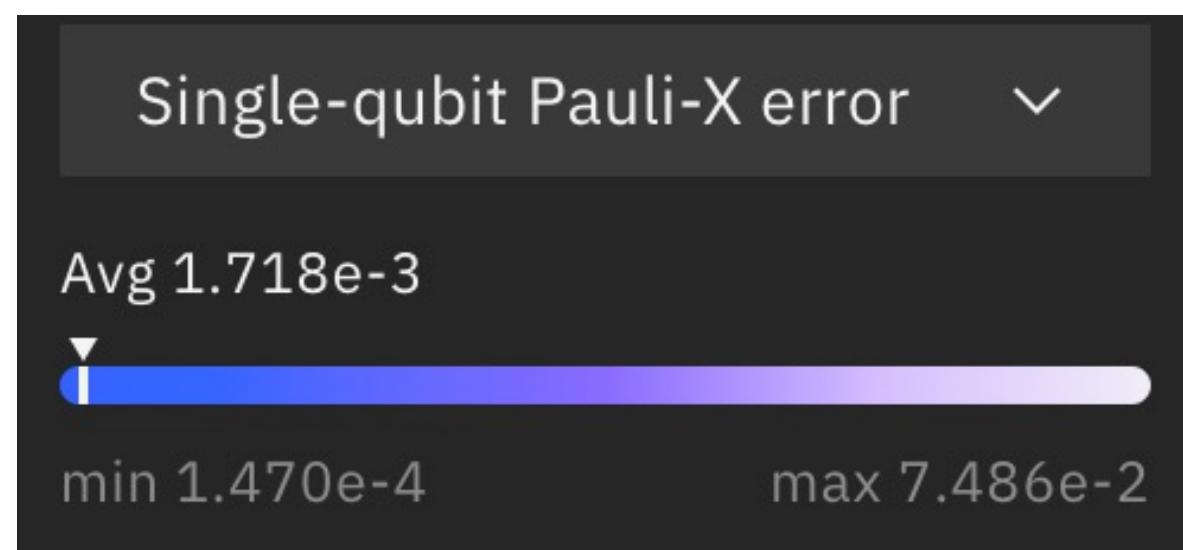
[DAC'22] QOC: Quantum On-Chip Training with Parameter Shift and Gradient Pruning

Outline

- Background
- QuantumNAS
- TorchQuantum Library
- Conclusion

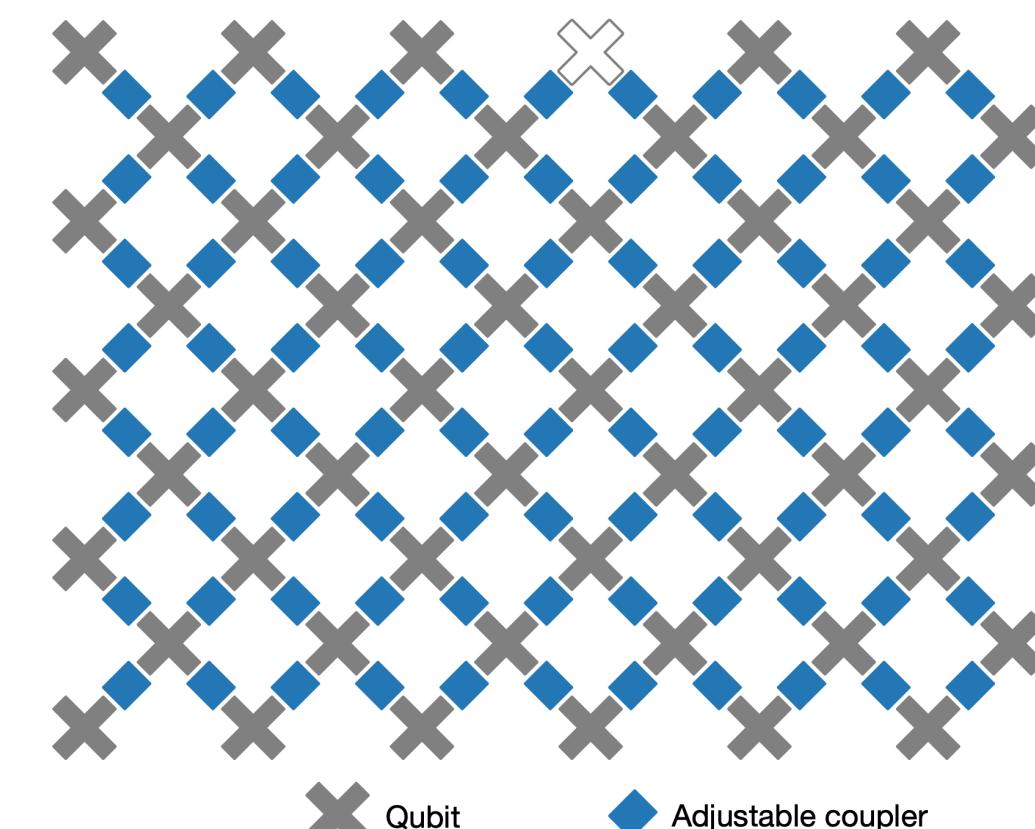
NISQ Era

- Noisy Intermediate-Scale Quantum (NISQ)
 - **Noisy**: qubits are sensitive to environment; quantum gates are unreliable
 - **Limited number** of qubits: tens to hundreds of qubits



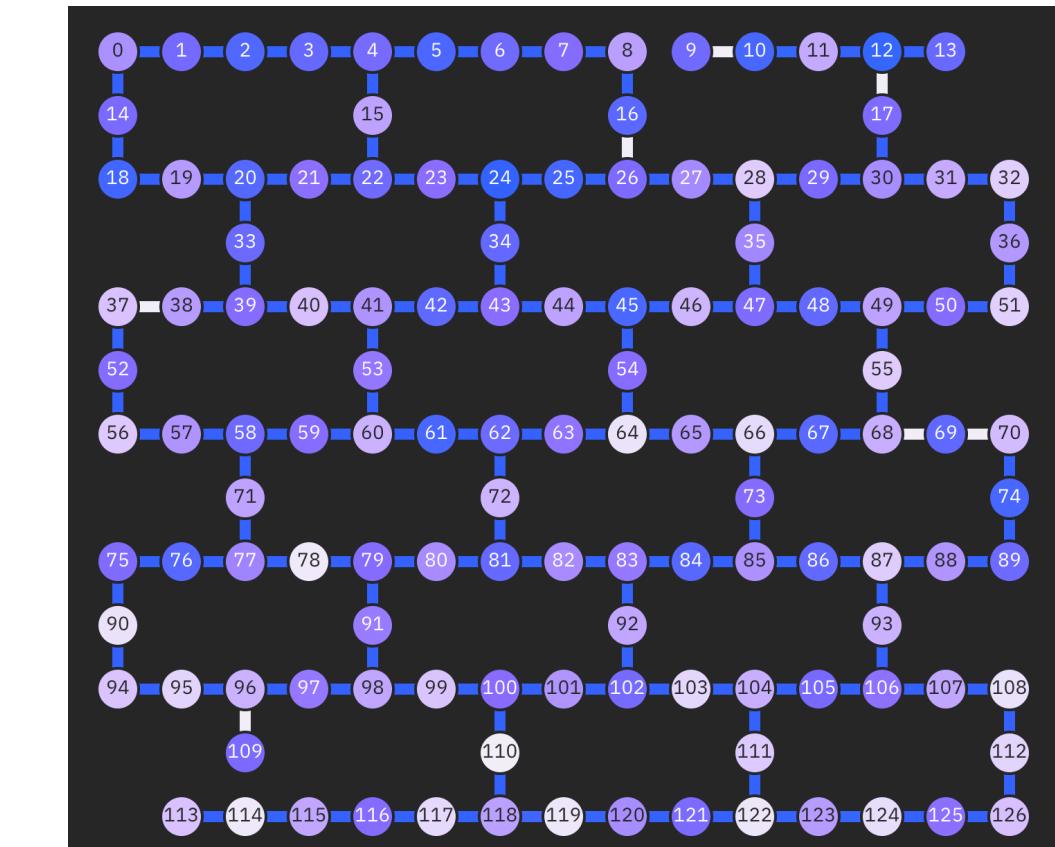
Gate Error Rate

<https://quantum-computing.ibm.com/>



Google Sycamore

<https://www.nature.com/articles/s41586-019-1666-5>

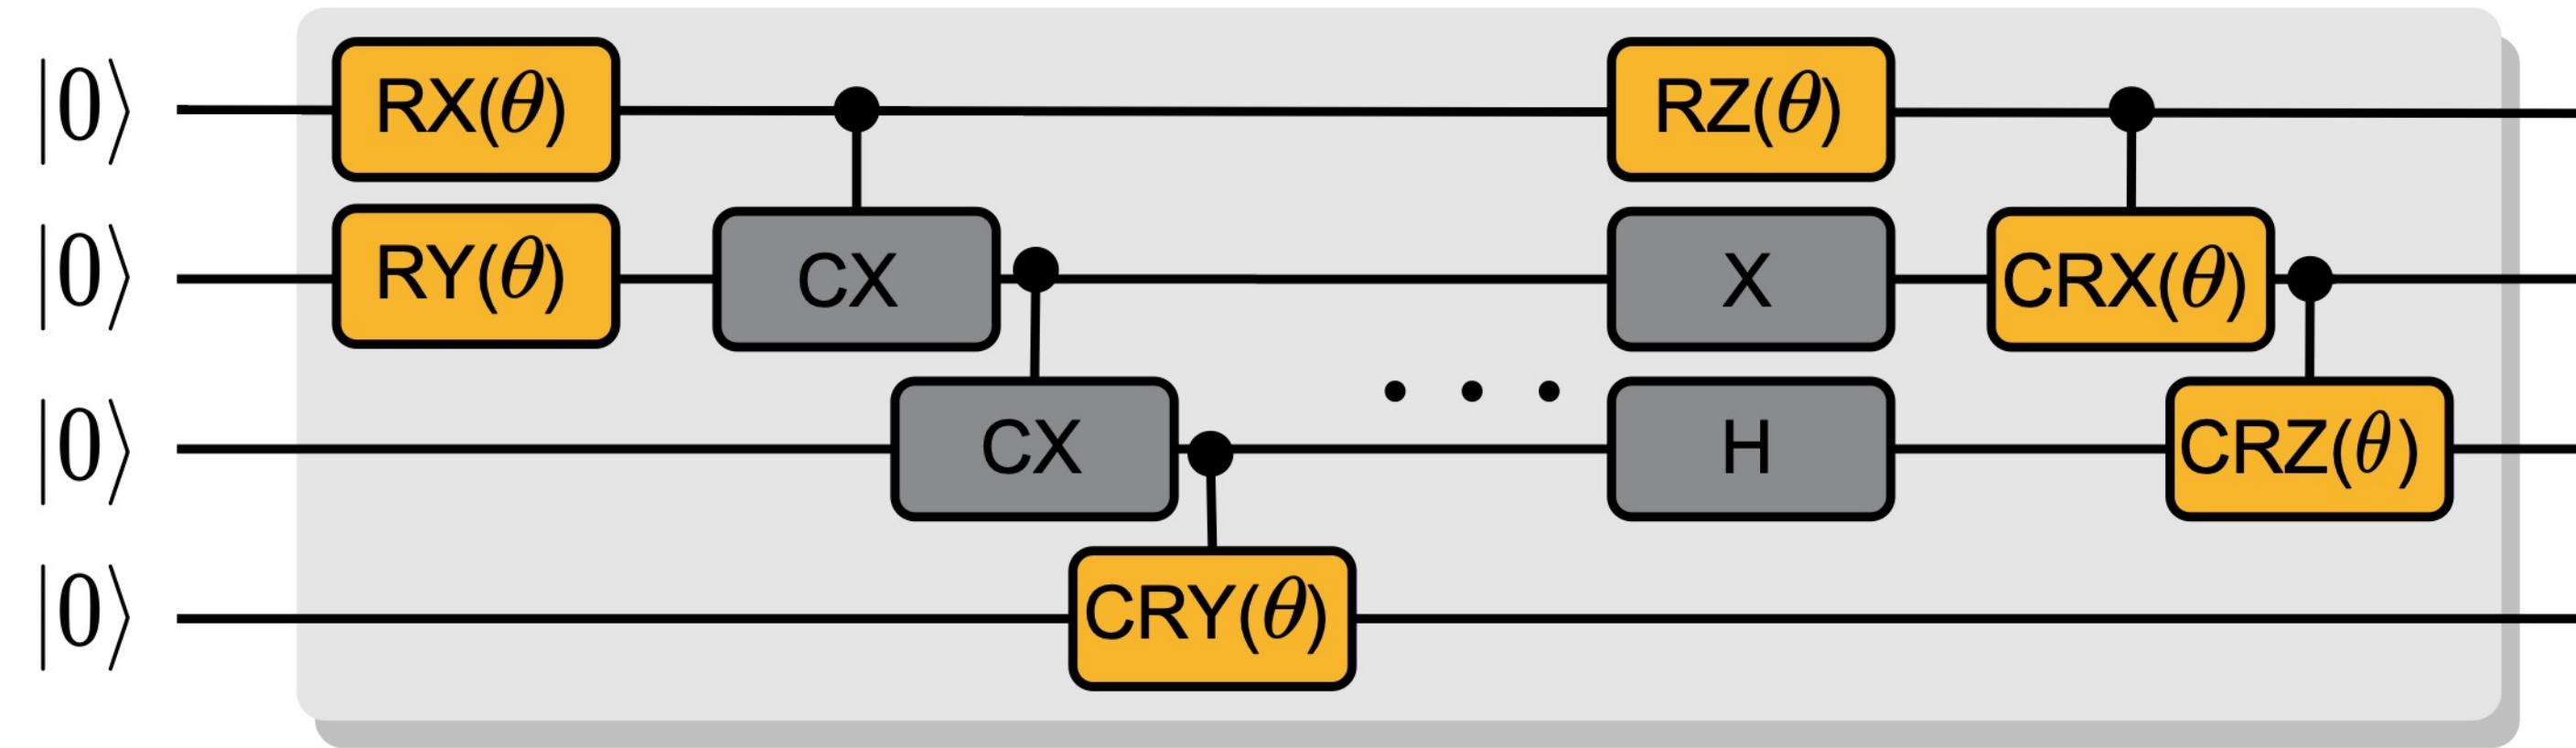


IBM Washington

<https://quantum-computing.ibm.com/>

Parameterized Quantum Circuits

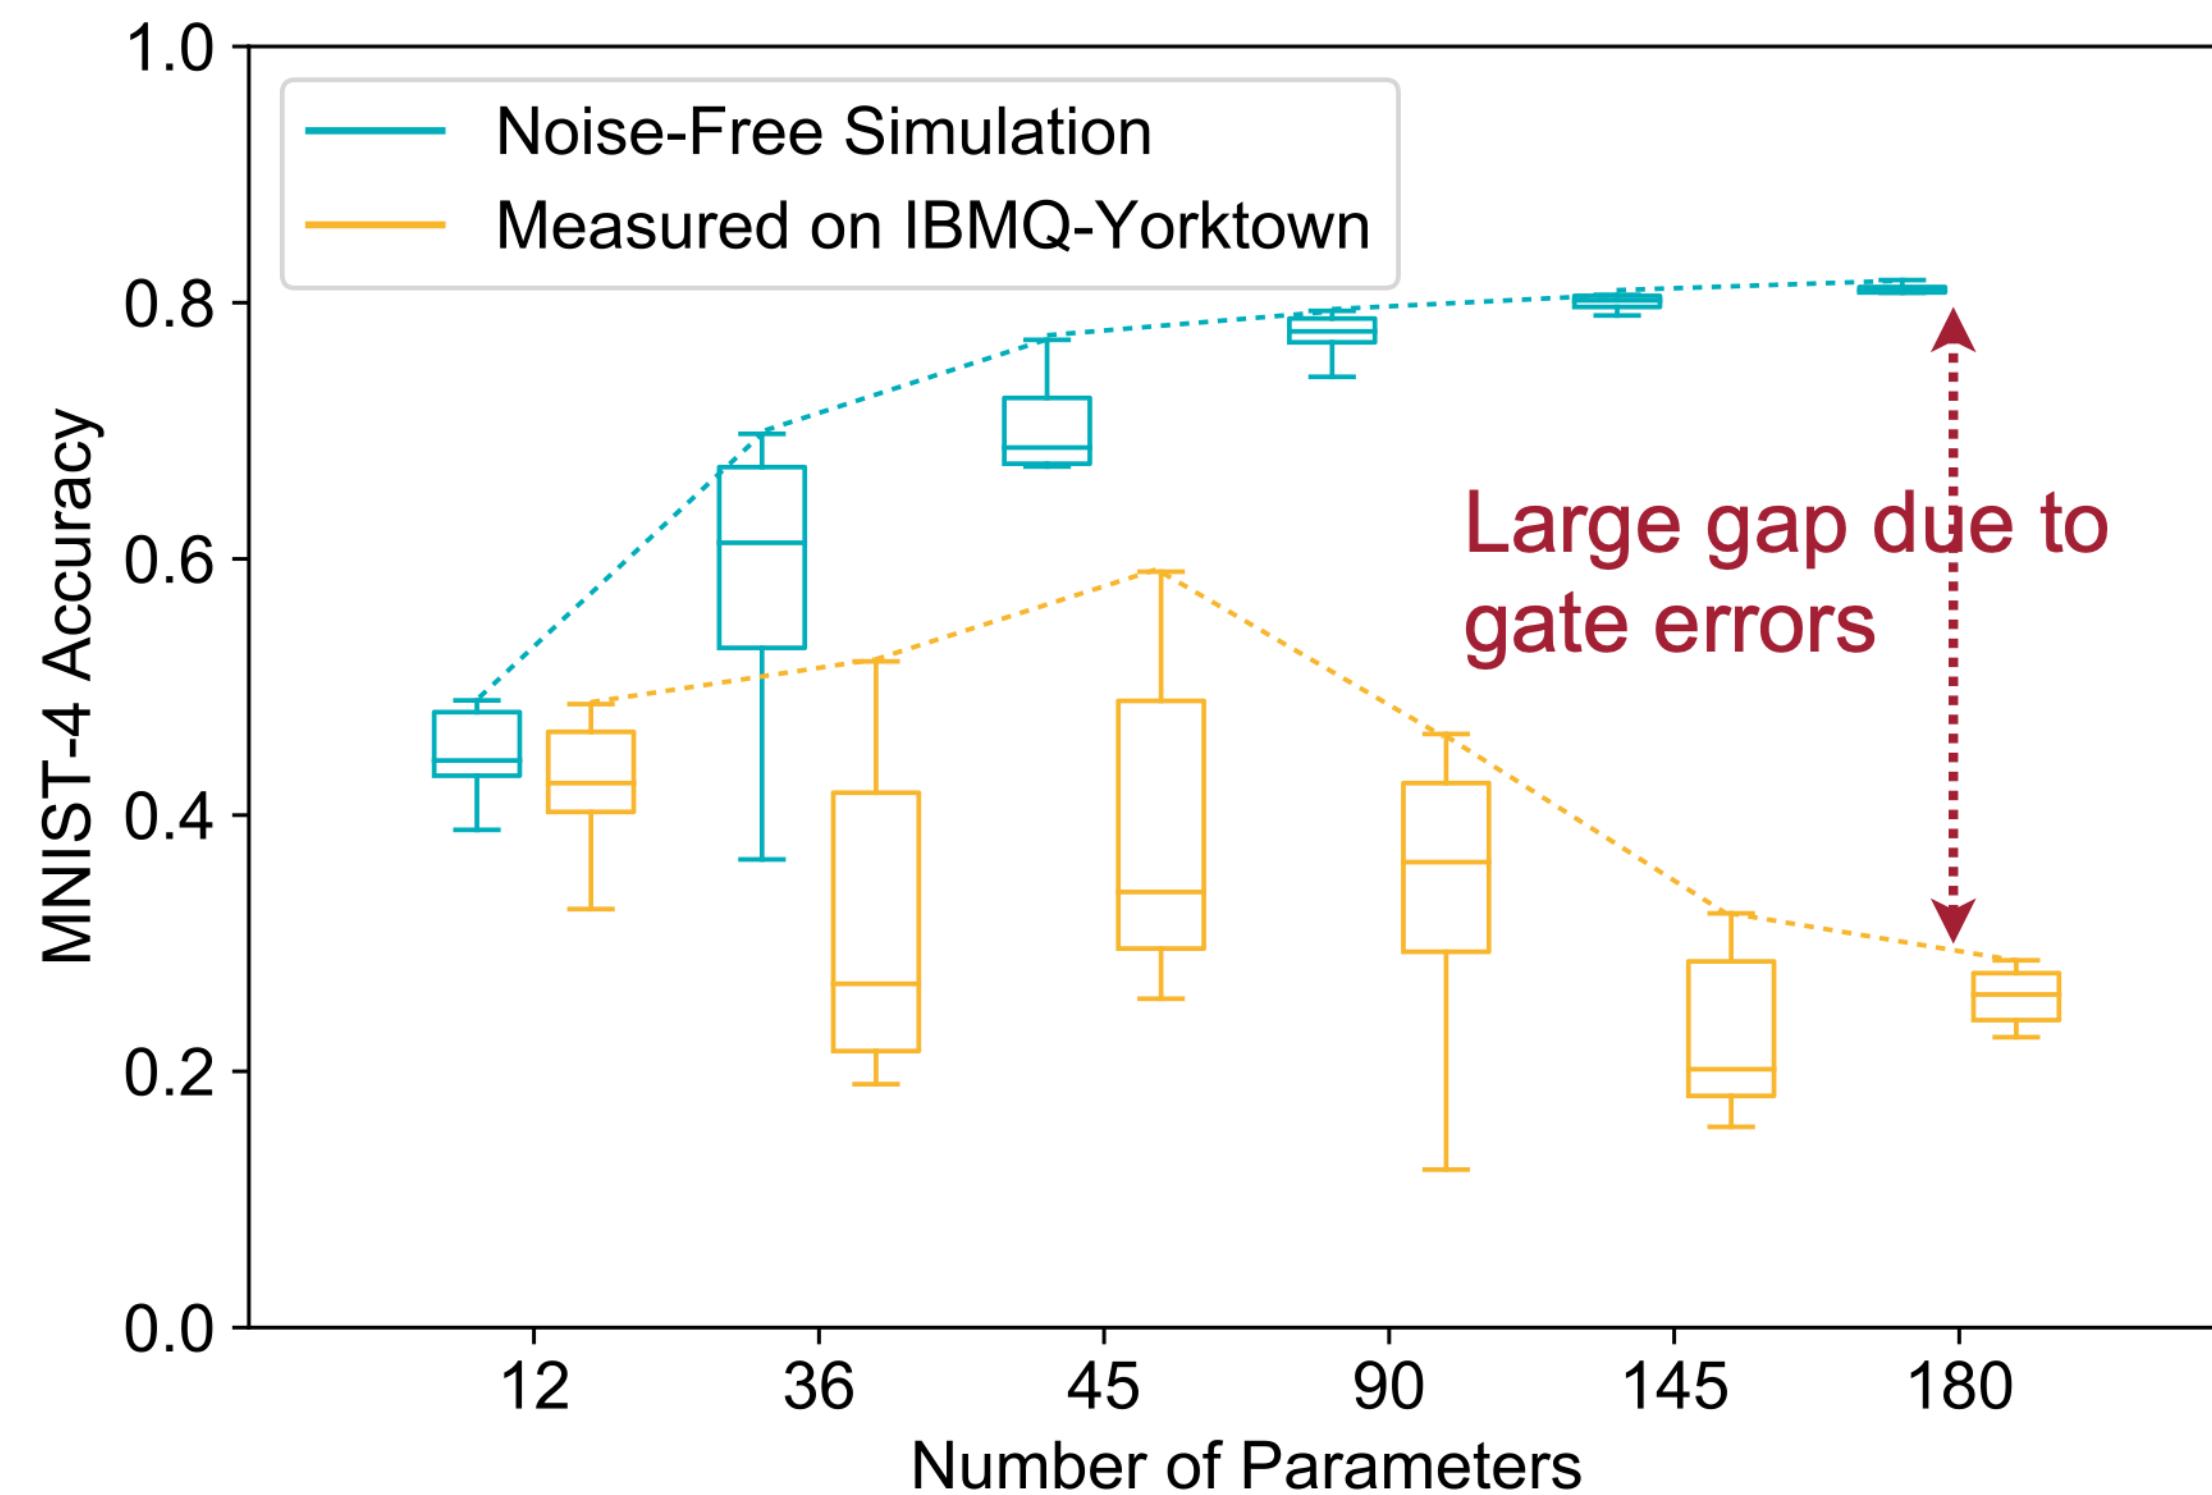
- Parameterized Quantum Circuits (PQC)
- Quantum circuit with fixed gates and parameterized gates



- PQCs are commonly used in **hybrid classical-quantum models** and show promises to achieve quantum advantage
 - Variational Quantum Eigensolver (VQE)
 - Quantum Neural Networks (QNN)
 - Quantum Approximate Optimization Algorithm (QAOA)

Challenges of PQC — Noise

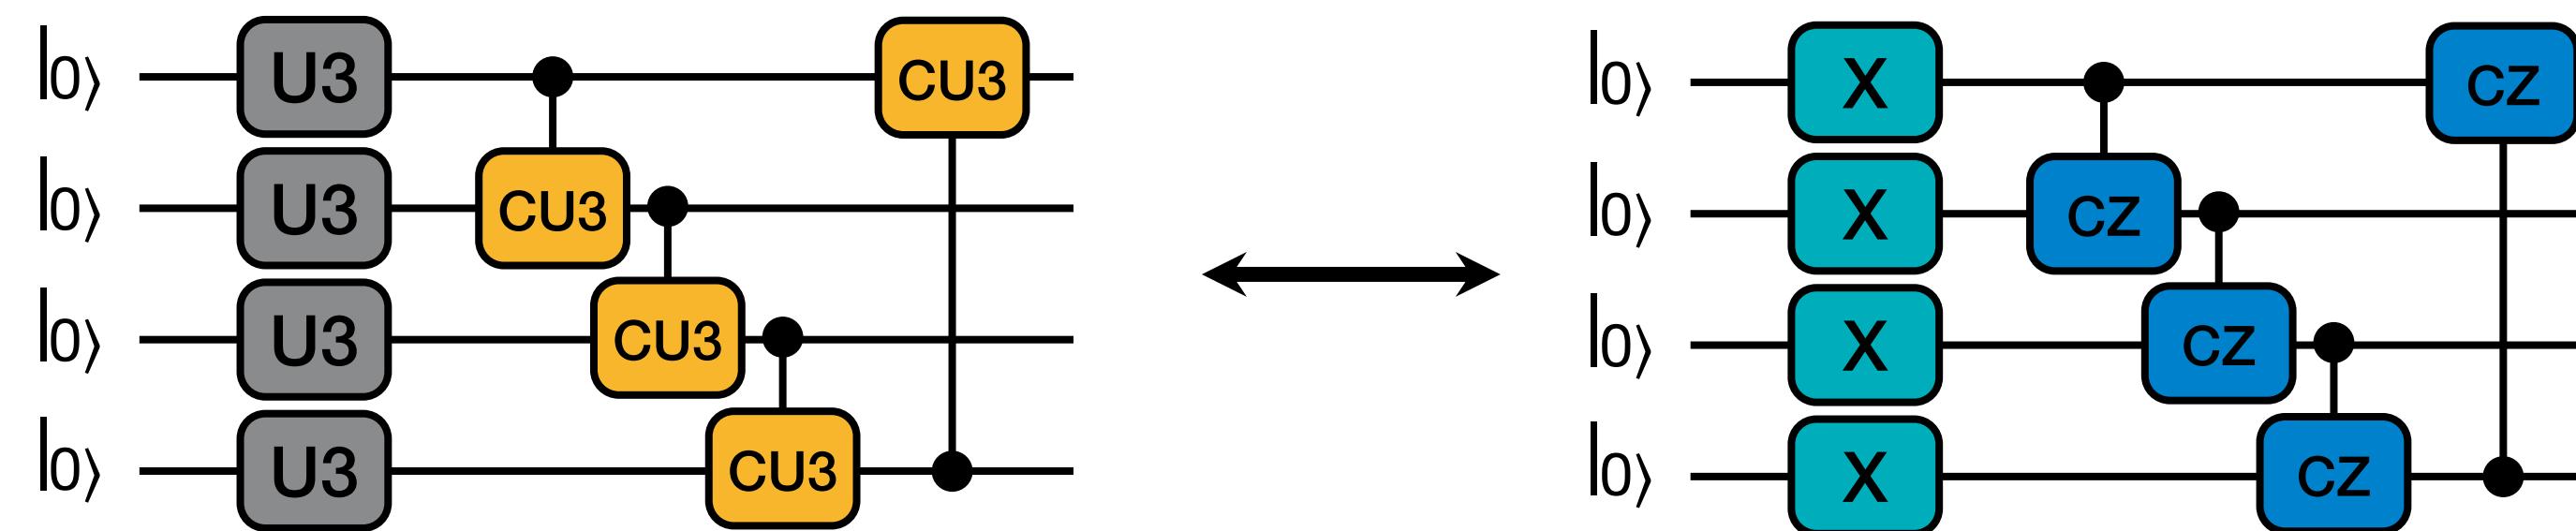
- Noise **degrades** PQC reliability
- More parameters increase the noise-free accuracy but degrade the measured accuracy
- Therefore, circuit architecture is critical



Challenges of PQC — Large Design Space

- Large design space for circuit architecture

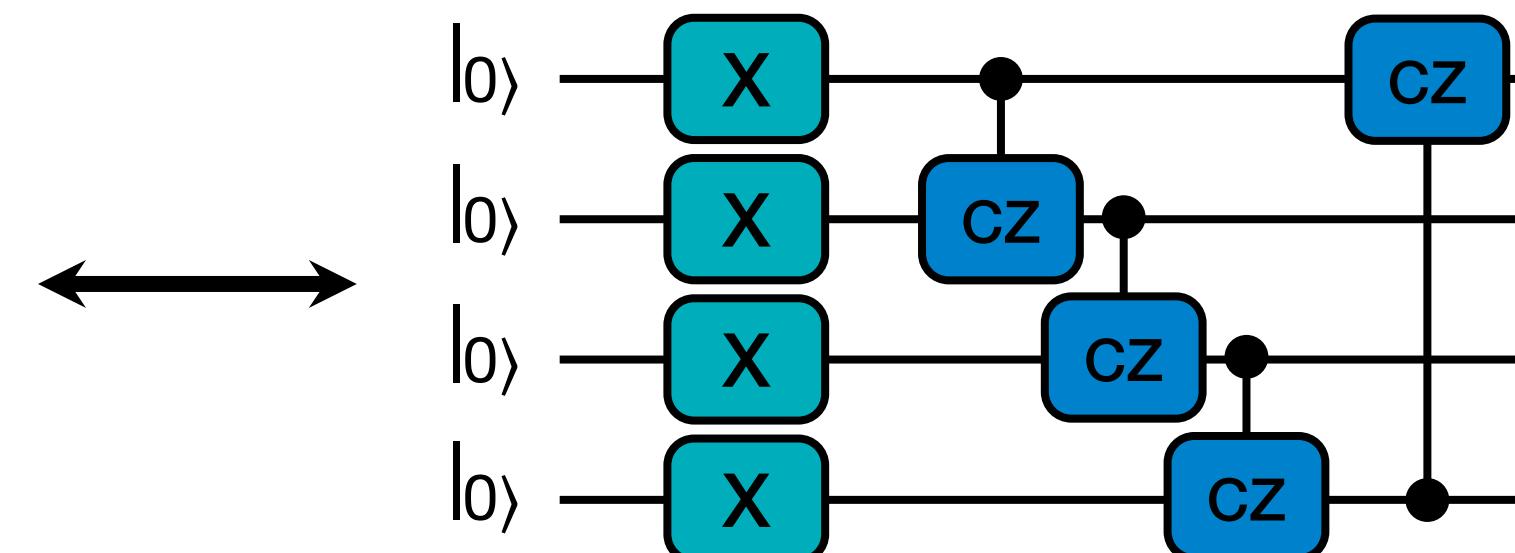
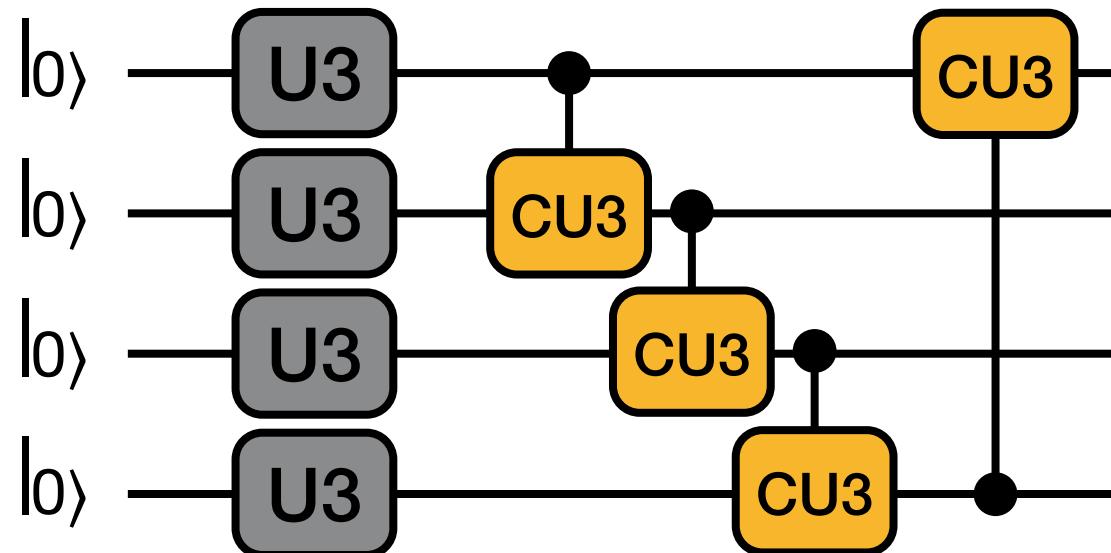
- Type of gates



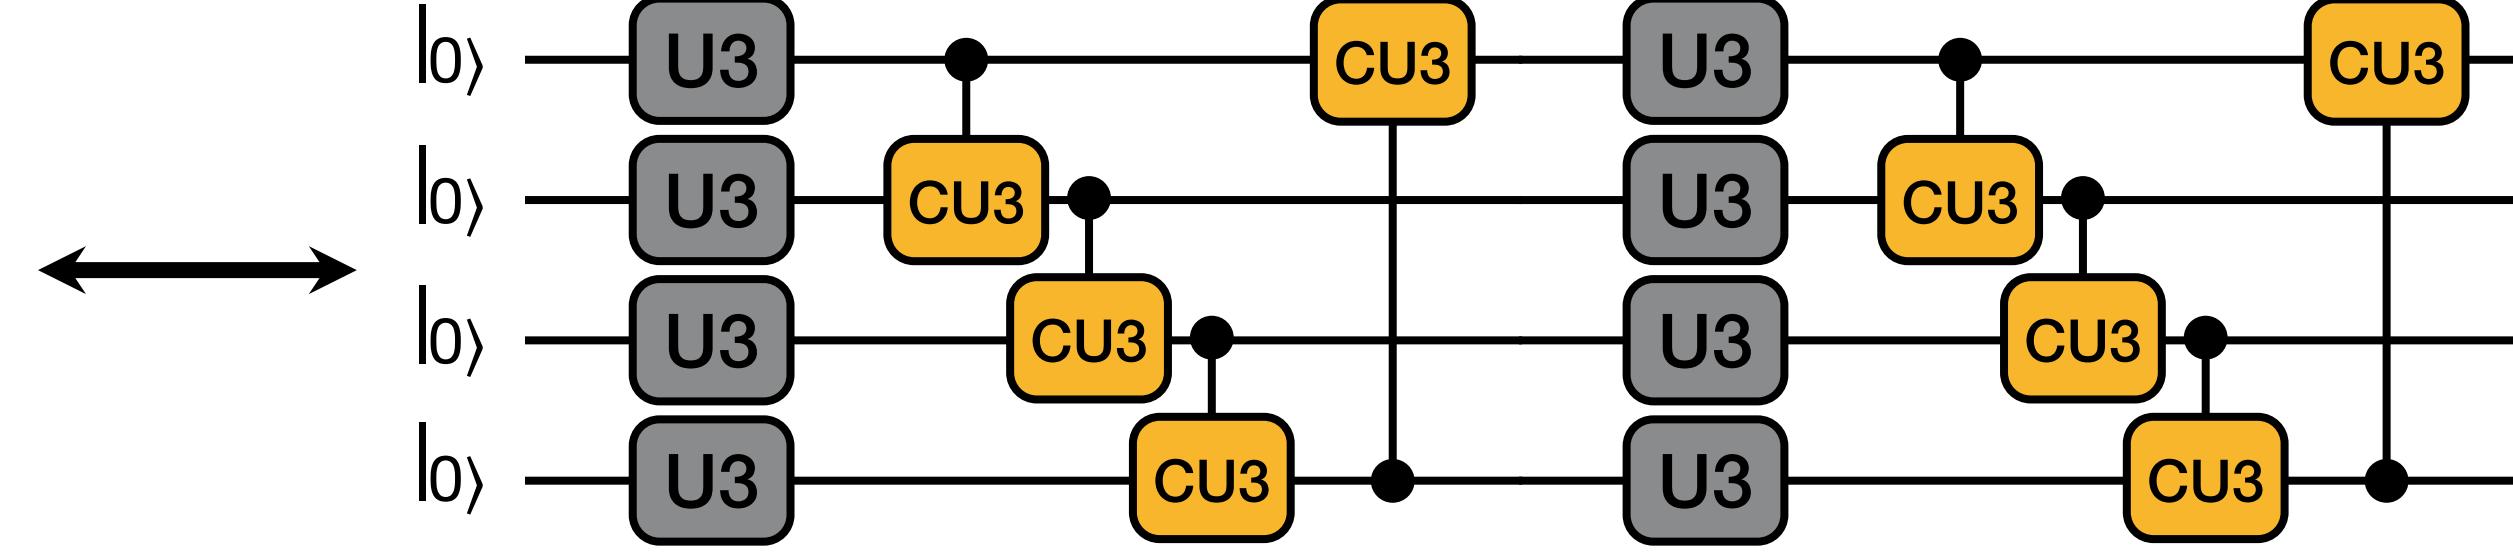
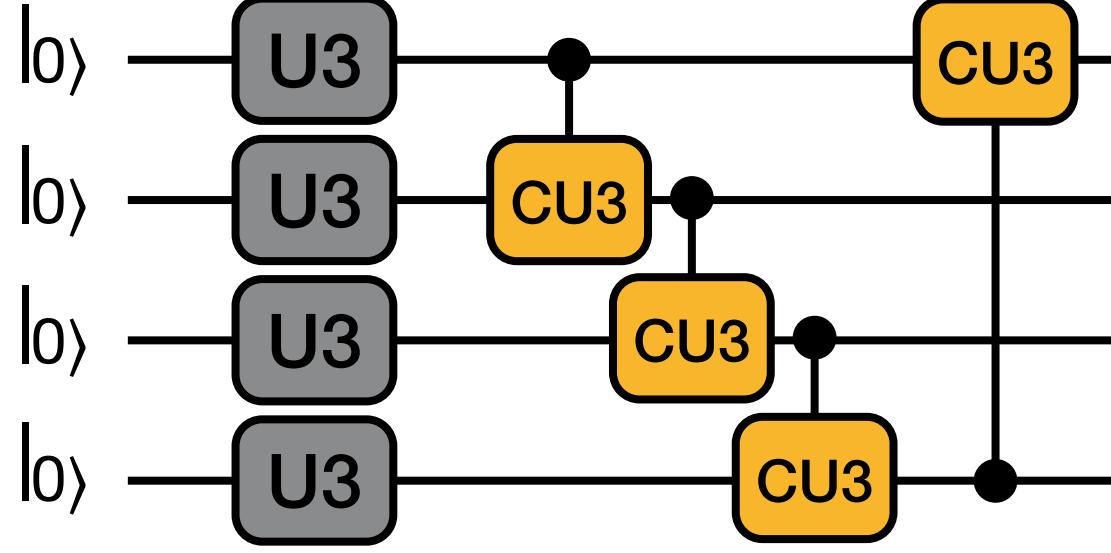
Challenges of PQC — Large Design Space

- Large design space for circuit architecture

- Type of gates



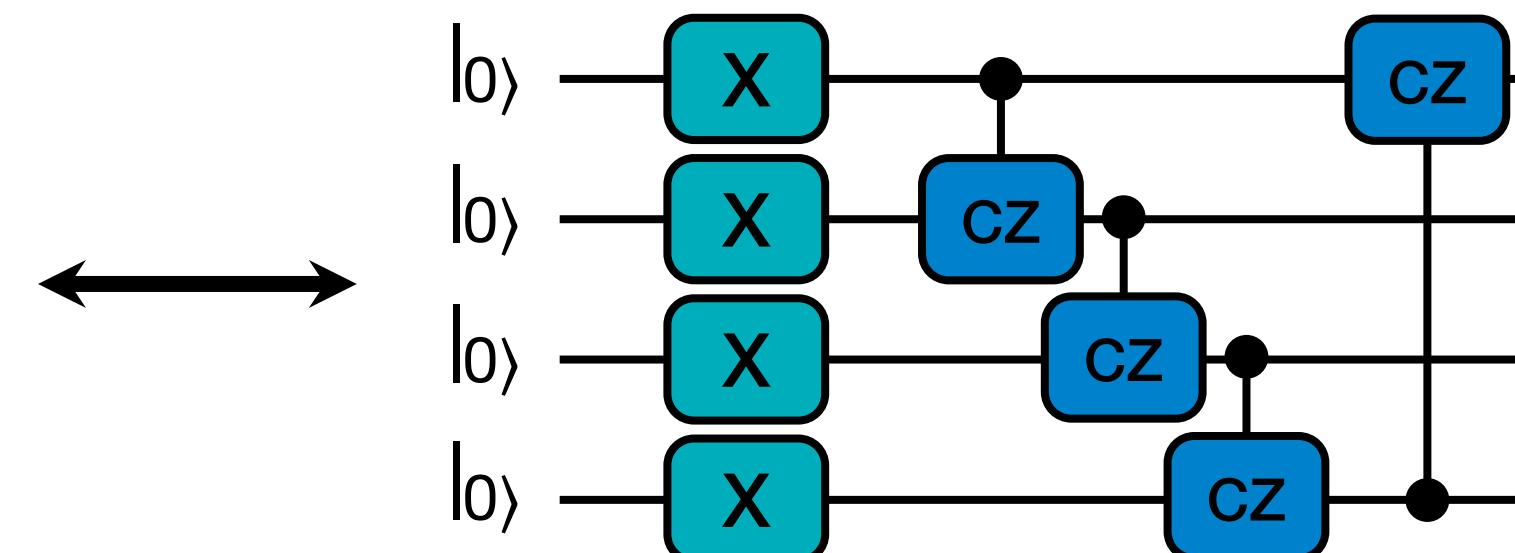
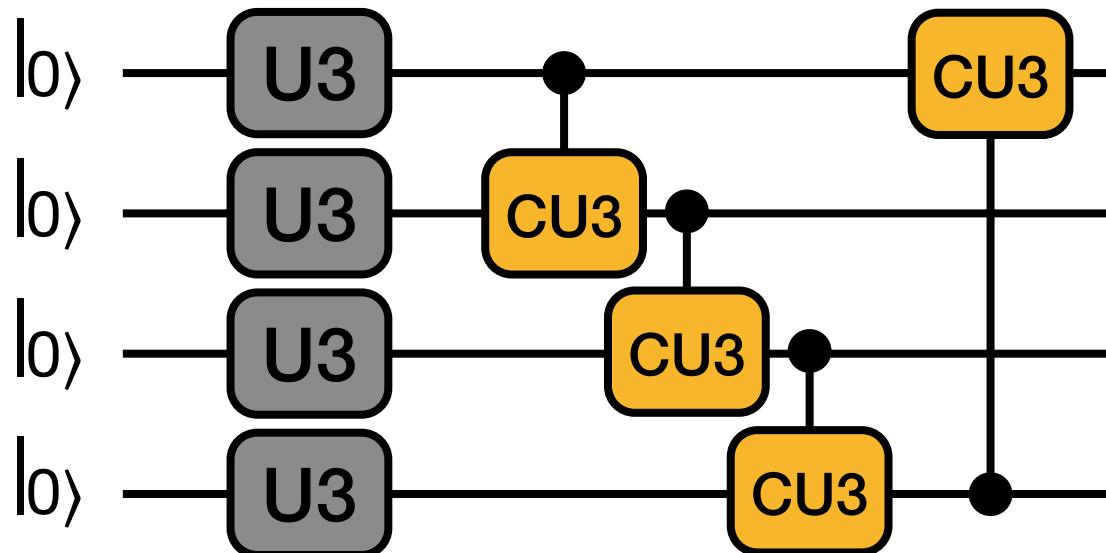
- Number of gates



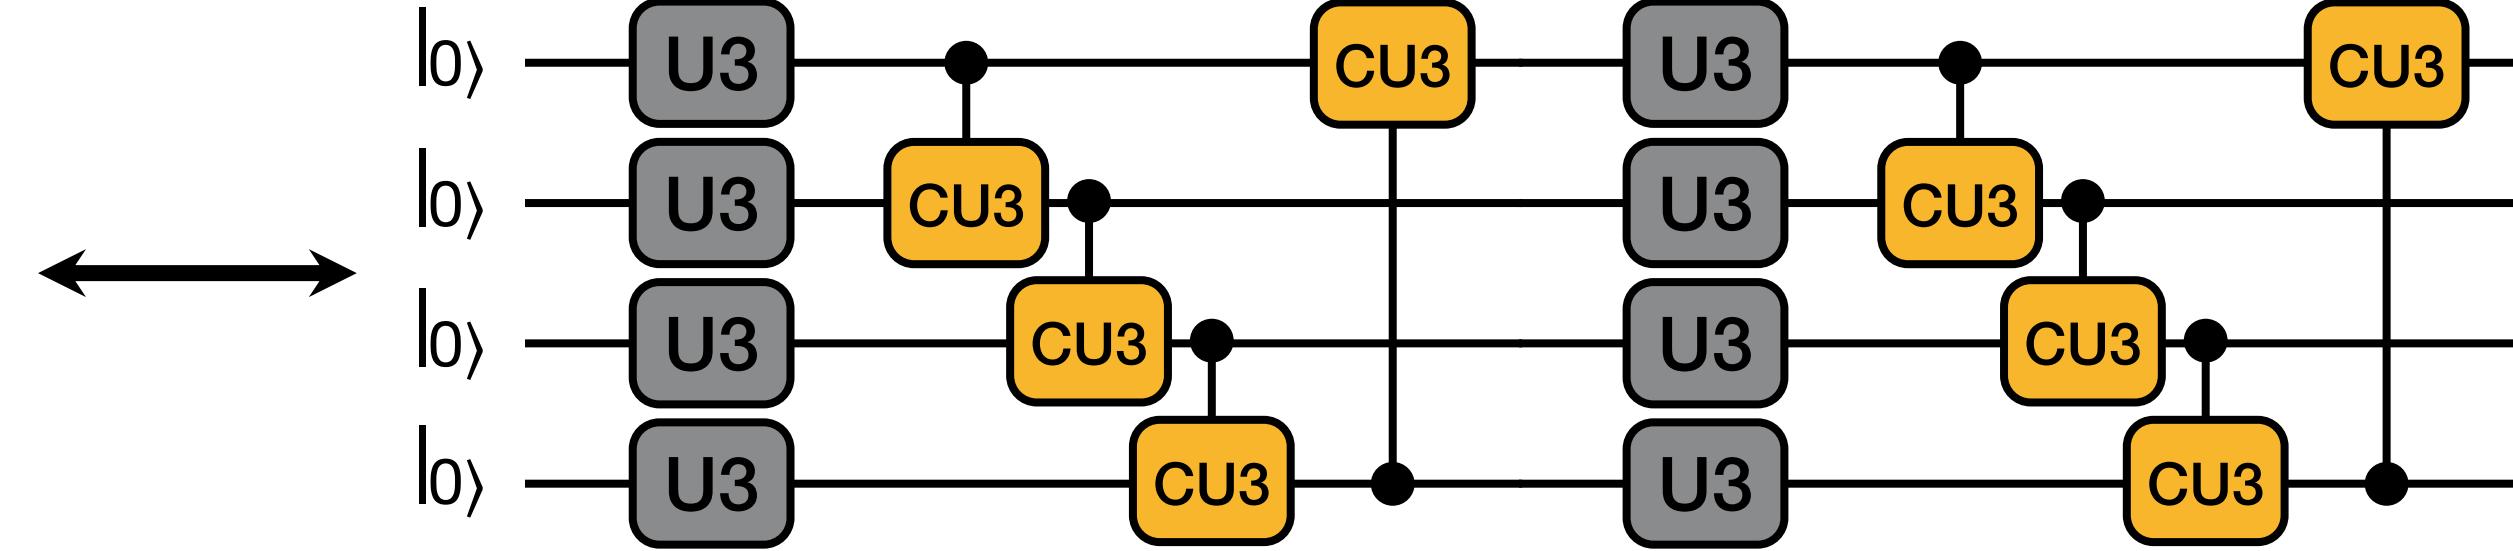
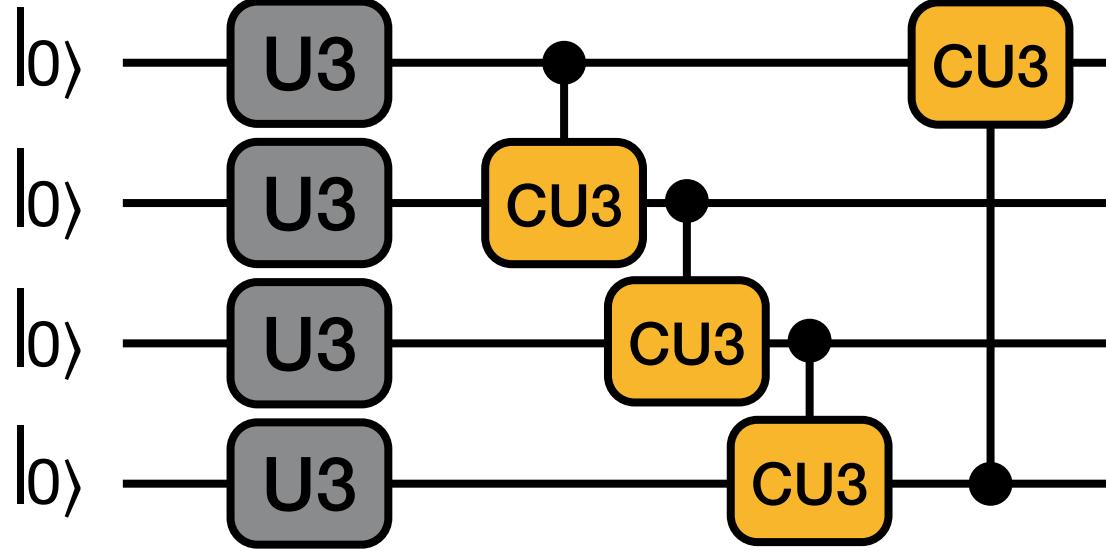
Challenges of PQC — Large Design Space

- Large design space for circuit architecture

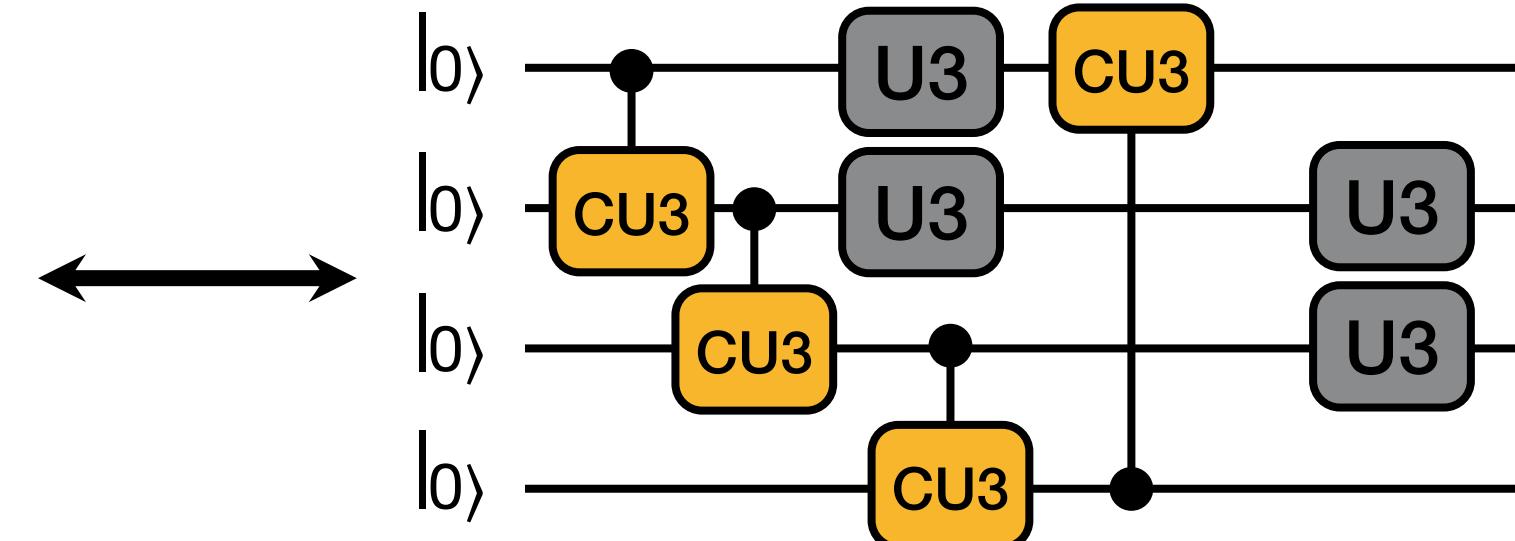
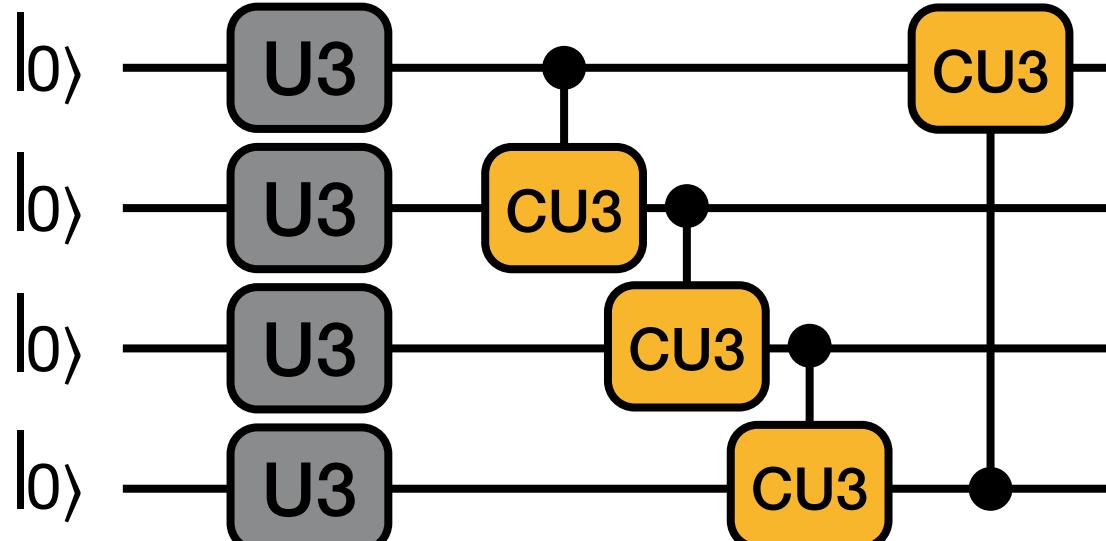
- Type of gates



- Number of gates



- Position of gates



Goal of QuantumNAS

Automatically & efficiently search for noise-robust quantum circuit

Train one “SuperCircuit”,
providing parameters to
many “SubCircuits”

Solve the challenge of large
design space

(1) Quantum noise feedback in
the search loop
(2) Co-search the circuit
architecture and qubit mapping

Solve the challenge of large
quantum noise

QuantumNAS

- SuperCircuit Construction and Training
- Noise-Adaptive Evolutionary Co-Search of SubCircuit and Qubit Mapping
- Train the Searched SubCircuit
- Iterative Quantum Gate Pruning

QuantumNAS

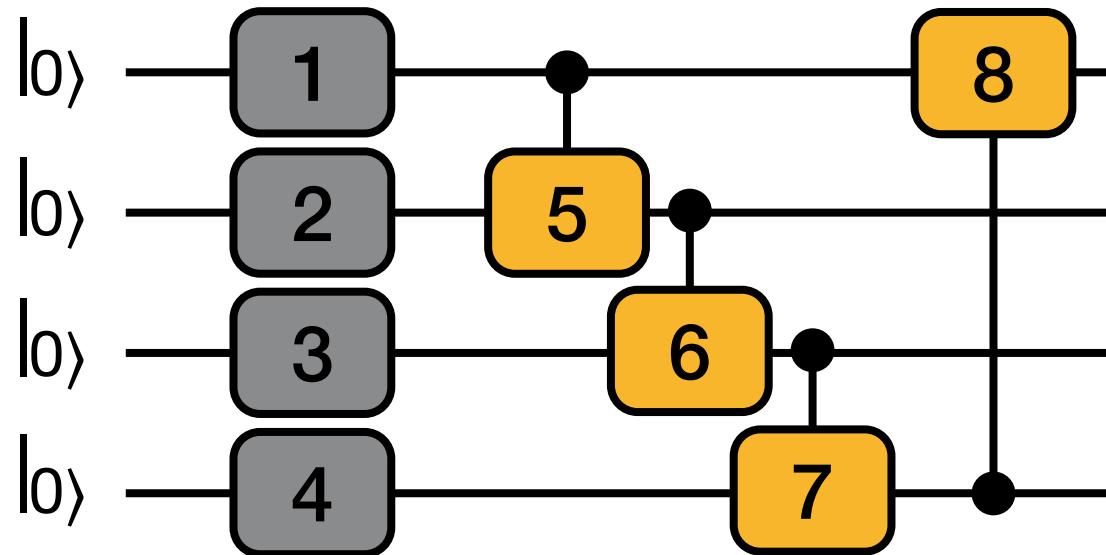
- SuperCircuit Construction and Training
- Noise-Adaptive Evolutionary Co-Search of SubCircuit and Qubit Mapping
- Train the Searched SubCircuit
- Iterative Quantum Gate Pruning

SuperCircuit & SubCircuit

- Firstly construct a design space. For example, a design space of maximum 4 U3 in the first layer and 4 CU3 gates in the second layer

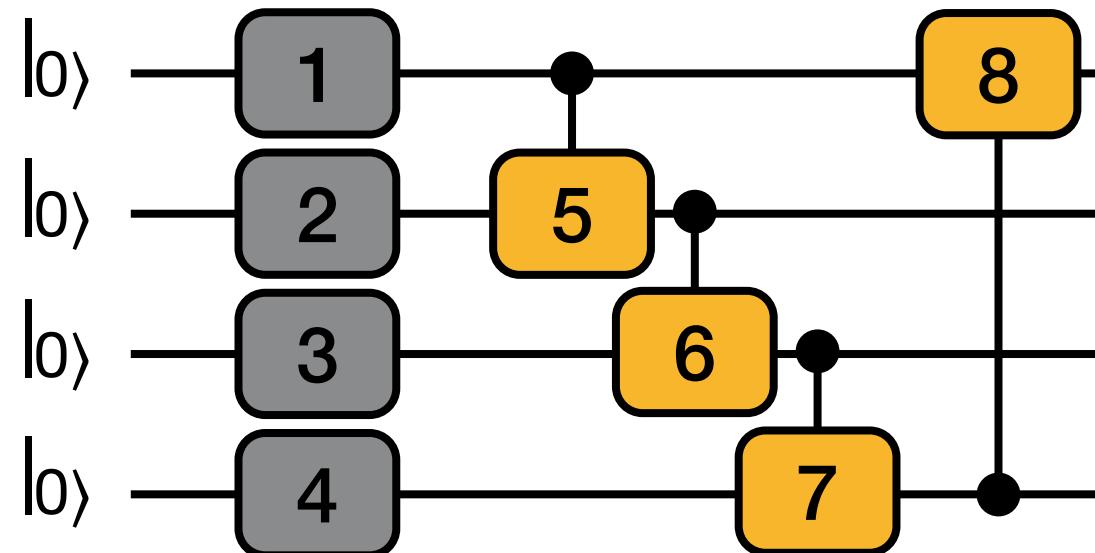
SuperCircuit & SubCircuit

- Firstly construct a design space. For example, a design space of maximum 4 U3 in the first layer and 4 CU3 gates in the second layer
- SuperCircuit: the circuit with the **largest** number of gates in the design space
 - Example: SuperCircuit in U3+CU3 space

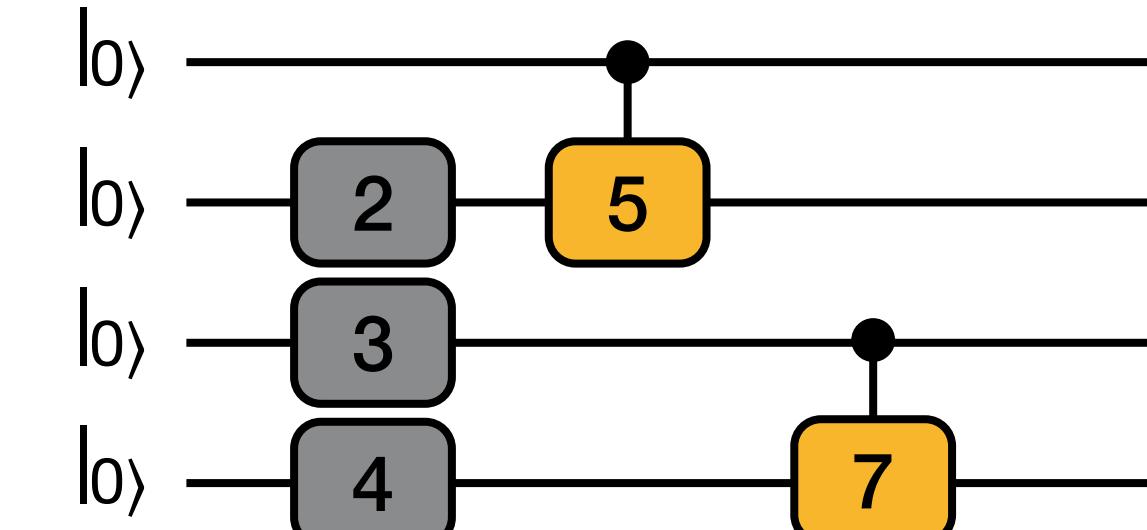
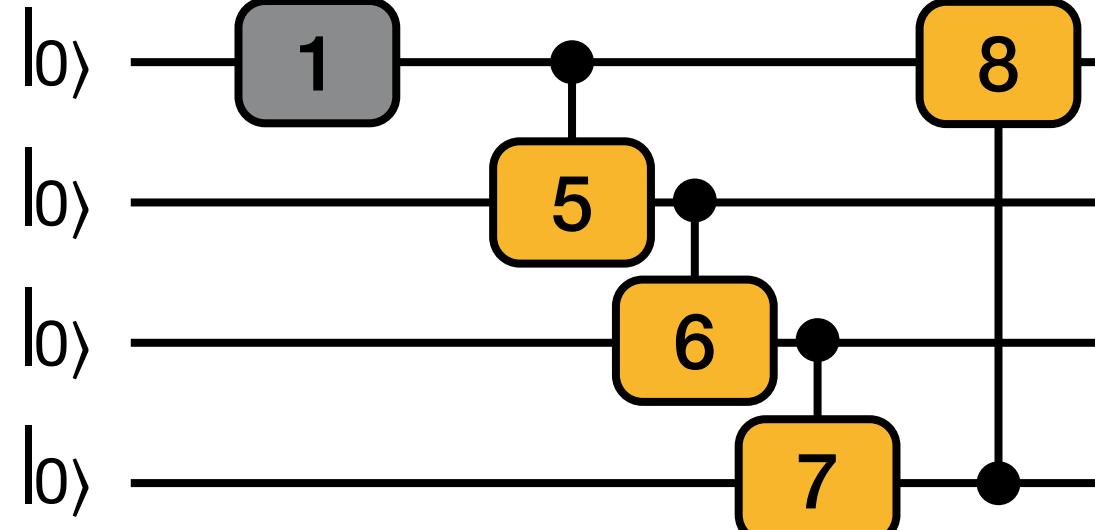
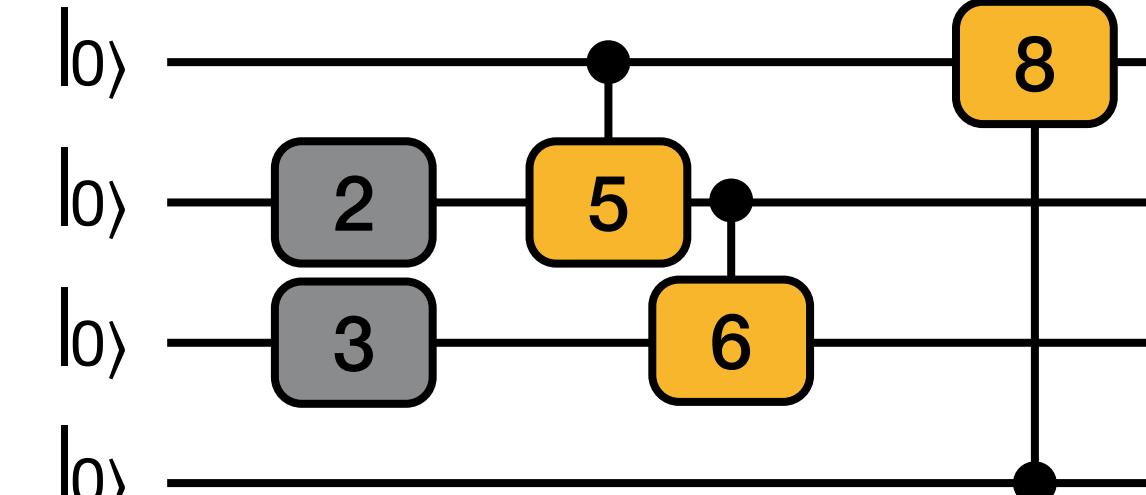


SuperCircuit & SubCircuit

- Firstly construct a design space. For example, a design space of maximum 4 U3 in the first layer and 4 CU3 gates in the second layer
- SuperCircuit: the circuit with the **largest** number of gates in the design space
 - Example: SuperCircuit in U3+CU3 space



- Each candidate circuit in the design space (called SubCircuit) is a **subset** of the SuperCircuit



SuperCircuit Construction

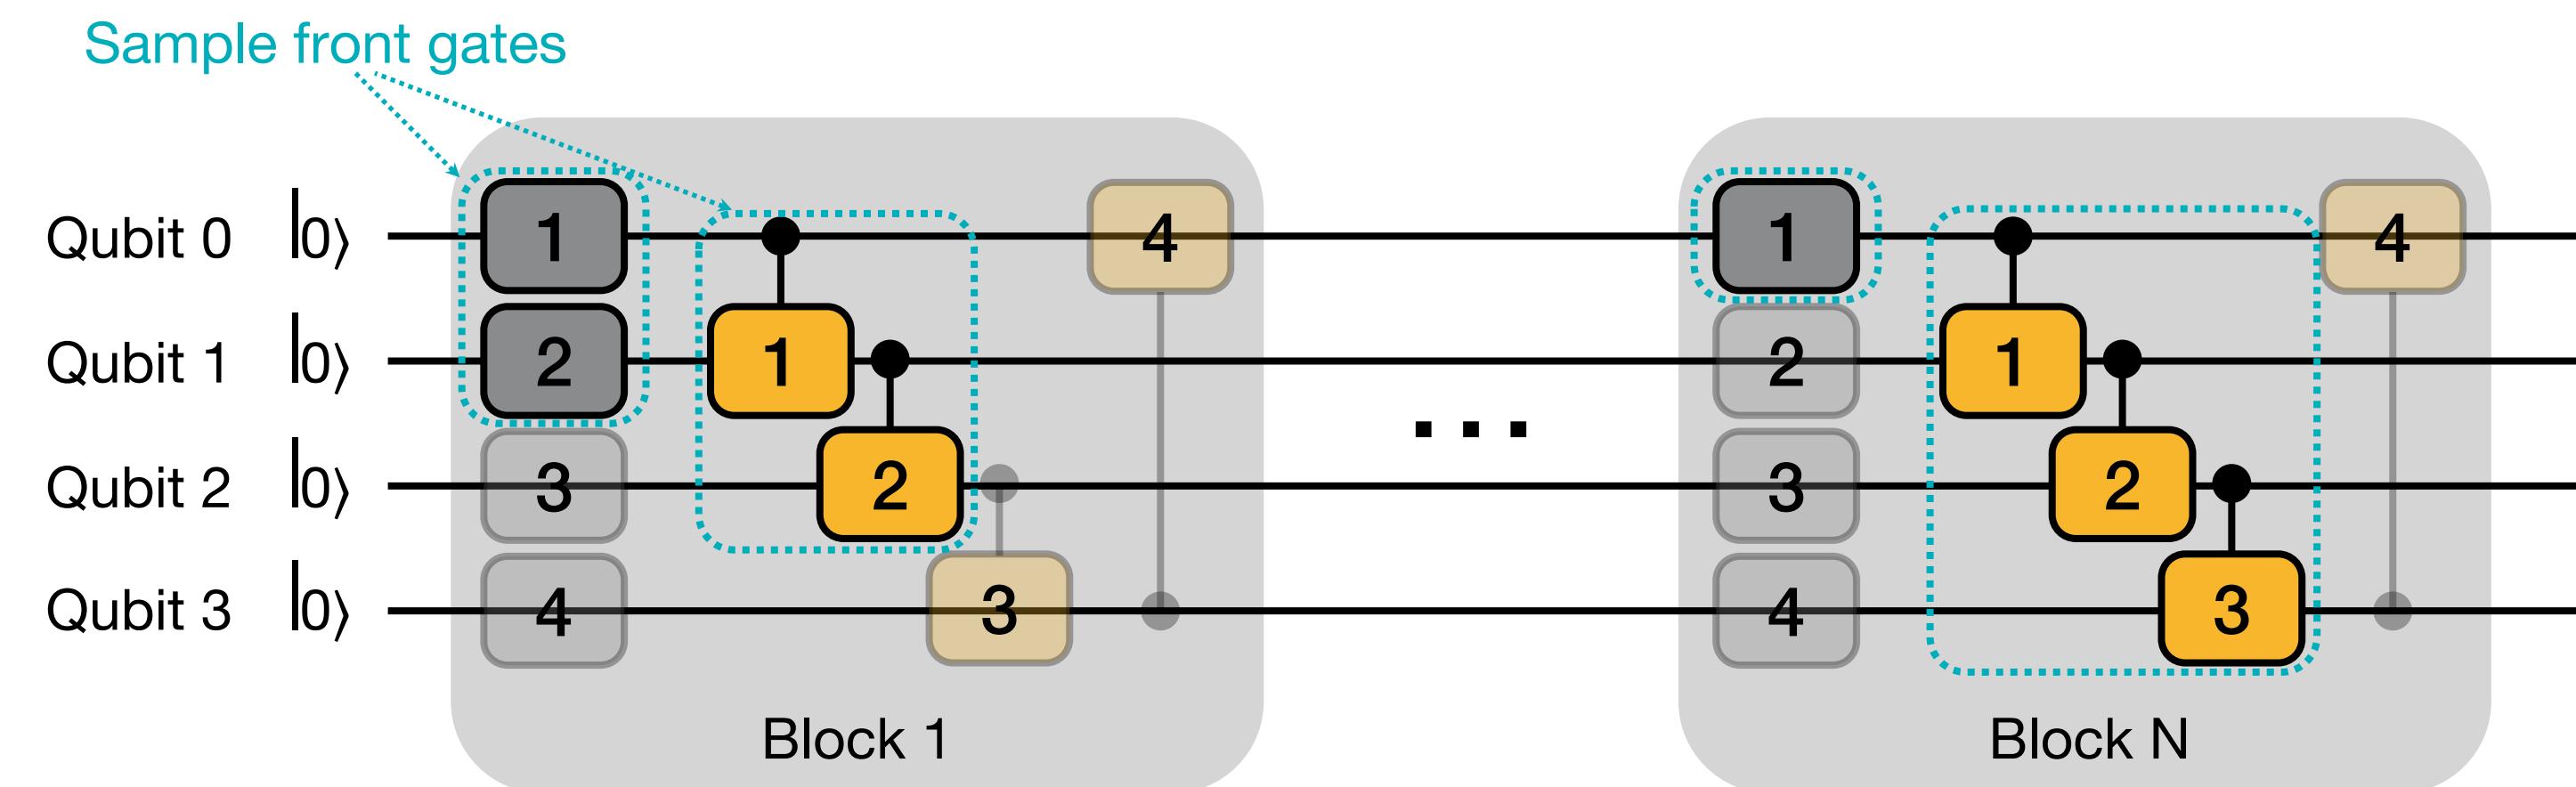
- Why use a SuperCircuit?
 - Enables **efficient** search of architecture candidates without training each
 - SubCircuit inherits parameters from SuperCircuit
 - With **inherited** parameters, we find some good SubCircuits, we find that they are **also good SubCircuits** with parameters **trained from-scratch** individually

SuperCircuit Training

- In one SuperCircuit Training step:
 - Sample a gate subset of SuperCircuit (a SubCircuit)
 - Front Sampling and Restricted Sampling
 - Only use the subset to perform the task and updates the parameters in the subset
 - Parameter updates are cumulative across steps

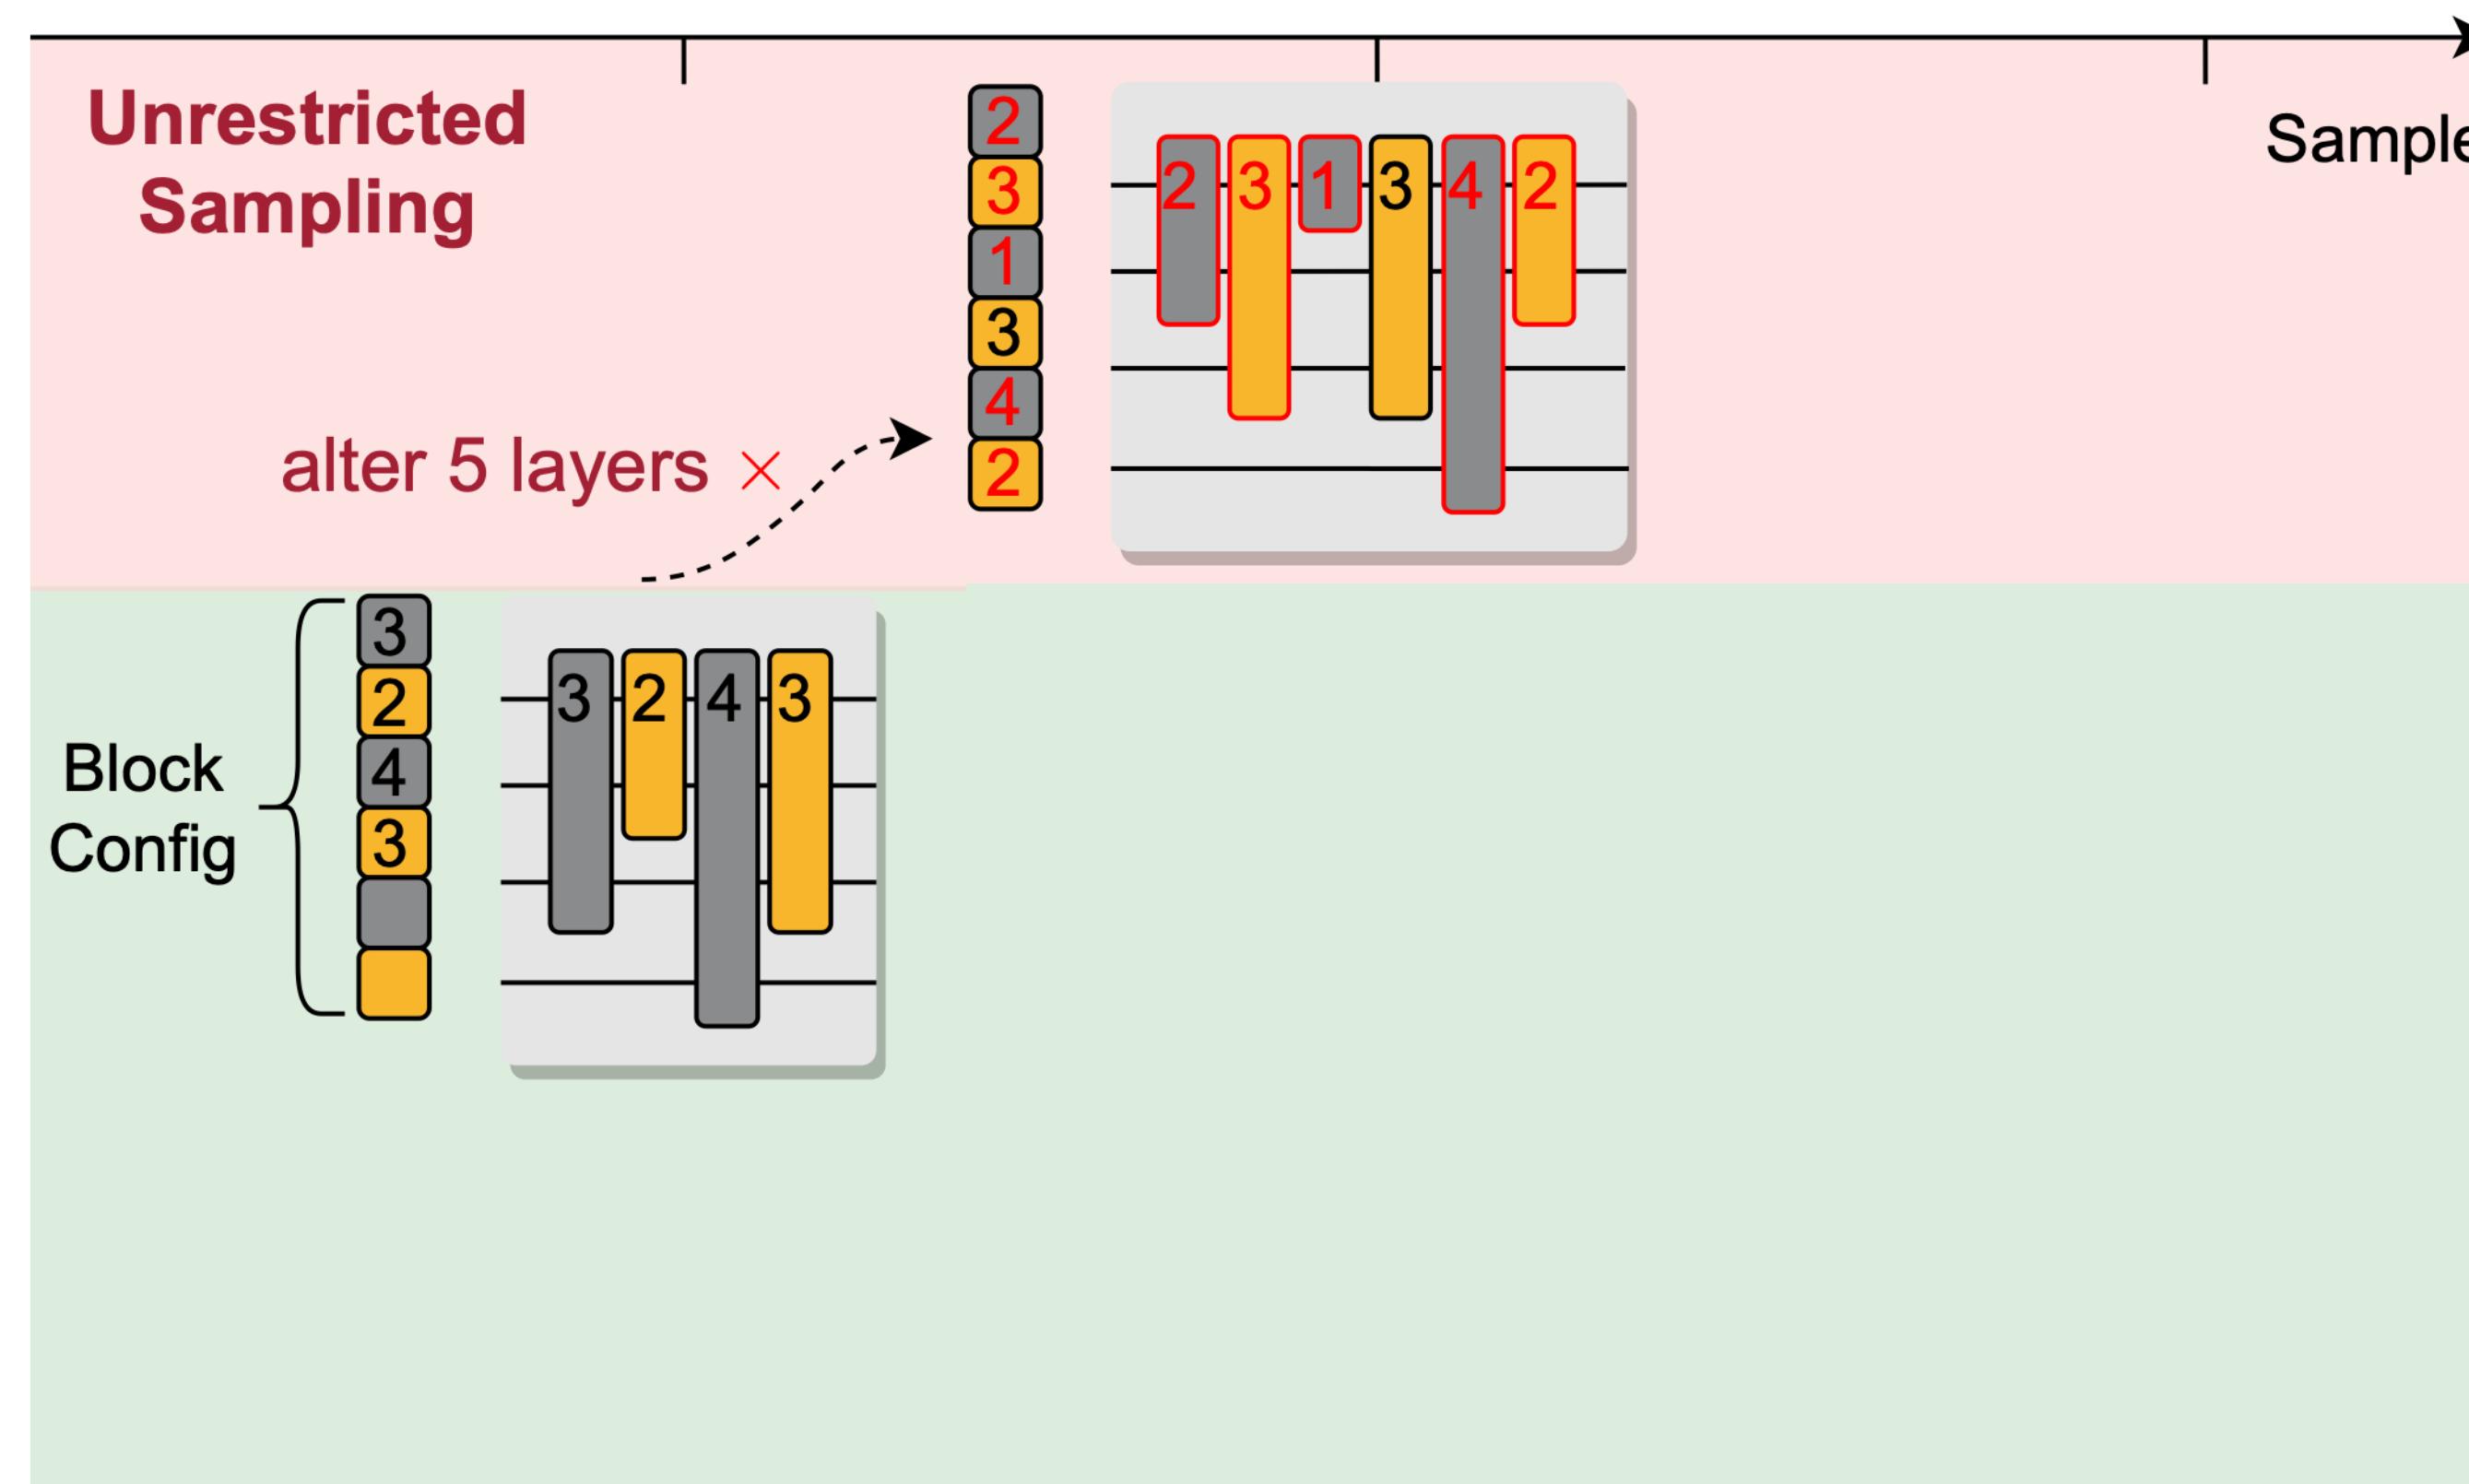
Front Sampling

- During sampling, we first sample total number of blocks, then sample gates within each block
 - Front sampling: Only the **front** several blocks and **front** several gates can be sampled to make SuperCircuit training more stable



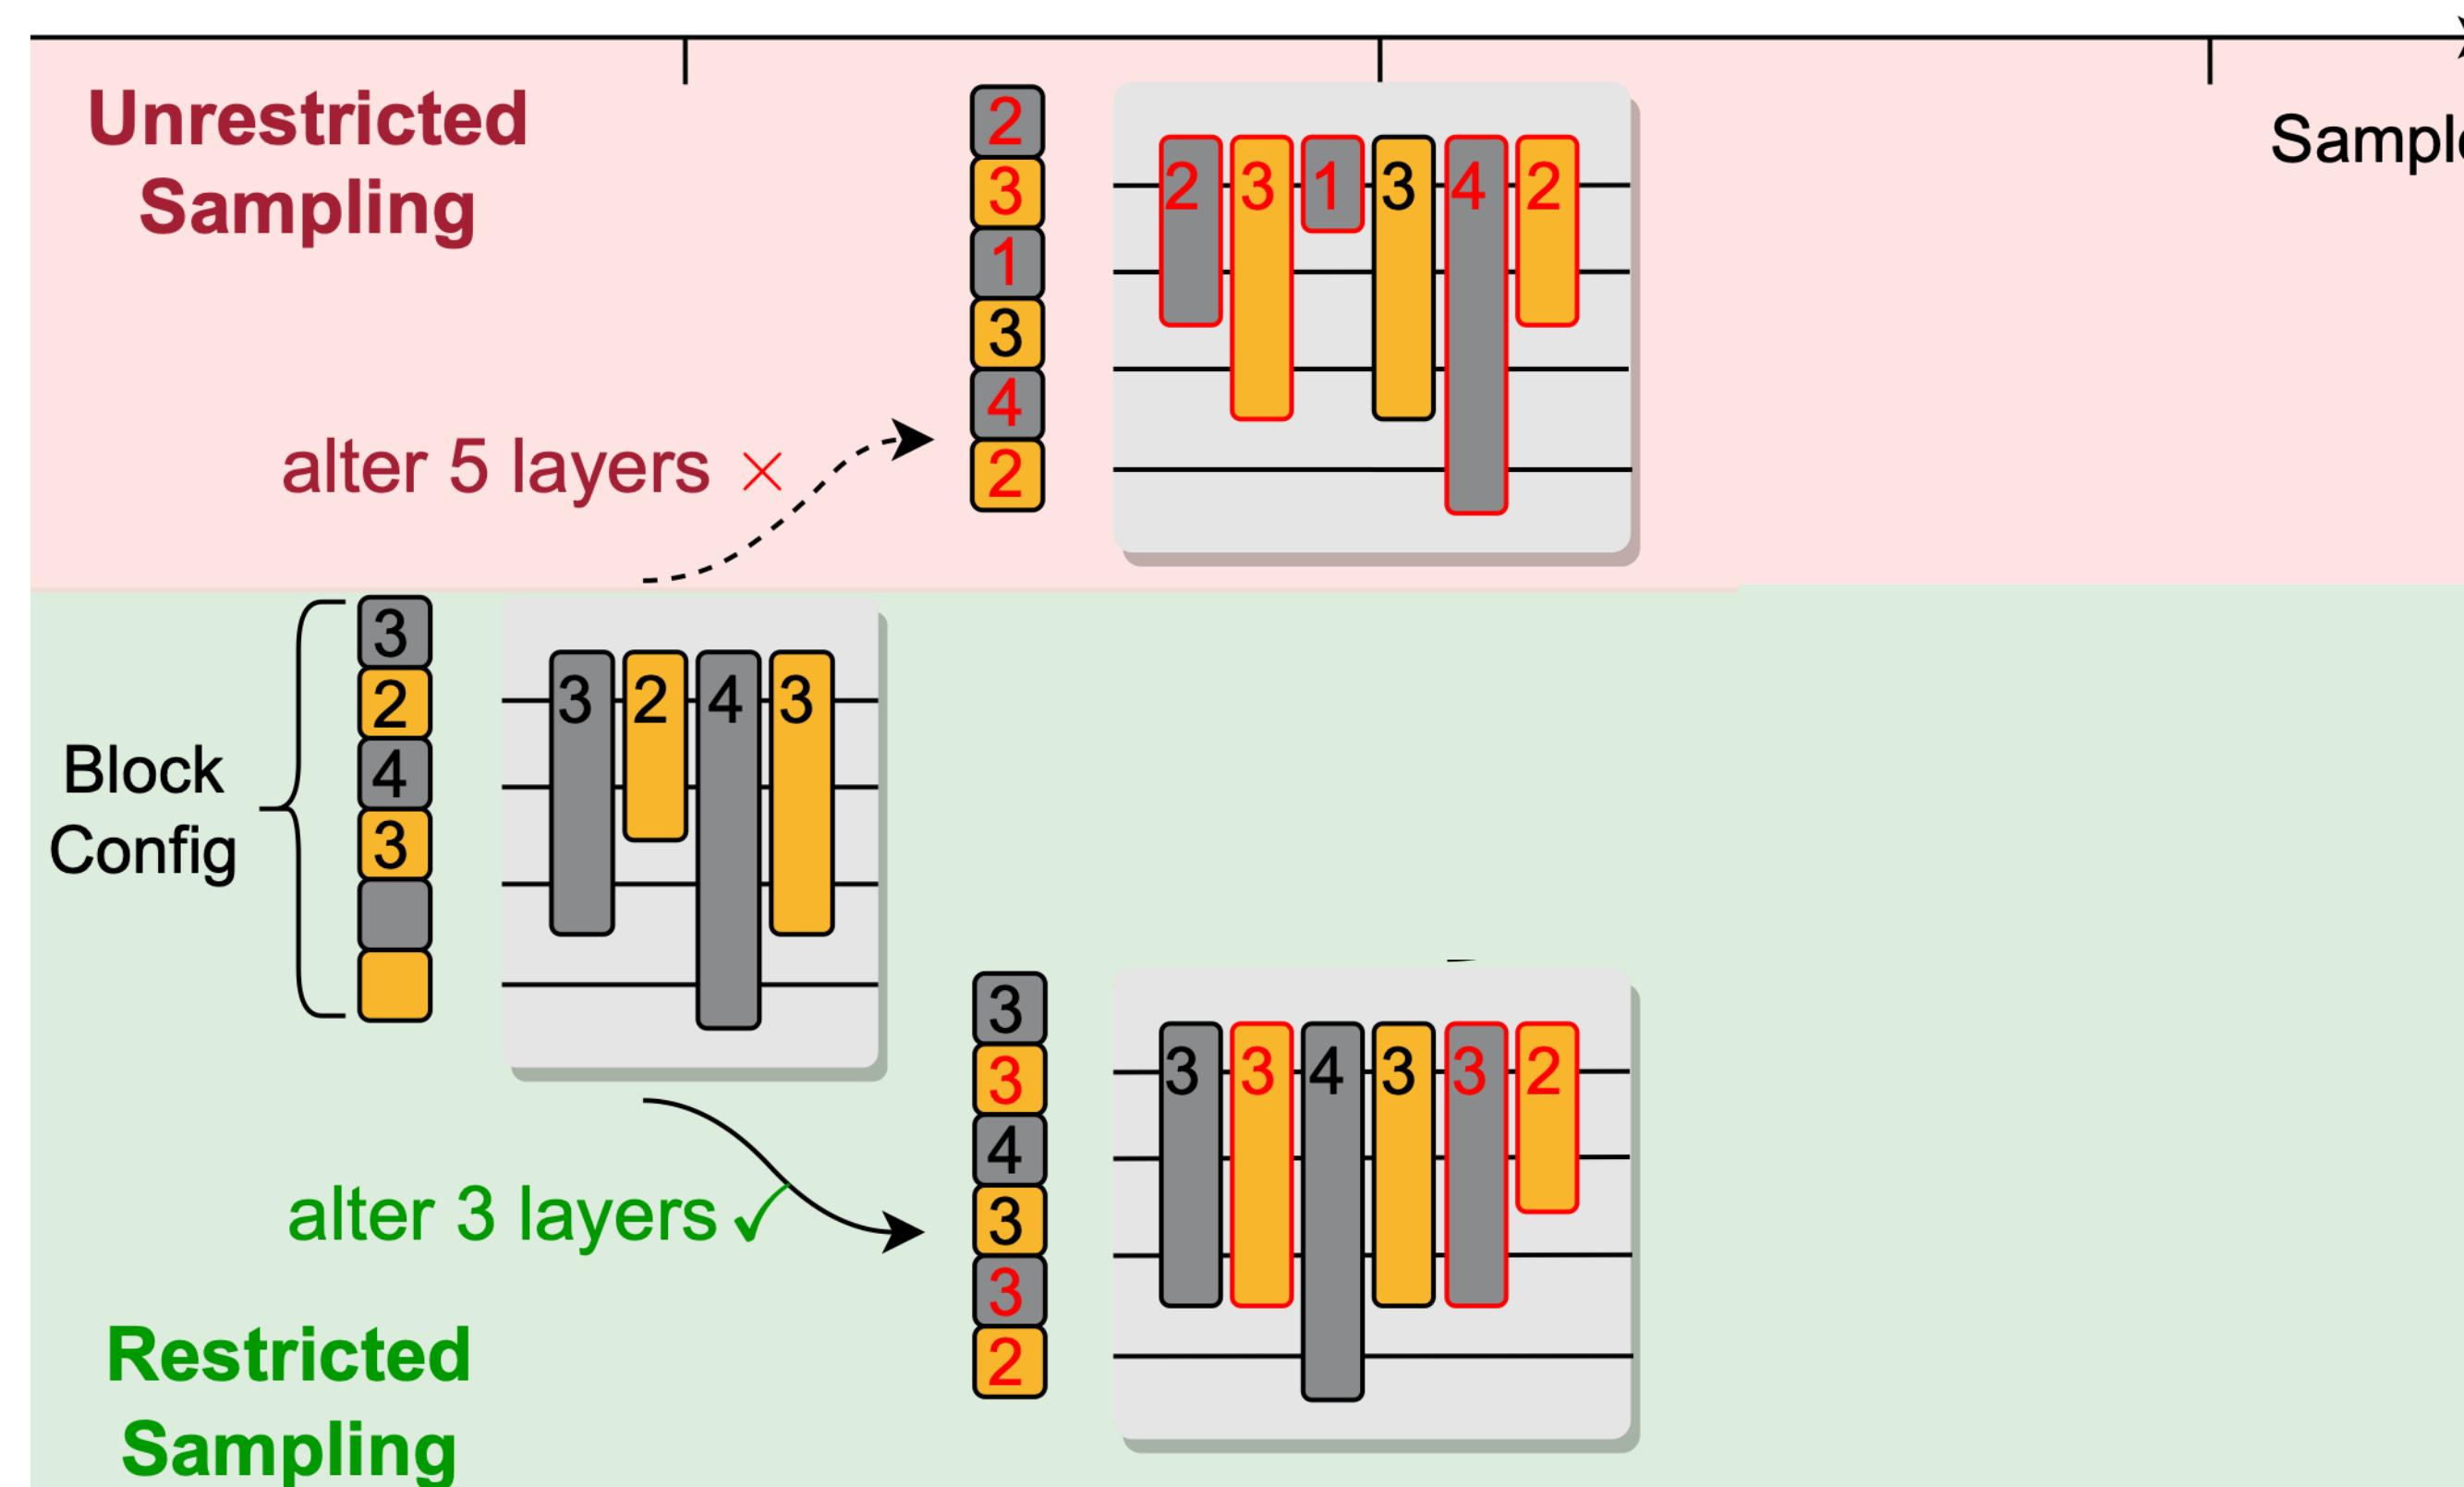
Restricted Sampling

- Restricted Sampling:
 - Restrict the difference between SubCircuits of two consecutive steps
 - For example: restrict to at most 4 different layers



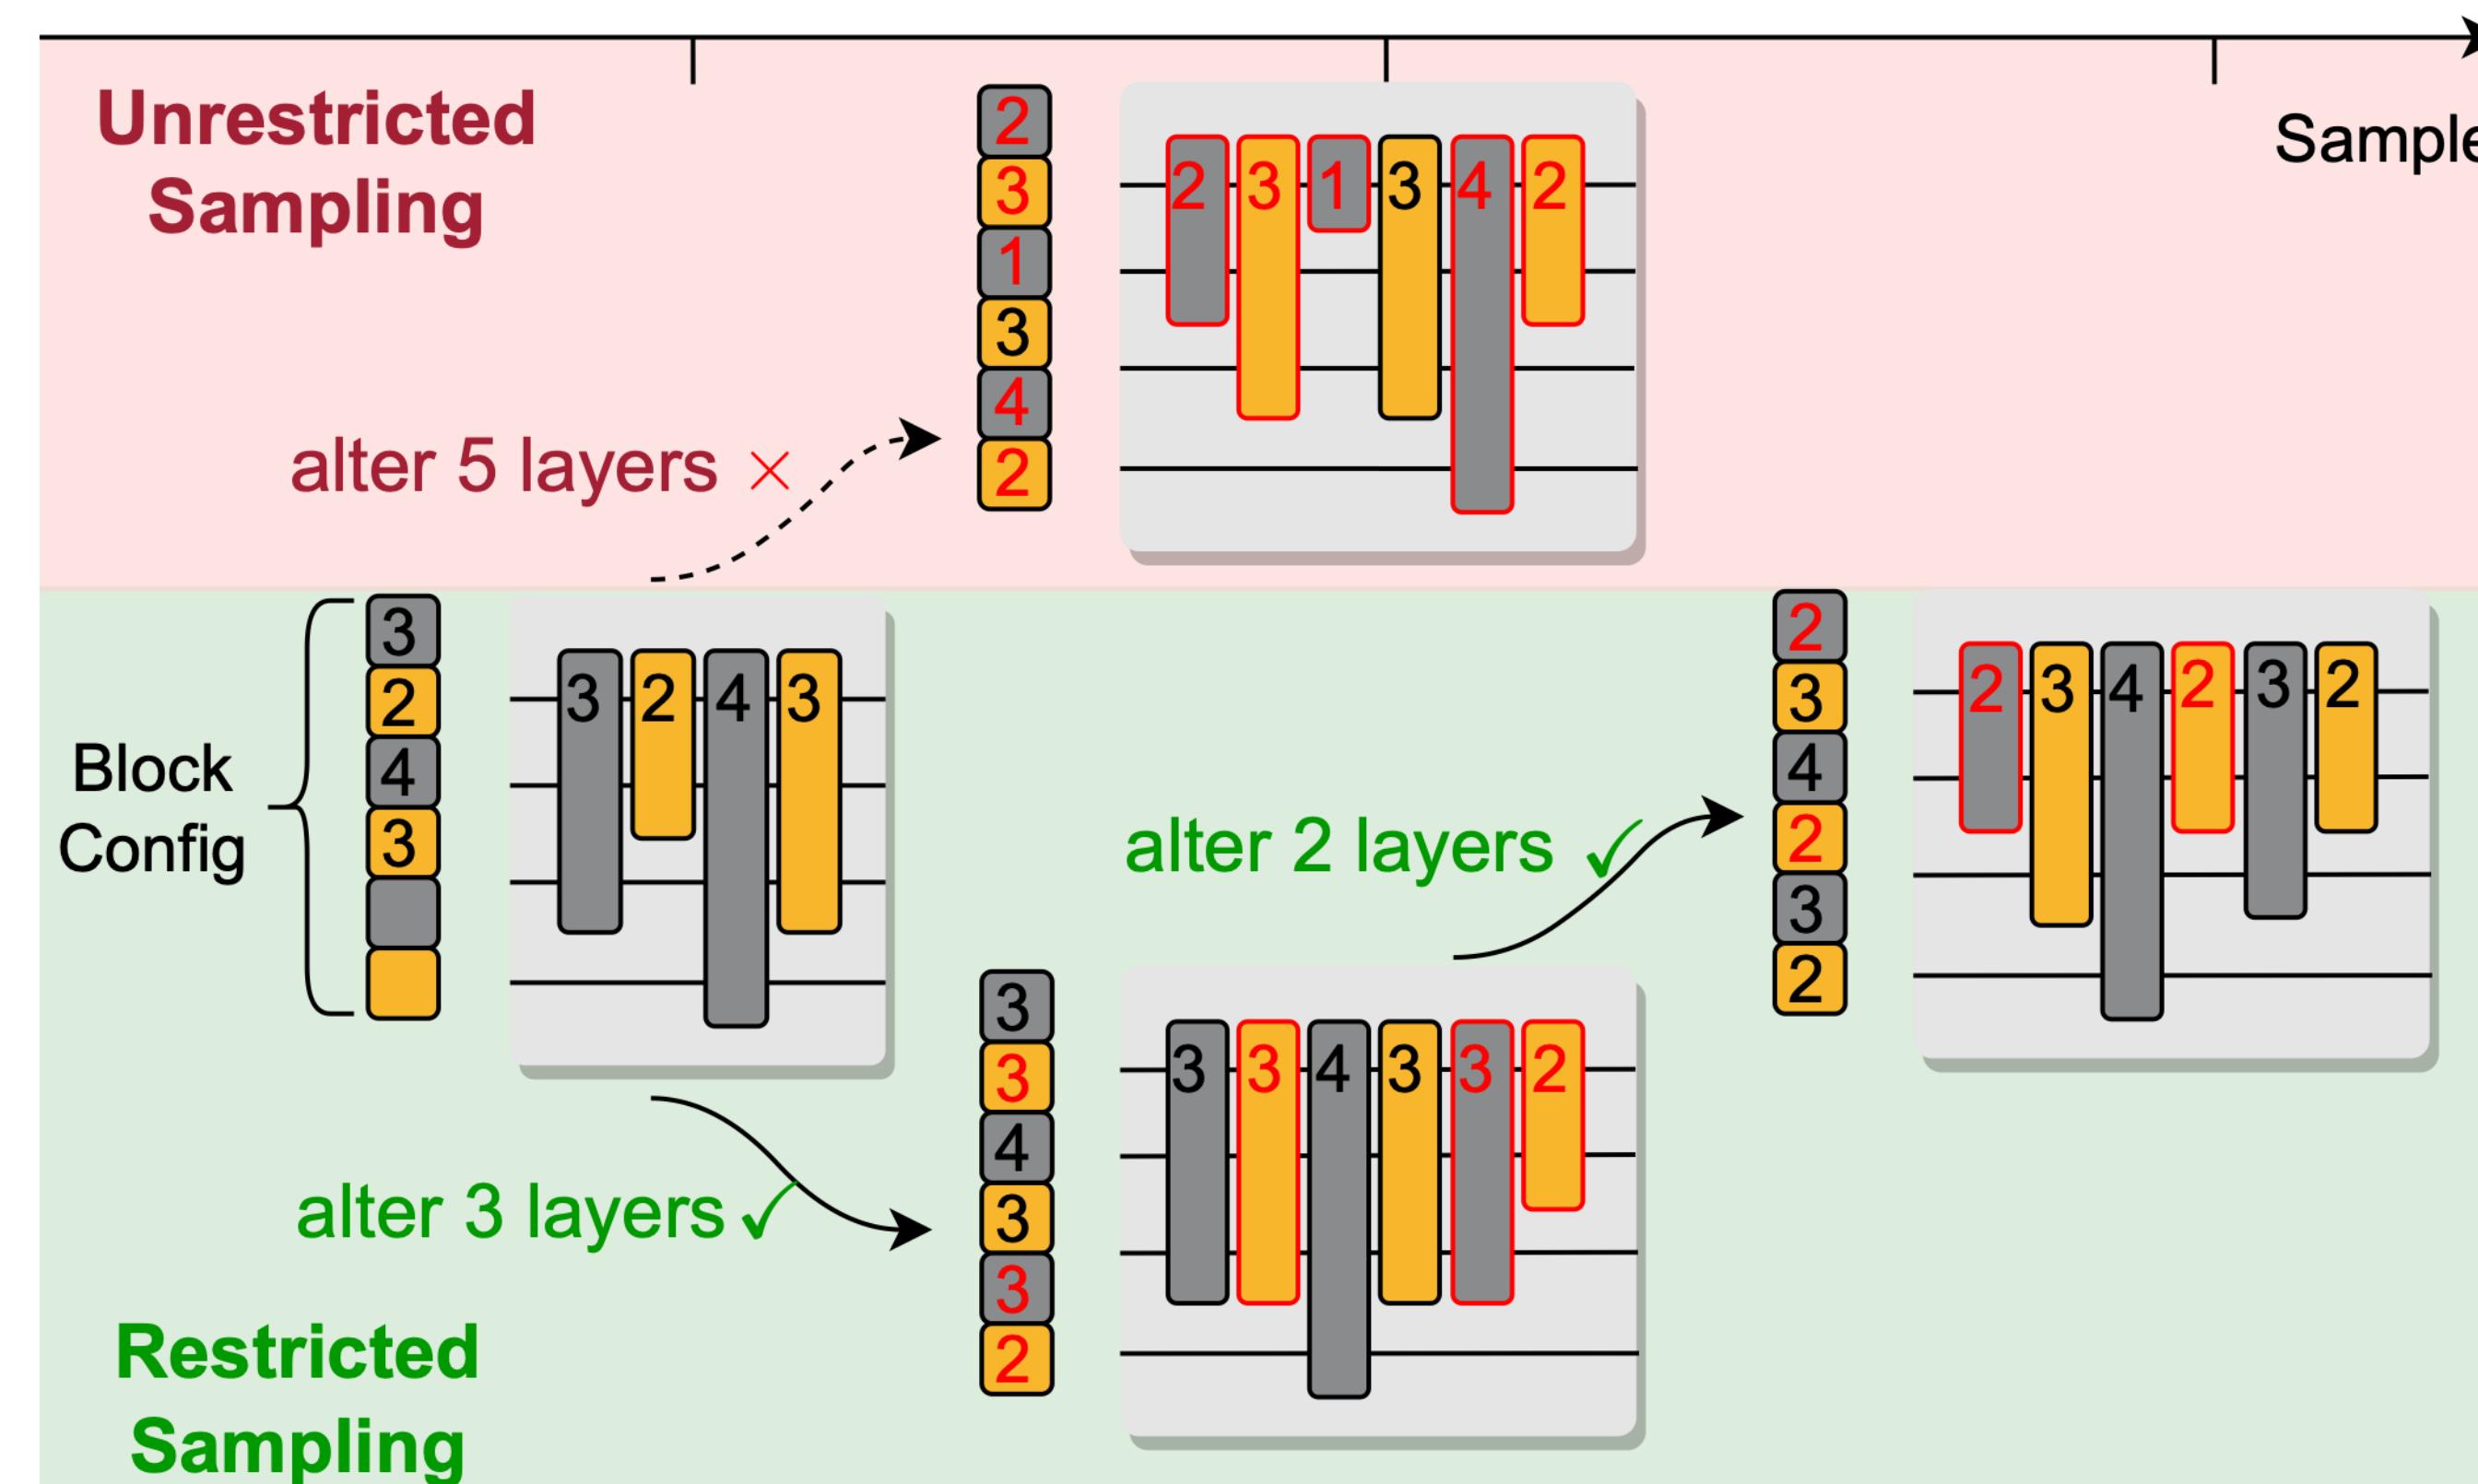
Restricted Sampling

- Restricted Sampling:
 - Restrict the difference between SubCircuits of two consecutive steps
 - For example: restrict to at most 4 different layers



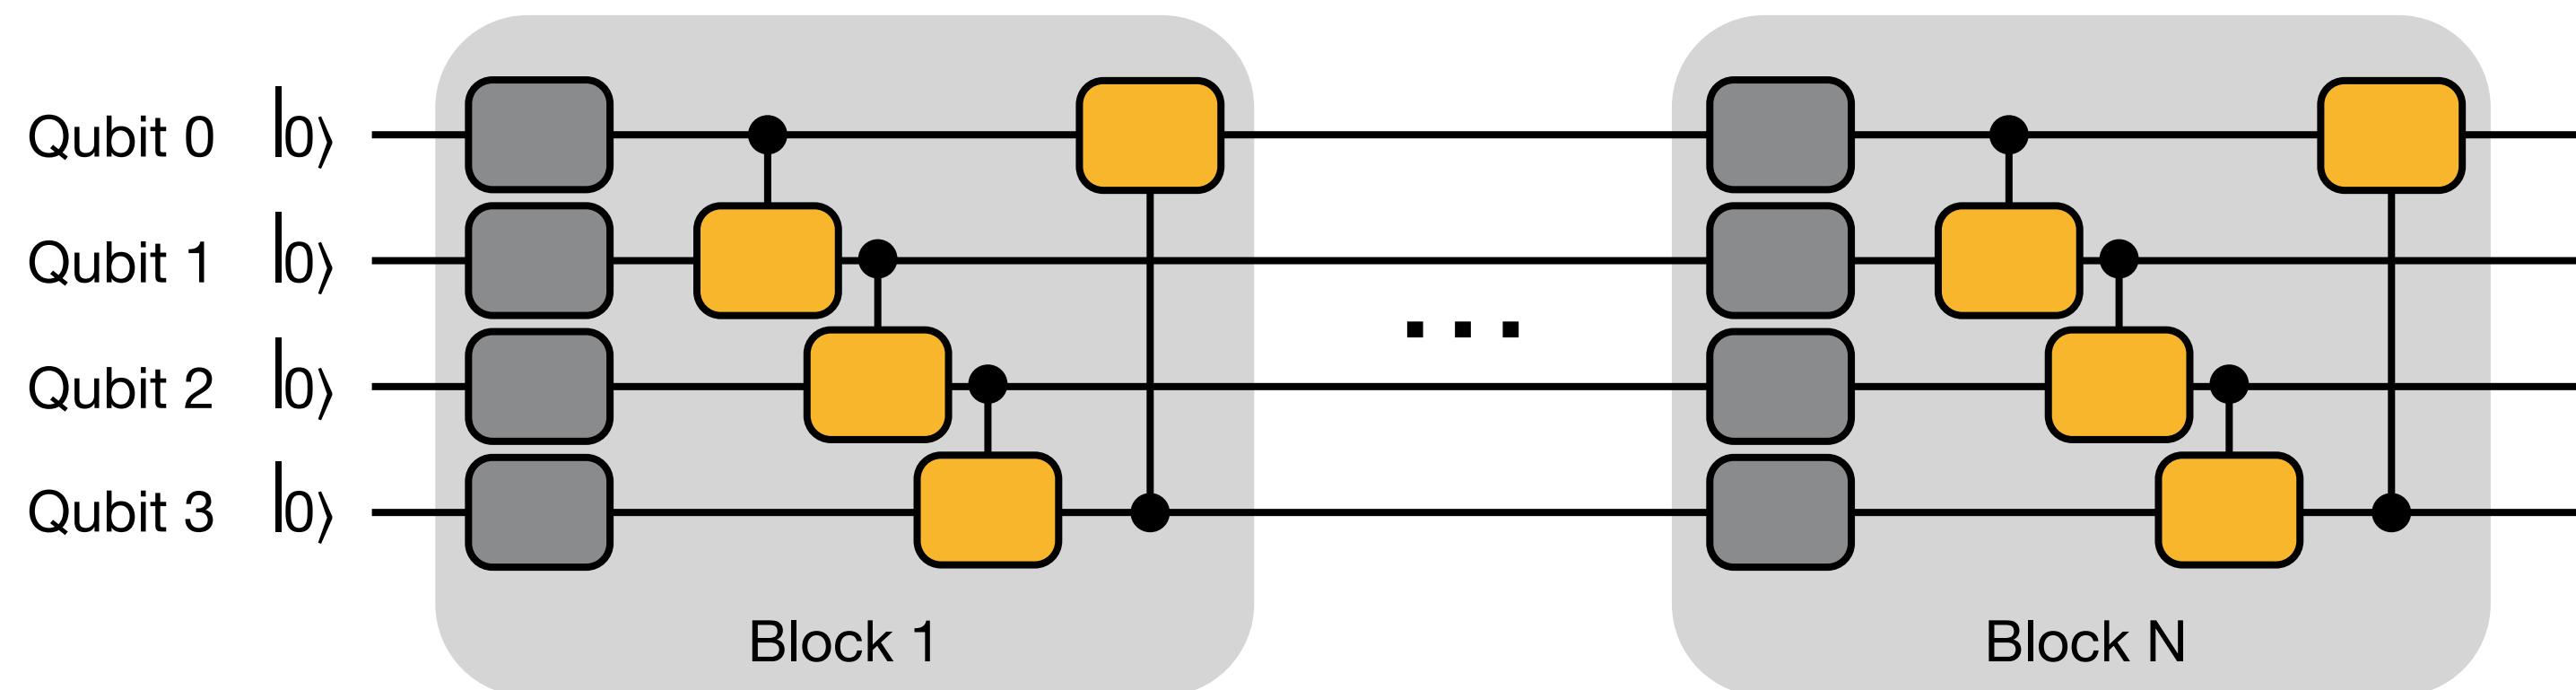
Restricted Sampling

- Restricted Sampling:
 - Restrict the difference between SubCircuits of two consecutive steps
 - For example: restrict to at most 4 different layers



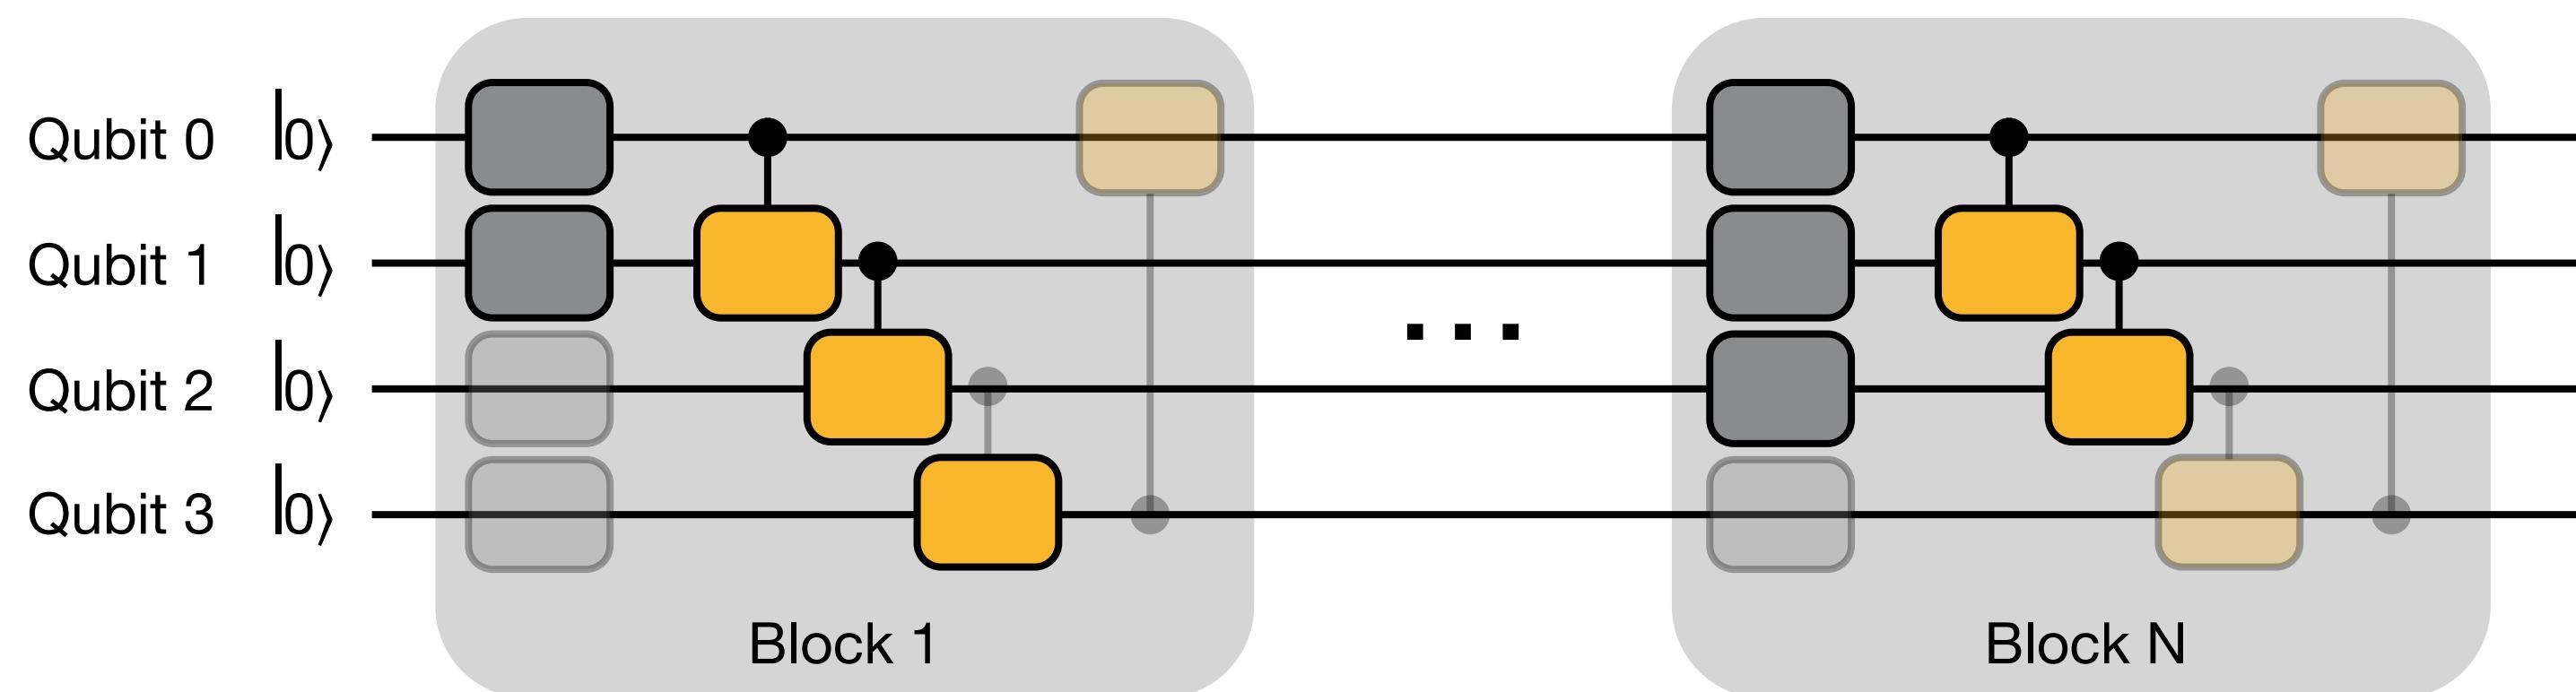
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



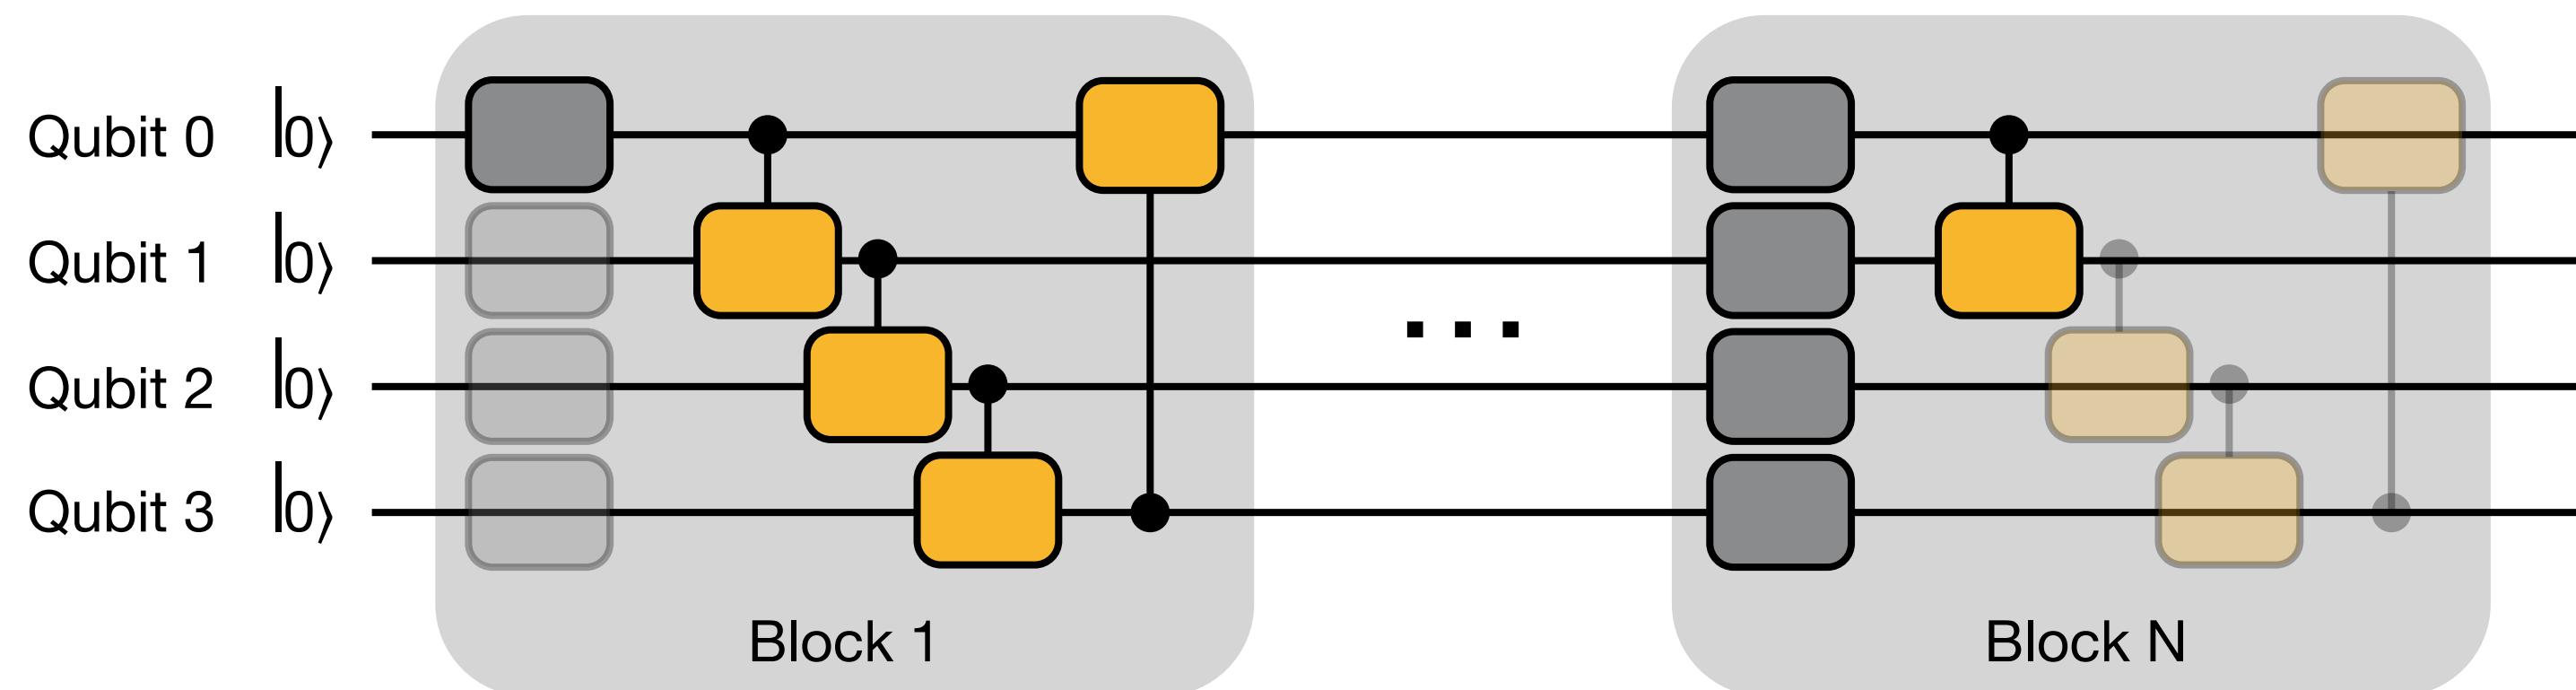
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



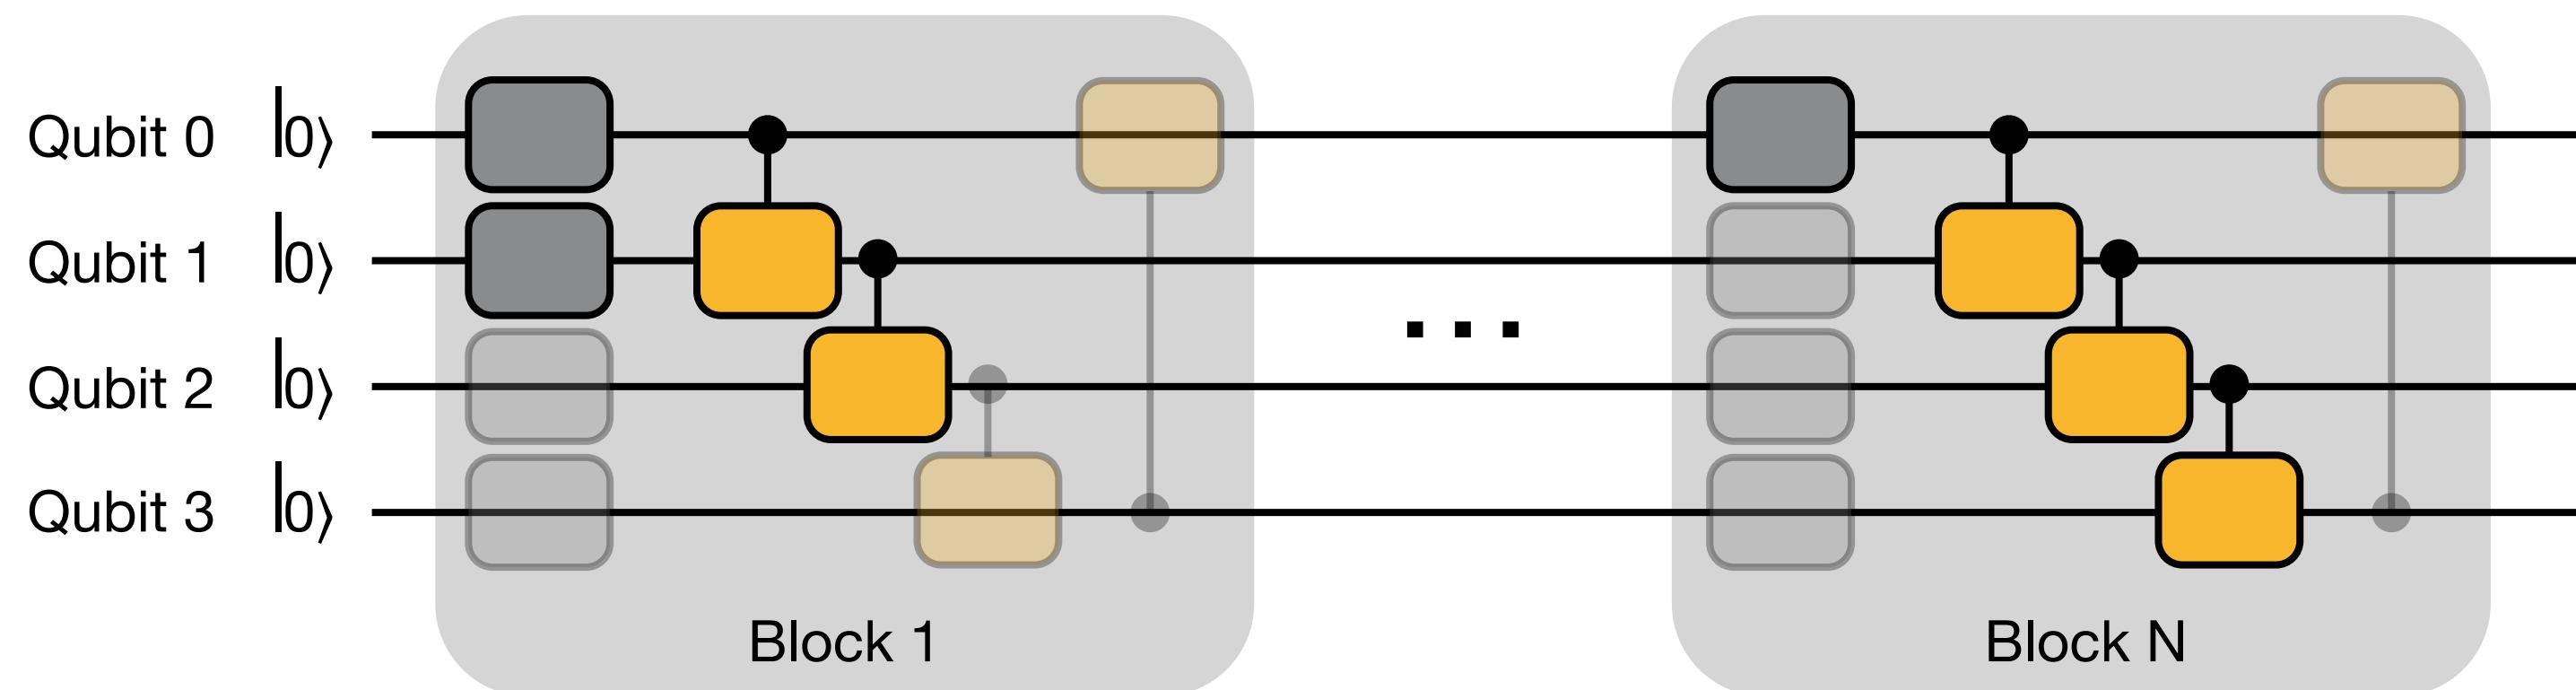
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



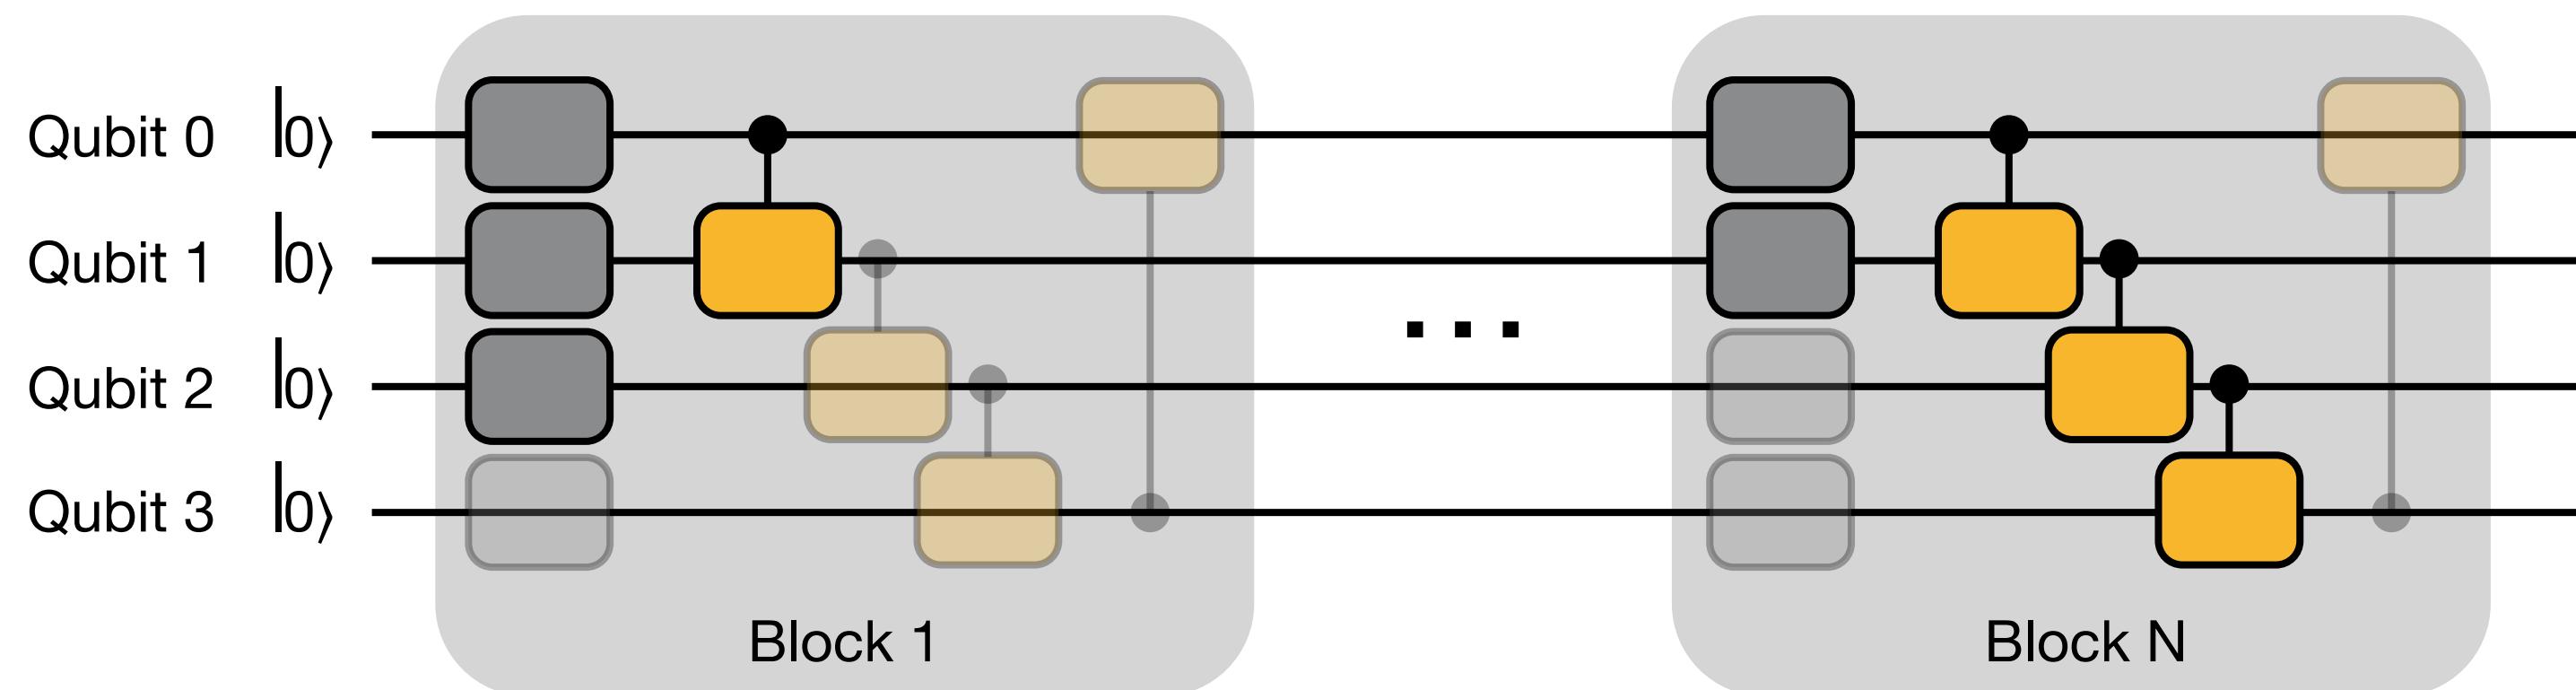
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



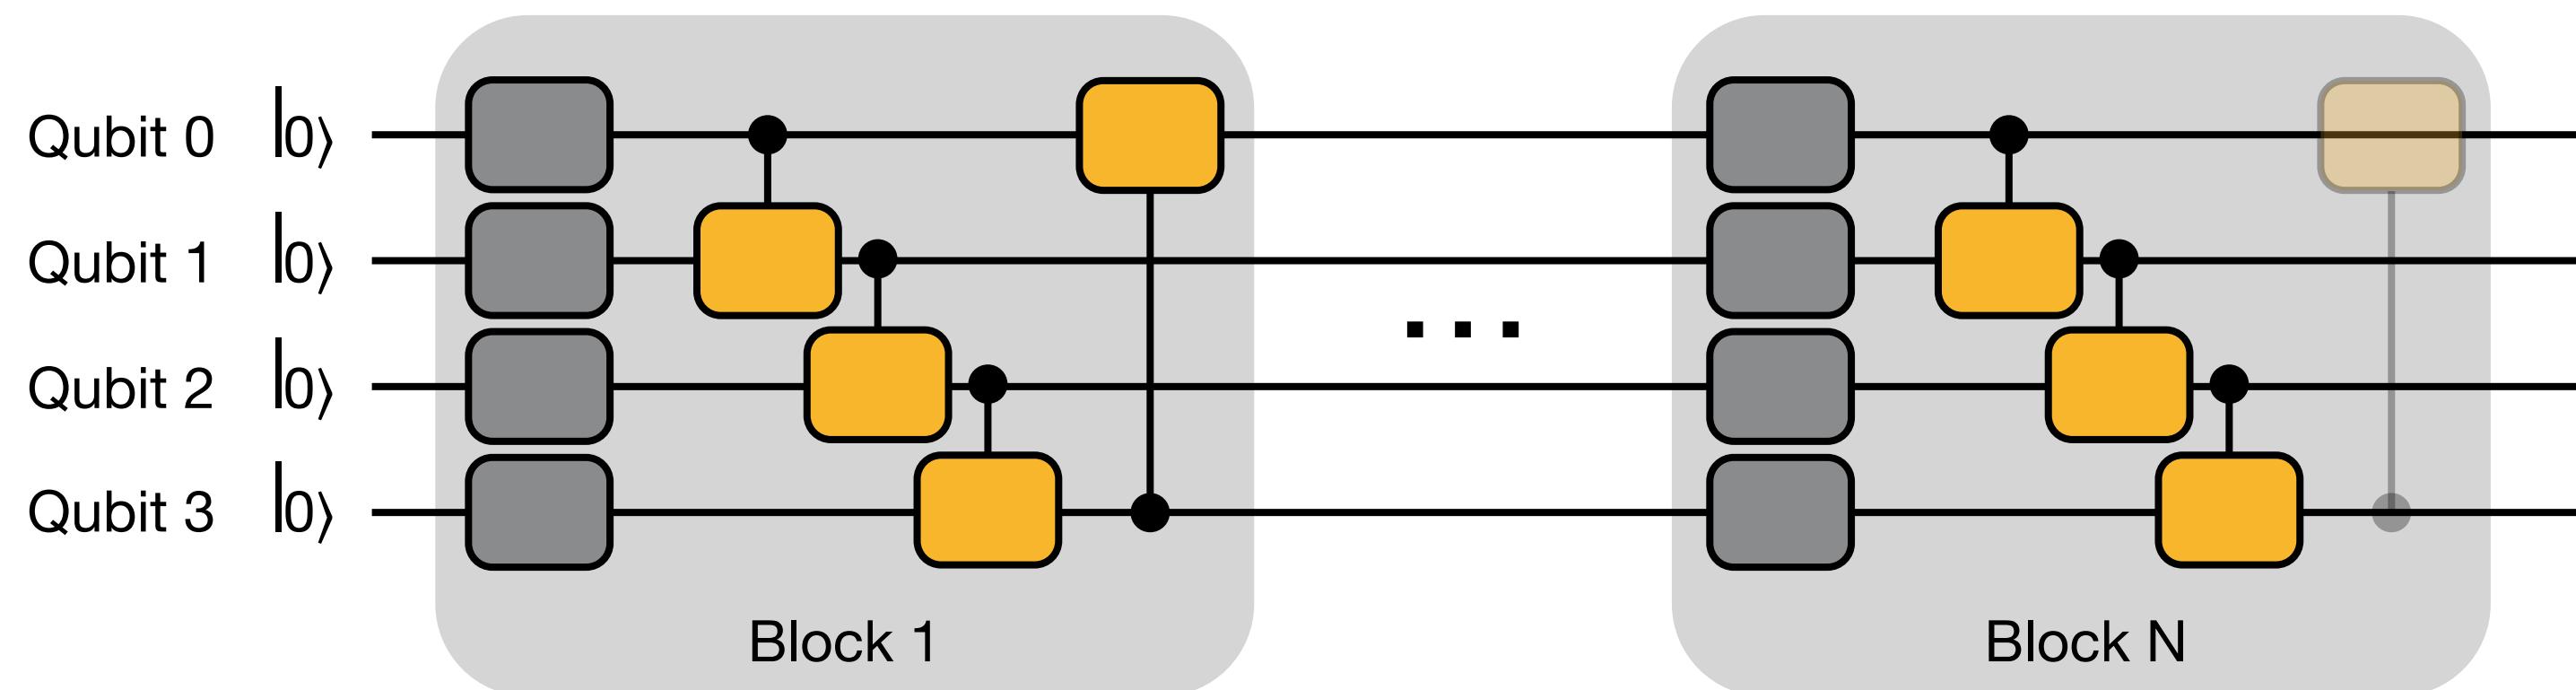
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



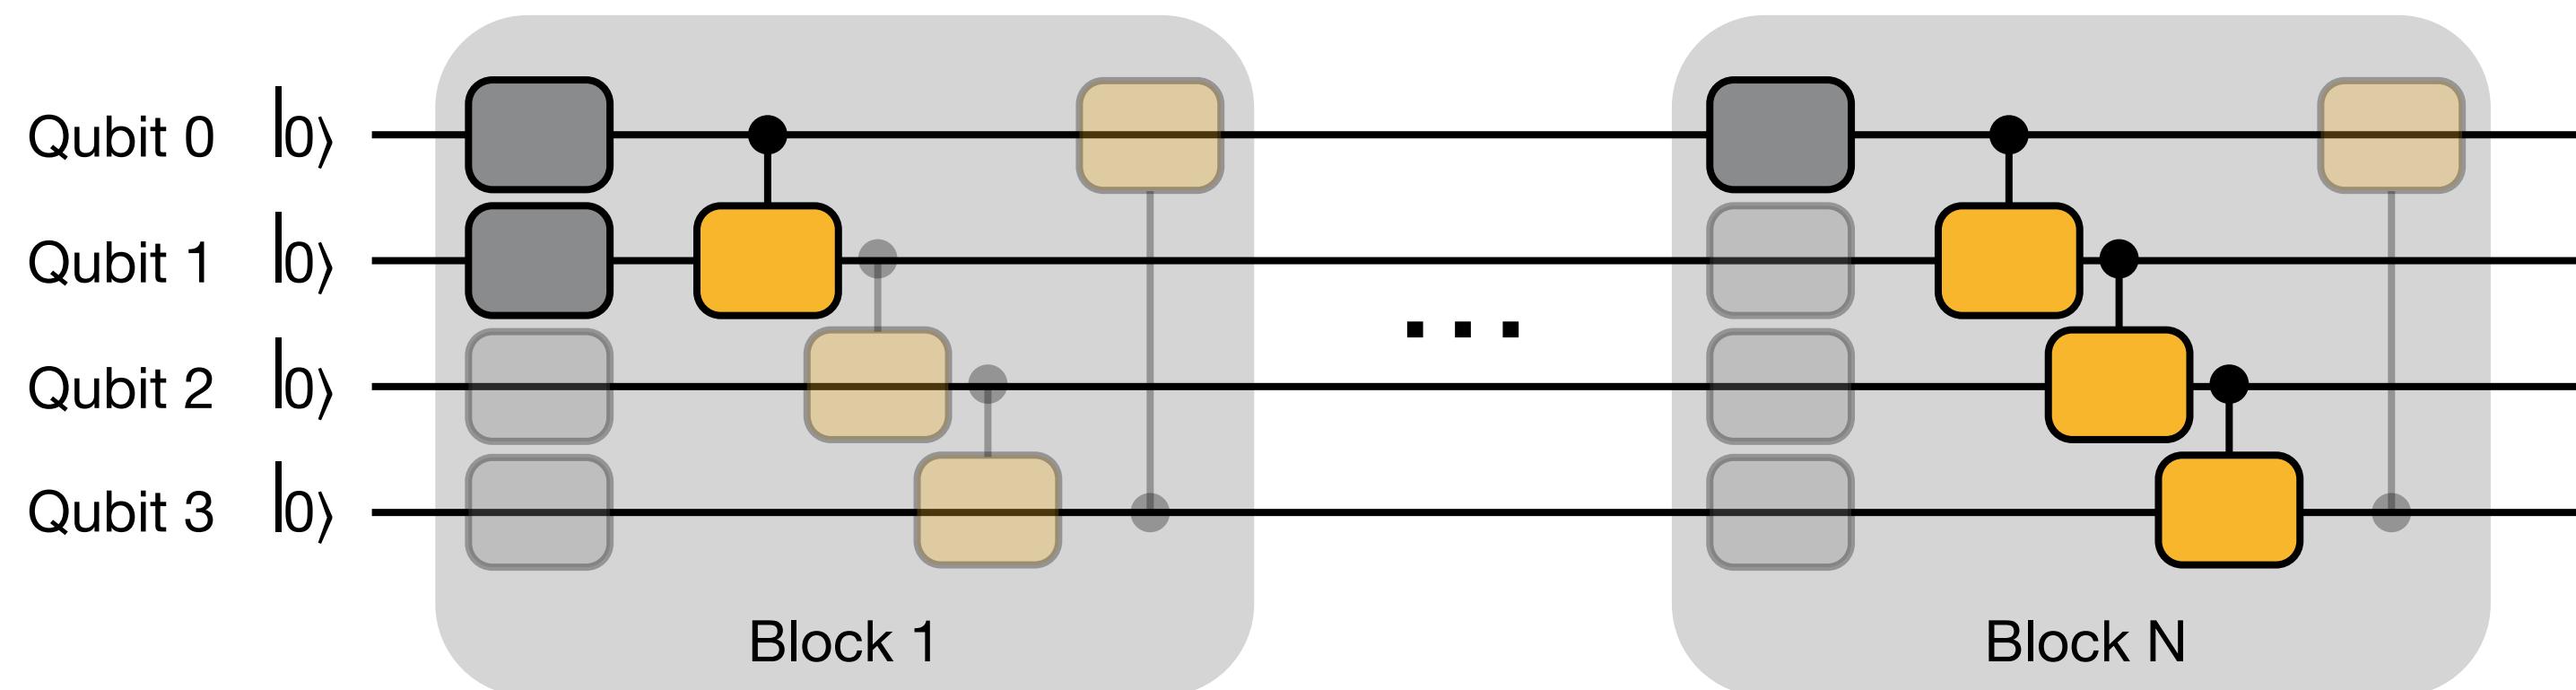
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



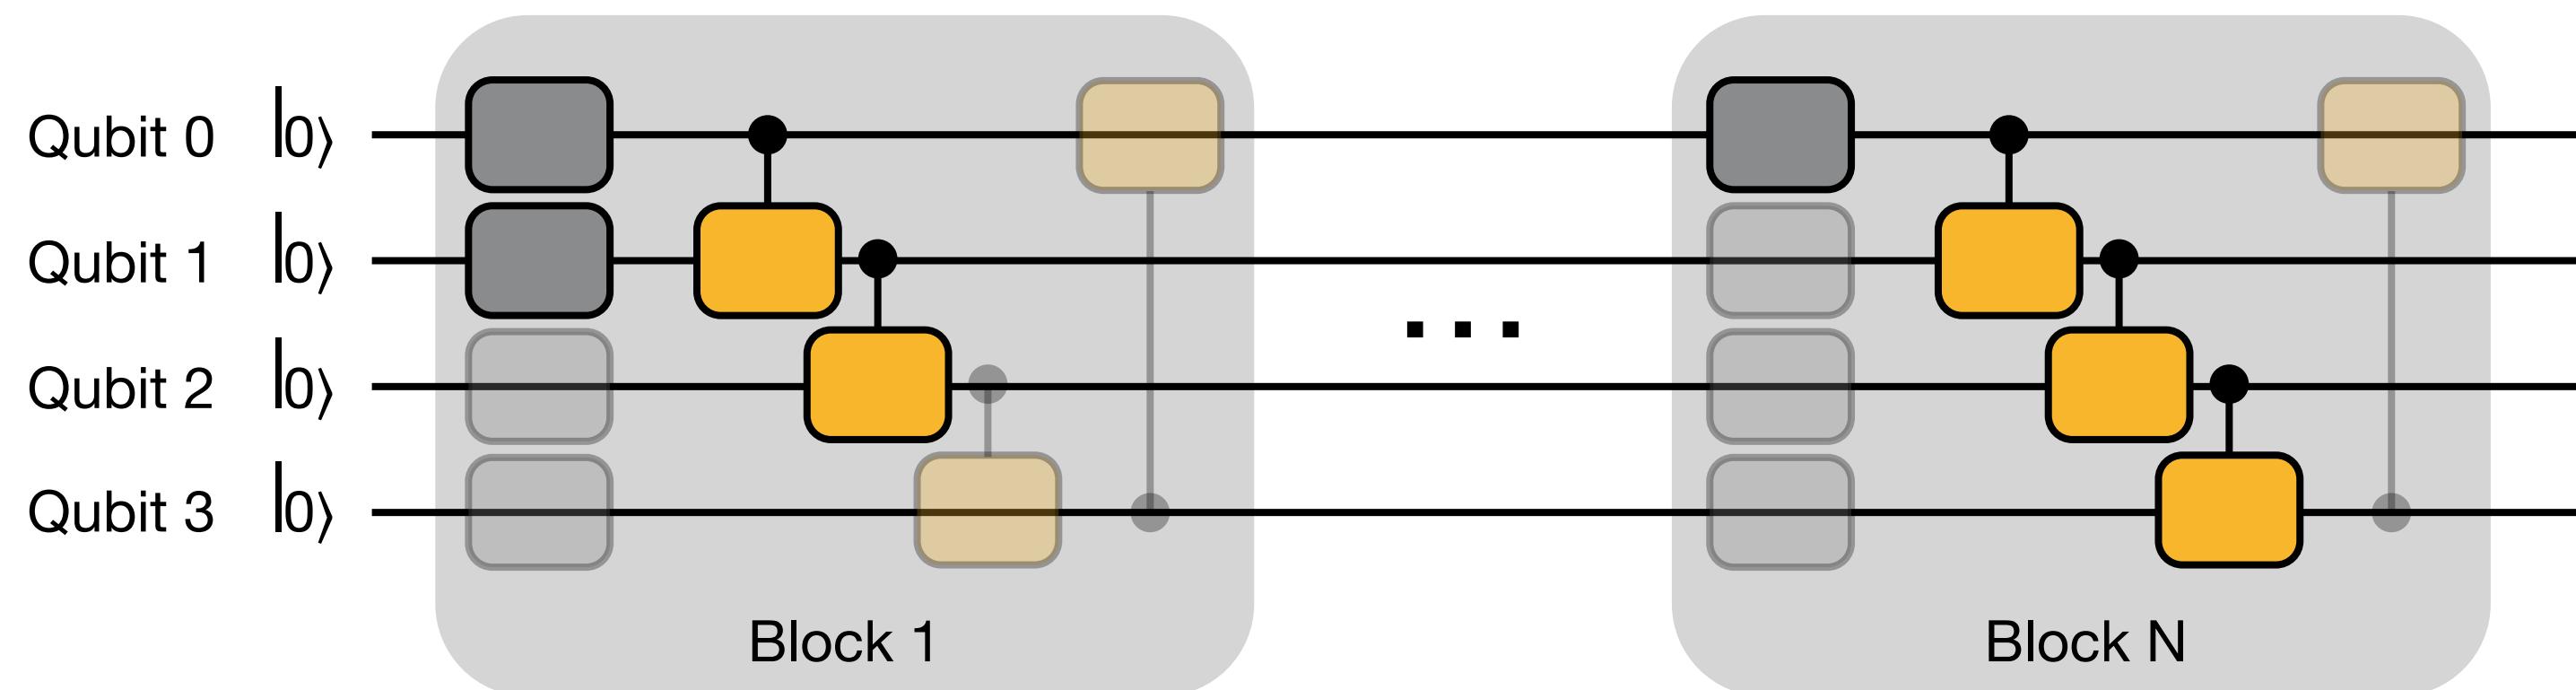
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



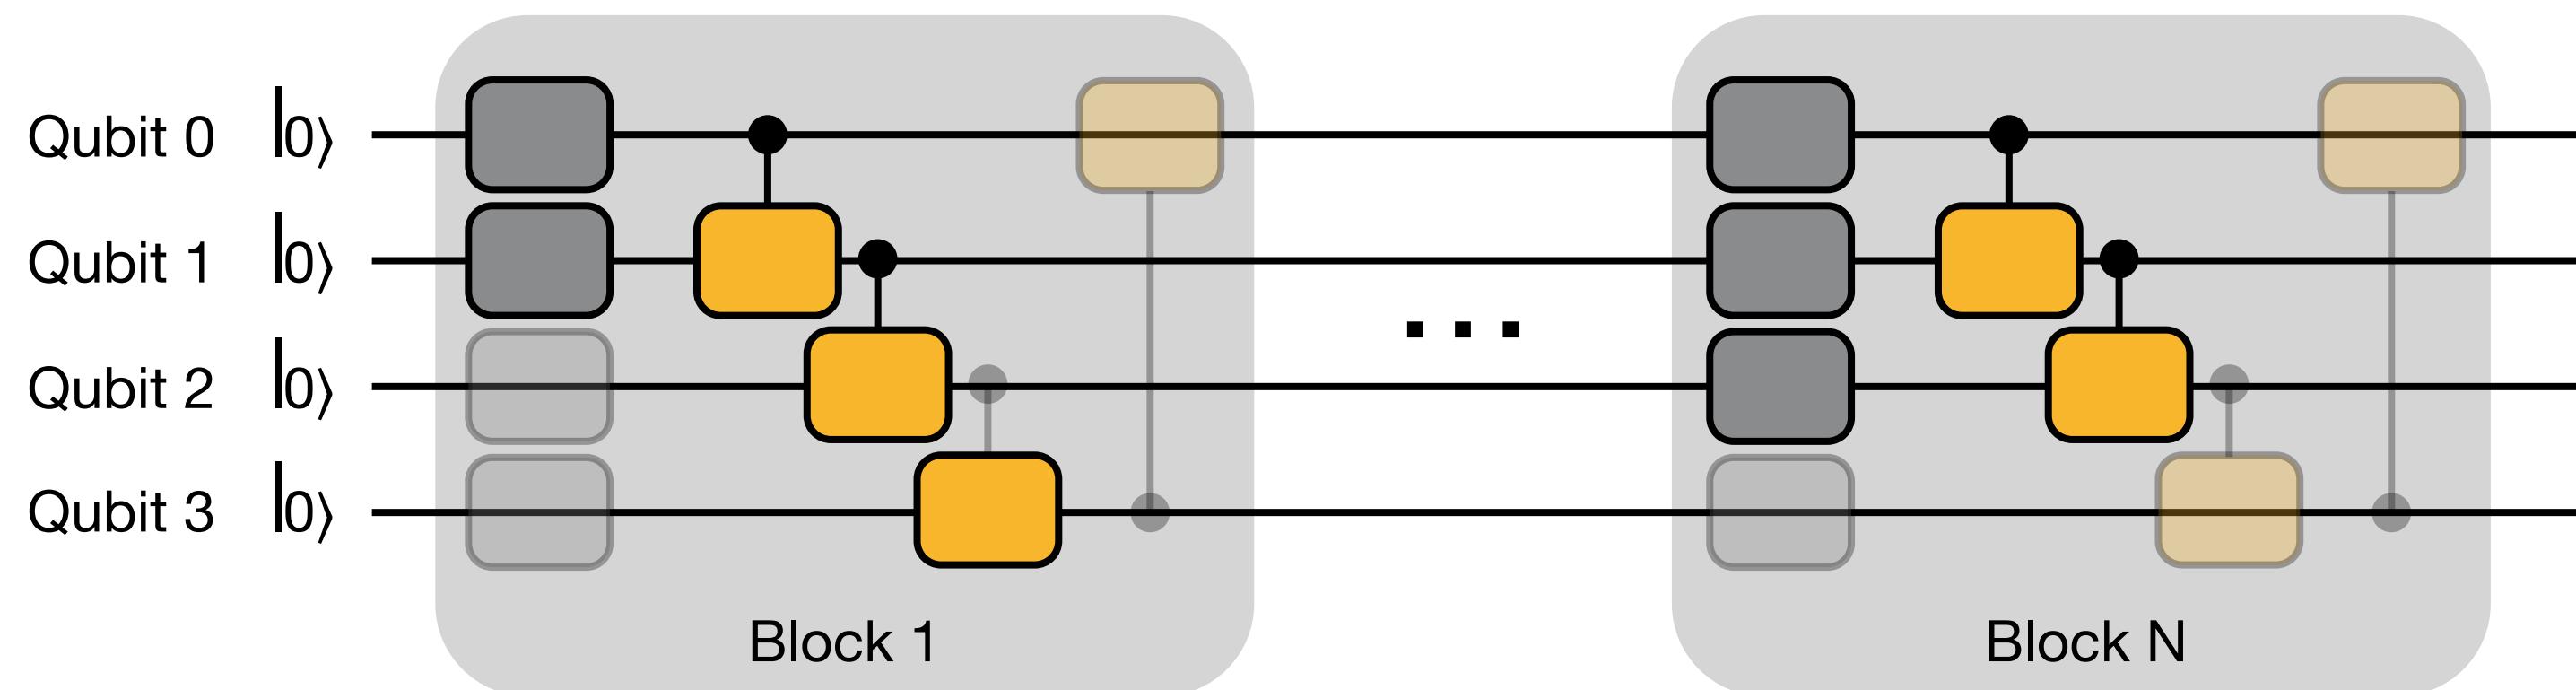
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



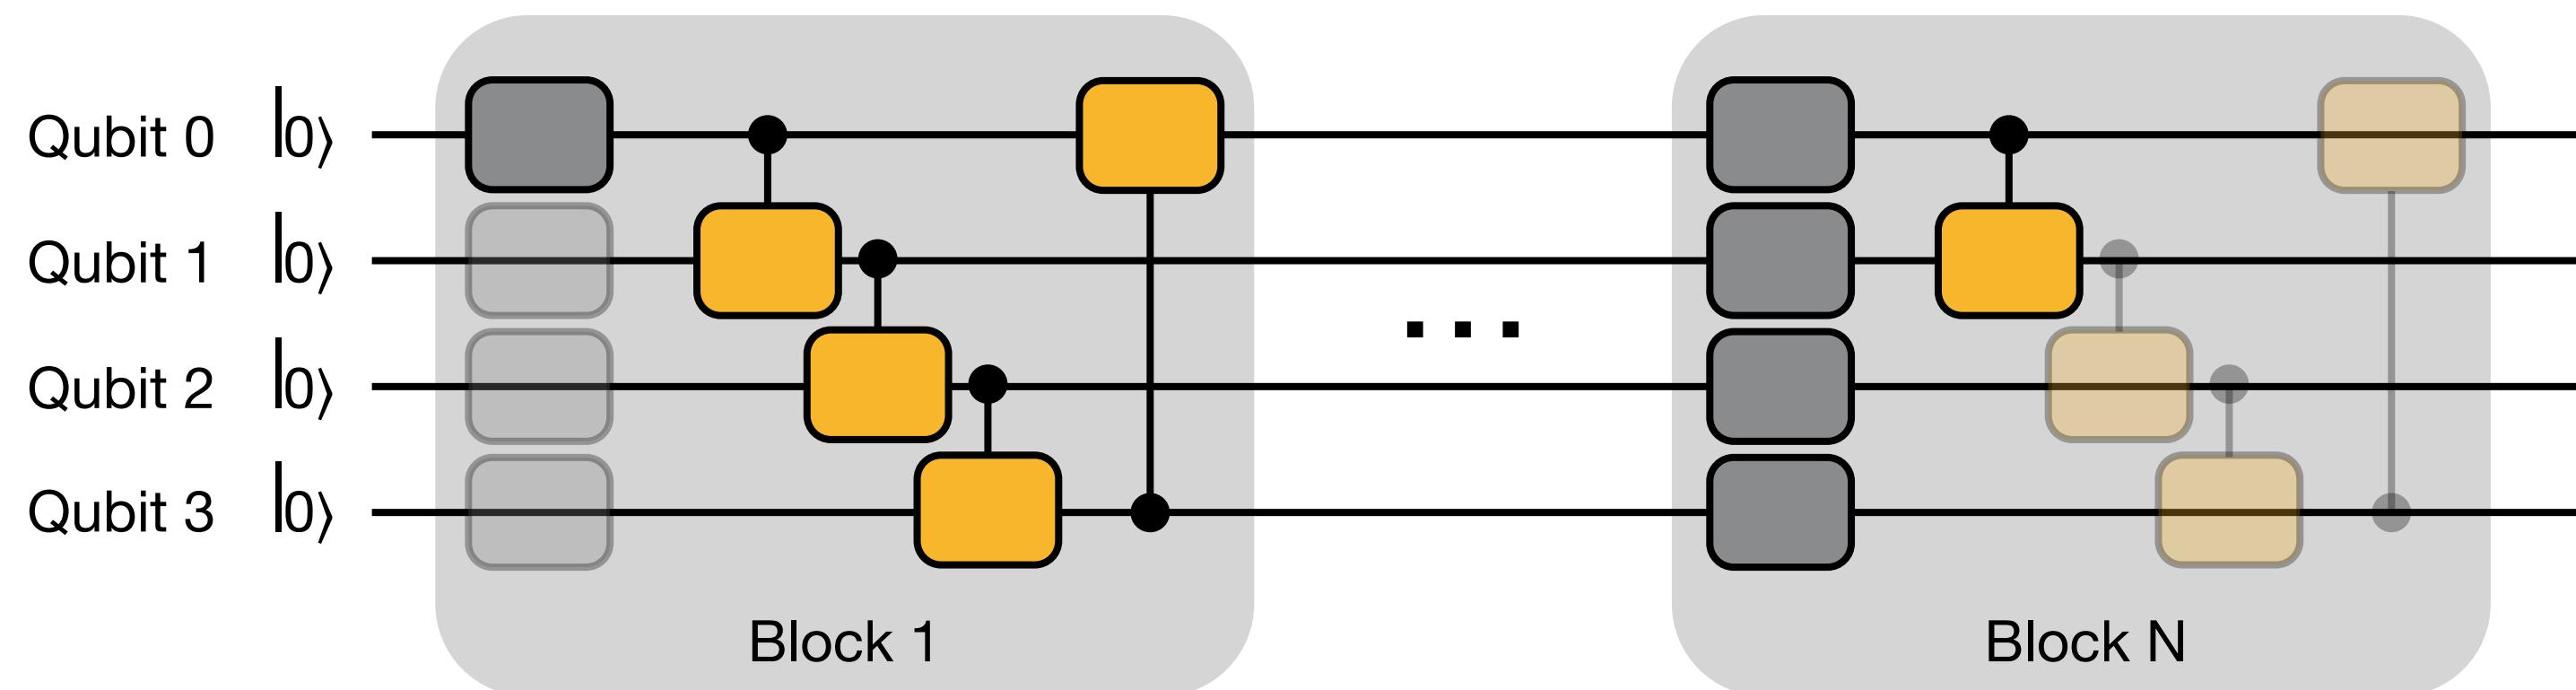
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



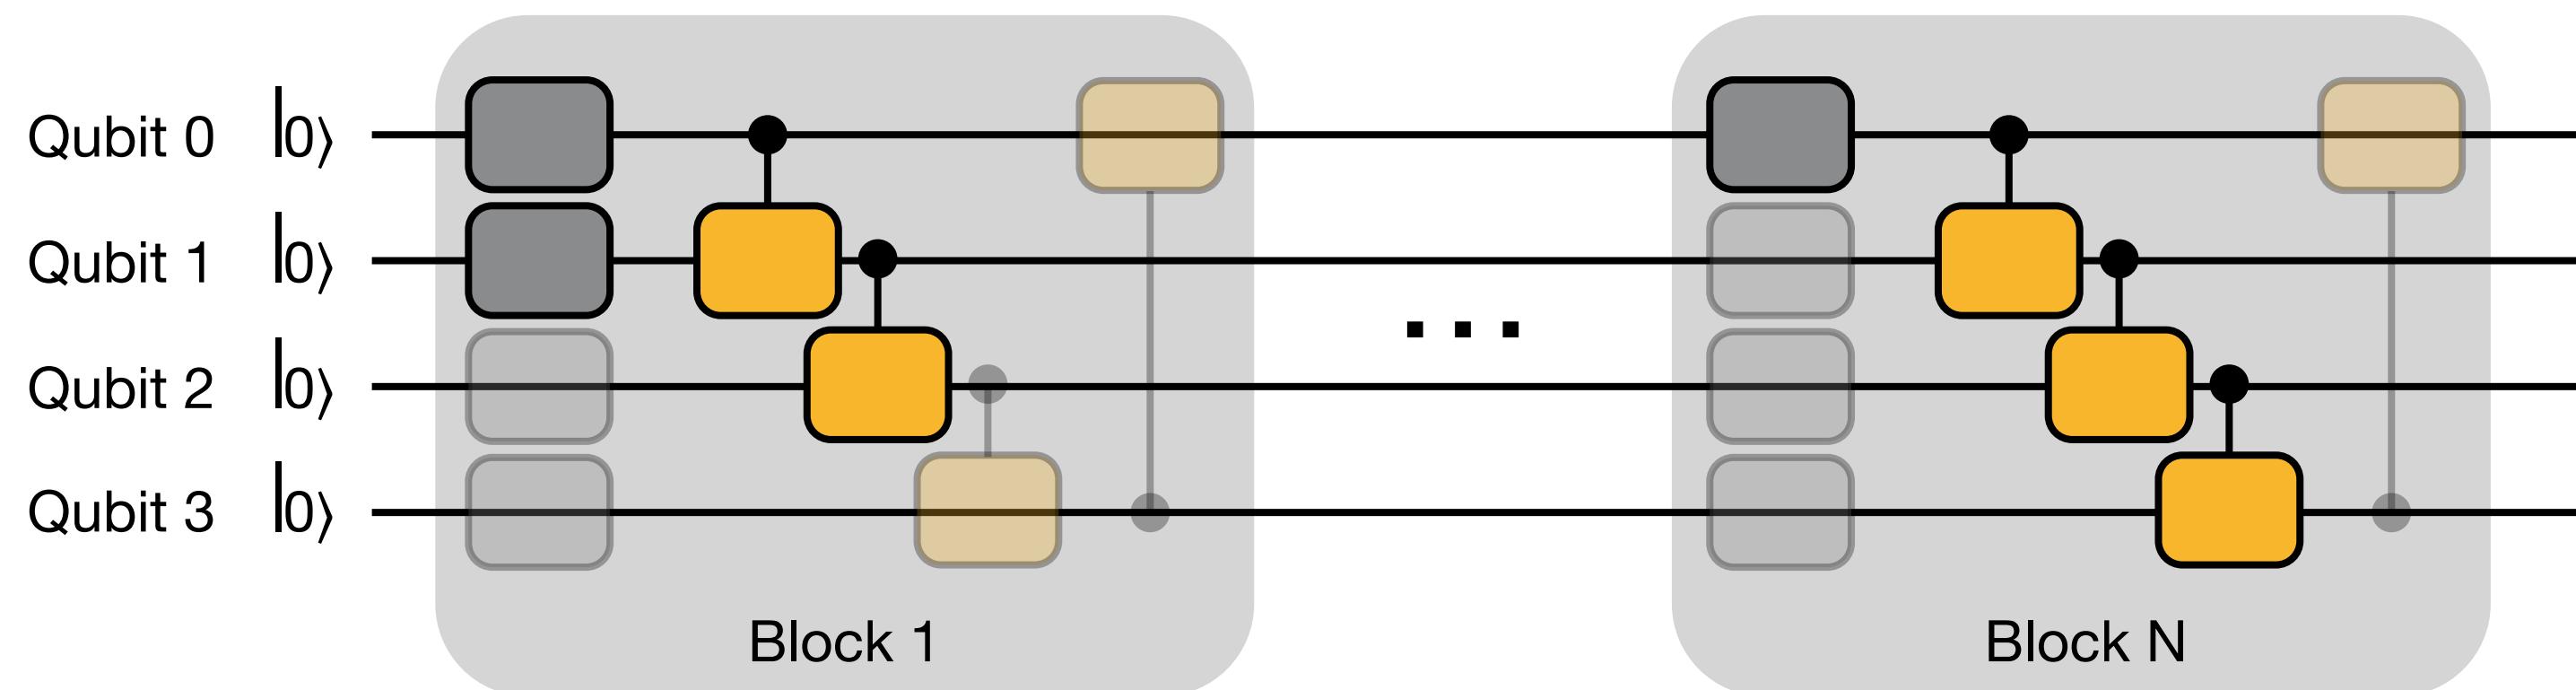
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



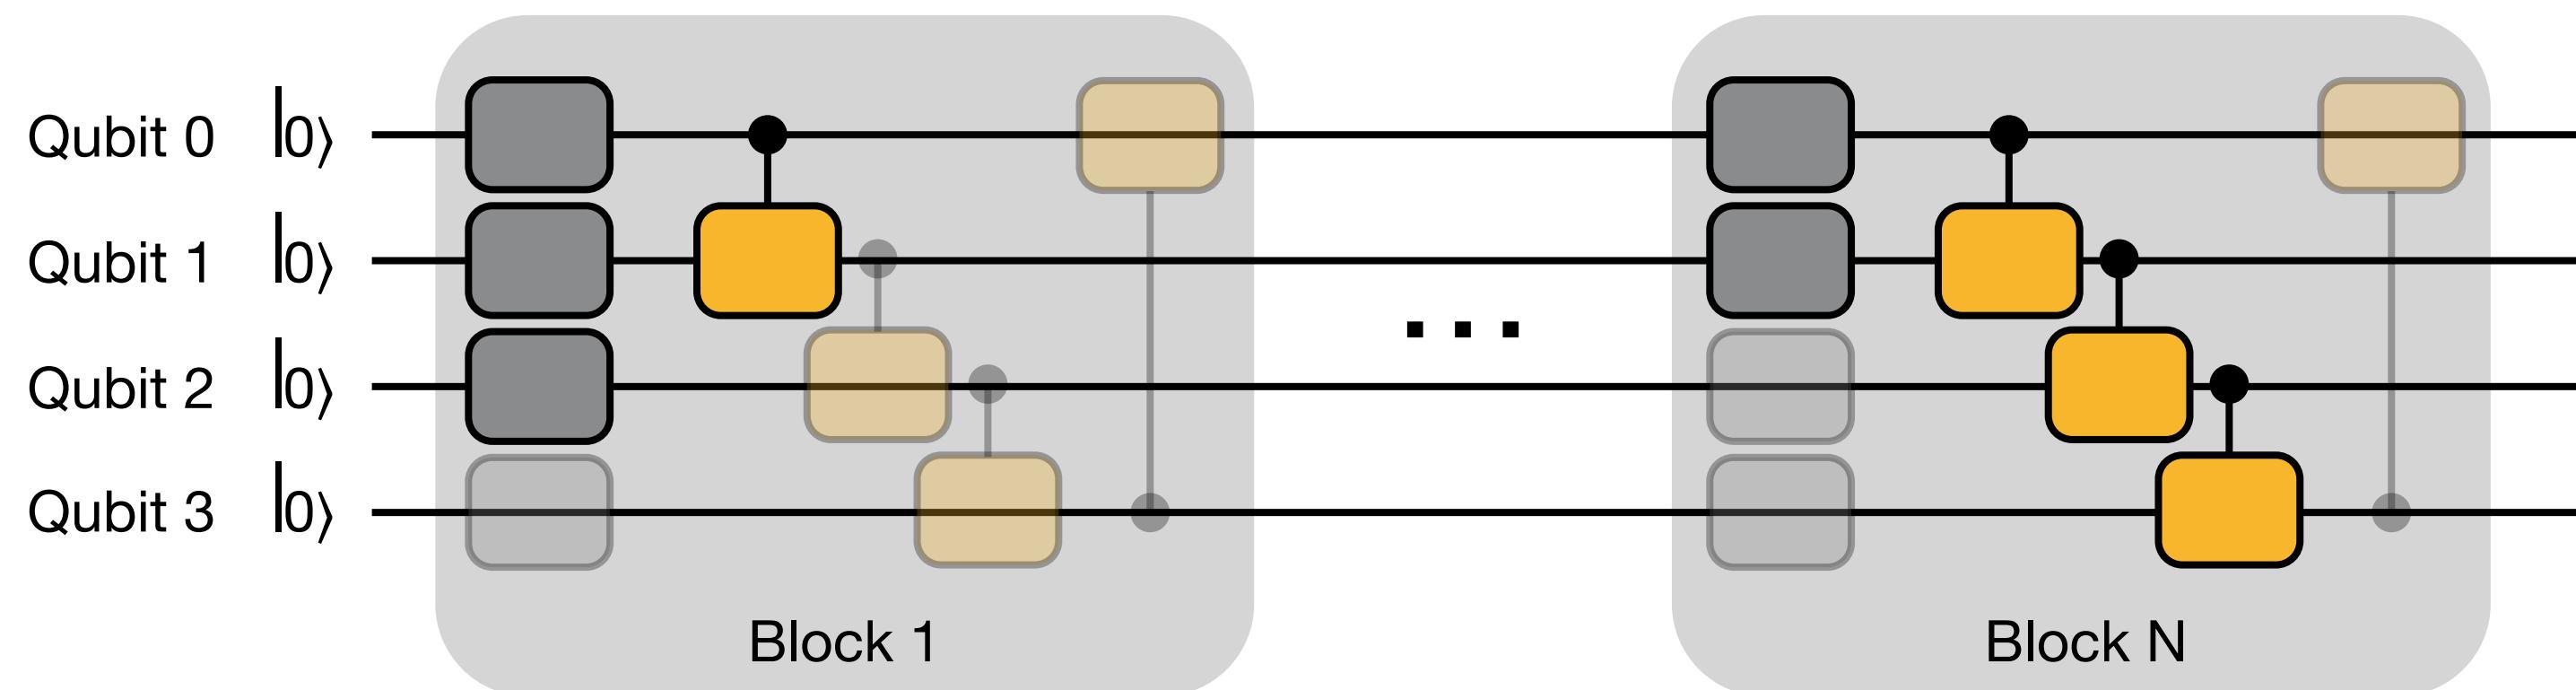
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



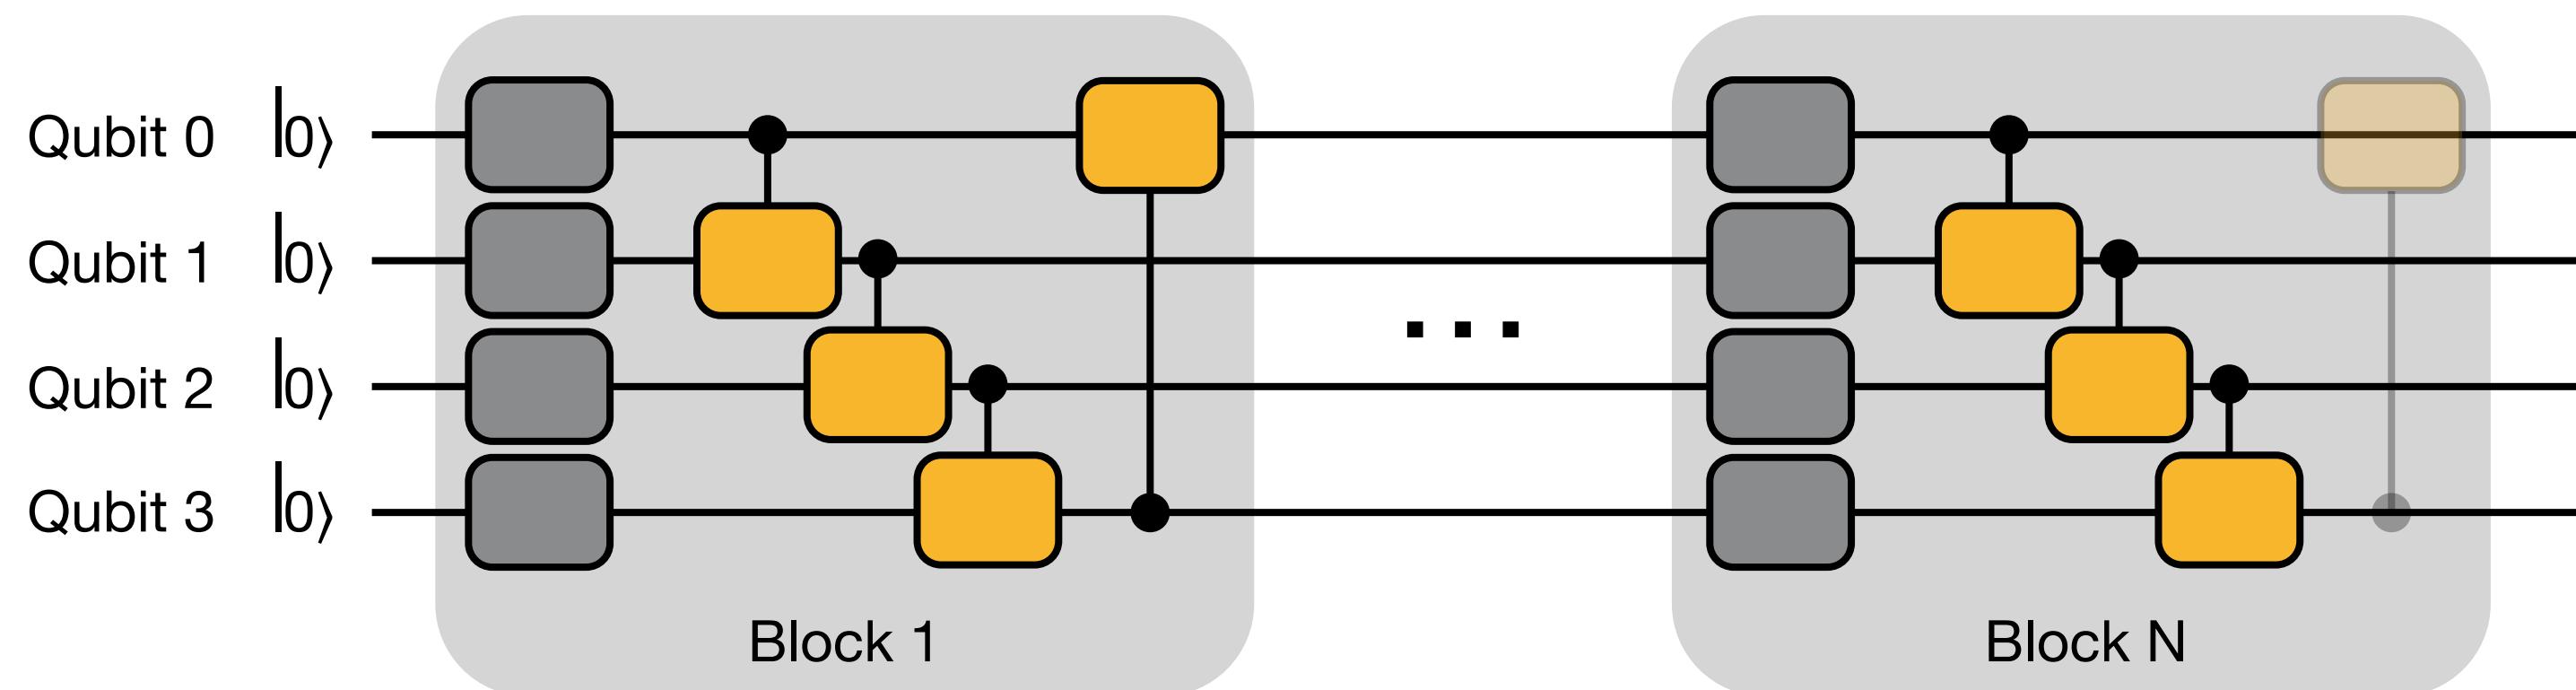
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



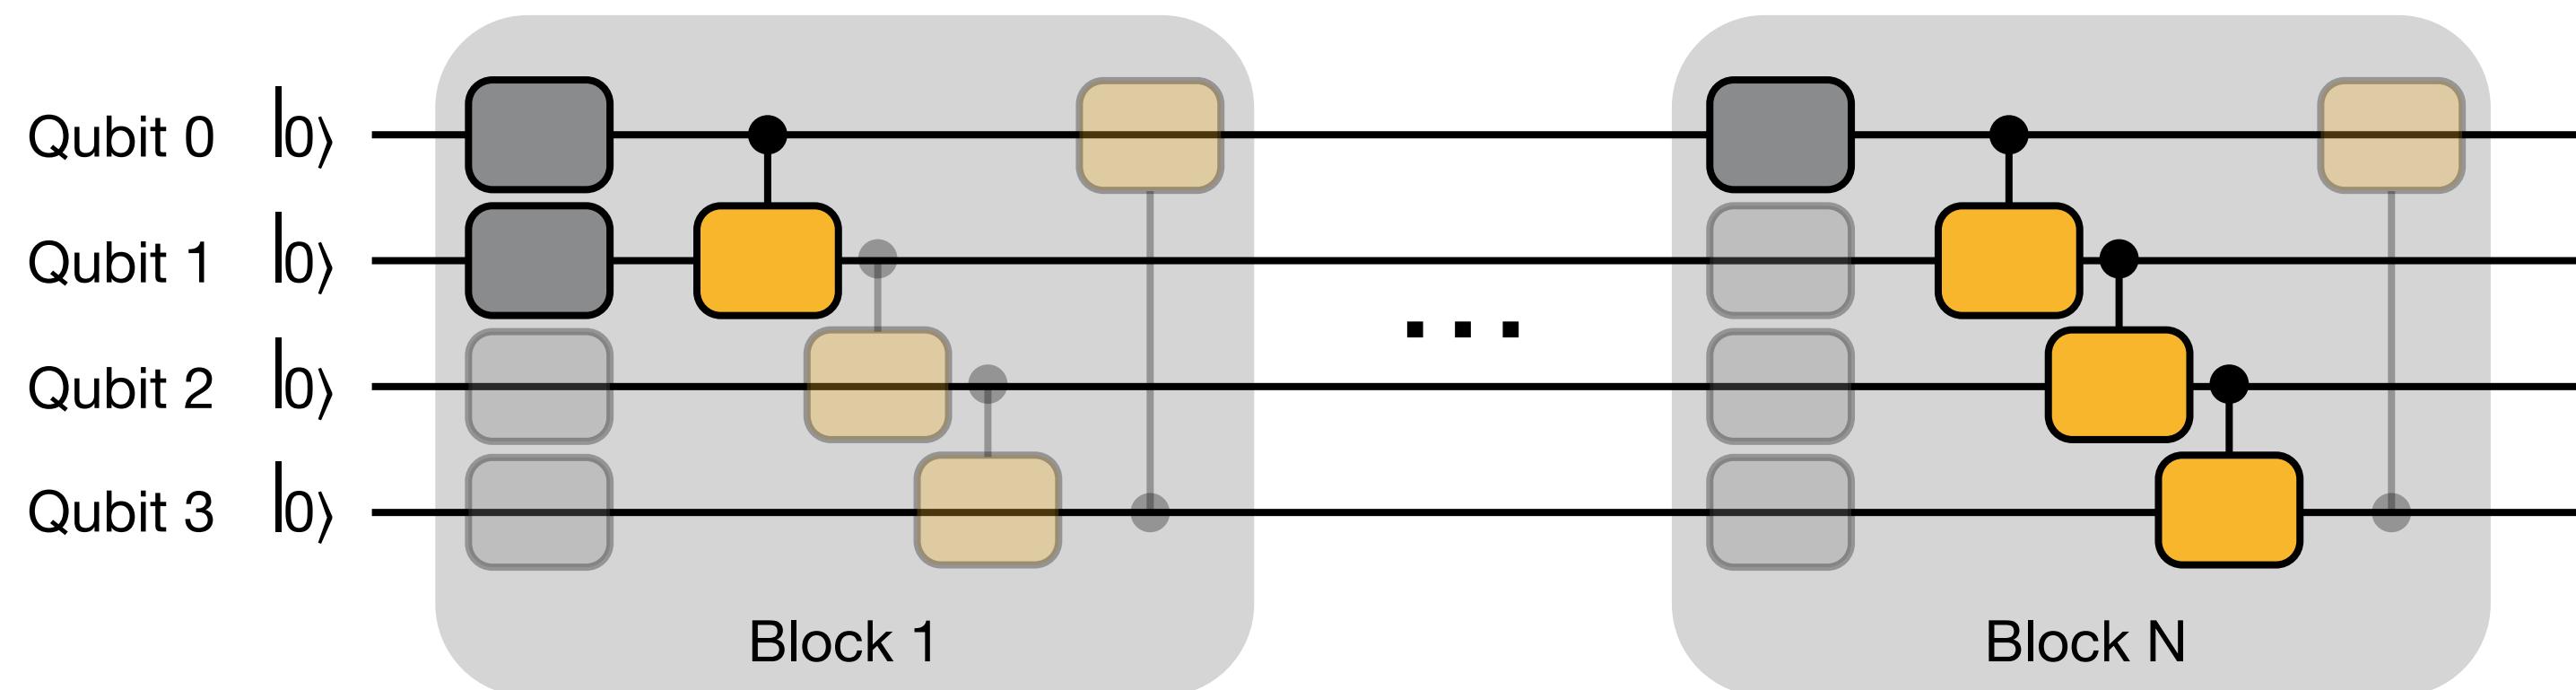
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train



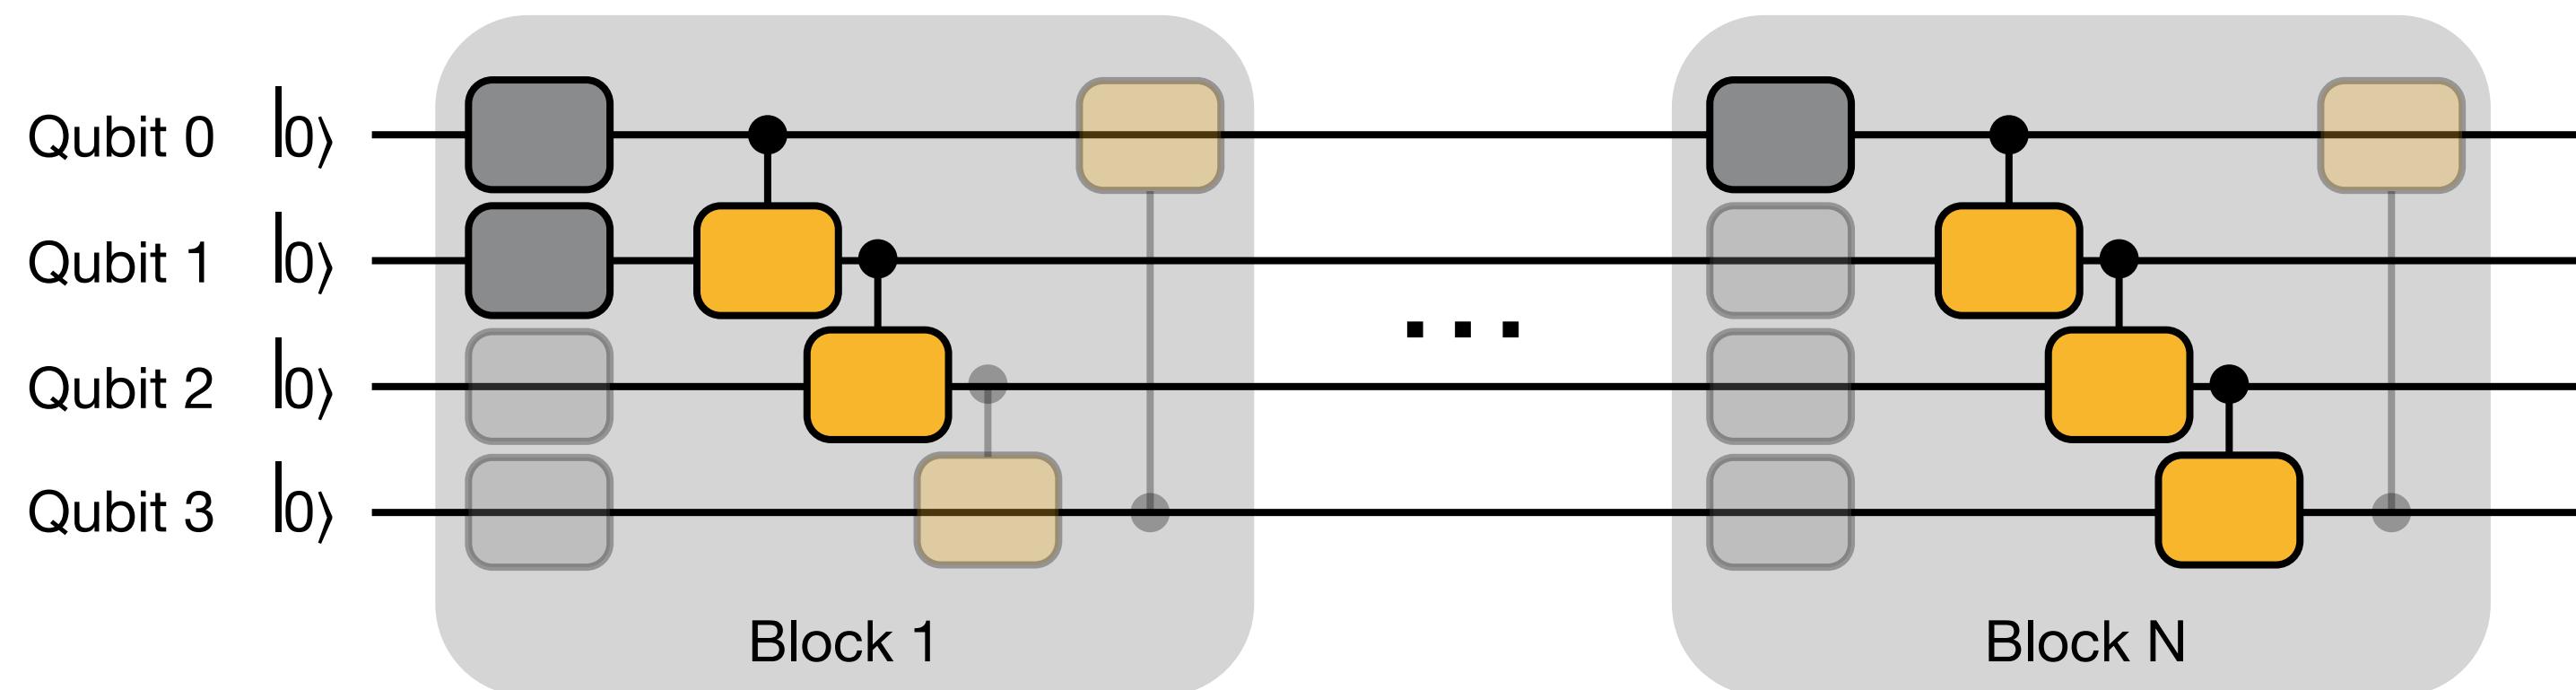
Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train

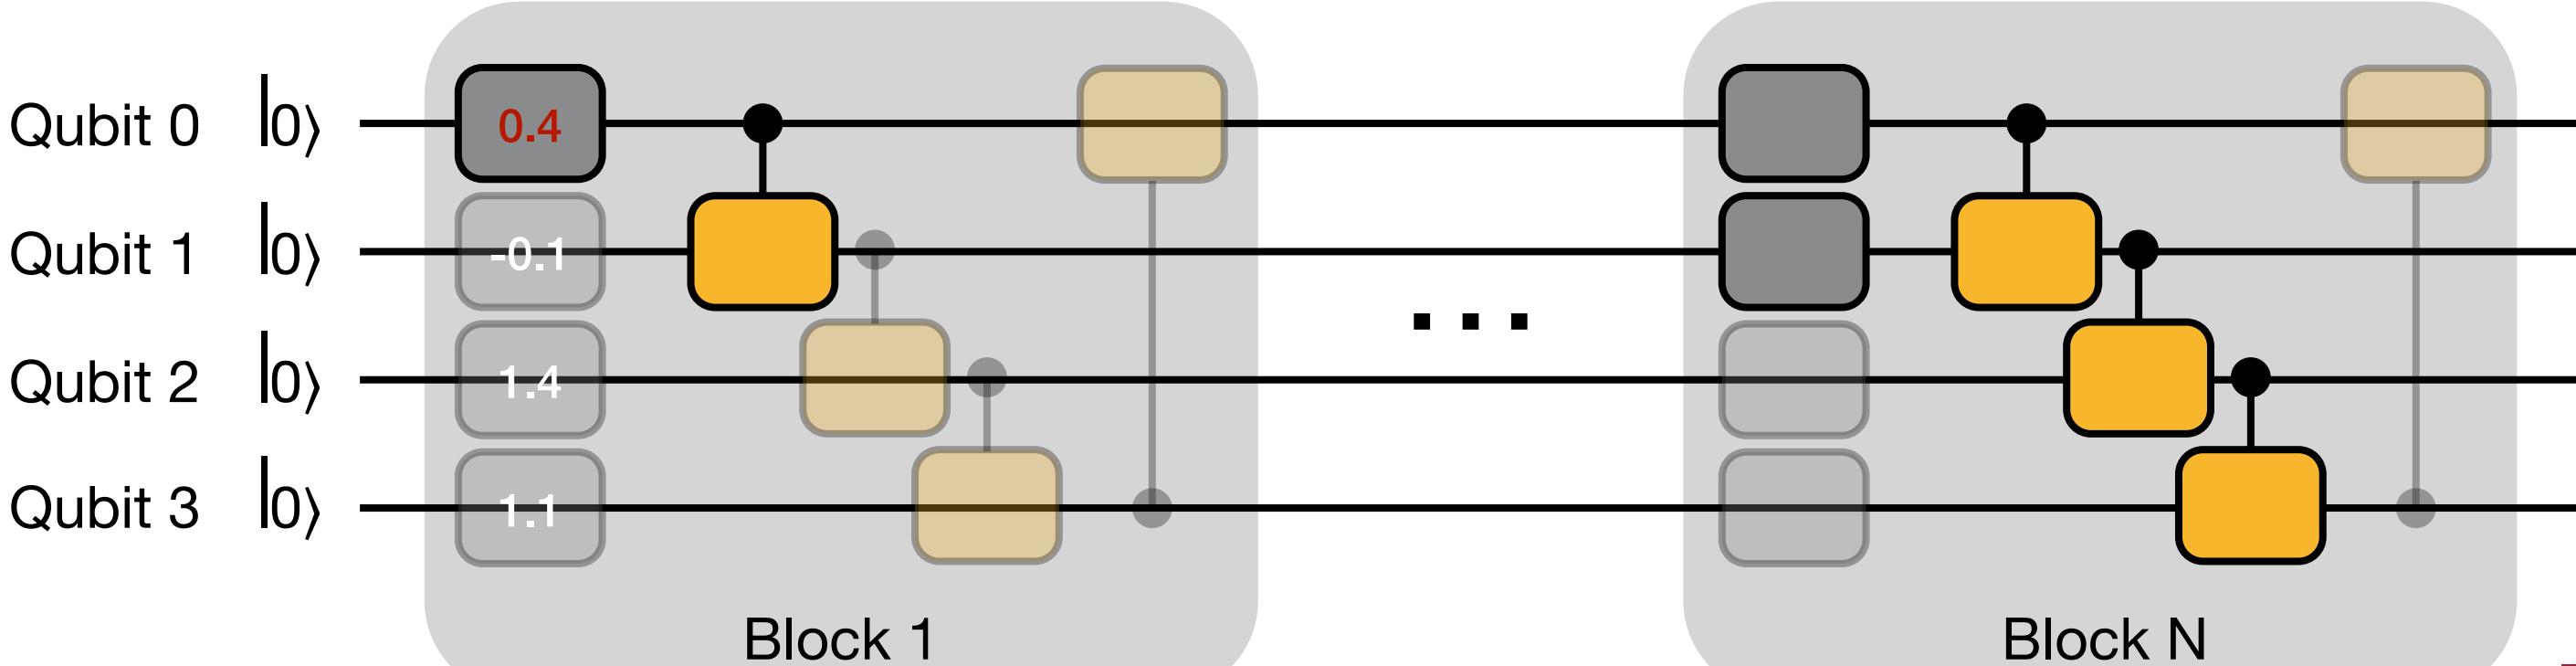
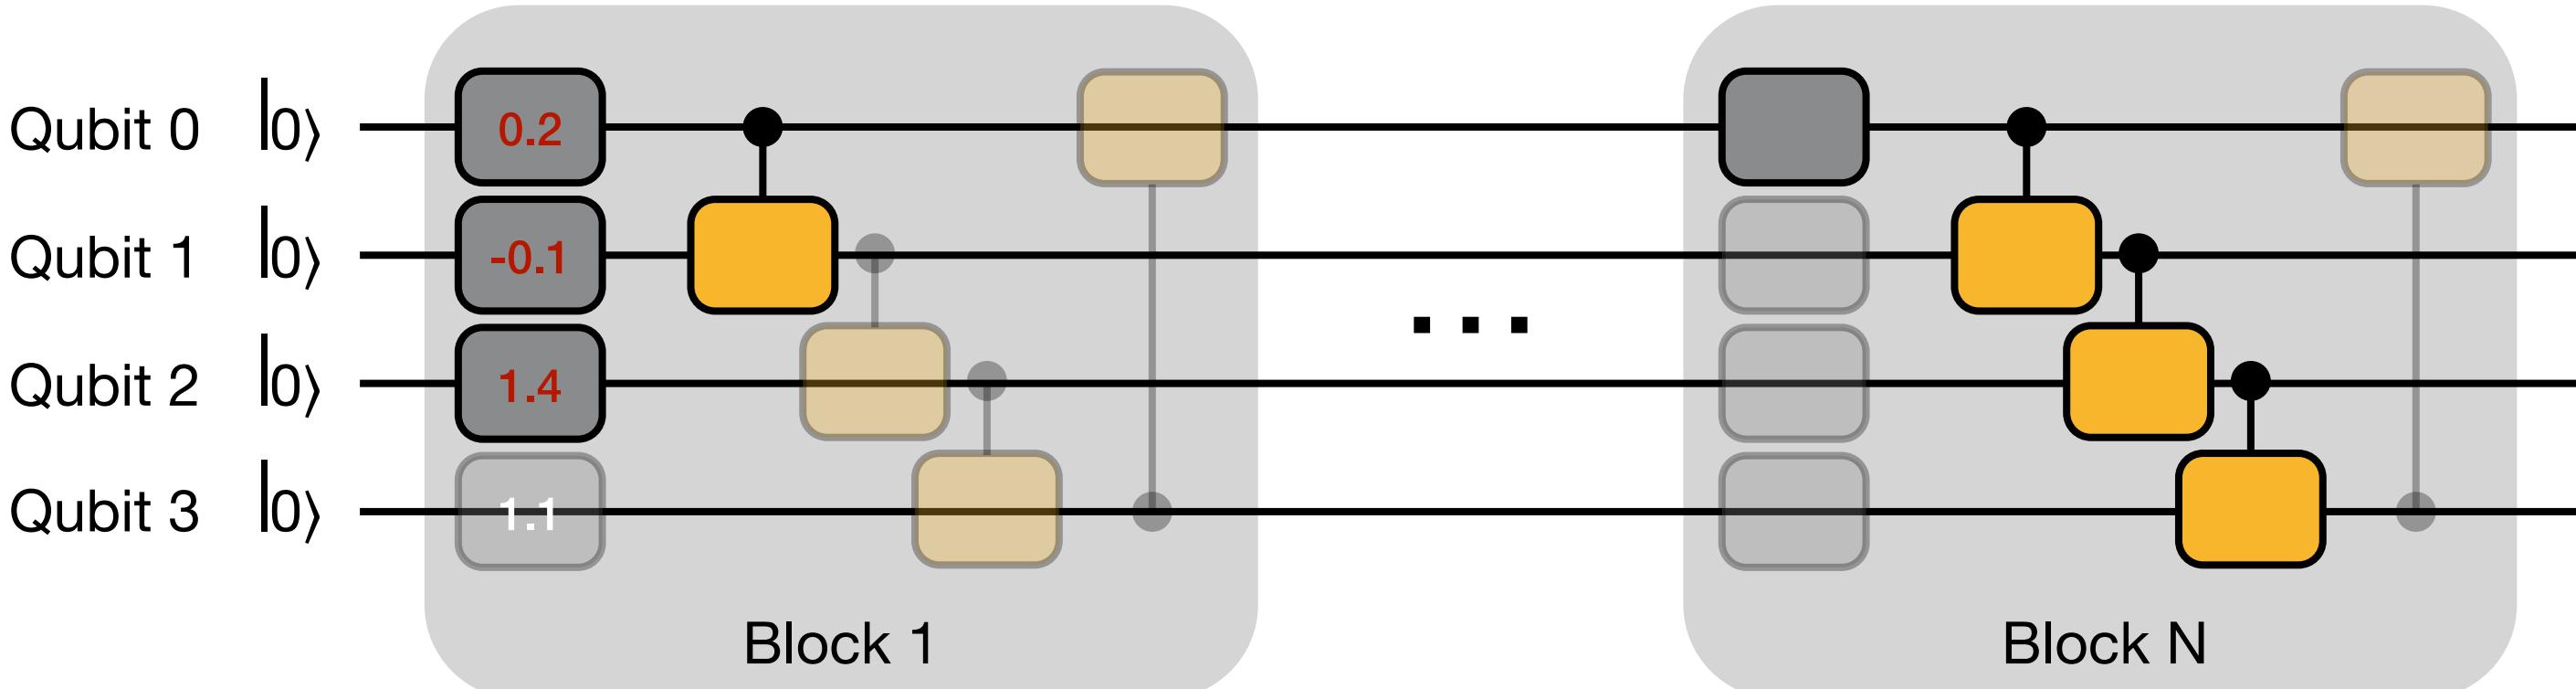
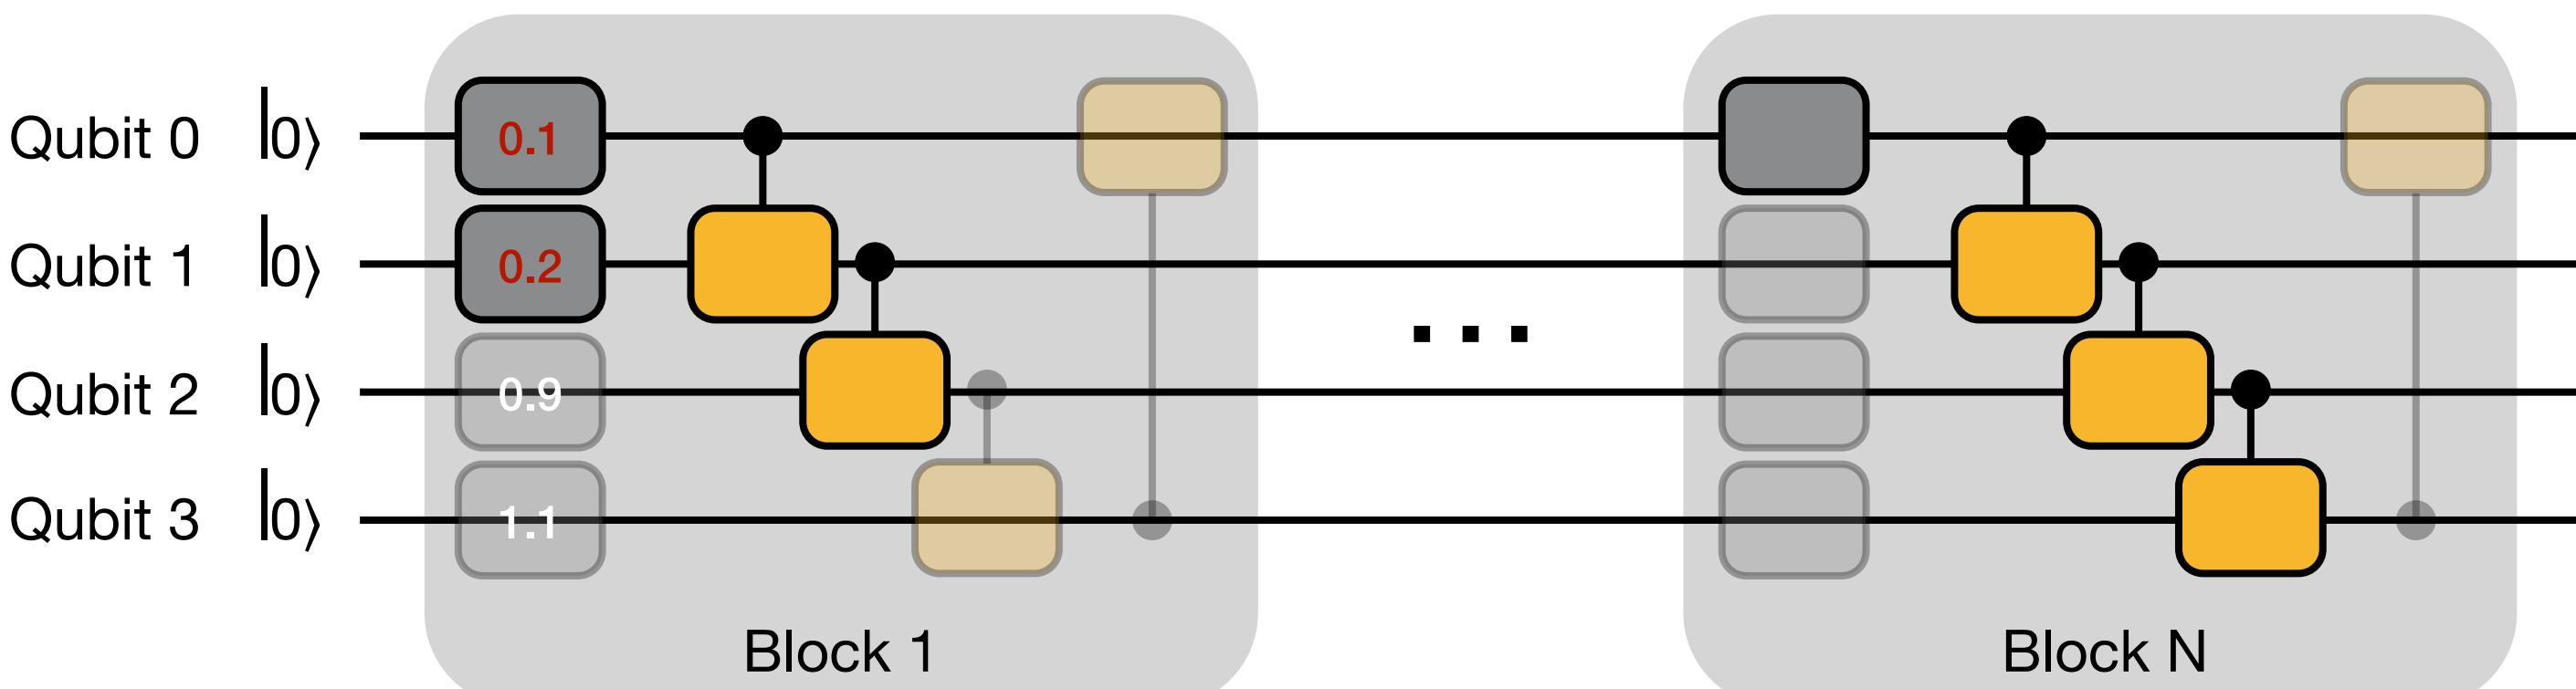


Train SuperCircuit for Multiple Steps

- In one SuperCircuit Training step: Sample and Train

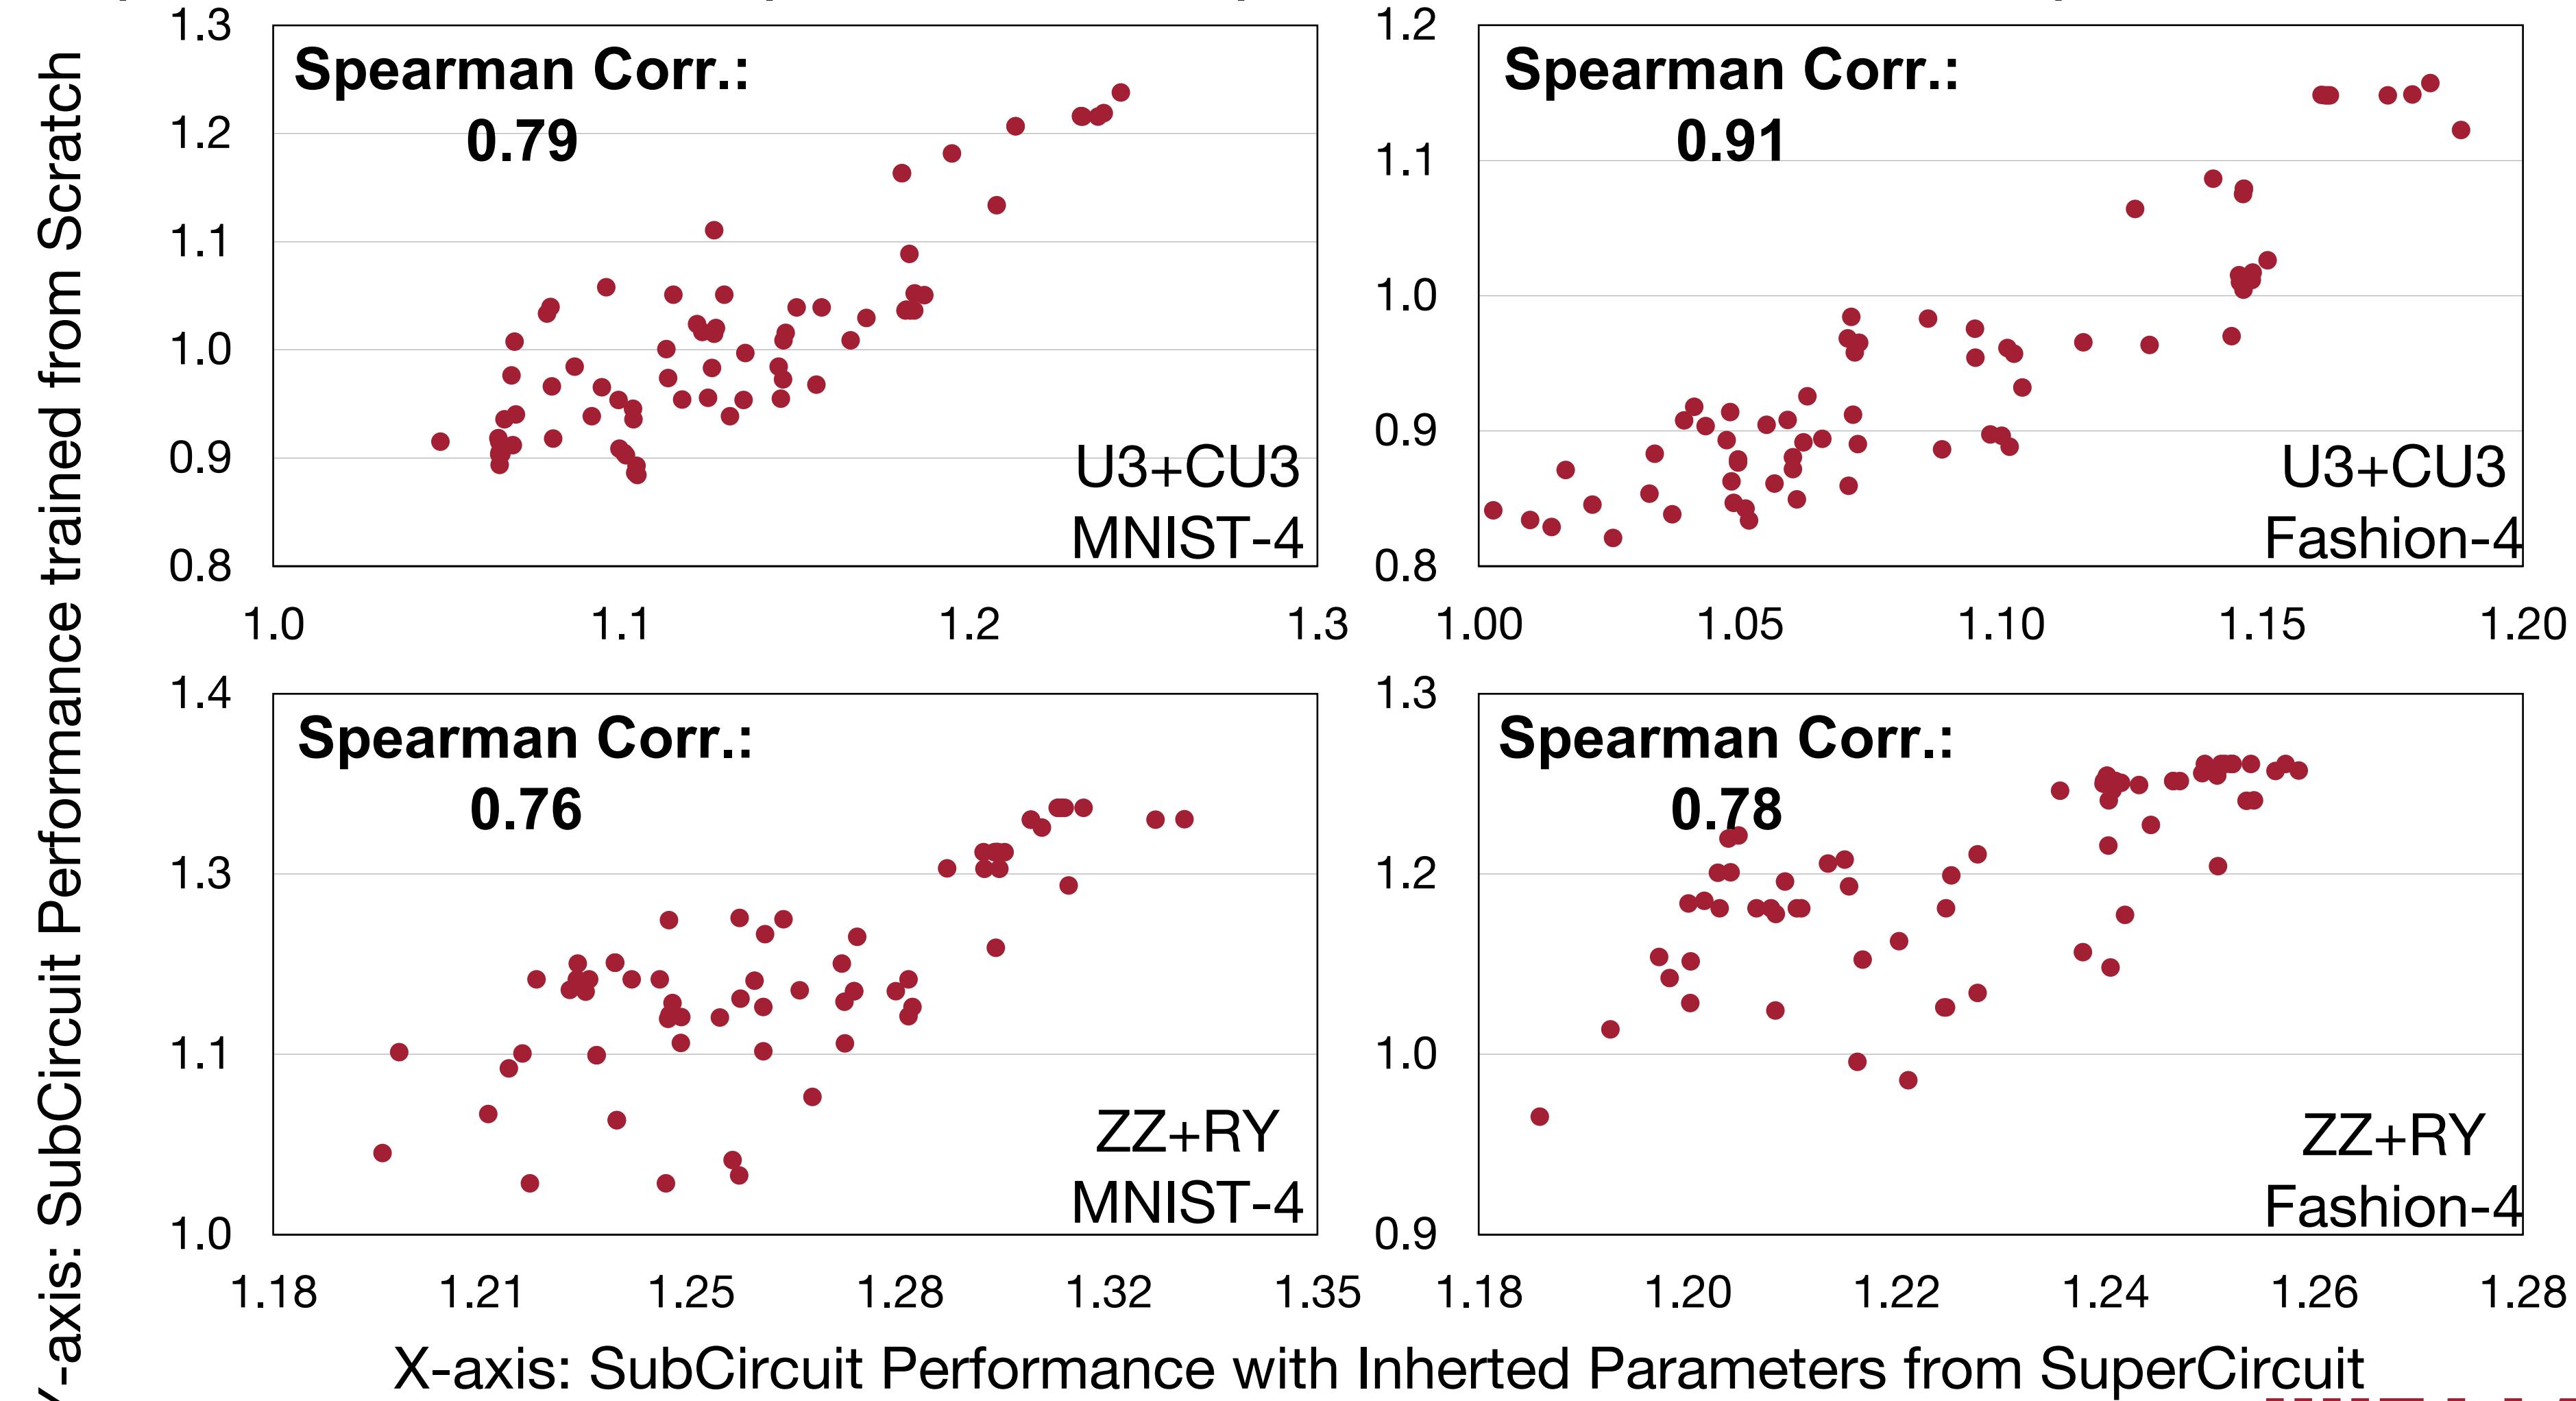


Train SuperCircuit for Multiple Steps



How Reliable is the SuperCircuit?

- Inherited parameters from SuperCircuit can provide accurate relative performance

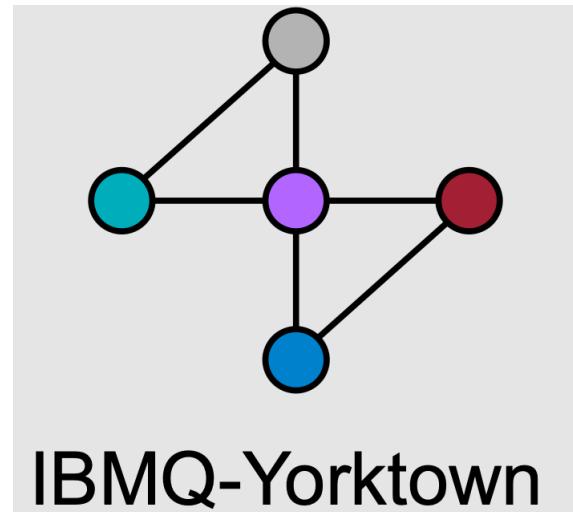


QuantumNAS

- SuperCircuit Construction and Training
- Noise-Adaptive Evolutionary Co-Search of SubCircuit and Qubit Mapping
- Train the Searched SubCircuit
- Iterative Quantum Gate Pruning

Noise-Adaptive Evolutionary Co-Search

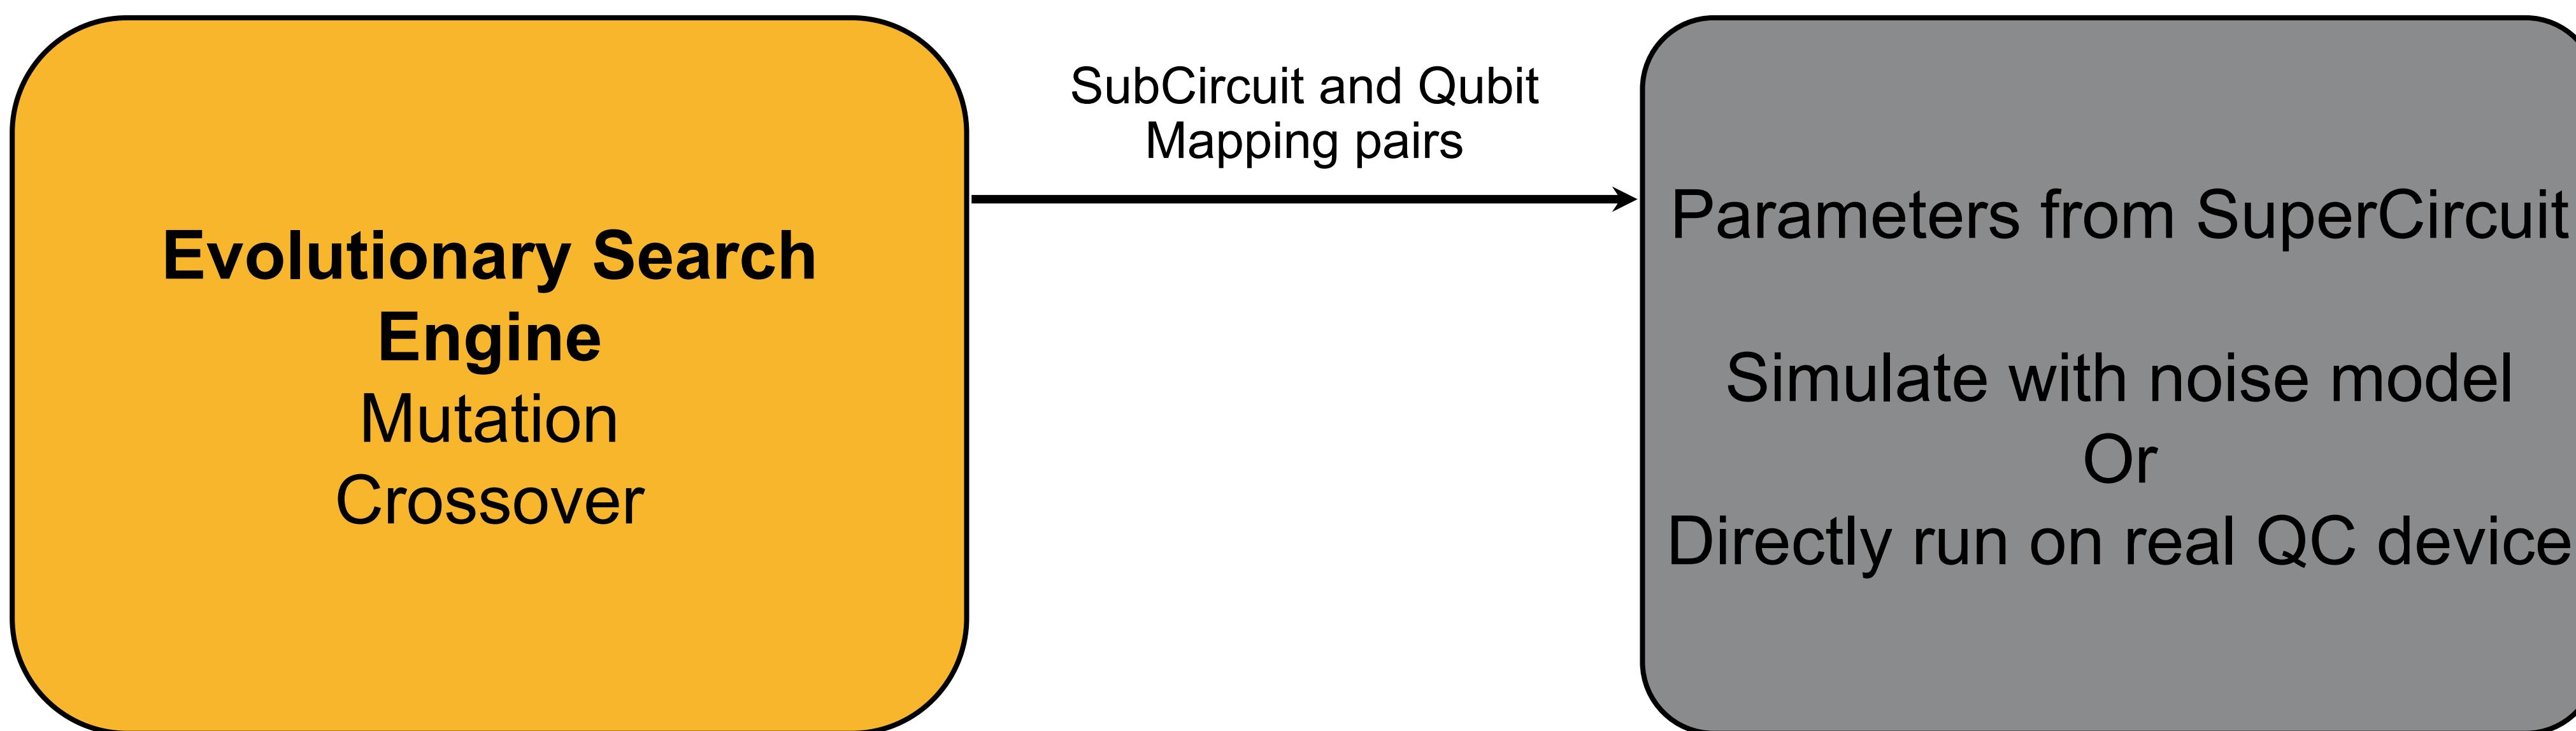
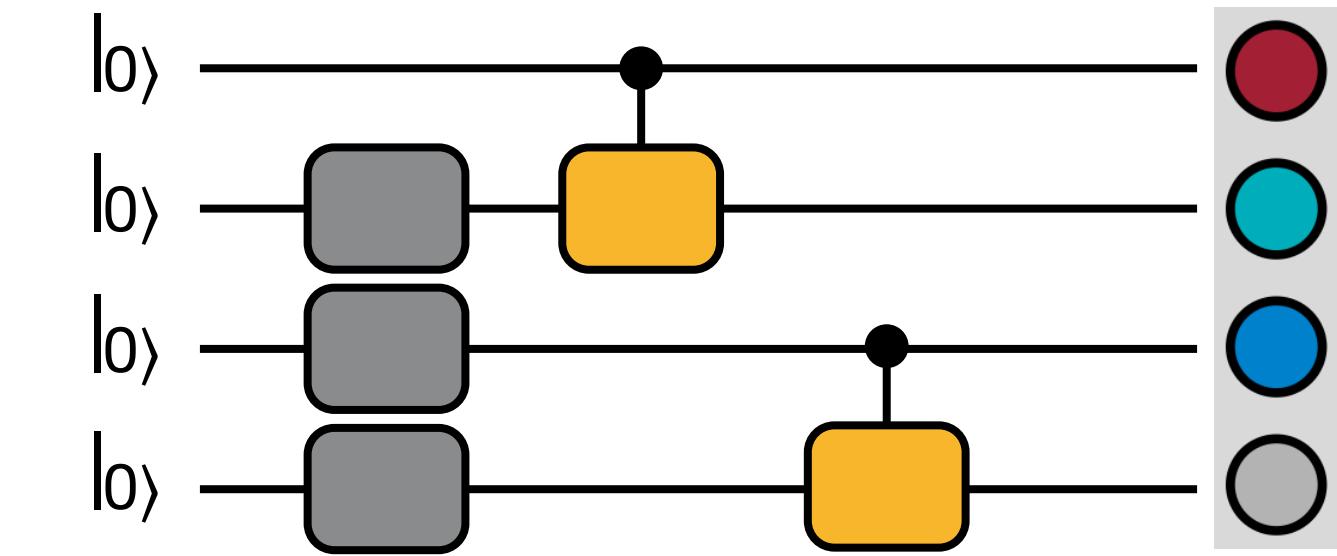
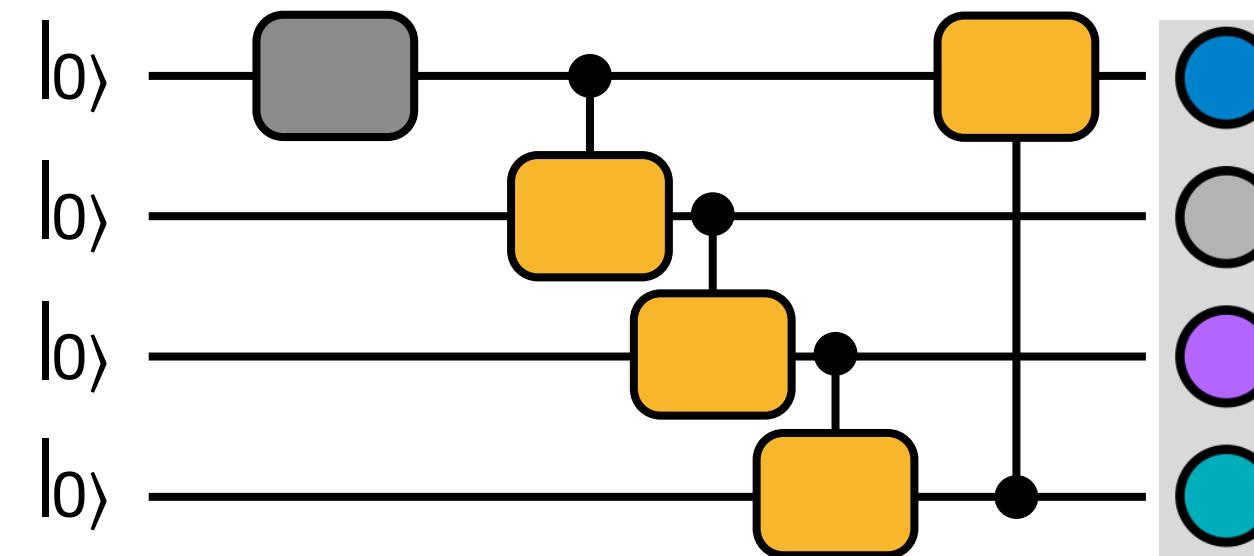
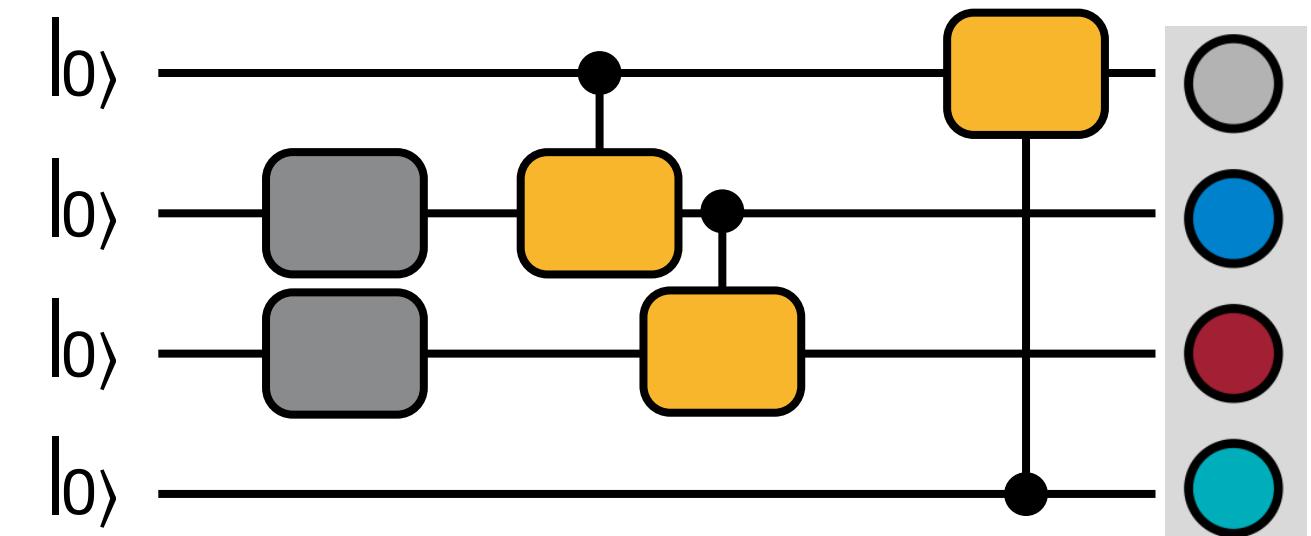
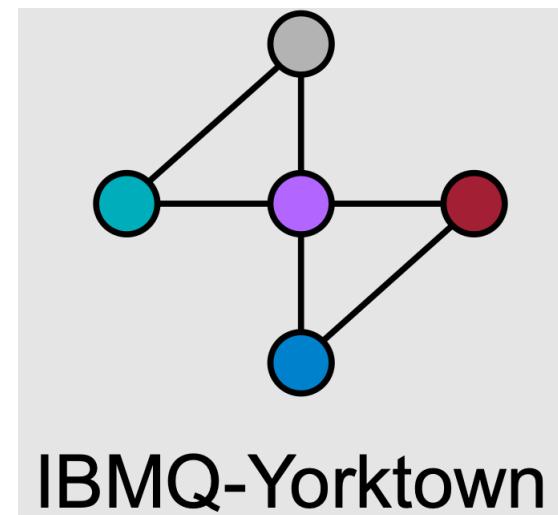
- Search the best SubCircuit and its qubit mapping on target device



IBMQ-Yorktown

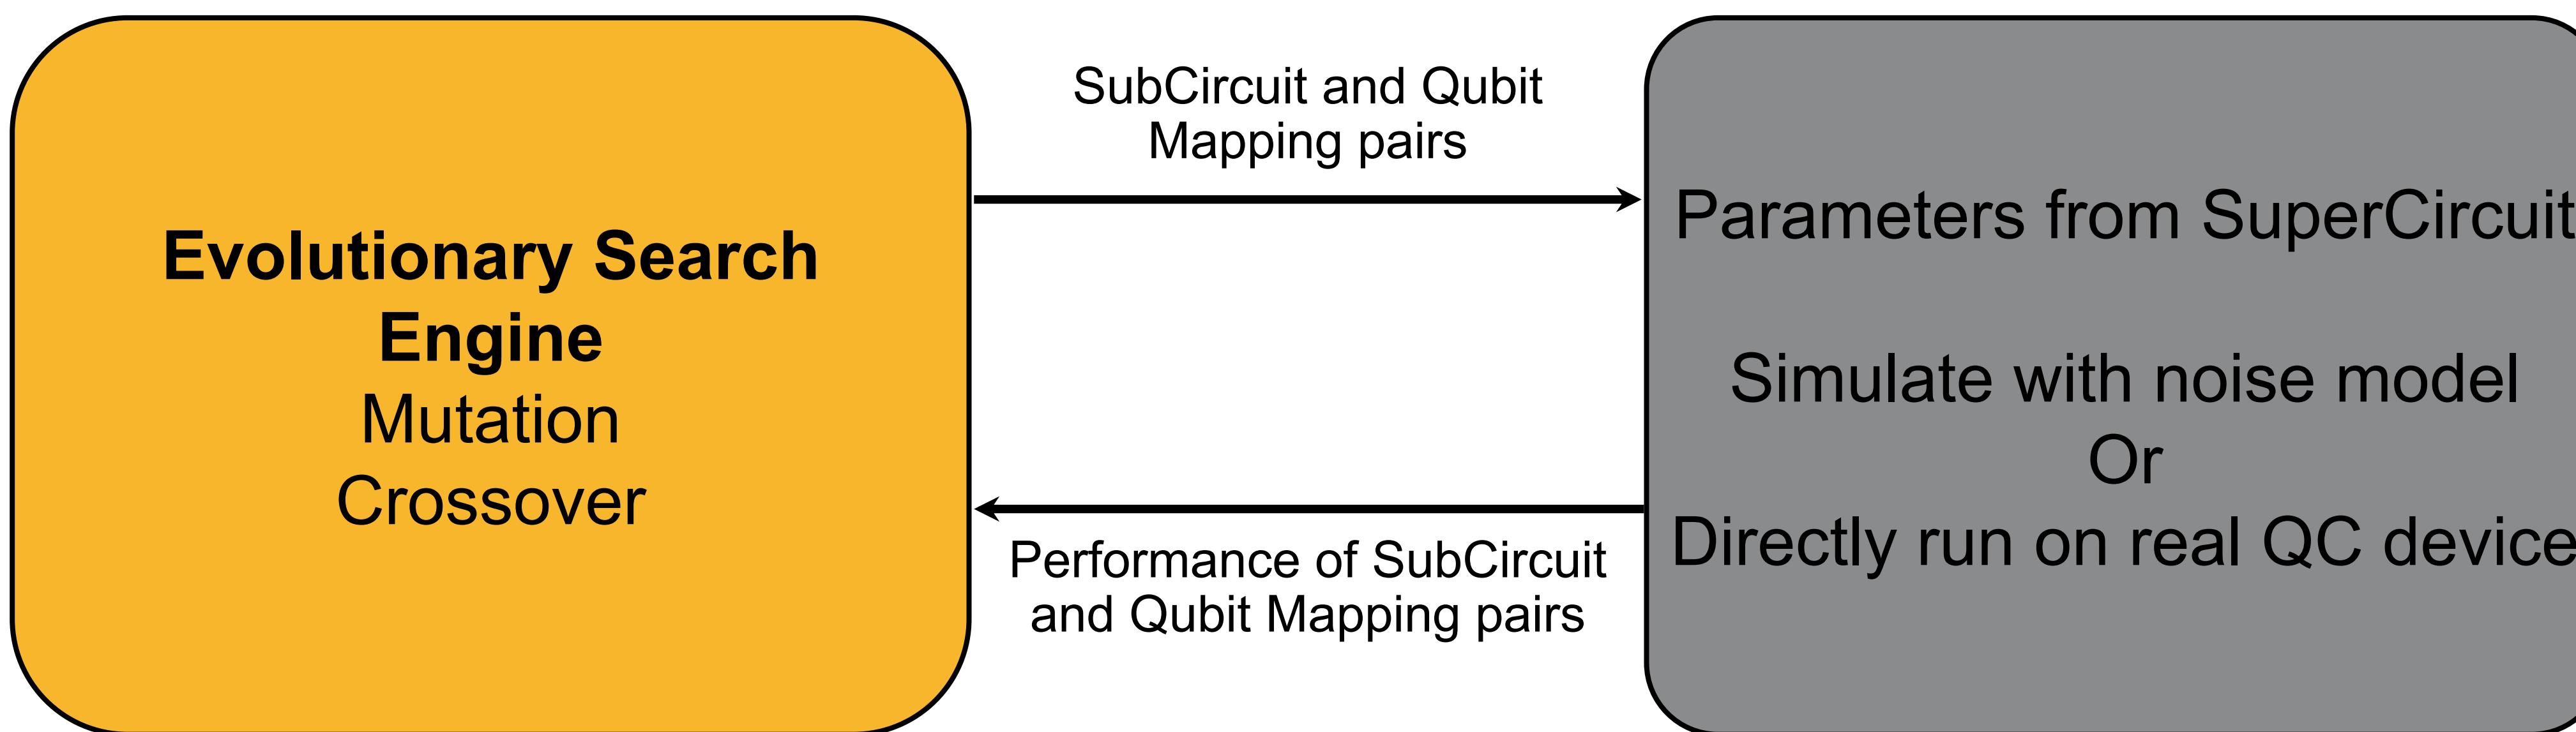
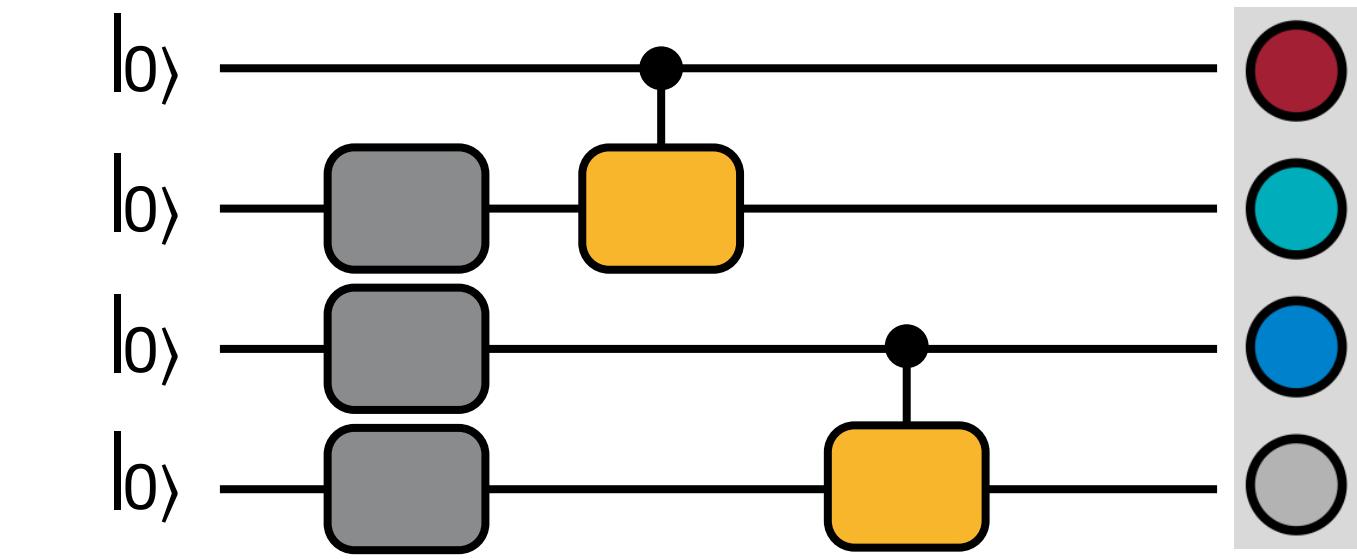
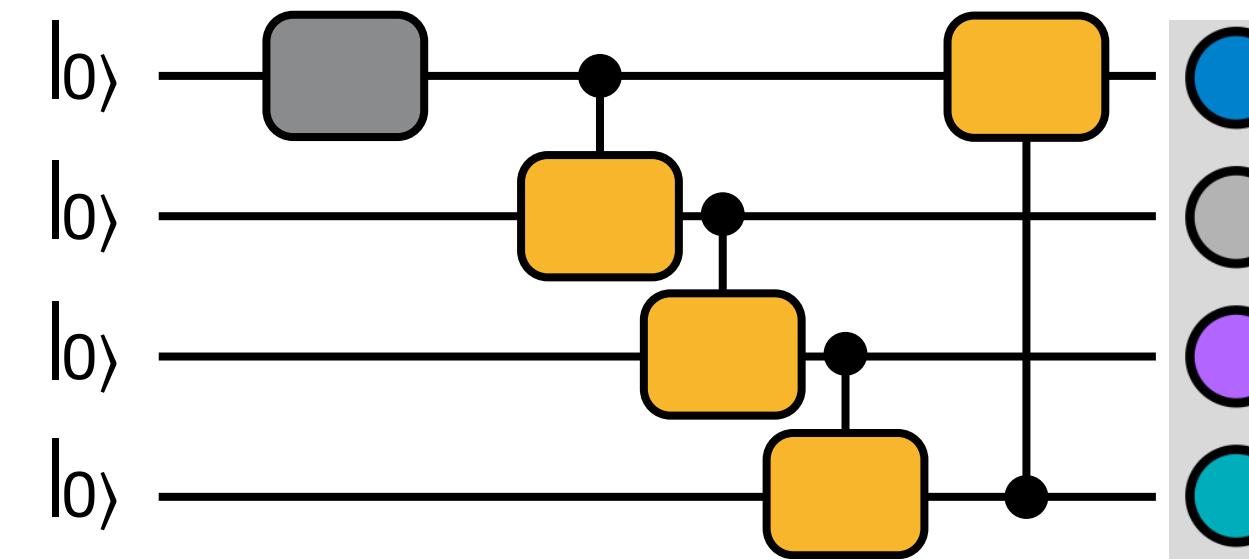
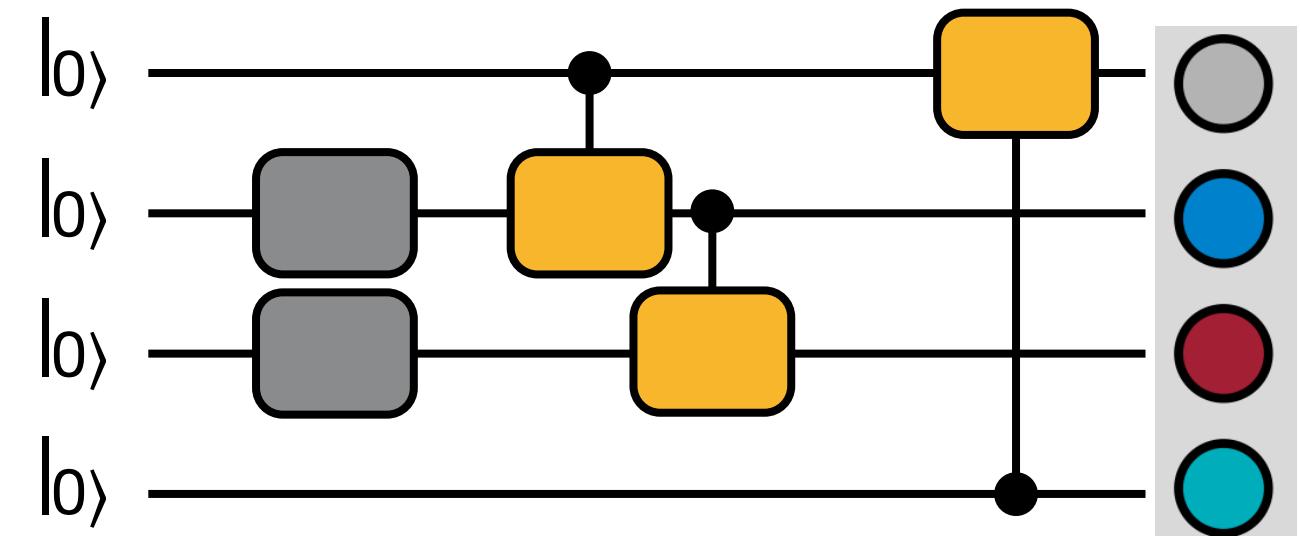
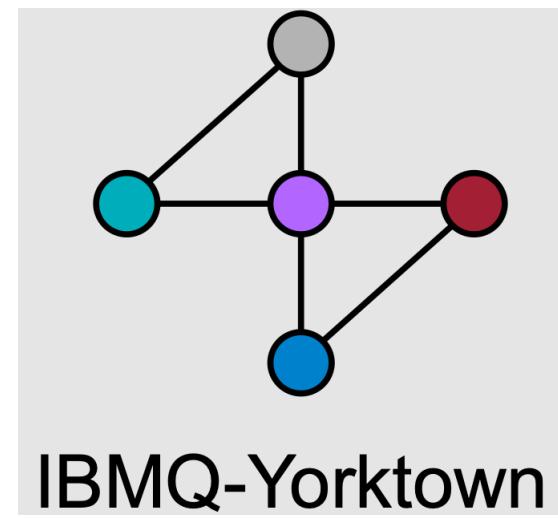
Noise-Adaptive Evolutionary Co-Search

- Search the best SubCircuit and its qubit mapping on target device



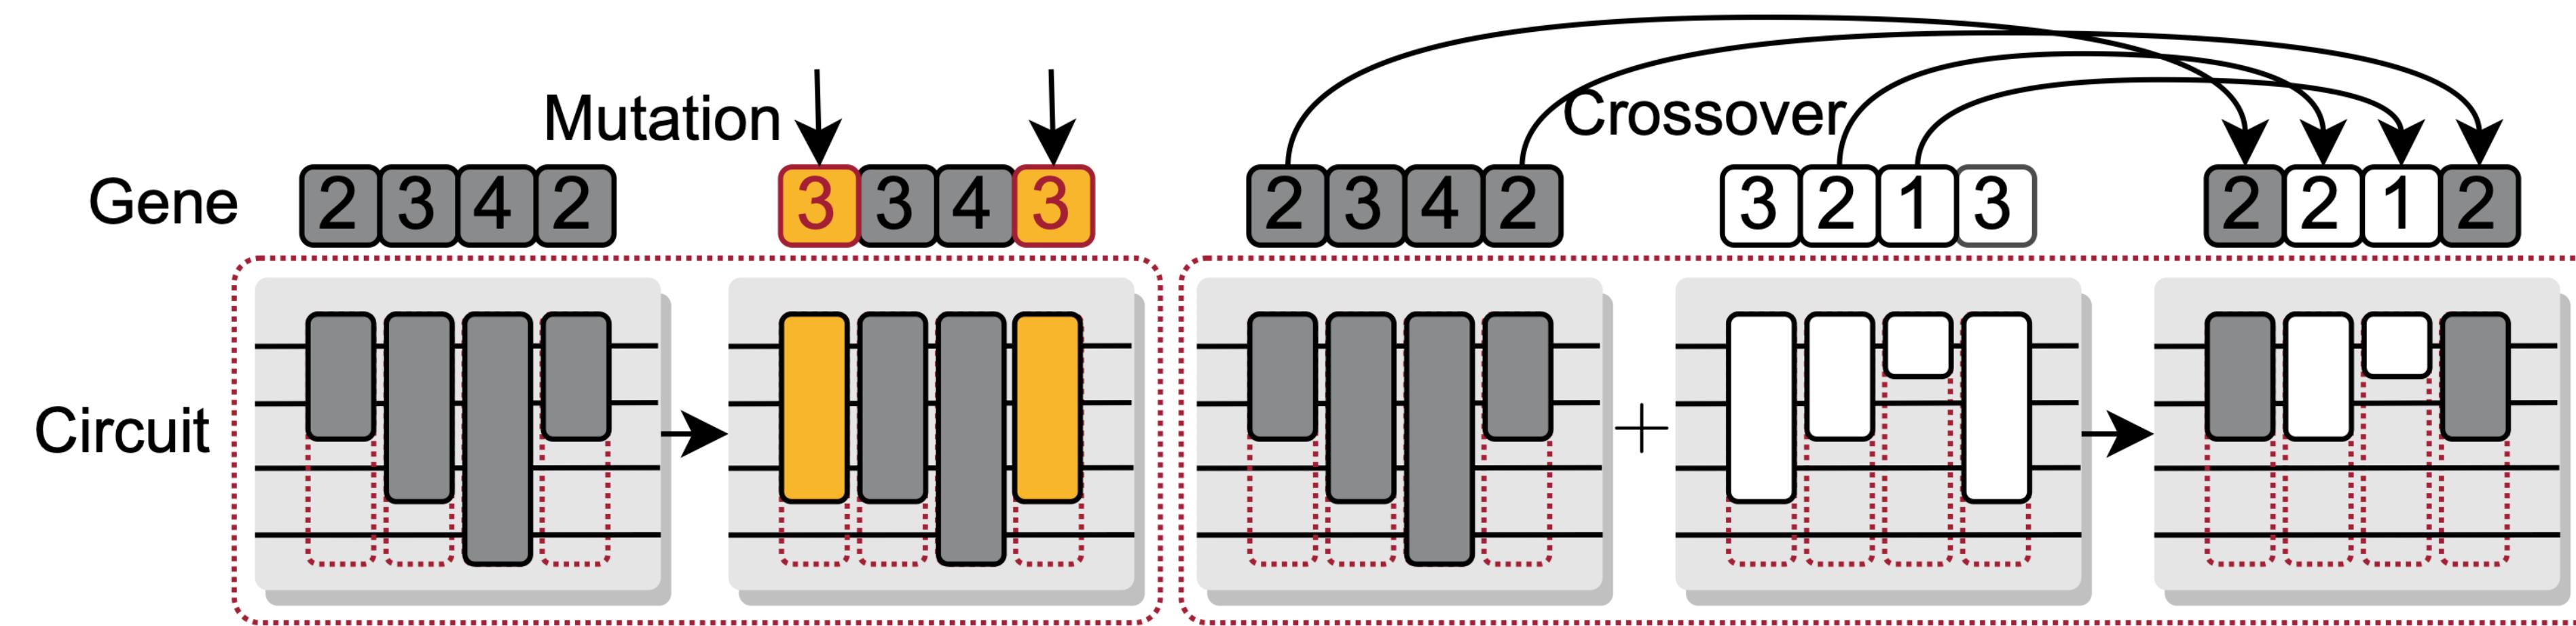
Noise-Adaptive Evolutionary Co-Search

- Search the best SubCircuit and its qubit mapping on target device



Mutation and Crossover

- Mutation and crossover create new SubCircuit candidates



QuantumNAS

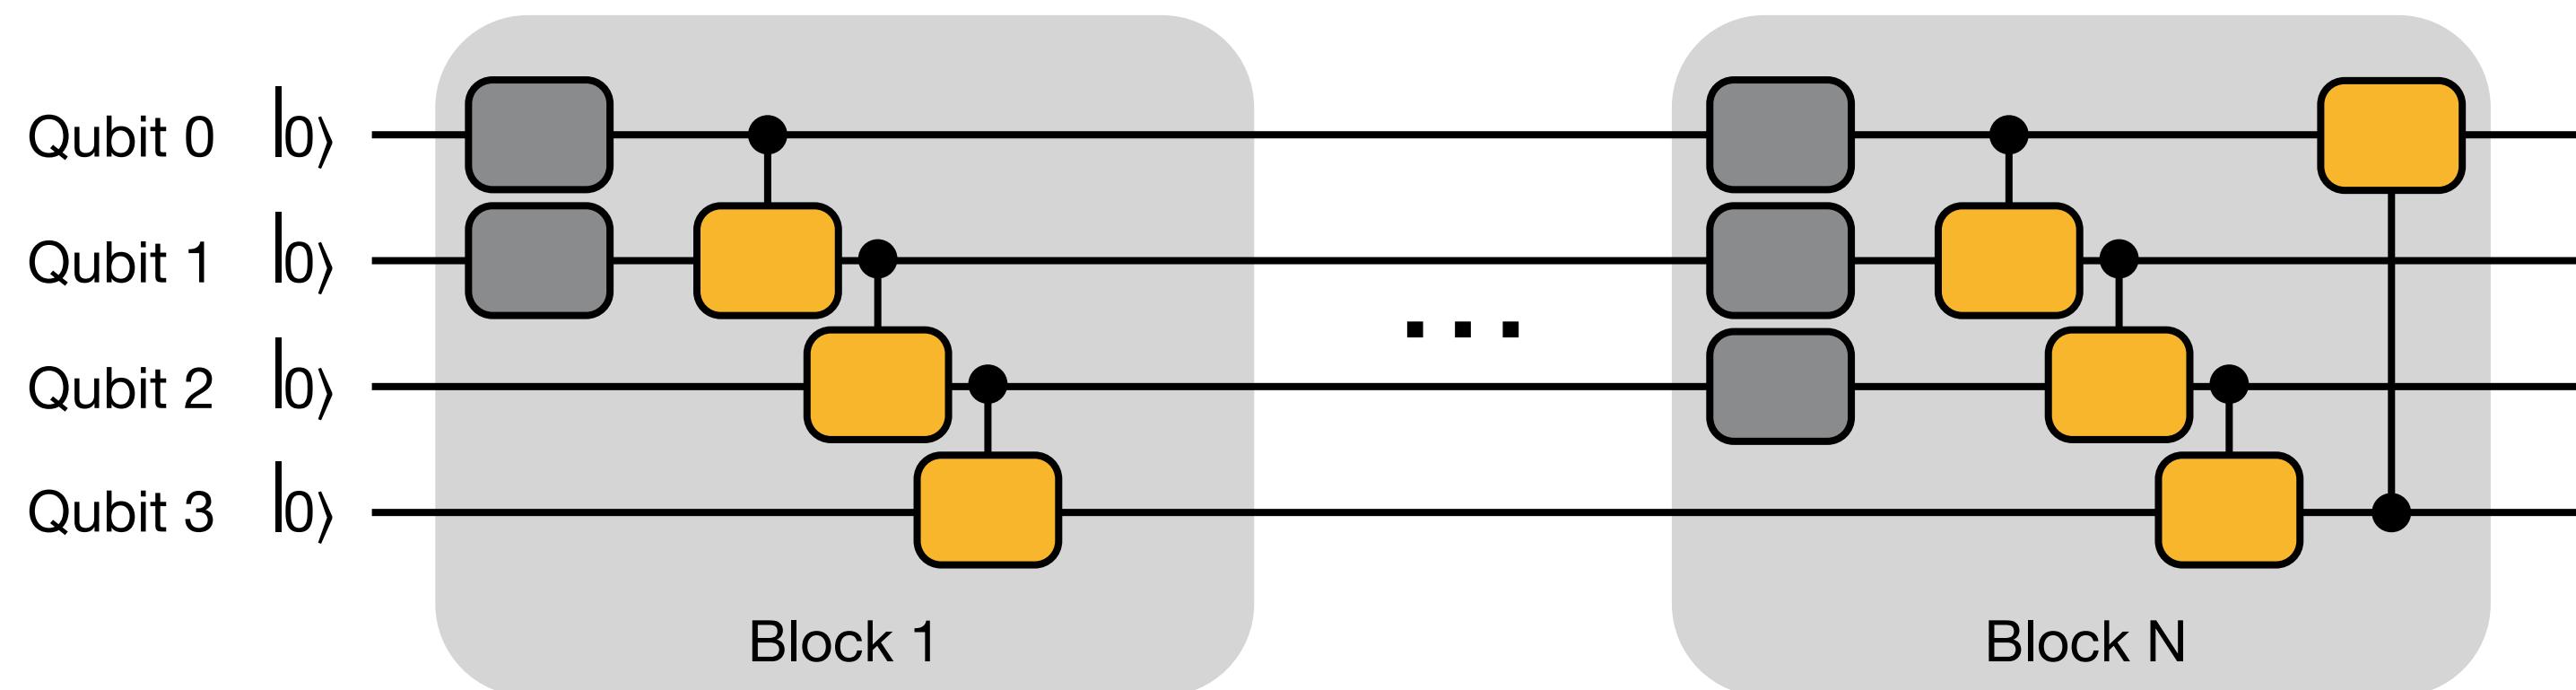
- SuperCircuit Construction and Training
- Noise-Adaptive Evolutionary Co-Search of SubCircuit and Qubit Mapping
- Train the Searched SubCircuit
- Iterative Quantum Gate Pruning

QuantumNAS

- SuperCircuit Construction and Training
- Noise-Adaptive Evolutionary Co-Search of SubCircuit and Qubit Mapping
- Train the Searched SubCircuit
- Iterative Quantum Gate Pruning

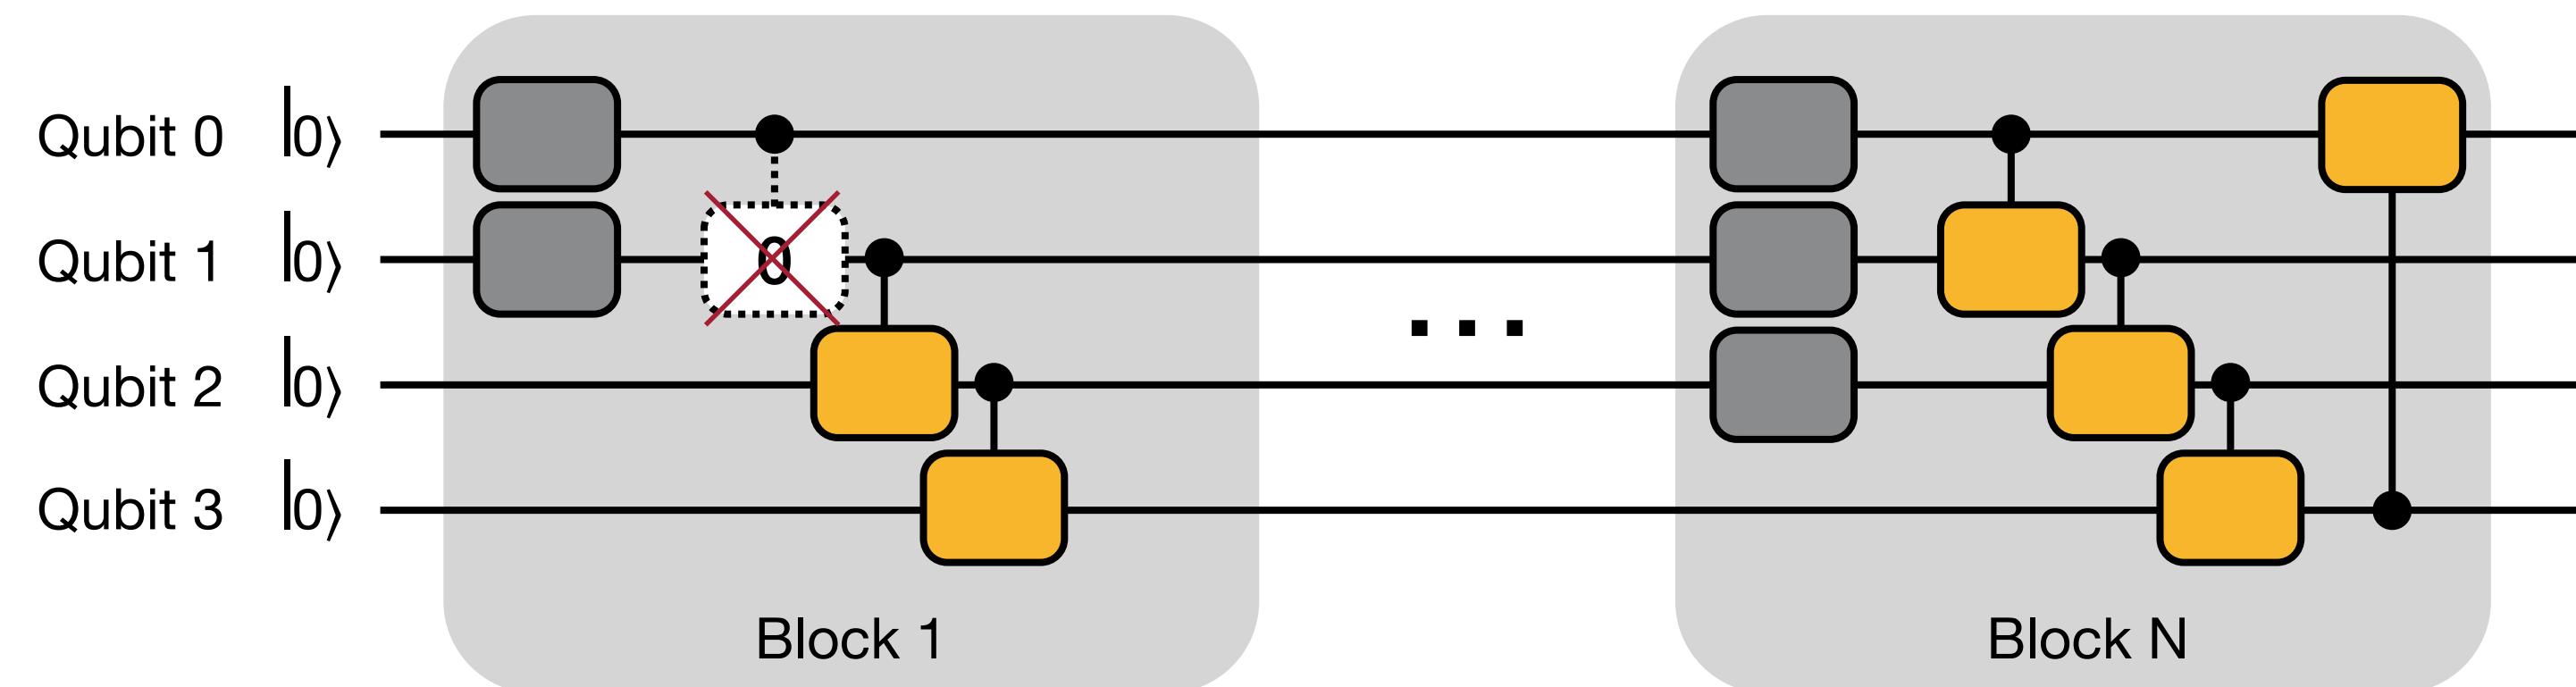
Iterative Pruning

- Some gates have parameters close to 0
 - Rotation gate with angle close to 0 has small impact on the results
 - Iteratively prune small-magnitude gates and fine-tune the remaining parameters



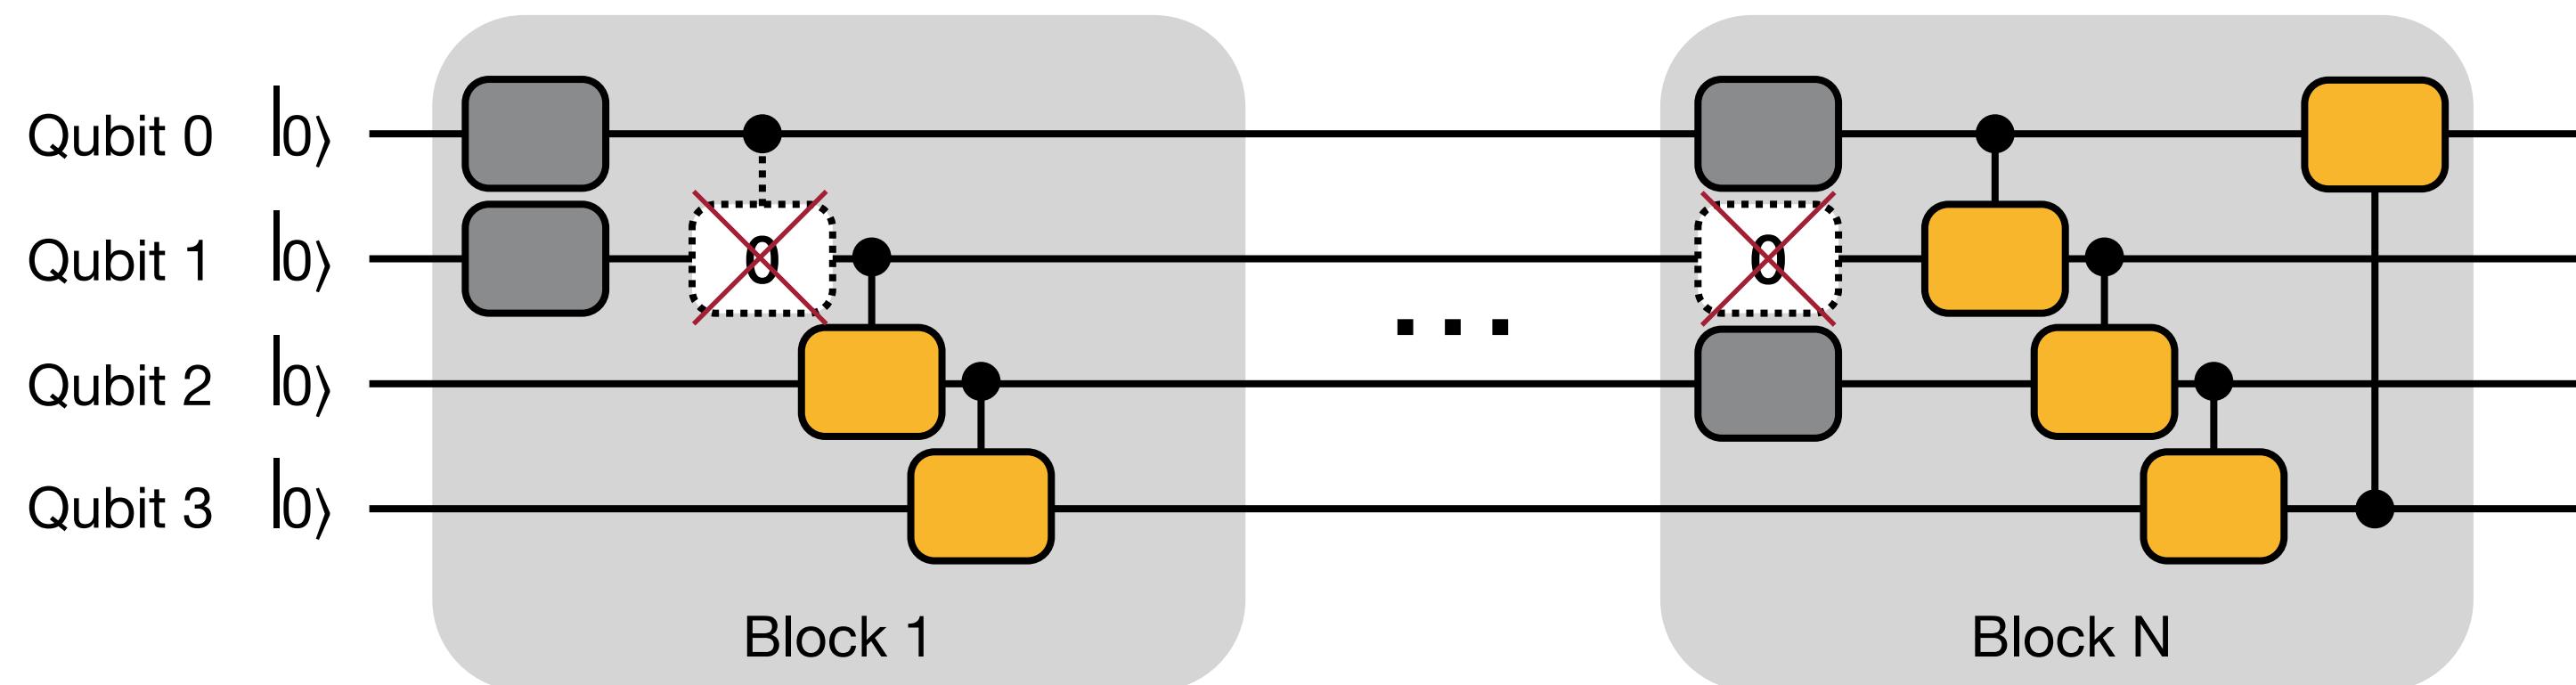
Iterative Pruning

- Some gates have parameters close to 0
 - Rotation gate with angle close to 0 has small impact on the results
 - Iteratively prune small-magnitude gates and fine-tune the remaining parameters



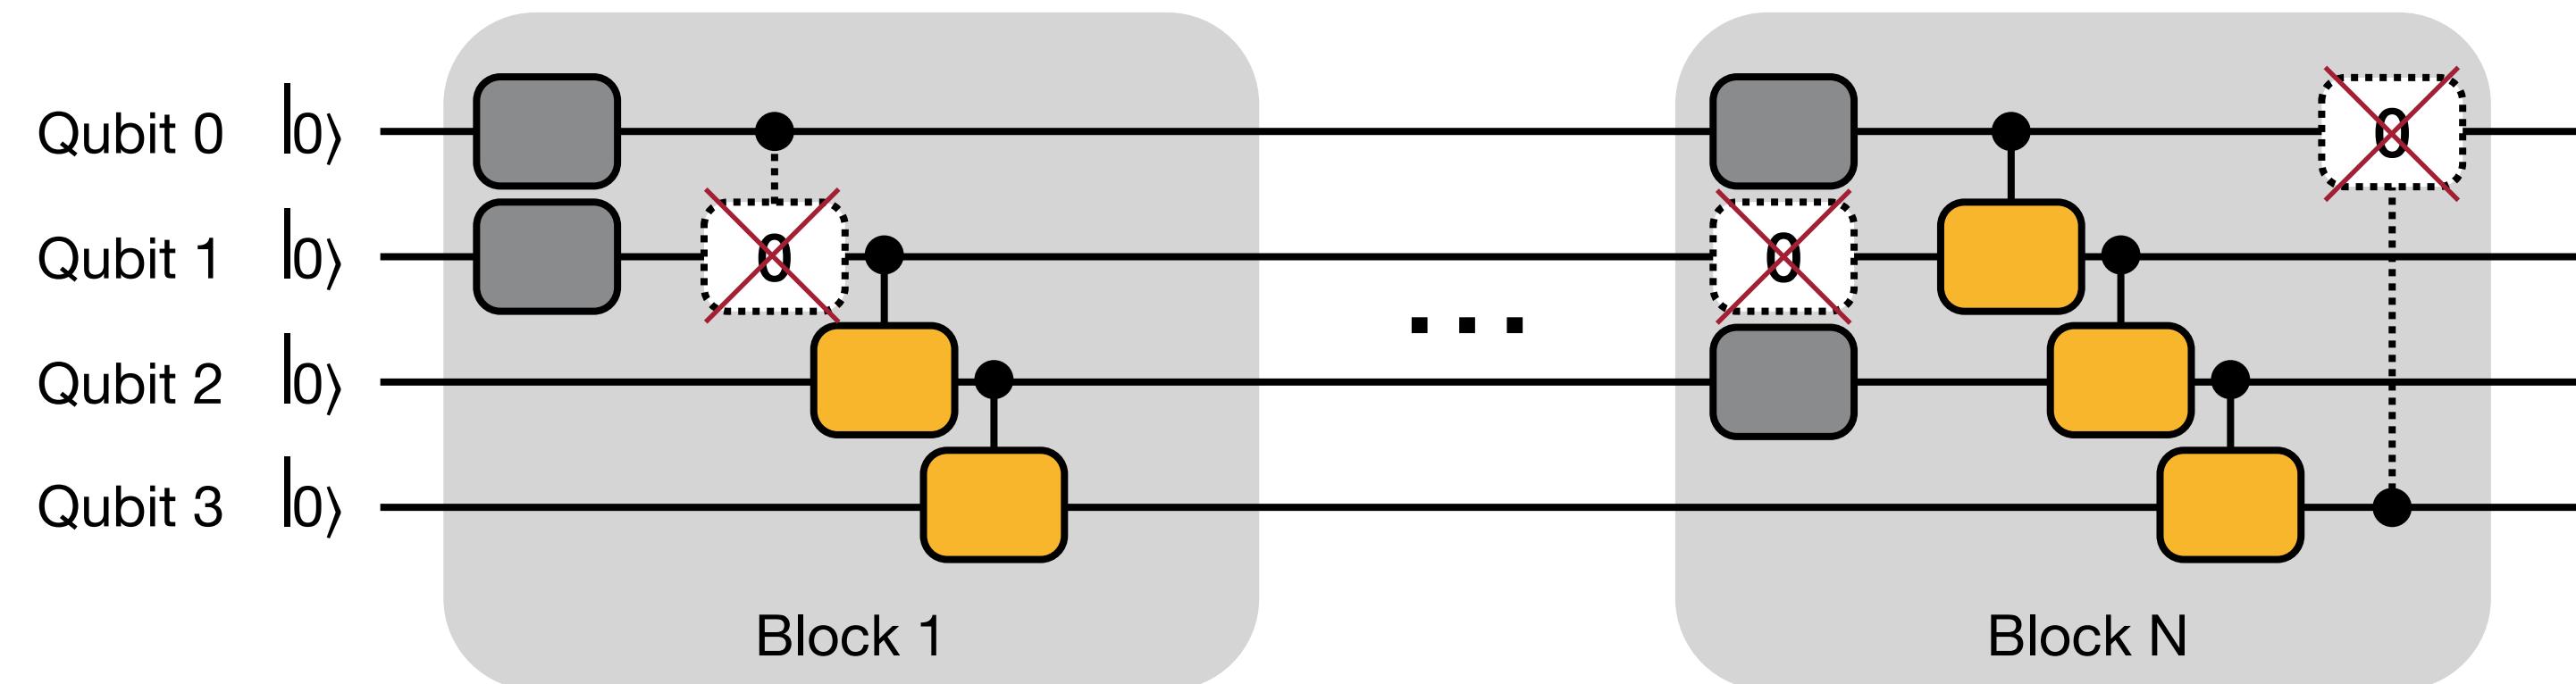
Iterative Pruning

- Some gates have parameters close to 0
 - Rotation gate with angle close to 0 has small impact on the results
 - Iteratively prune small-magnitude gates and fine-tune the remaining parameters



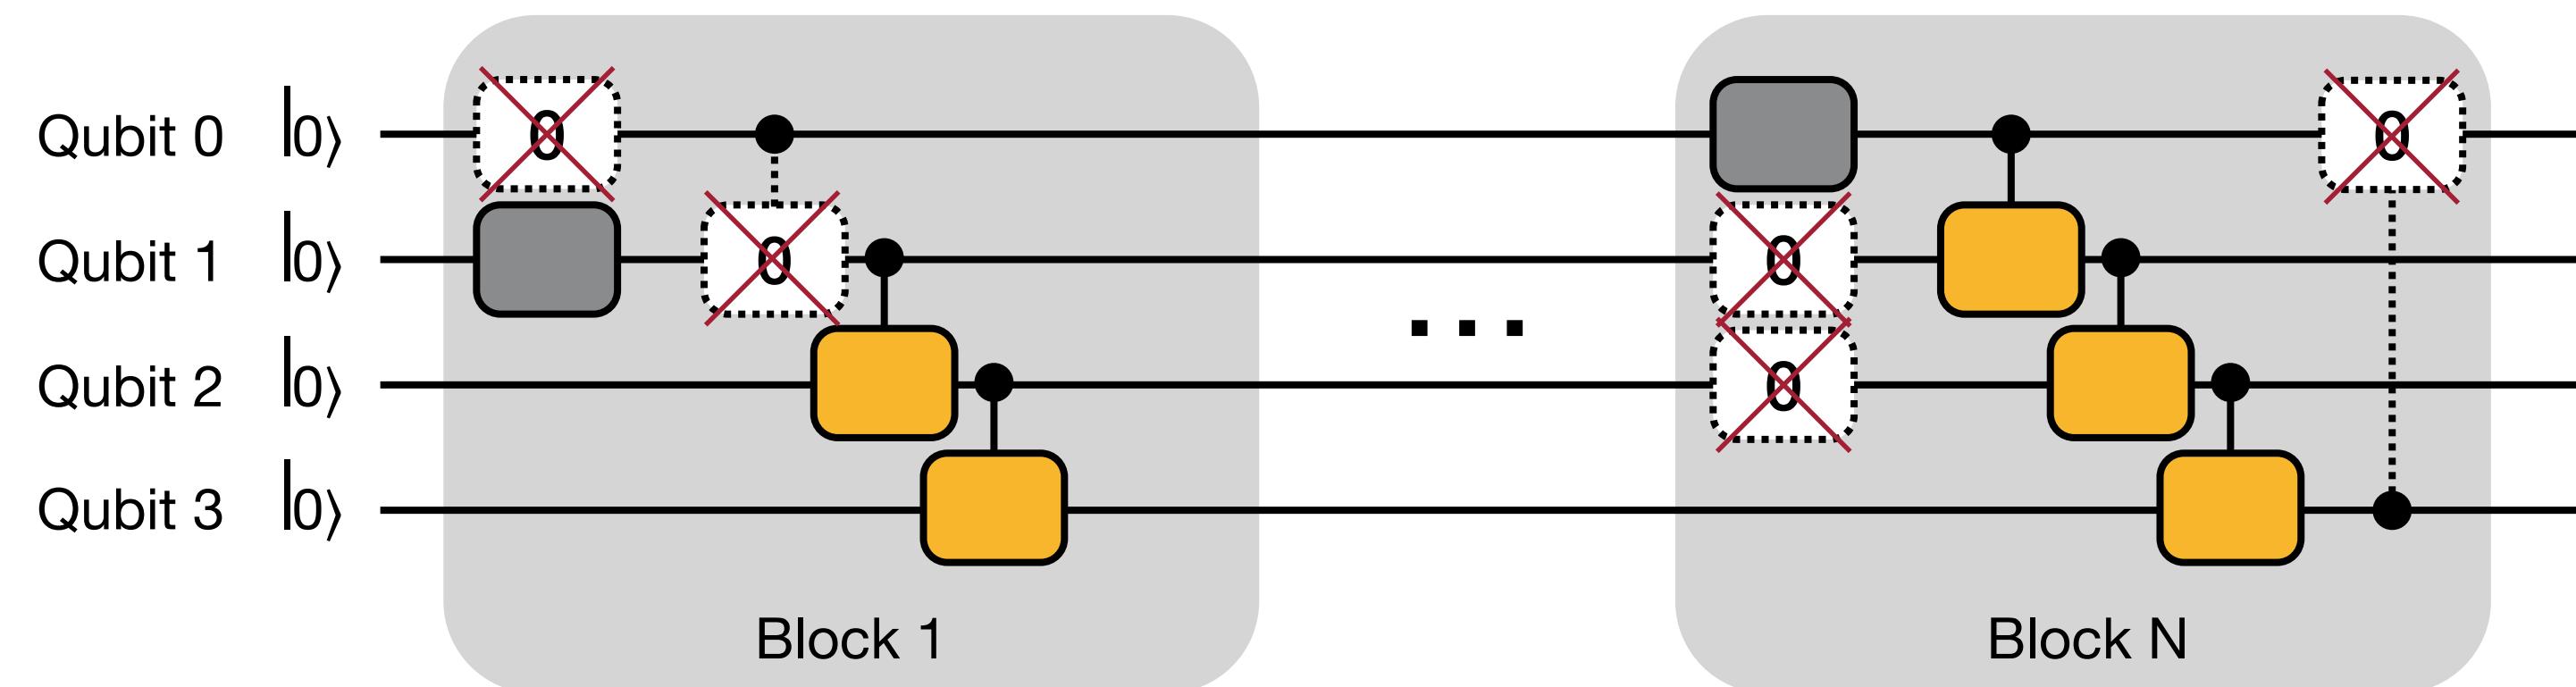
Iterative Pruning

- Some gates have parameters close to 0
 - Rotation gate with angle close to 0 has small impact on the results
 - Iteratively prune small-magnitude gates and fine-tune the remaining parameters



Iterative Pruning

- Some gates have parameters close to 0
 - Rotation gate with angle close to 0 has small impact on the results
 - Iteratively prune small-magnitude gates and fine-tune the remaining parameters

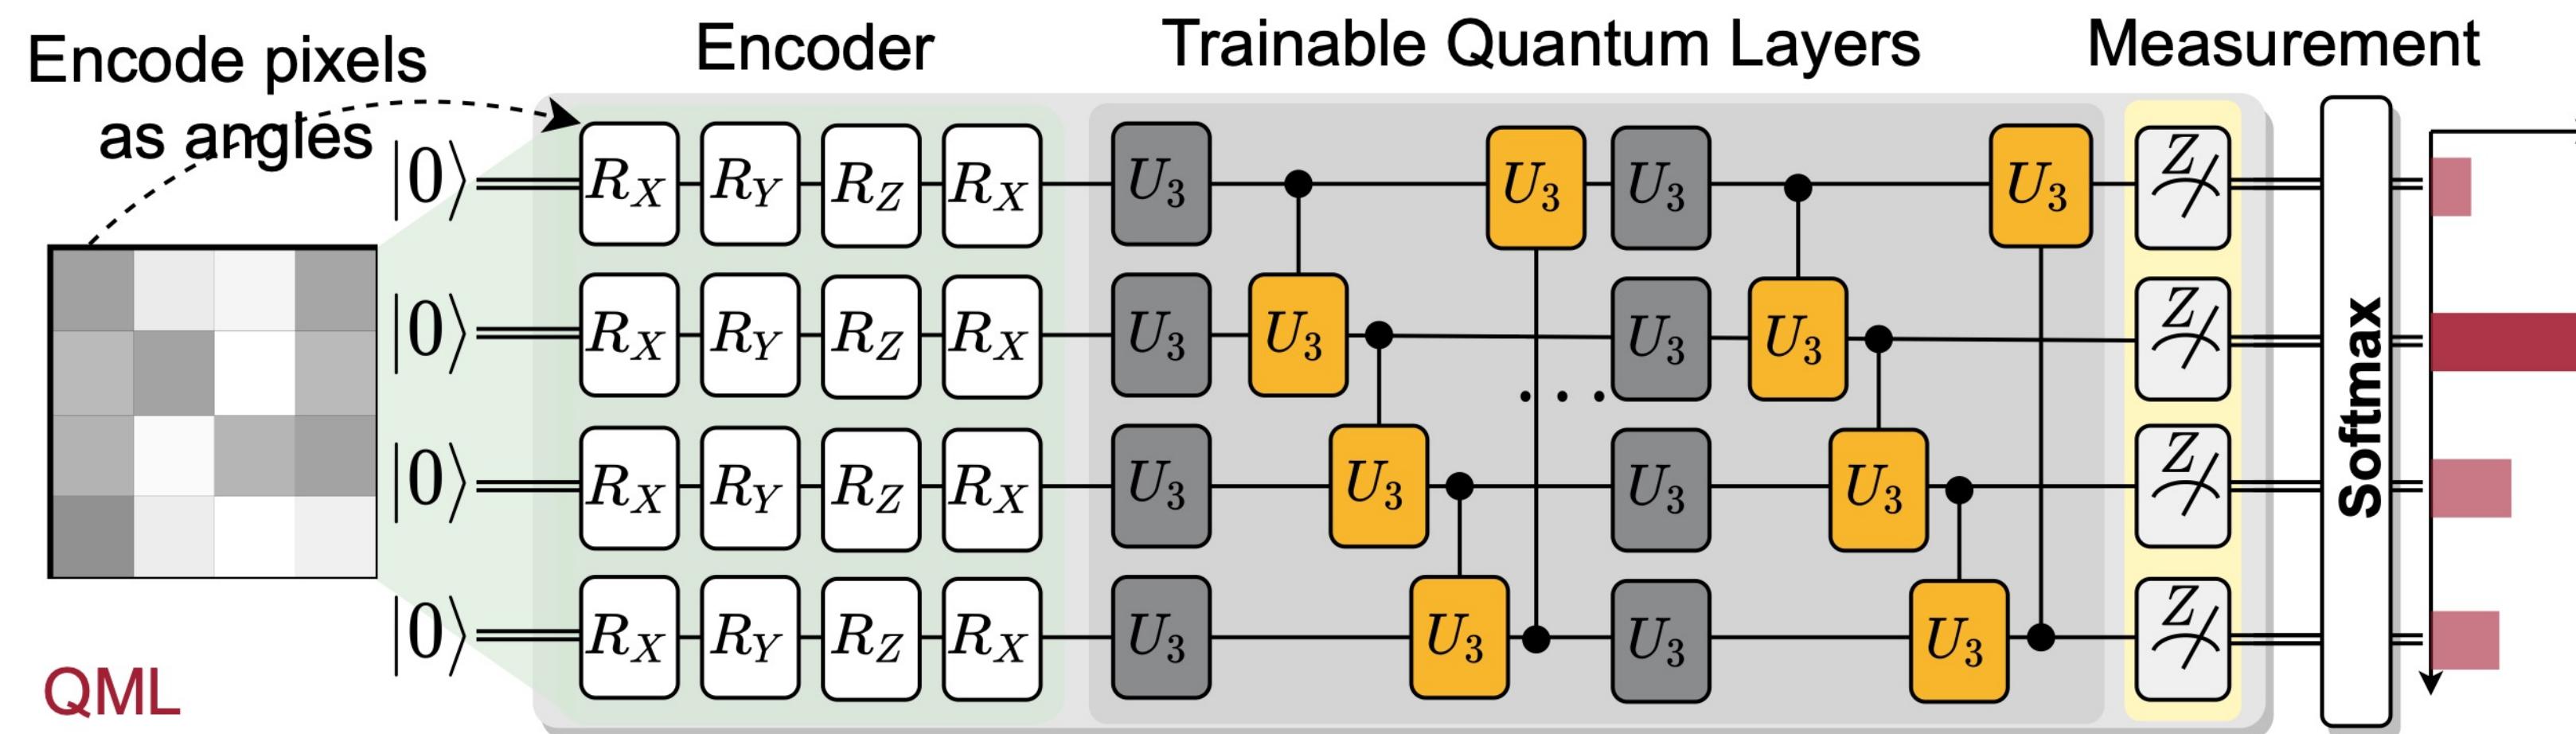


Evaluation Setups: Benchmarks and Devices

- Benchmarks
 - QML classification tasks: MNIST 10-class, 4-class, 2-class, Fashion 4-class, 2-class, Vowel 4-class
 - VQE task molecules: H₂, H₂O, LiH, CH₄, BeH₂
- Quantum Devices
 - IBMQ
 - #Qubits: 5 to 65
 - Quantum Volume: 8 to 128

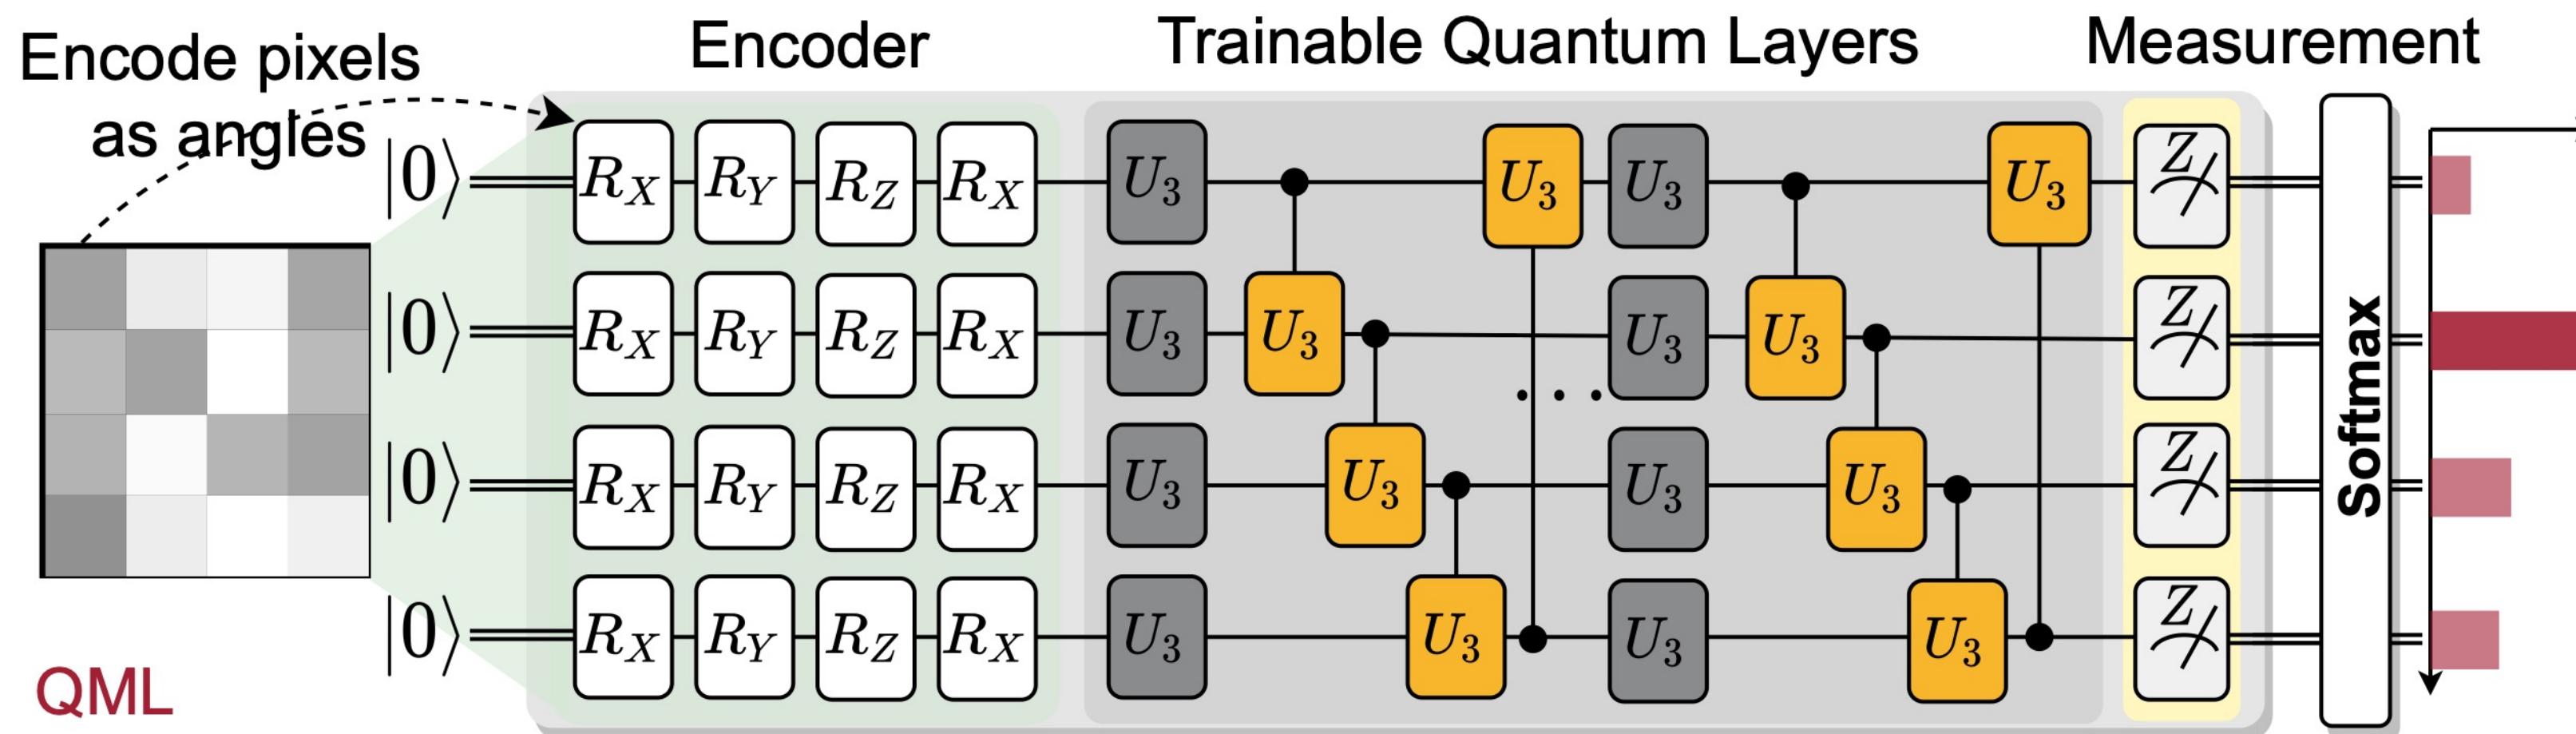
Benchmarks: QNN and VQE

- Quantum Neural Networks: classification

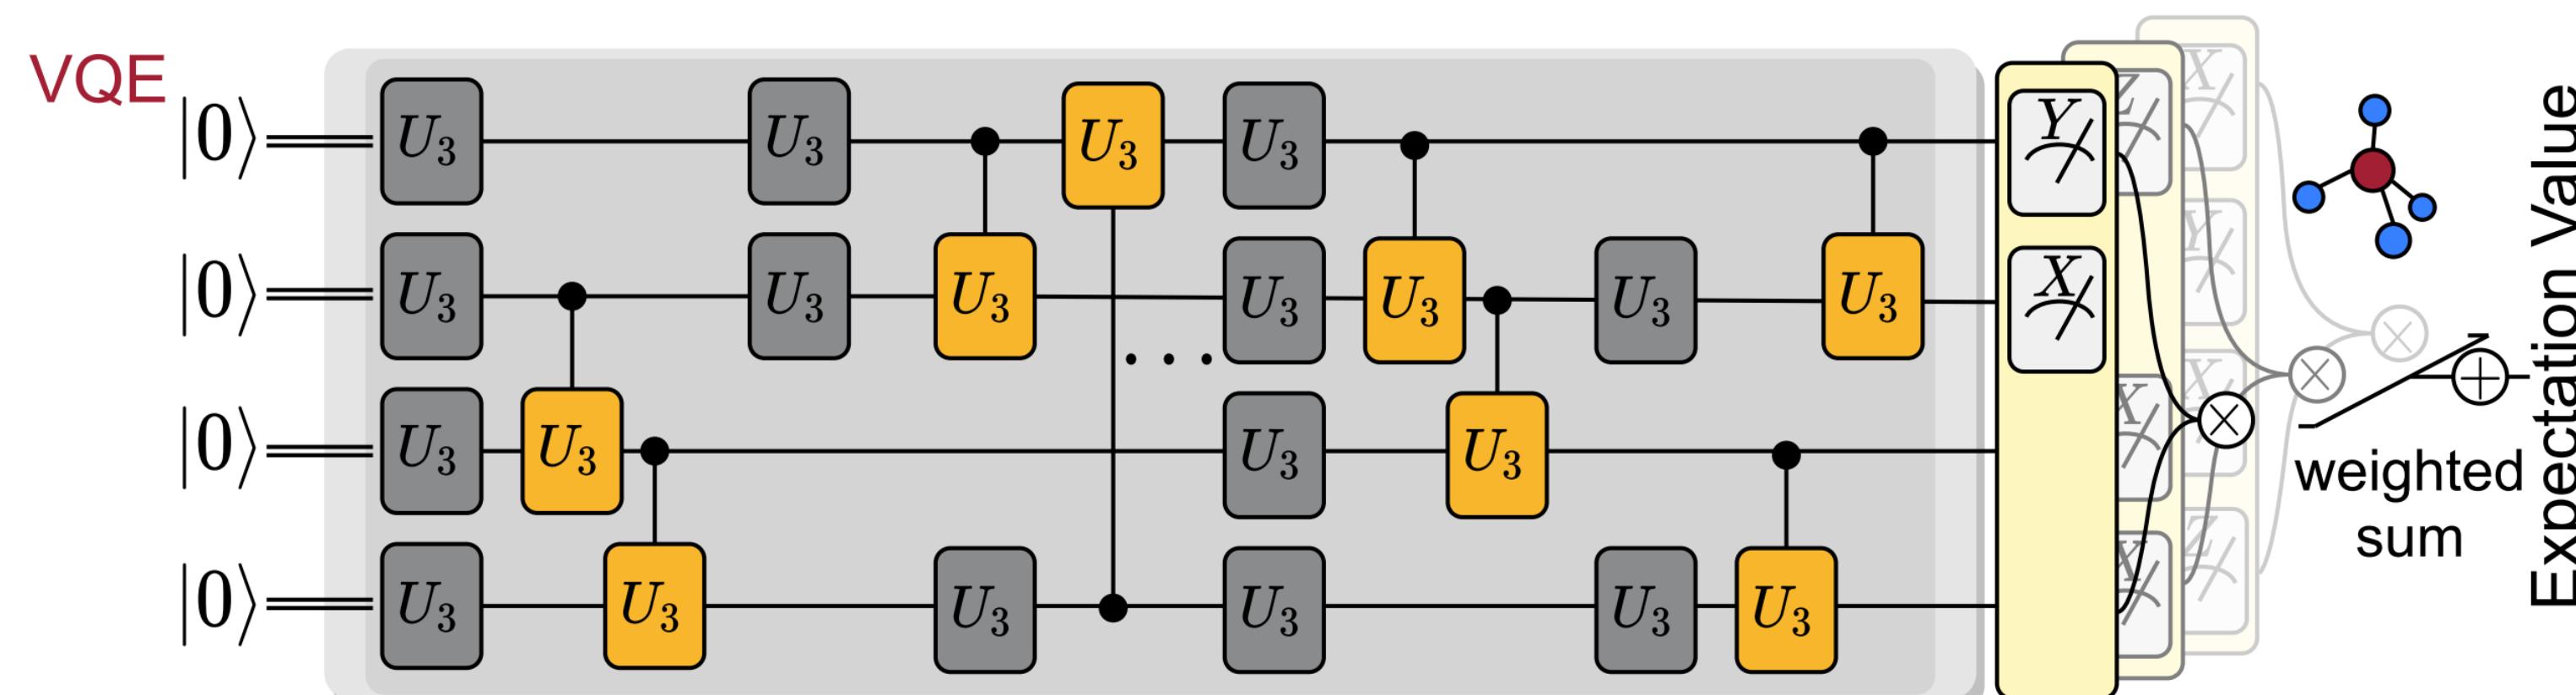


Benchmarks: QNN and VQE

- Quantum Neural Networks: classification

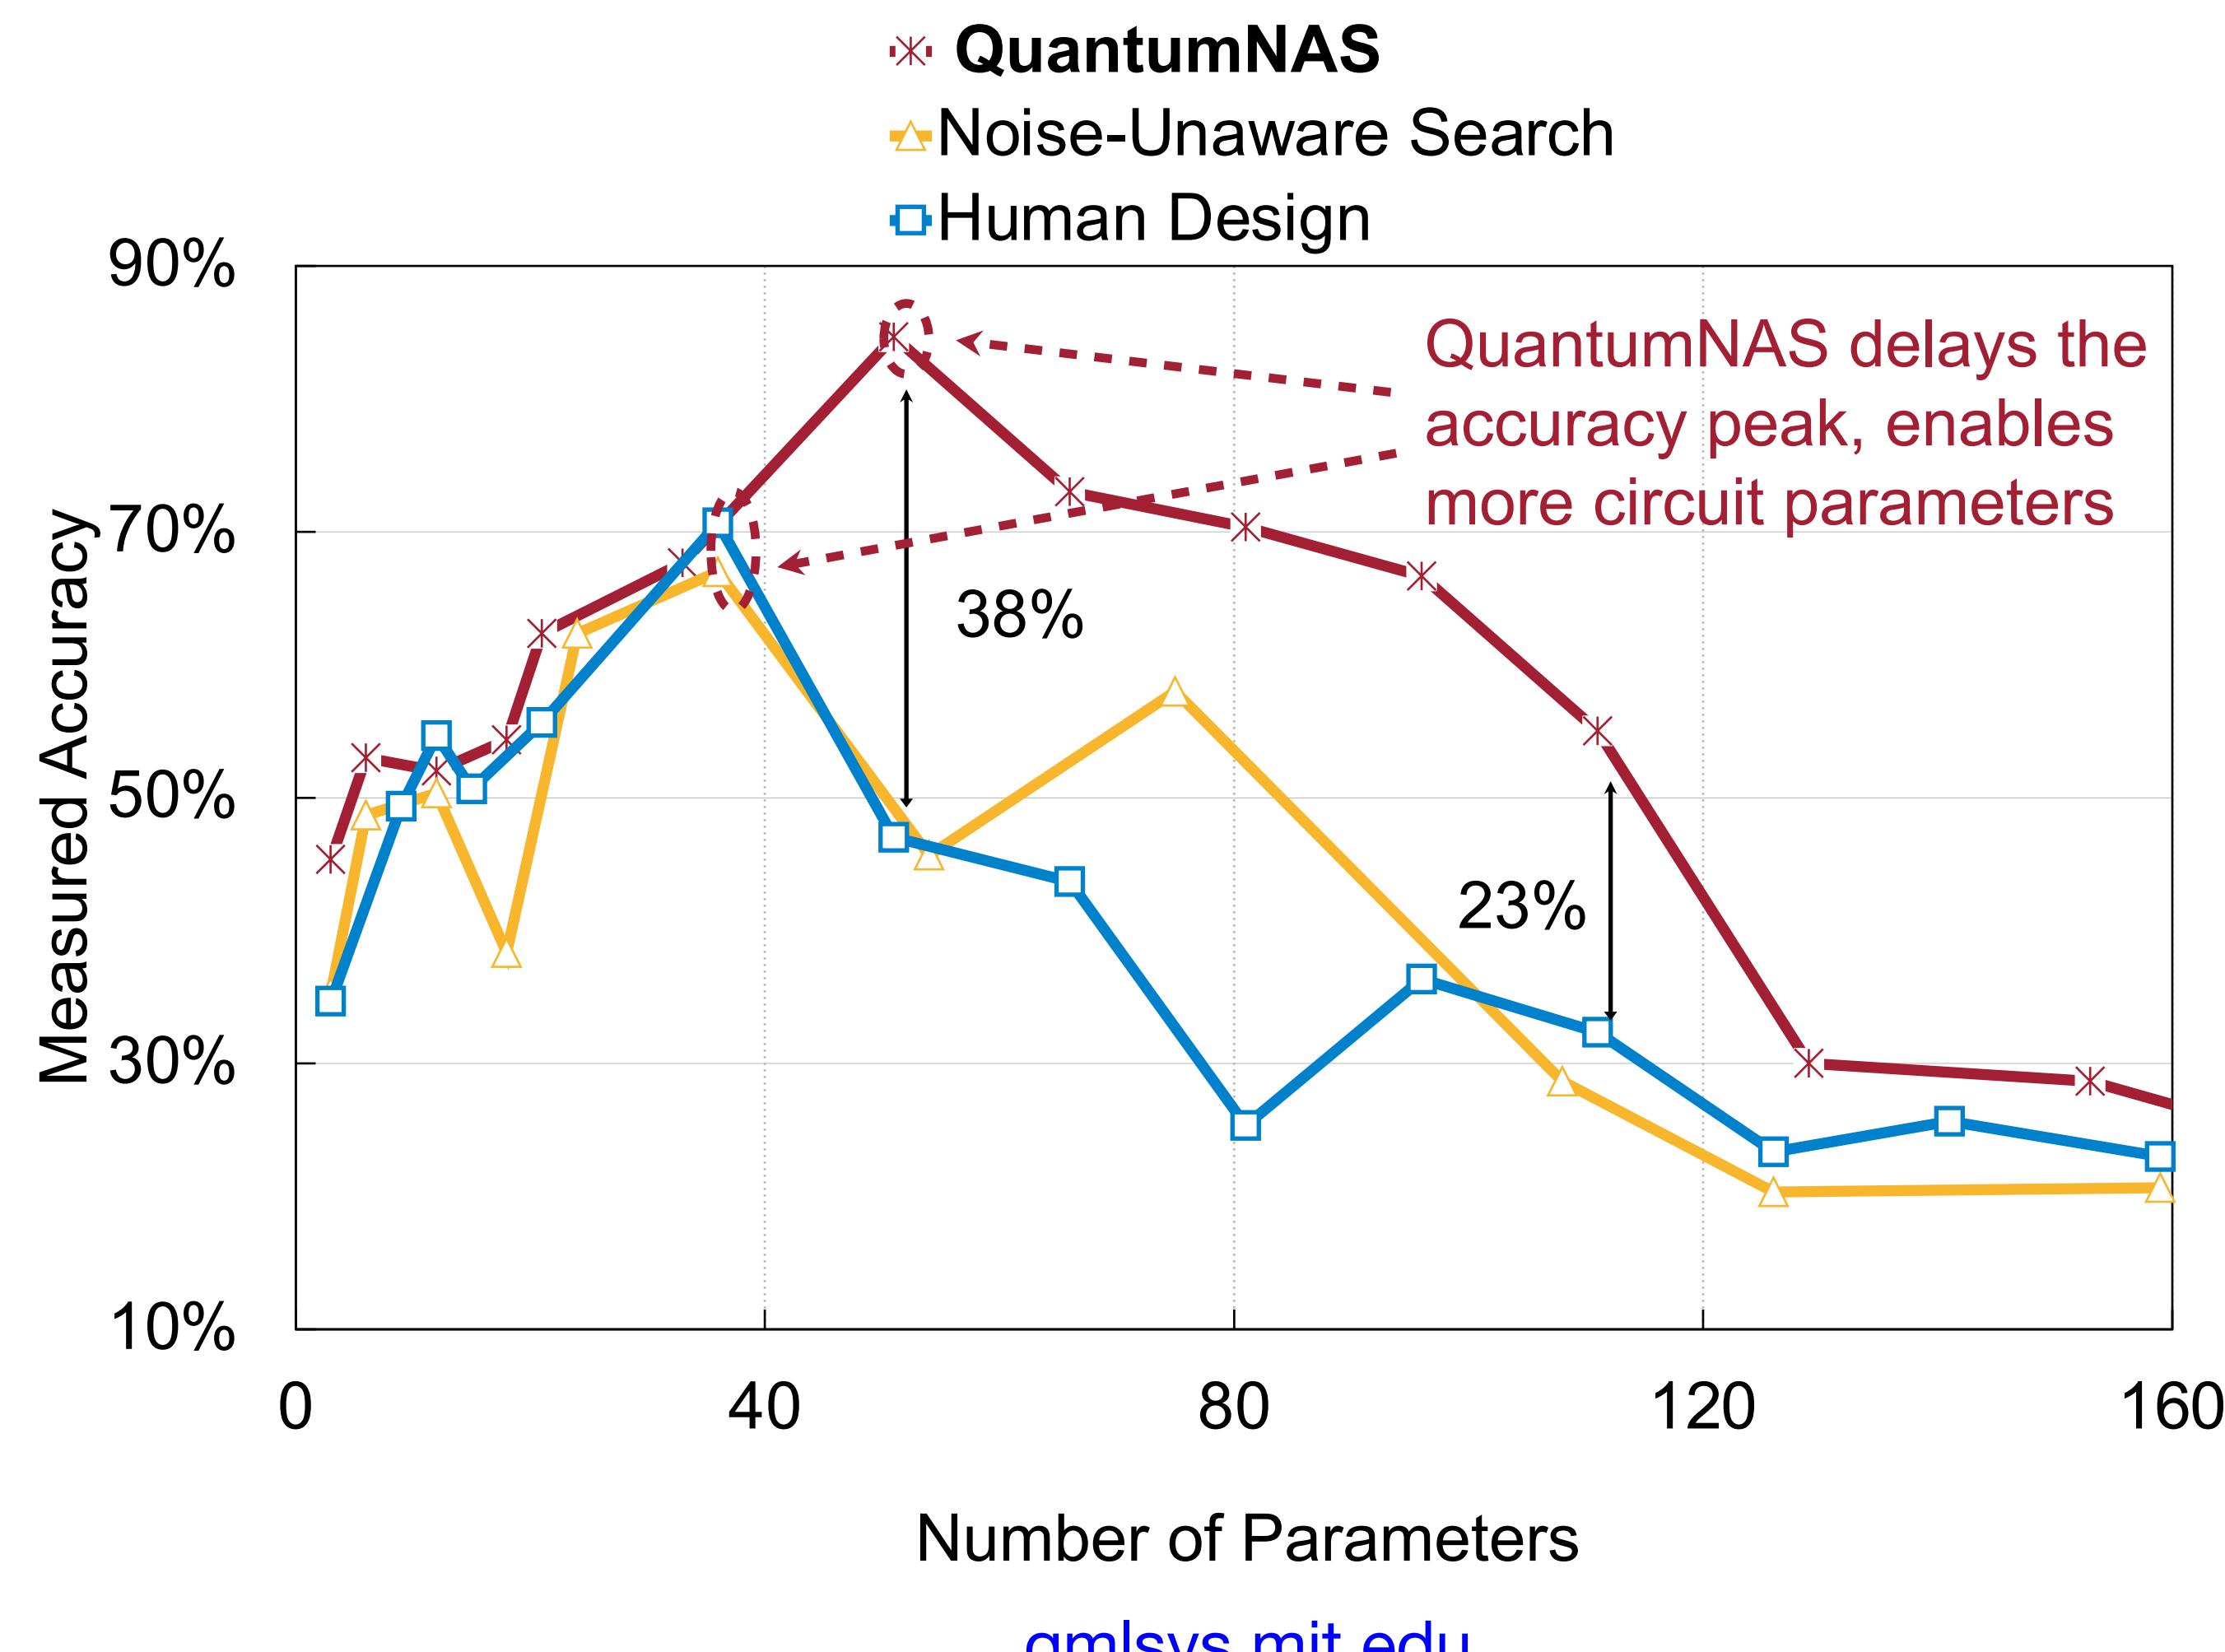


- Variational Quantum Eigensolver: finds the ground state energy of molecule Hamiltonian



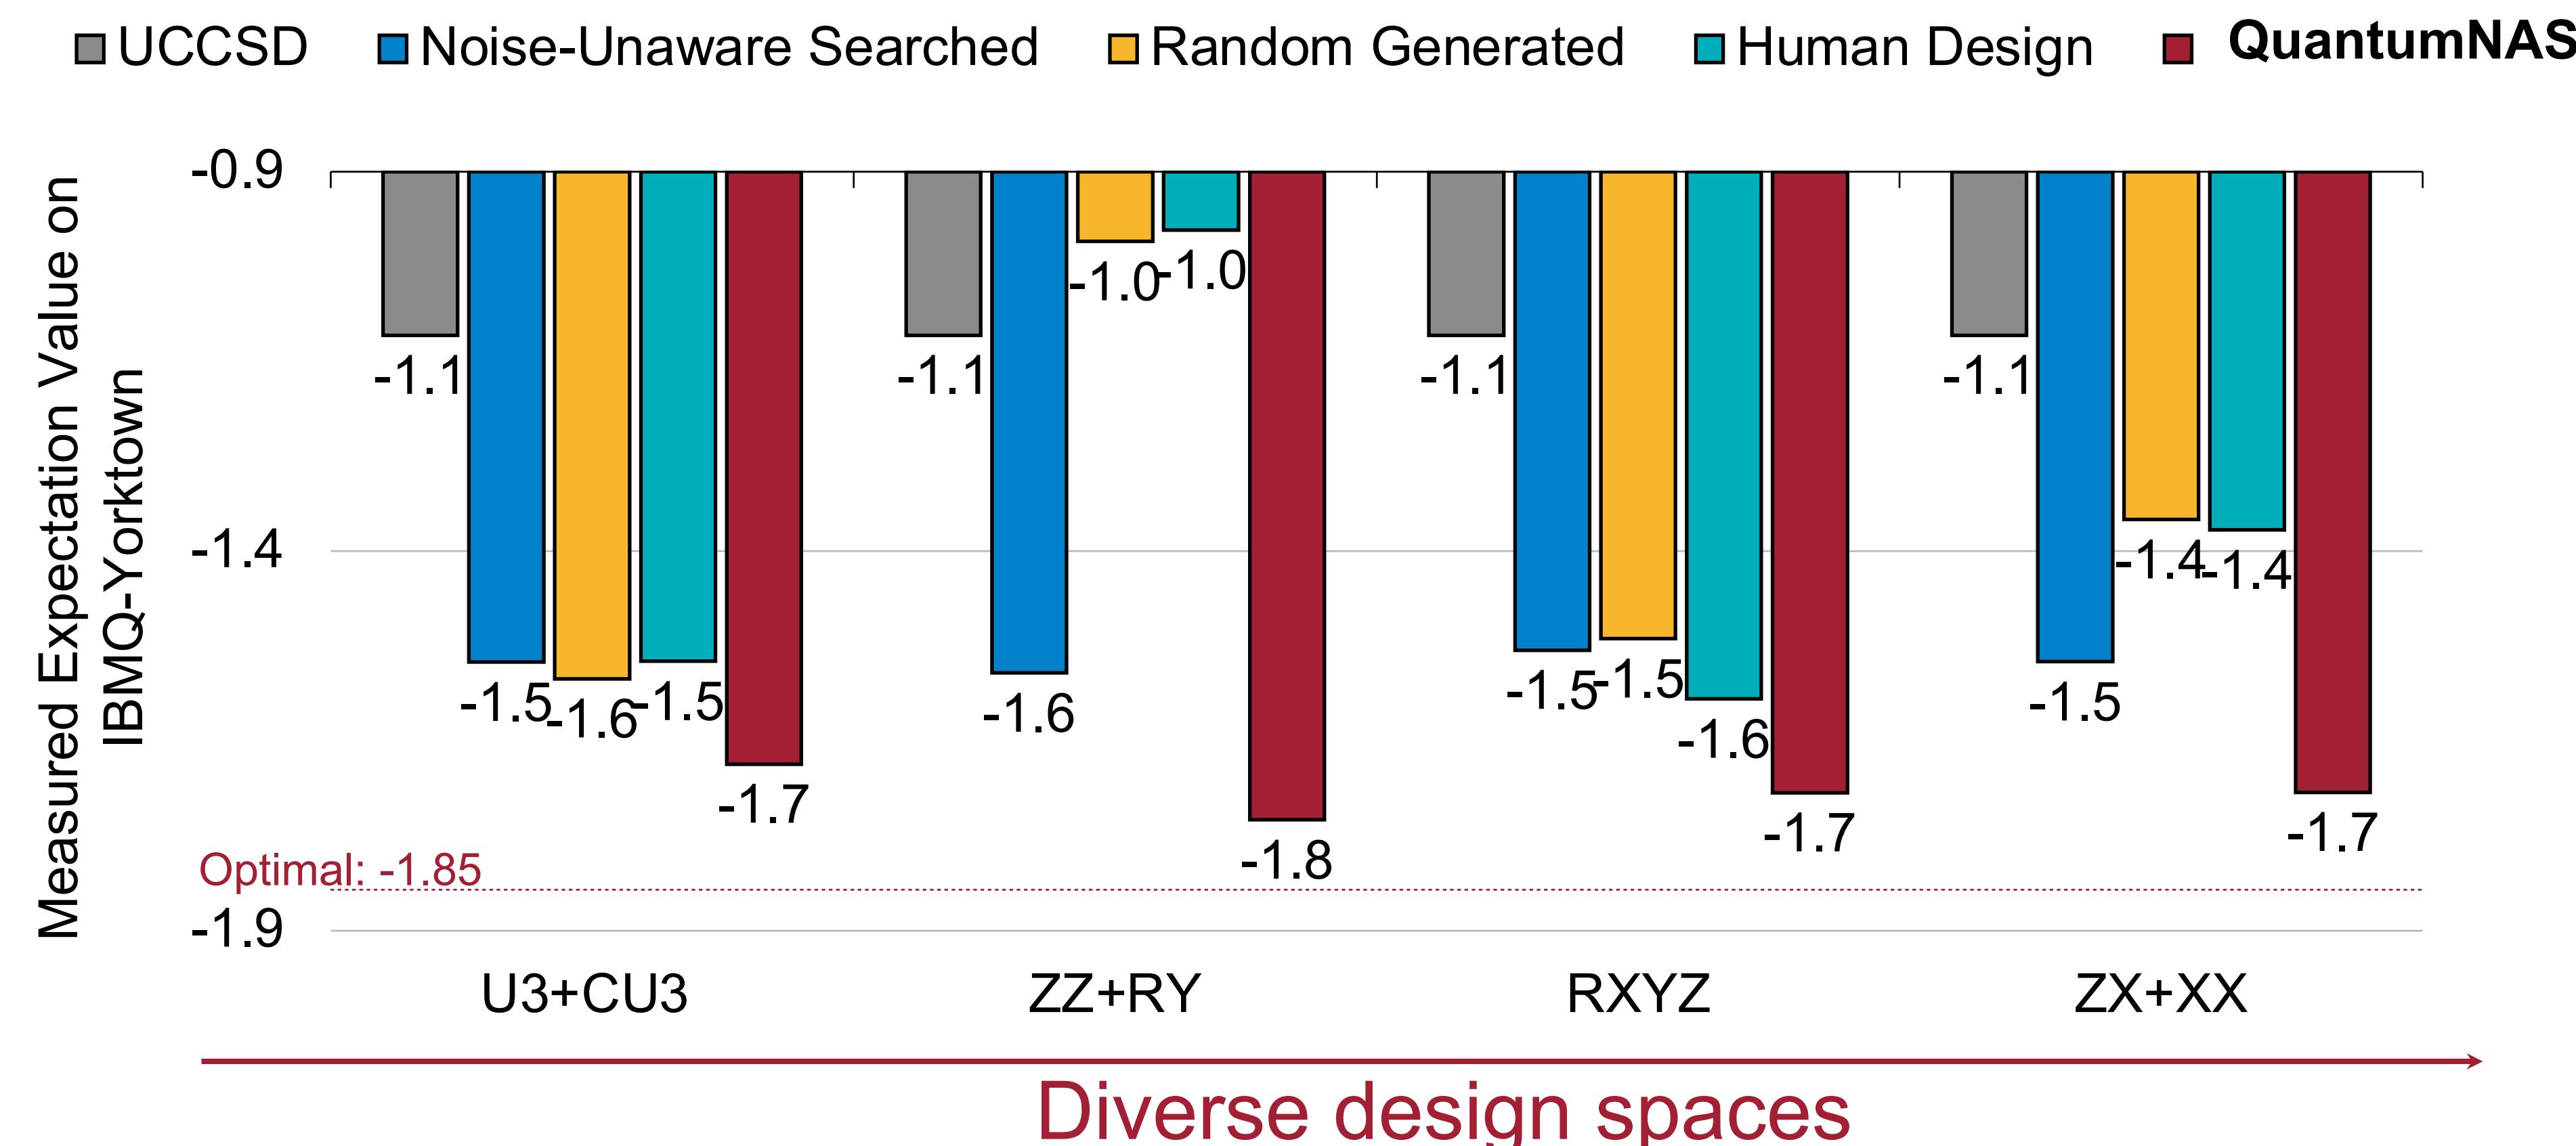
QML Results

- 4-classification: MNIST-4 U3+CU3 on IBMQ-Yorktown



Consistent Improvements on Diverse Design Spaces

- H2 in different design spaces on IBMQ-Yorktown



Scalable to Large #Qubits

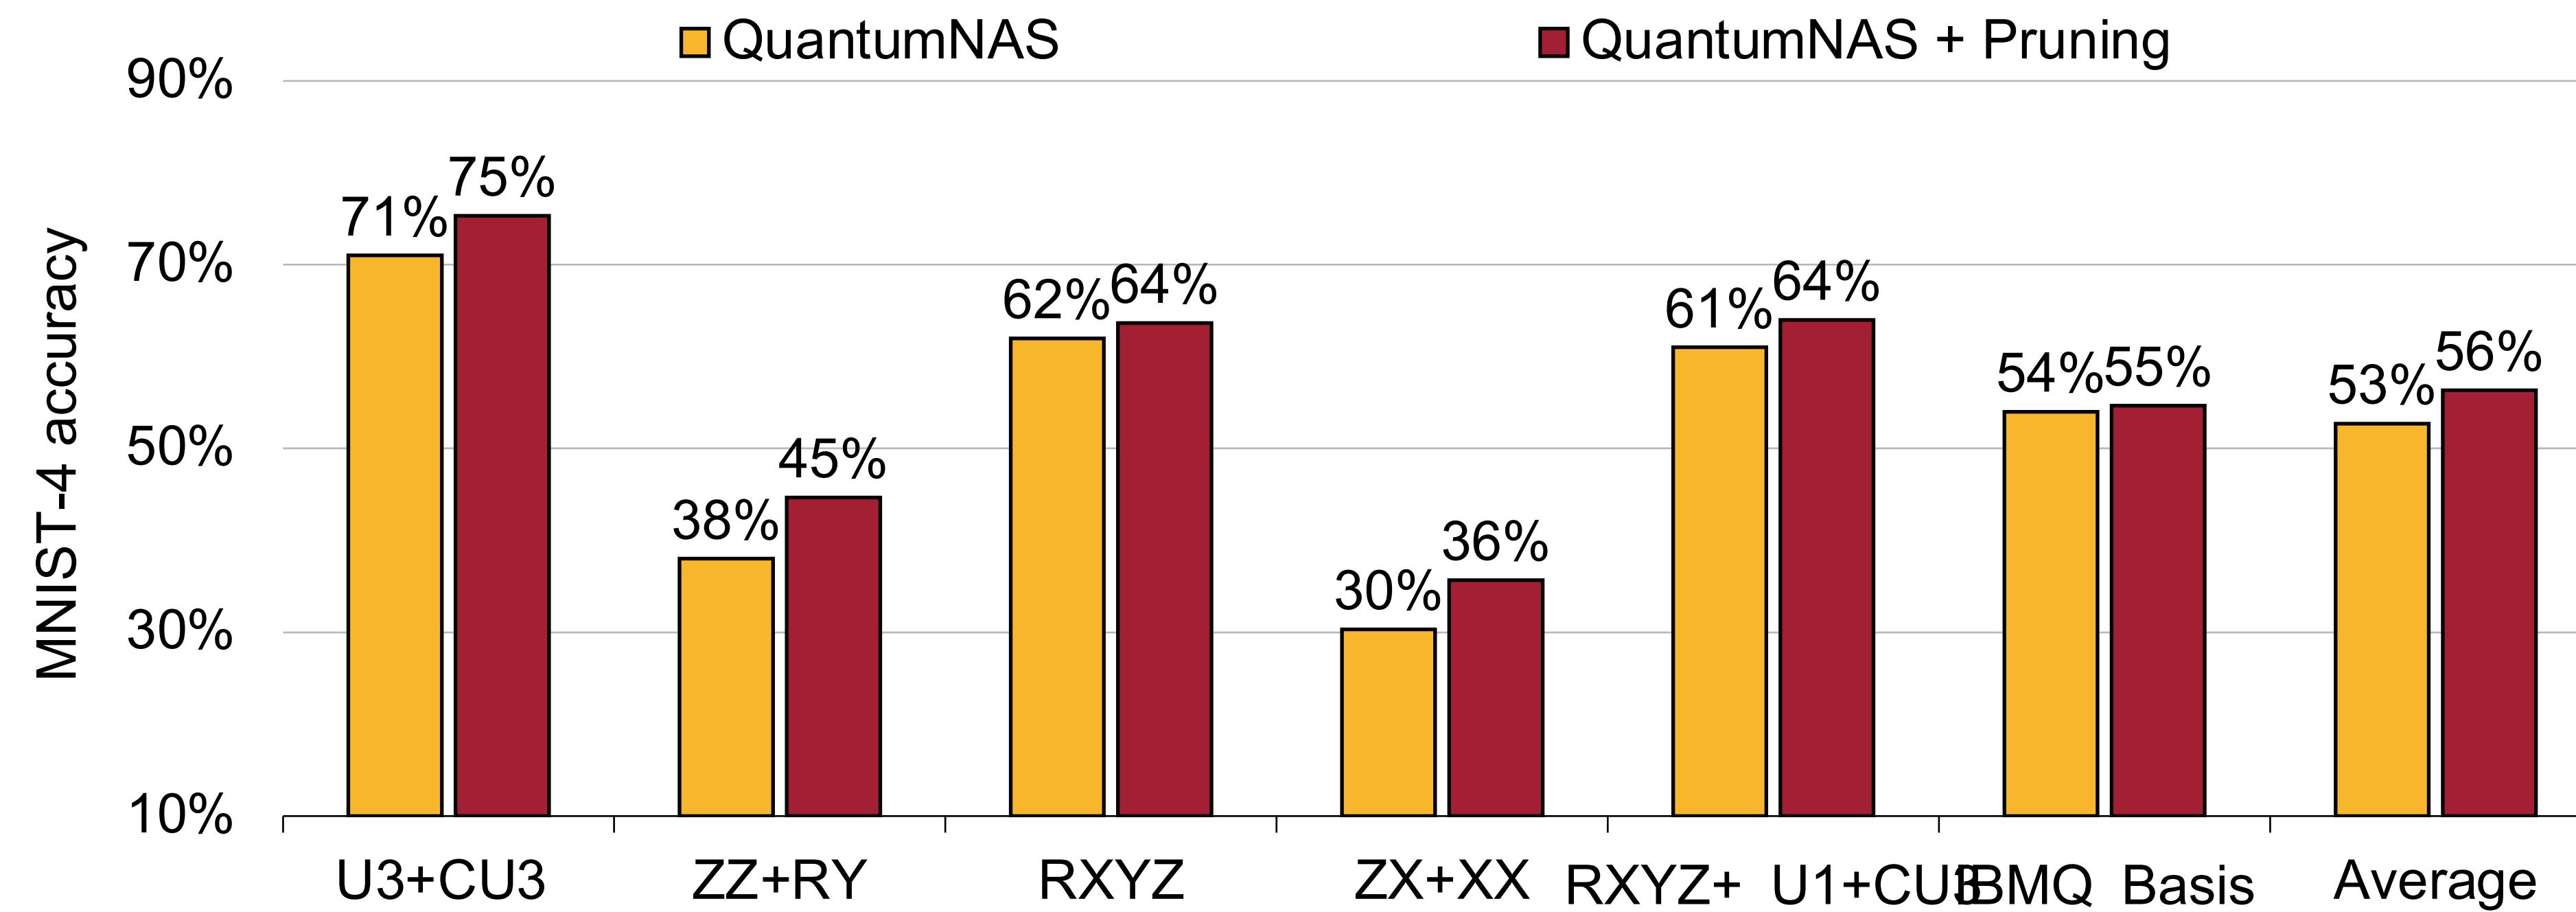
- On large devices
- MNIST-10 accuracy

More Qubits

Method	Noise-Unaware Searched	Random	Human	QuantumNAS
Melbourne (15Q, 8QV, use 15Q)	11%	10%	15%	32%
Guadalupe (16Q, 32QV, use 16Q)	14%	12%	10%	15%
Montreal (27Q, 128QV, use 21Q)	13%	7%	14%	16%
Manhattan (65Q, 32QV, use 21Q)	11%	11%	15%	18%

Effectiveness of Quantum Gate Pruning

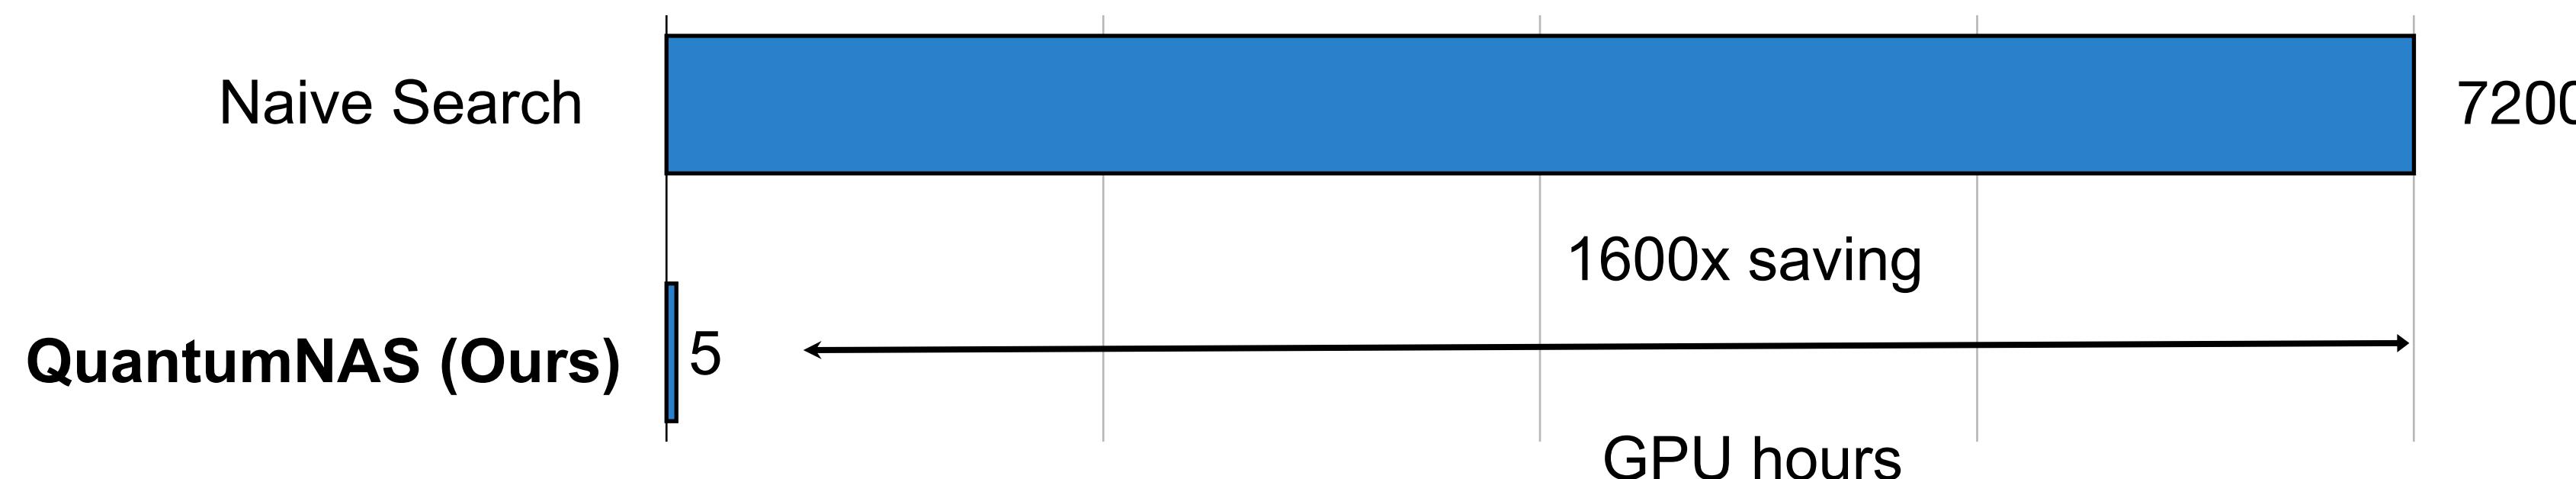
- For MNIST-4, Quantum gate pruning improves accuracy by 3% on average



Time Cost

- On 1 Nvidia Titan RTX 2080 ti GPU

#qubits	Step	SuperCircuit Training	Noise-Adaptive Co-search	SubCircuit Training	Deployment on Real QC
4 Qubits		0.5h	3h	0.5h	0.5h
15 Qubits		5h	5h	5h	1h
21 Qubits		20h	10h	15h	1h





Torch
Quantum

Open-source: TorchQuantum

- TorchQuantum — An open-source library for differentiable quantum simulation

<https://github.com/mit-han-lab/torchquantum>

TorchQuantum

- Features
 - Easy construction and simulation of quantum circuits in **PyTorch**
 - **Dynamic** computation graph for easy debugging
 - **Gradient** support via autograd
 - **Batch** mode inference and training on CPU/GPU.
 - Easy **deployment** on real quantum devices such as IBMQ
 - Easy **hybrid** classical-quantum model construction
 - (coming soon) **pulse-level** simulation
 - Tutorials, videos and example projects of QML and using ML to optimize quantum computer system problems
- Will be on **QCE** (Quantum Computing Engineering) tutorials

Open-source: TorchQuantum

- Statevector

```
_state = torch.zeros(2 ** self.n_wires, dtype=C_DTYPE)
_state[0] = 1 + 0j
```

- Quantum Gates

```
'cnot': torch.tensor([[1, 0, 0, 0],
                      [0, 1, 0, 0],
                      [0, 0, 0, 1],
                      [0, 0, 1, 0]], dtype=C_DTYPE),
```

Open-source: TorchQuantum

- Quantum Gates

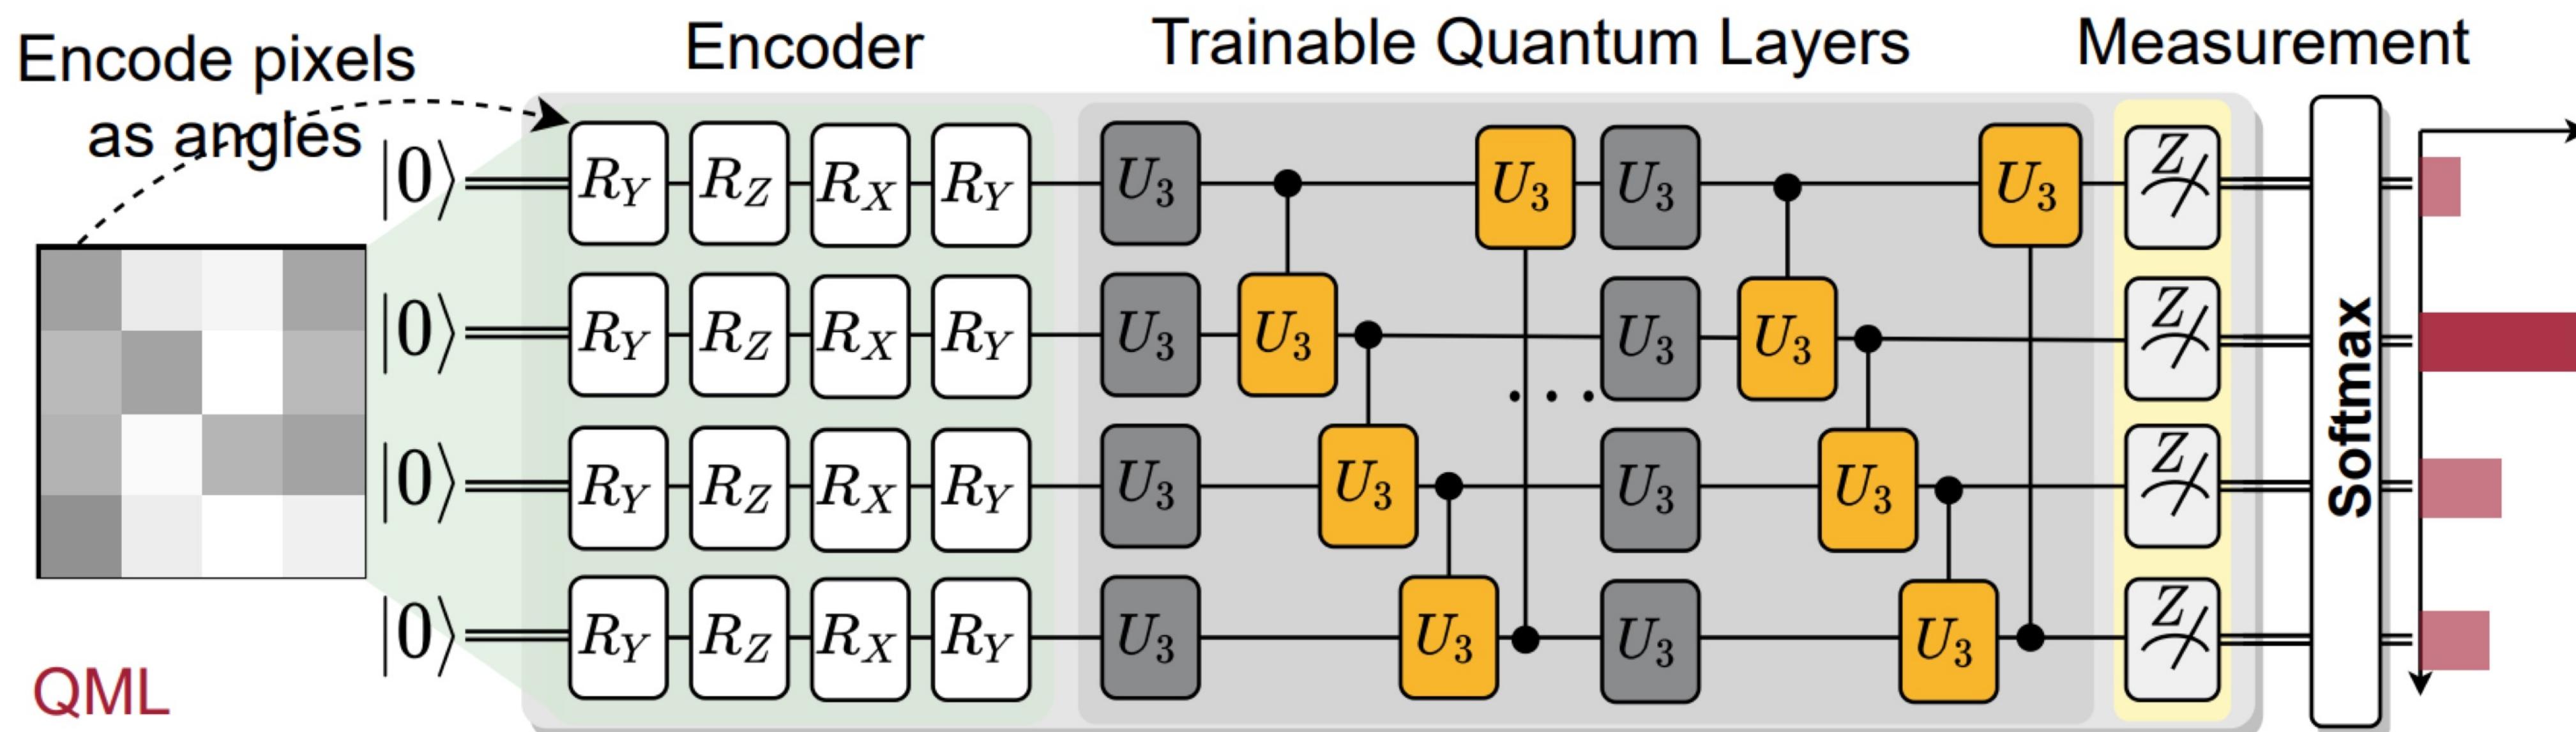
```
def crx_matrix(params):
    theta = params.type(C_DTYPE)
    co = torch.cos(theta / 2)
    jsi = 1j * torch.sin(-theta / 2)

    matrix = torch.tensor([[1, 0, 0, 0],
                          [0, 1, 0, 0],
                          [0, 0, 0, 0],
                          [0, 0, 0, 0]], dtype=C_DTYPE, device=params.device
                         ).unsqueeze(0).repeat(co.shape[0], 1, 1)
    matrix[:, 2, 2] = co[:, 0]
    matrix[:, 2, 3] = jsi[:, 0]
    matrix[:, 3, 2] = jsi[:, 0]
    matrix[:, 3, 3] = co[:, 0]

    return matrix.squeeze(0)
```

- Matrix-vector multiplication: torch.einsum and torch.bmm

MNIST Example with QNN



MNIST Example

Initialize a quantum device

```
import torch.nn as nn
import torch.nn.functional as F
import torchquantum as tq
import torchquantum.functional as tqf

class QFCModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.n_wires = 4
        self.q_device = tq.QuantumDevice(n_wires=self.n_wires)
        self.measure = tq.MeasureAll(tq.PauliZ)
```

Specify encoder gates

```
self.encoder_gates = [tqf.rx] * 4 + [tqf.ry] * 4 + \
                     [tqf.rz] * 4 + [tqf.rx] * 4
self.rx0 = tq.RX(has_params=True, trainable=True)
self.ry0 = tq.RY(has_params=True, trainable=True)
self.rz0 = tq.RZ(has_params=True, trainable=True)
self.crx0 = tq.CRX(has_params=True, trainable=True)
```

Specify trainable gates

MNIST Example

Reset statevector

```
def forward(self, x):
    bsz = x.shape[0]
    # down-sample the image
    x = F.avg_pool2d(x, 6).view(bsz, 16)

    # reset qubit states
    self.q_device.reset_states(bsz)
```

Encode classical pixels

```
# encode the classical image to quantum domain
for k, gate in enumerate(self.encoder_gates):
    gate(self.q_device, wires=k % self.n_wires, params=x[:, k])
```

Apply the trainable gates

```
# add some trainable gates (need to instantiate ahead of time)
self.rx0(self.q_device, wires=0)
self.ry0(self.q_device, wires=1)
self.rz0(self.q_device, wires=3)
self.crx0(self.q_device, wires=[0, 2])
```

Apply some non-trainable gates

```
# add some more non-parameterized gates (add on-the-fly)
tqf.hadamard(self.q_device, wires=3)
tqf.sx(self.q_device, wires=2)
tqf.cnot(self.q_device, wires=[3, 0])
tqf.qubitunitary(self.q_device0, wires=[1, 2], params=[[1, 0, 0, 0],
                                                       [0, 1, 0, 0],
                                                       [0, 0, 0, 1j],
                                                       [0, 0, -1j, 0]])
```

Measure to get classical values

```
# perform measurement to get expectations (back to classical domain)
x = self.measure(self.q_device).reshape(bsz, 2, 2)

# classification
x = x.sum(-1).squeeze()
x = F.log_softmax(x, dim=1)

return x
```

MNIST Example

Initialize a QiskitProcessor

```
# then try to run on REAL QC
backend_name = 'ibmqx2'
print(f"\nTest on Real Quantum Computer {backend_name}")
processor_real_qc = QiskitProcessor(use_real_qc=True,
                                     backend_name=backend_name)
model.set_qiskit_processor(processor_real_qc)
valid_test(dataflow, 'test', model, device, qiskit=True)
```

The 'forward' will be run on real QC

Examples and tutorials

- Tutorial Colab and videos



**TorchQuantum Tutorials
Opening**



Hanrui Wang
MIT HAN Lab



**TorchQuantum Tutorials
Quanvolutional Neural Network**

Zirui Li, Hanrui Wang
MIT HAN Lab



MIT HAN LAB



MIT HAN LAB

Thank you for listening!

- Take home
 - **TorchQuantum**: fast open-source **library** for quantum simulation
 - **QuantumNAS**: framework to search for **noise-robust** circuit architecture



Torch
Quantum

<https://github.com/mit-han-lab/torchquantum>

qmlsys.mit.edu

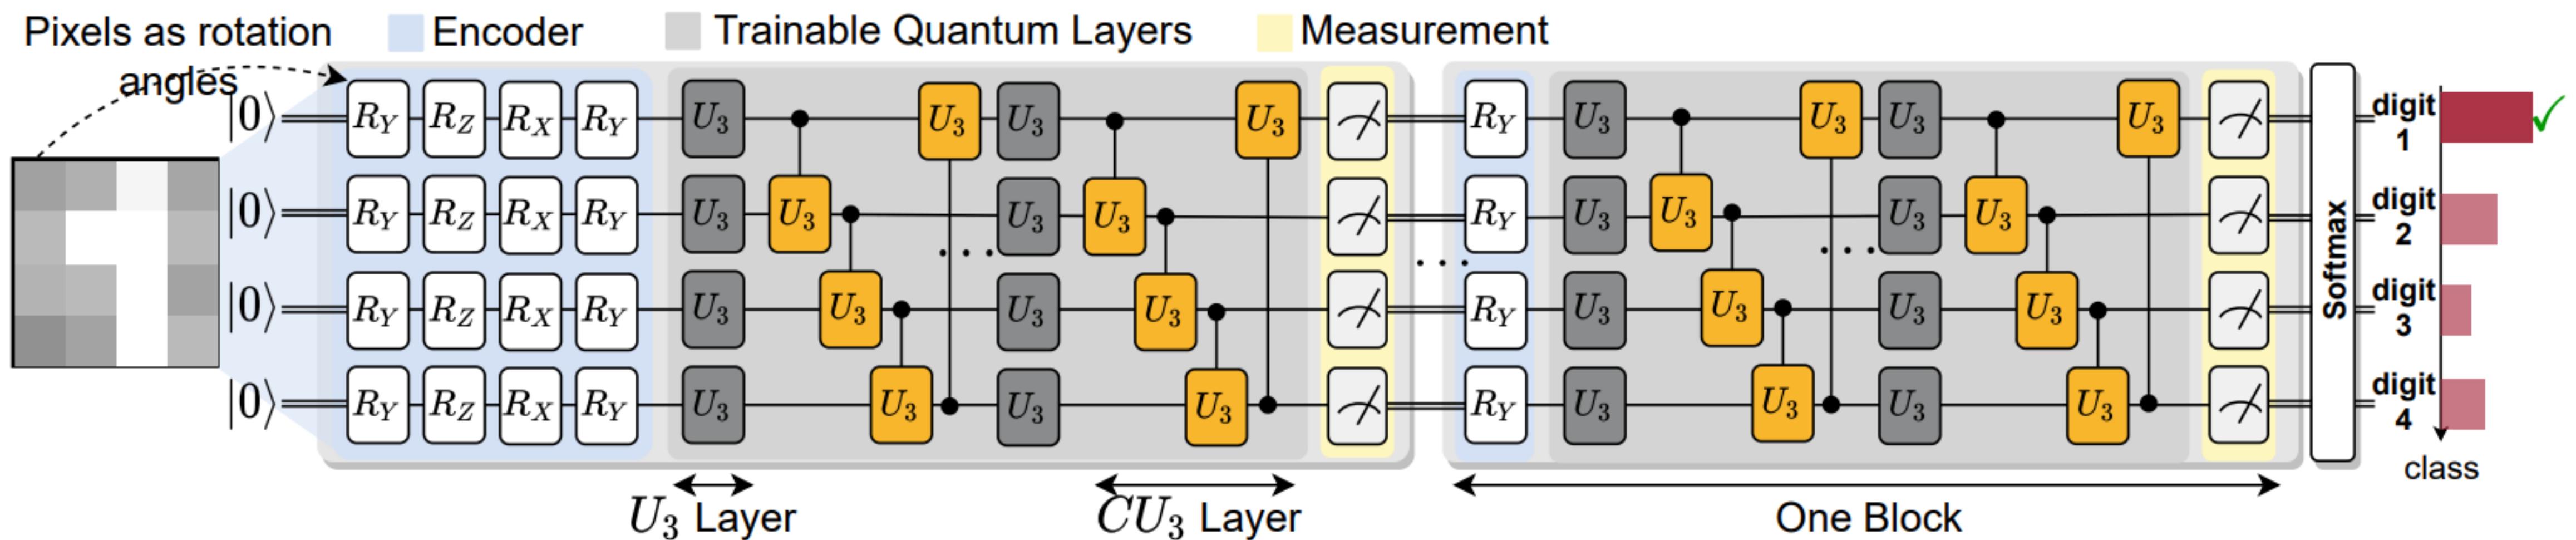


qmlsys.mit.edu

MIT HAN LAB

QuantumNAT

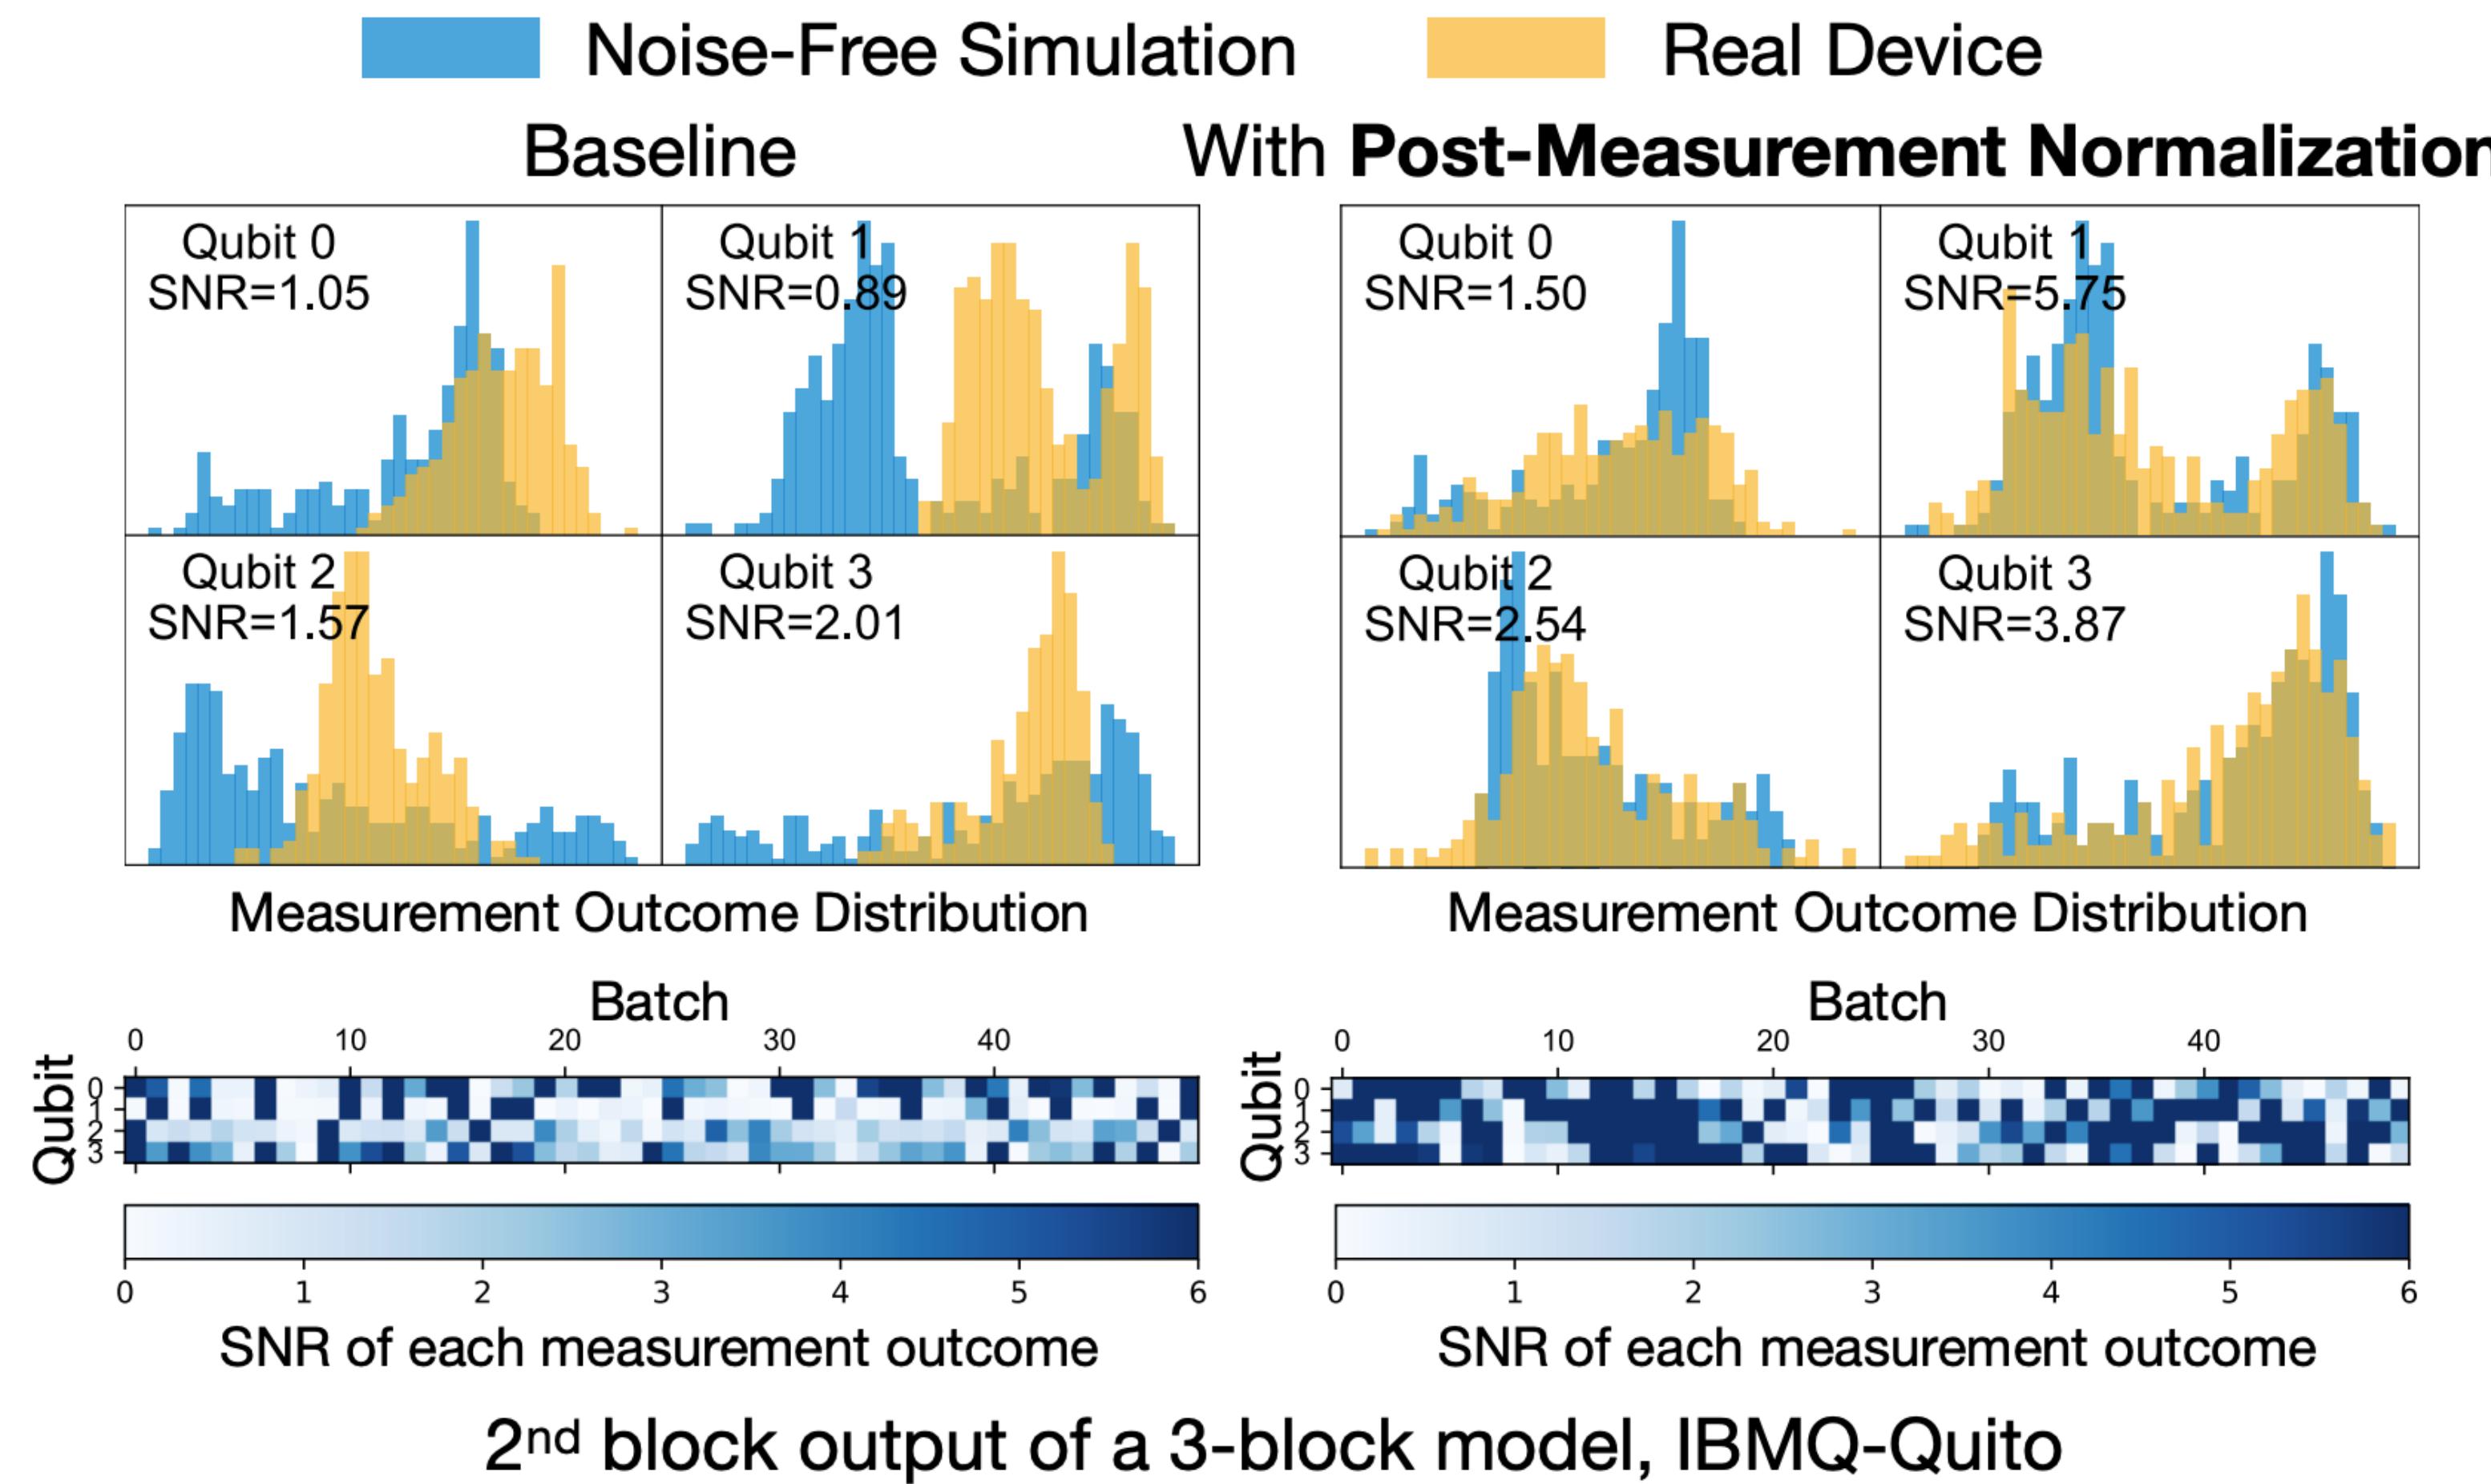
- QuantumNAS: find the circuit architecture robust to noise
- QuantumNAT: further make the parameter robust to noise



[DAC'22] Wang et, at, QuantumNAT: Quantum Noise-Aware Training with Noise Injection, Quantization and Normalization

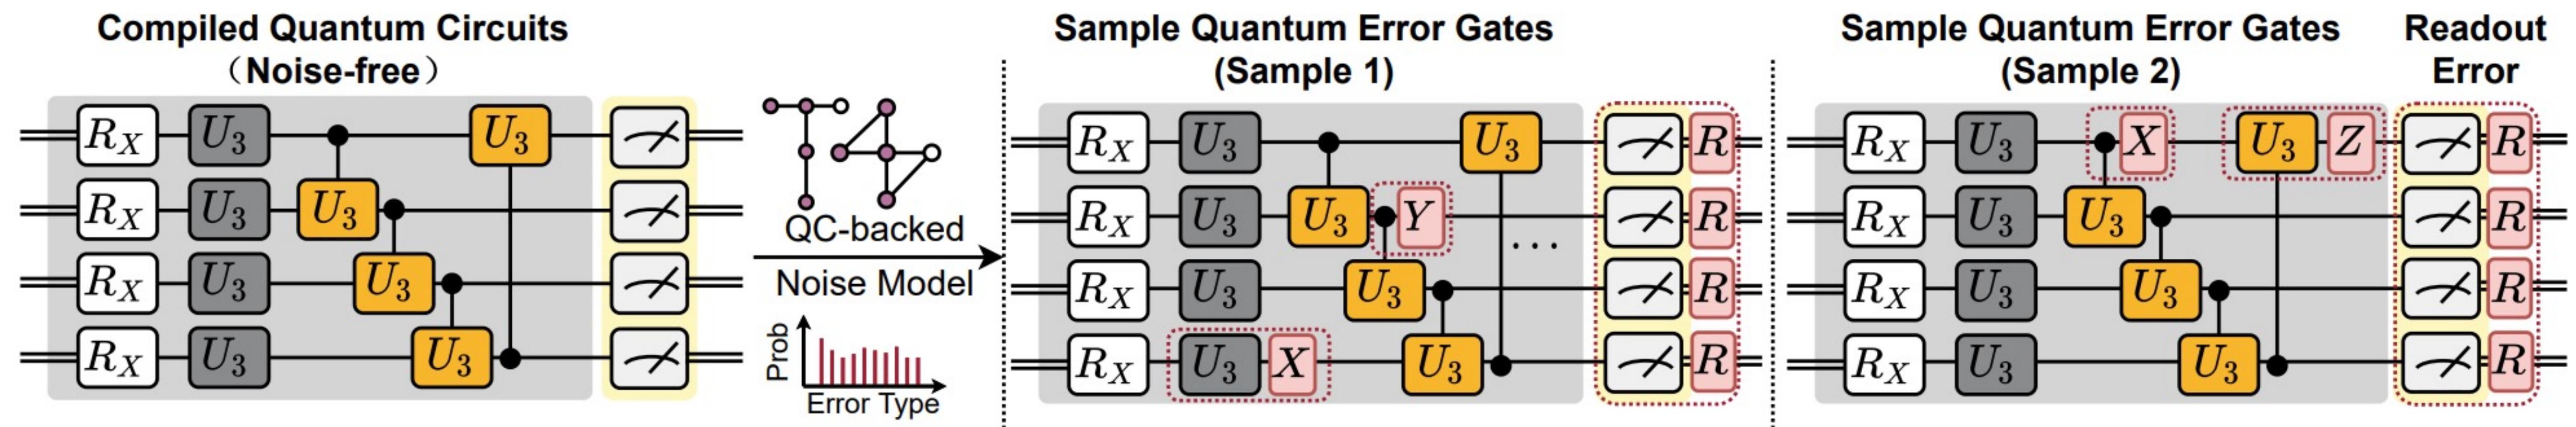
Post-Measurement Normalization

- Normalize the measurement outcomes on the batch dimension



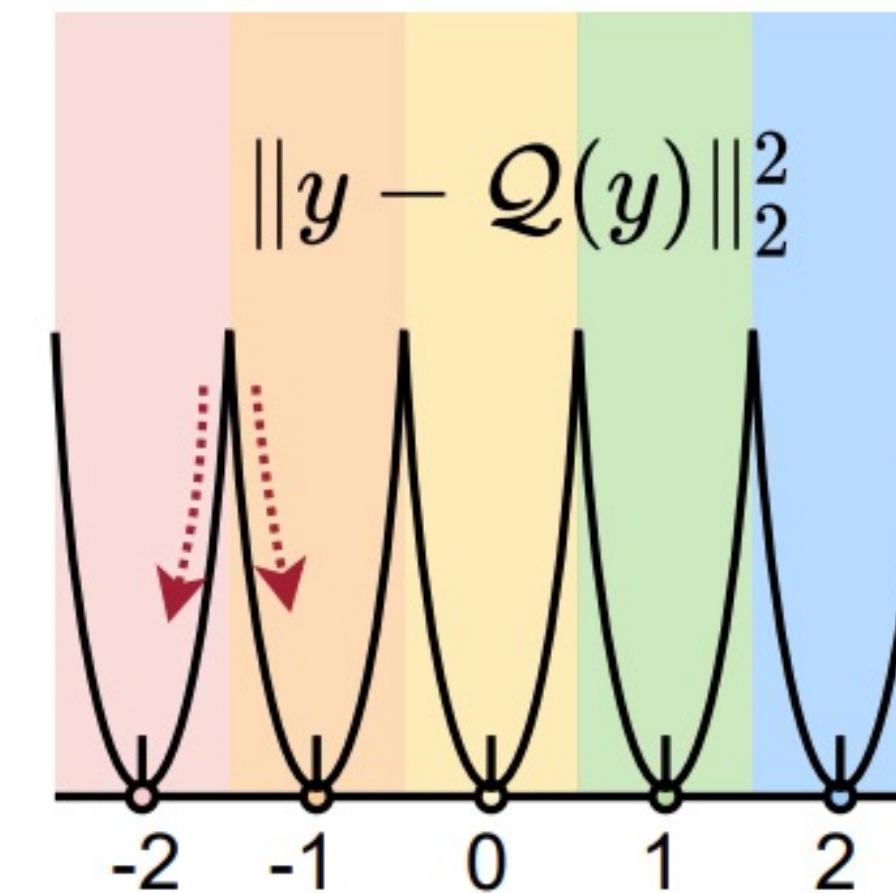
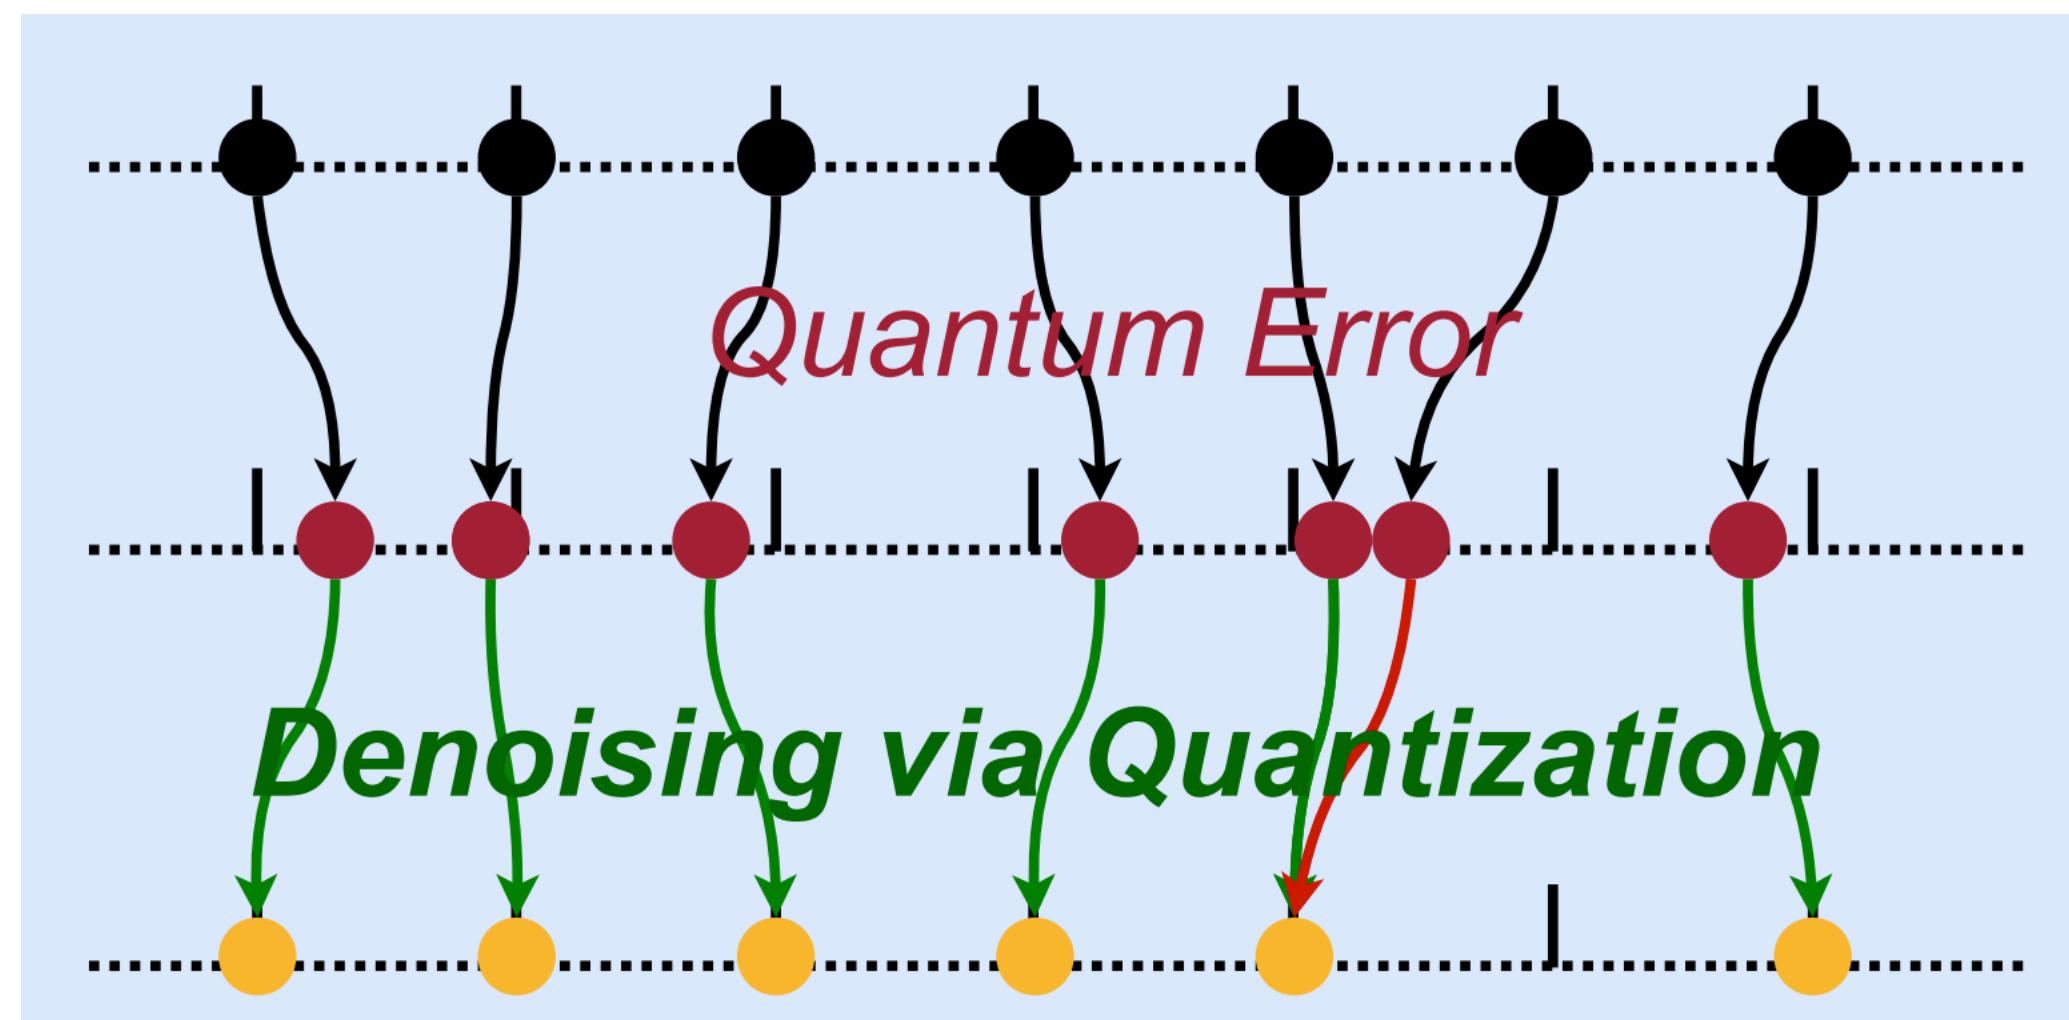
Noise Injection

- Insert noise gate during training, according to the noise model
- For each step, sample new positions for noise gates



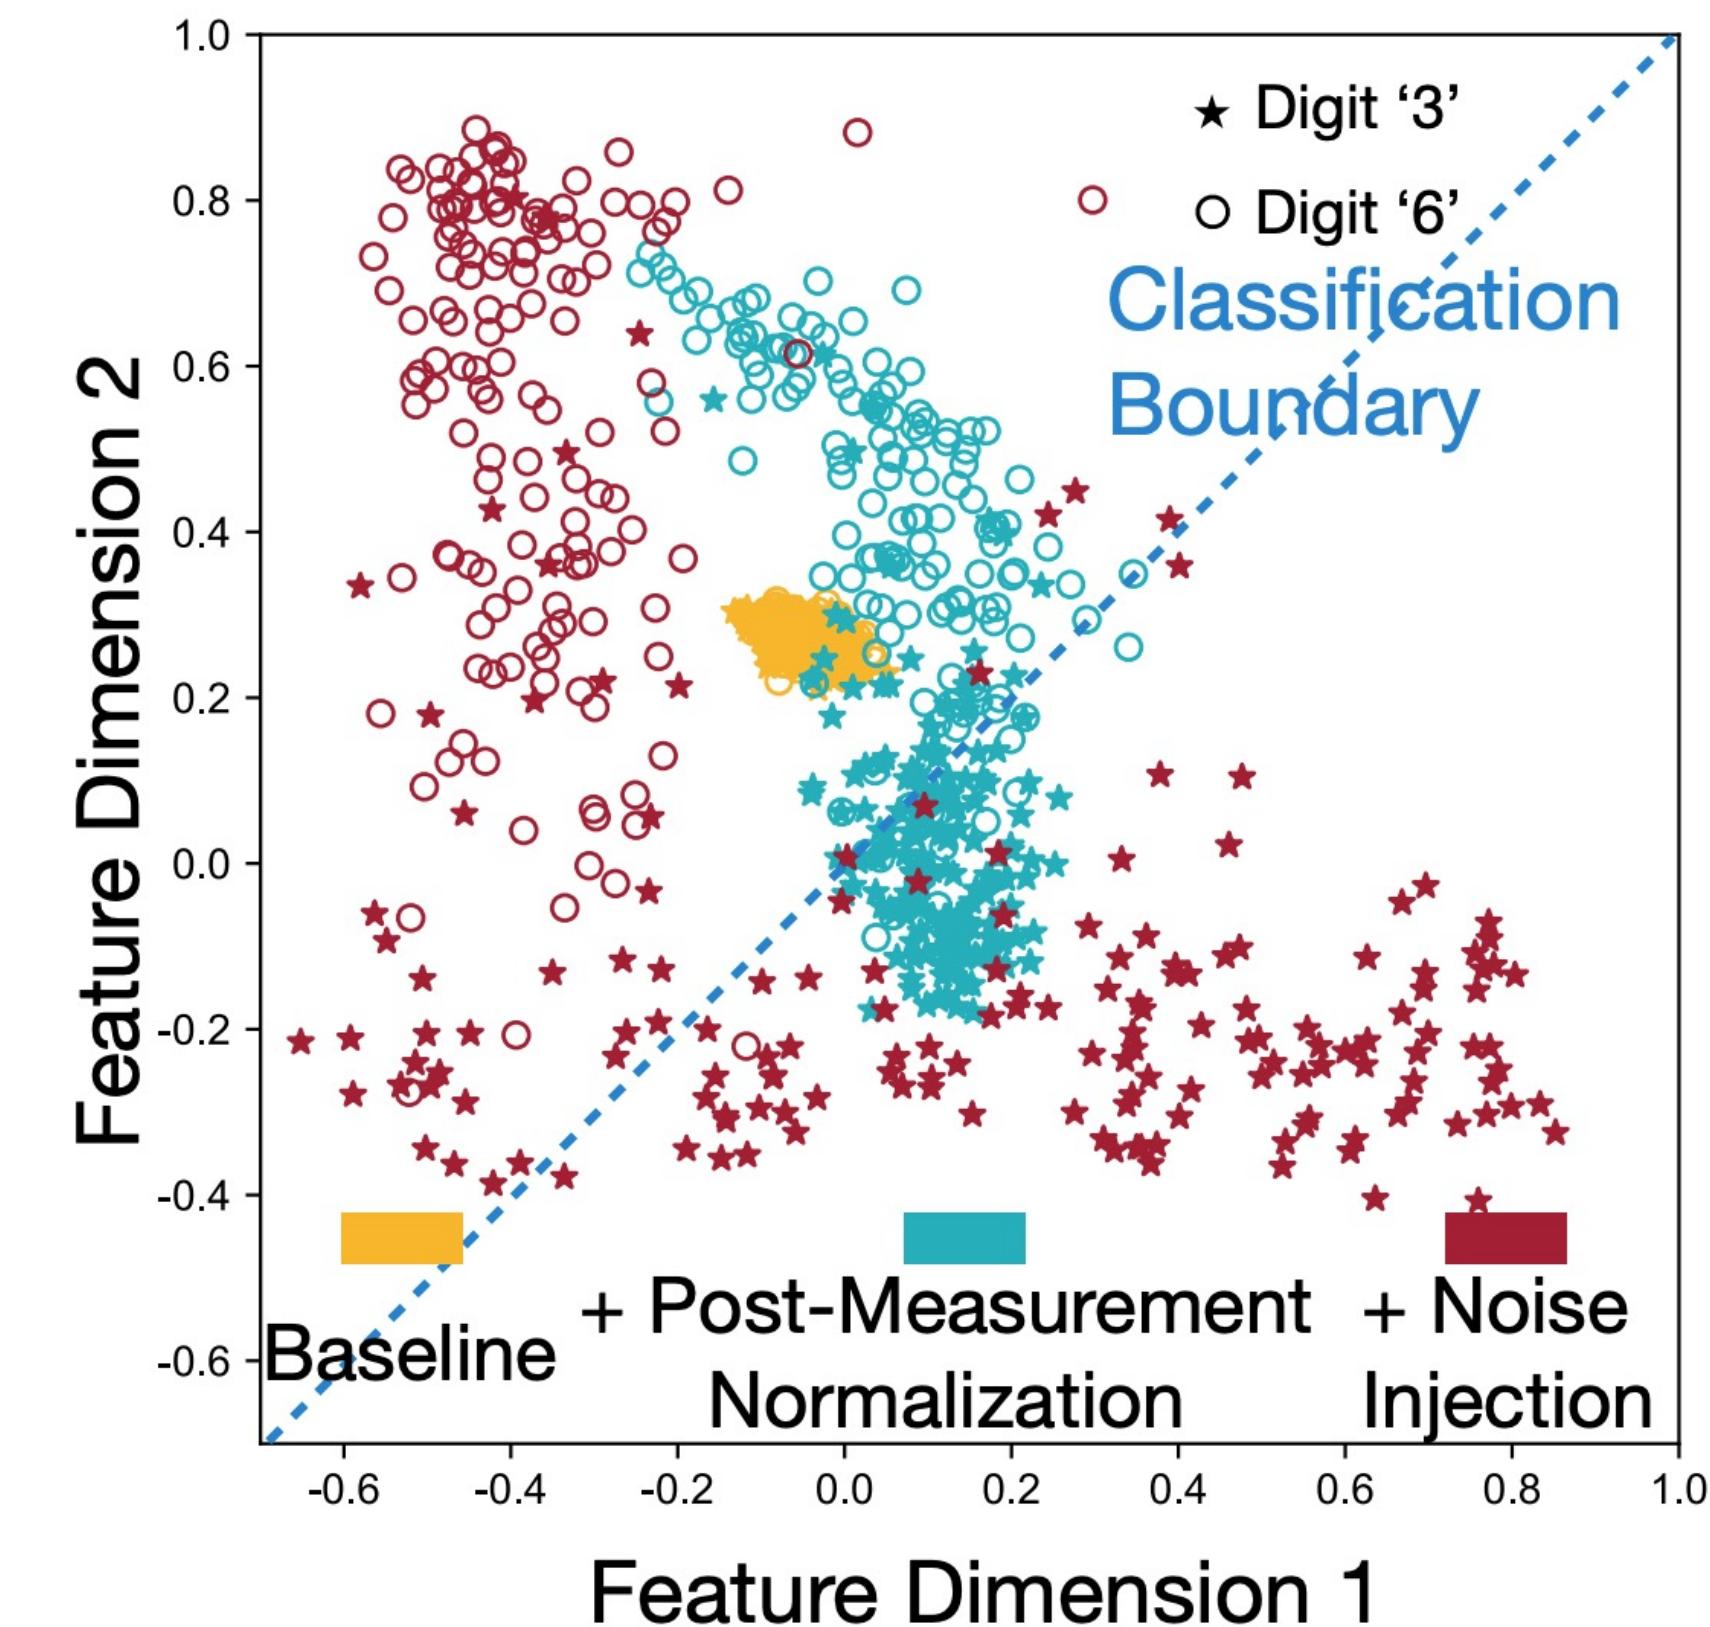
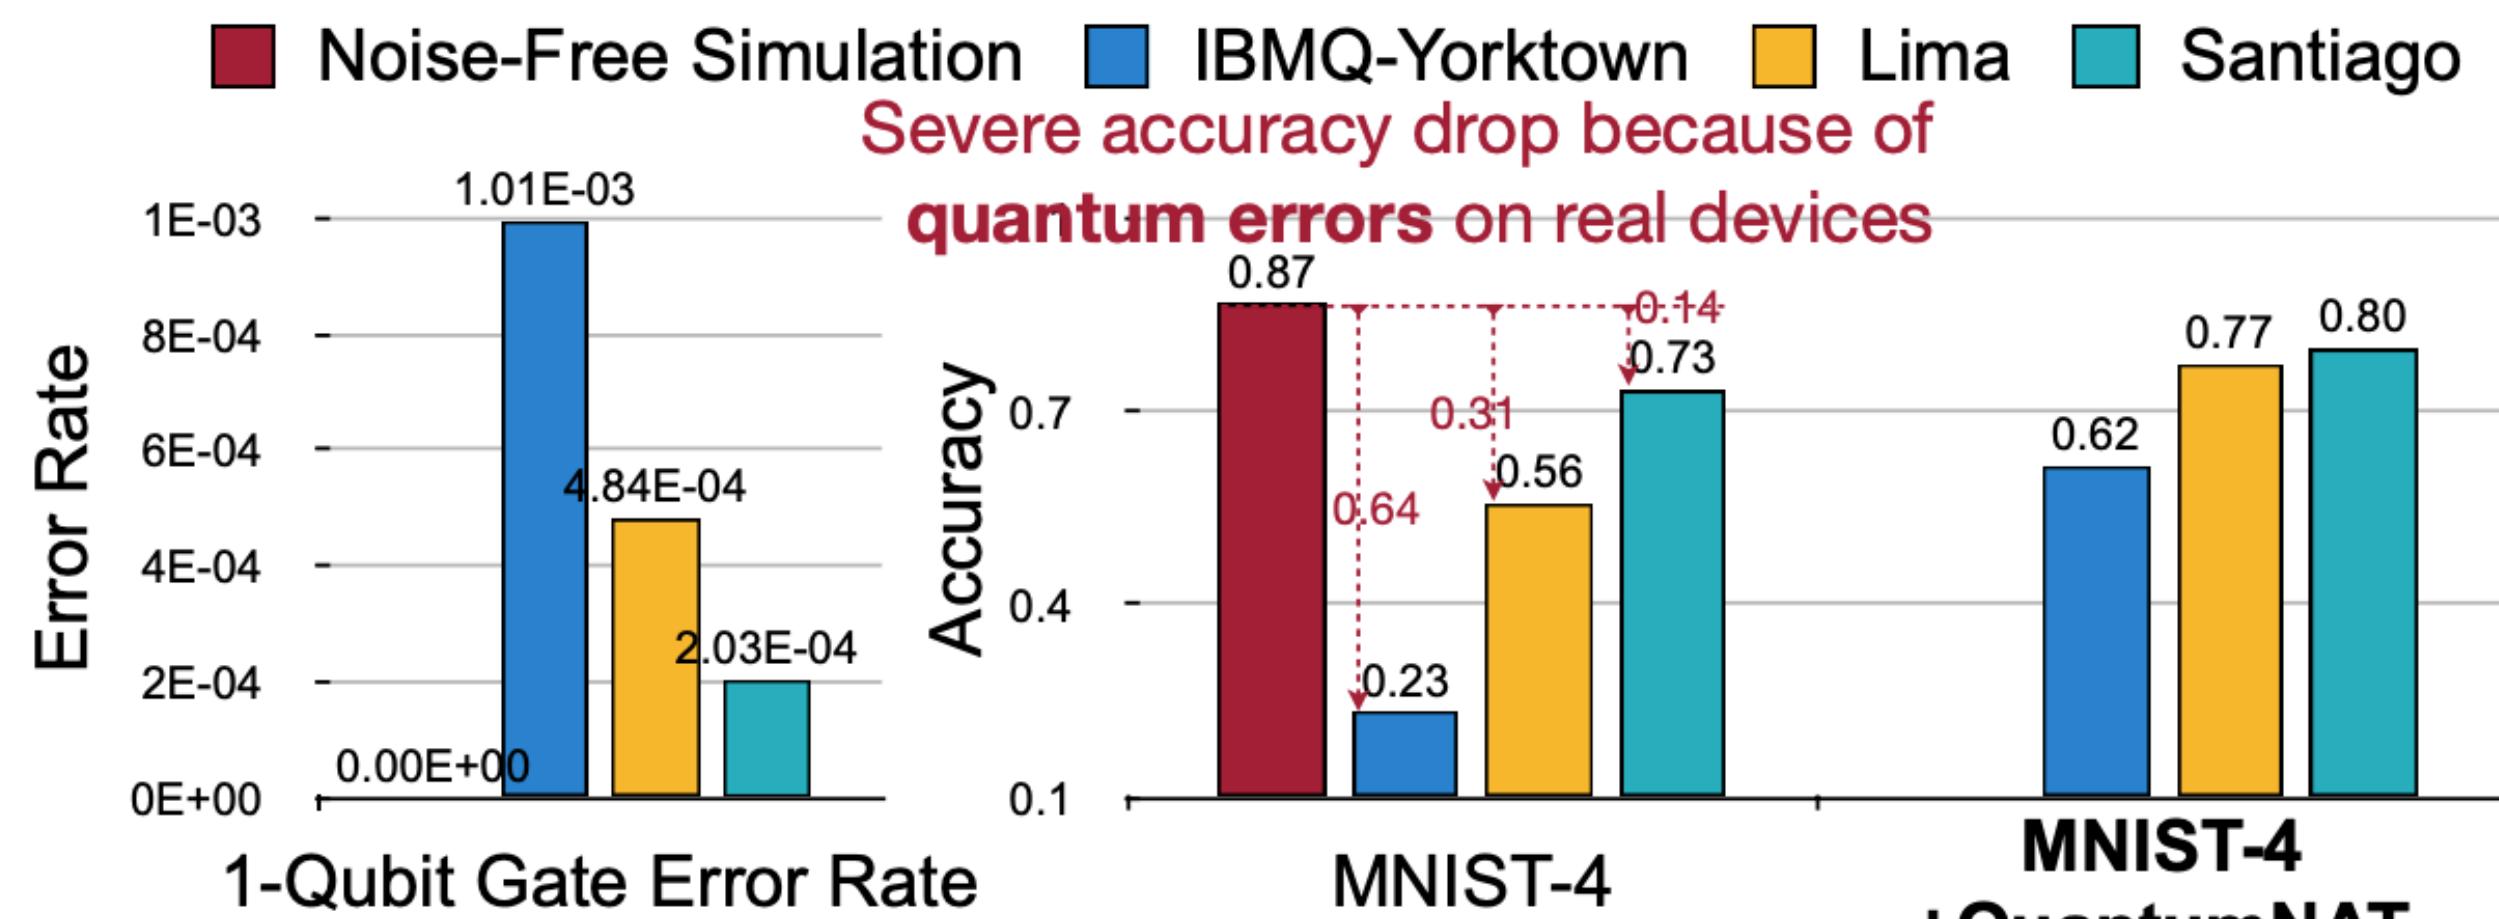
Post-Measurement Quantization

- Quantization provides denoising effects
- Quadratic penalty loss to encourage measurement outcomes close to quantization centroids

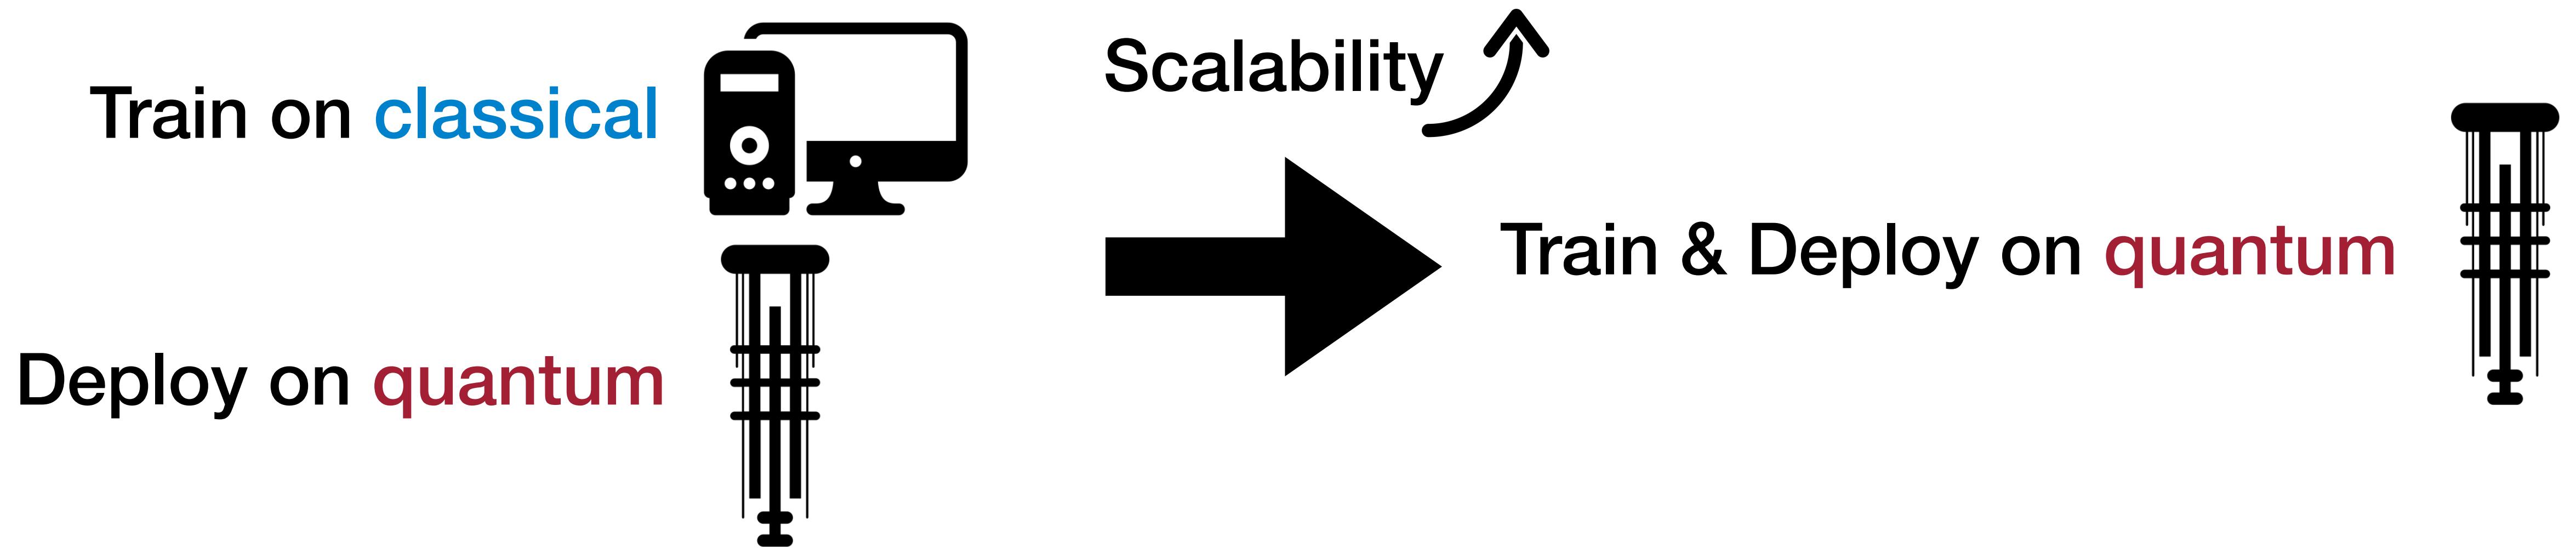


Evaluation

- On classification tasks with QNN



QOC (Quantum On-Chip Training)



[DAC'22] Wang et, at, QOC: Quantum On-Chip Training with Parameter Shift and Gradient Pruning